

Algorithmic Problem Solving

Johan Sannemo

2020

This version of the book is a preliminary draft. Expect to find typos and other mistakes. If you do, please report them to `jsannemo+book@jsannemo.se`. A number of sections and chapters are also unfinished, and a number of problems are not yet uploaded to the judge – this are known issues.

Note: the linked problems are sometimes available on Kattis (<https://open.kattis.com/problems/PROBLEMID>) and sometimes on Kodsport.dev (<https://kodsport.dev/problems/PROBLEMID>). In this particular version, you should try the first one for most chapters.

Contents

Preface	ix
Reading this Book	xi
 I Preliminaries	 1
1 Algorithms and Problems	3
1.1 Computational Problems	3
1.2 Algorithms	5
1.3 Programming Languages	8
1.4 Pseudo Code	9
1.5 The KS.Dev Online Judge	10
 2 Programming in C++	 15
2.1 Development Environments	16
2.2 Hello World!	17
2.3 Variables and Types	20
2.4 Input and Output	25
2.5 Operators	26
2.6 If Statements	28
2.7 For Loops	31
2.8 While Loops	32
2.9 Functions	33
2.10 Structures	37
2.11 Arrays	40
2.12 Lambdas	42
2.13 The Preprocessor	43
2.14 Template	44

3	The C++ Standard Library	47
3.1	vector	47
3.2	Iterators	50
3.3	queue	51
3.4	stack	52
3.5	priority_queue	53
3.6	set and map	54
3.7	Math	55
3.8	Algorithms	57
3.9	Strings	59
3.10	Input/Output	60
4	Implementation Problems	65
5	Time Complexity	83
5.1	The Complexity of Insertion Sort	83
5.2	Asymptotic Notation	86
5.3	NP-complete problems	92
5.4	Other Types of Complexities	92
5.5	The Importance of Constant Factors	92
5.6	Additional Exercises	93
5.7	Chapter Notes	94
6	Data Structures	97
6.1	Dynamic Arrays	97
6.2	Stacks	101
6.3	Queues	102
6.4	Priority Queues	103
6.5	Bitsets	108
6.6	Hash Tables	110
7	Recursion	117
7.1	Recursive Definitions	117
7.2	The Time Complexity of Recursive Functions	120
7.3	Choice	121
7.4	Multidimensional Recursion	126
7.5	Recursion vs. Iteration	127

8	Graph Theory	131
8.1	Graphs	131
8.2	Representing Graphs	135
8.3	Breadth-First Search	138
8.4	Depth-First Search	143
8.5	Trees	146
II	Basics	147
9	Brute Force	149
9.1	Optimization Problems	149
9.2	Generate and Test	150
9.3	Backtracking	154
9.4	Fixing Parameters	162
9.5	Meet in the Middle	165
9.6	Chapter Notes	170
10	Greedy Algorithms	171
10.1	Change-making Problem	171
10.2	Optimal Substructure	172
10.3	Locally Optimal Choices	173
10.4	Scheduling	175
10.5	Huffman Coding	178
11	Dynamic Programming	181
11.1	Best Path in a DAG	181
11.2	Dynamic Programming	183
11.3	Multidimensional DP	186
11.4	Subset DP	188
11.5	Digit DP	190
11.6	Standard Problems	193
12	Divide and Conquer	201
12.1	Inductive Constructions	201
12.2	Merge Sort	208
12.3	Binary Search	210
12.4	Karatsuba's algorithm	217

12.5	Chapter Notes	219
13	Data Structures	221
13.1	Disjoint Sets	221
13.2	Range Queries	224
13.3	Chapter Notes	230
14	Graph Algorithms	231
14.1	Breadth-First Search	231
14.2	Depth-First Search	236
14.3	Weighted Shortest Path	239
14.4	Minimum Spanning Tree	242
14.5	Chapter Notes	244
15	Maximum Flows	247
15.1	Flow Networks	247
15.2	Edmonds-Karp	249
15.3	Applications of Flows	252
15.4	Chapter Notes	254
16	Strings	255
16.1	Tries	255
16.2	String Matching	260
16.3	Chapter Notes	265
17	Combinatorics	267
17.1	The Addition and Multiplication Principles	267
17.2	Permutations	270
17.3	Ordered Subsets	276
17.4	Binomial Coefficients	277
17.5	The Principle of Inclusion and Exclusion	286
17.6	The Pigeon Hole Principle	288
17.7	Invariants	288
17.8	Monovariants	290
17.9	Chapter Notes	295

18 Game Theory	297
18.1 Mathematical Techniques	297
18.2 The Graph Game	299
19 Number Theory	303
19.1 Divisibility	303
19.2 Prime Numbers	310
19.3 The Euclidean Algorithm	319
19.4 Modular Arithmetic	329
19.5 Chinese Remainder Theorem	332
19.6 Euler's totient function	335
20 Competitive Programming Strategy	343
20.1 IOI	343
20.2 ICPC	346
21 Papers	351
21.1 Paper 3	352
III Advanced Topics	355
22 Data Structures	357
22.1 Self-Balancing Trees	357
22.2 Persistent Data Structures	357
22.3 Heavy-Light Decomposition	357
23 Combinatorics	359
23.1 Convolutions	359
24 Strings	361
24.1 Hashing	361
24.2 Dynamic Hashing	370
A Discrete Mathematics	371
A.1 Logic	372
A.2 Sets and Sequences	374
A.3 Sums and Products	377

CONTENTS

Hints	381
Solutions	383
Bibliography	387
Index	390

Preface

Algorithmic problem solving is the art of formulating efficient methods that solve problems of a mathematical nature. From the many numerical algorithms developed by the ancient Babylonians to the founding of graph theory by Euler, algorithmic problem solving has been a popular intellectual pursuit during the last few thousand years. For a long time, it was a purely mathematical endeavor with algorithms meant to be executed by hand. During the recent decades algorithmic problem solving has evolved. What was mainly a topic of research became a mind sport known as competitive programming. As a sport algorithmic problem solving rose in popularity with the largest competitions attracting tens of thousands of programmers. While its mathematical counterpart has a rich literature, there are only a few books on algorithms with a strong problem solving focus.

The purpose of this book is to contribute to the literature of algorithmic problem solving in two ways. First of all, it tries to fill in some holes in existing books. Many topics in algorithmic problem solving lack any treatment at all in the literature – at least in English books. Much of the content is instead documented only in blog posts and solutions to problems from various competitions. While this book attempts to rectify this, it is not to detract from those sources. Many of the best treatments of an algorithmic topic I have seen are as part of a well-written solution to a problem. However, there is value in completeness and coherence when treating such a large area. Secondly, I hope to provide another way of learning the basics of algorithmic problem solving by helping the reader build an intuition for problem solving. A large part of this book describes techniques using worked-through examples of problems. These examples attempt not only to describe the manner in which a problem is solved, but to give an insight into how a thought process might be guided to yield the insights necessary to arrive at a solution.

This book is different from pure programming books and most other algorithm textbooks. Programming books are mostly either in-depth studies of a specific programming language or describe various programming paradigms. A single language is used in this book – C++. The text on C++ exists for the

sole purpose of enabling those readers without prior programming experience to implement the solutions to algorithm problems. Such a treatment is necessarily minimal and teach neither good coding style nor advanced programming concepts. Algorithm textbooks teach primarily algorithm analysis, basic algorithm design, and some standard algorithms and data structures. They seldom include as much problem solving as this book does. The book also falls somewhere between the practical nature of a programming book and the heavy theory of algorithm textbooks. This is in part due to the book's dual nature of being not only about algorithmic problem solving, but also competitive programming to some extent. As such there is more real code and efficient C++ implementations of algorithms included compared to most algorithm books.

Acknowledgments. First and foremost, thanks to Per Austrin who provided much valuable advice and feedback during the writing of this book. Thanks to Simon and Mårten who have competed with me for several years as Omogen Heap. A lot of the knowledge in this book has its roots in you. Finally, thanks to several others who have read through drafts and caught numerous mistakes of my own.

Reading this Book

This book consists of three parts. The first part contains some preliminary background, such as algorithm analysis and programming in C++. With an undergraduate education in computer science most of these chapters are probably familiar to you. It is recommended that you at least skim through the first part since the remainder of the book assumes you know the contents of the preliminary chapters.

The second part makes up most of the material in the book. Some of it should be familiar if you have taken a course in algorithms and data structures. The take on those topics is a bit different compared to an algorithms course. We therefore recommend that you read through even the parts you feel familiar with – in particular those on the basic problem solving paradigms, i.e. brute force, greedy algorithms, dynamic programming and divide & conquer. The chapters in this part are structured so that a chapter builds upon only the preliminaries and previous chapters to the largest extent possible.

In the third part you will find the advanced topics. These are extensions of the topics from the second part. This part is less cohesive, with few dependencies between chapters. You can to a larger degree choose what topics you wish to study, though most of them depend on several of the chapters from the basics.

At the end of the book you can find an appendix with some mathematical background, together with hints and solutions for selected exercises.

When reading this book, know that every problem and technique was chosen with care; every step on the way in a solution added to provide value. Sometimes, this can make the book feel boring – a solution can take a long time tracing out the intuition behind some small step, or show partial solutions that are unused in the final result. At other times, missing a single sentence can leave you with a crucial gap in your knowledge. I have tried to make sure that every sentence written is important; when the book is long-winding, trust that it is useful, and when difficult, endure to make sure you attain the deep understanding I hope this book will be able to provide.

Similarly, the exercises are meant as attempts for you to construct some crucial knowledge on your own. There may be fewer end-of-chapter exercises

than you might be used to in a textbook, and more exercises inlined in chapters. This is because we expect you to solve **all** exercises as part of the reading of the book. Sometimes, the text after an exercise will assume that you read and solved the exercise. The lecture analogue would be the lecturer pausing to ask the class a question; only giving an answer if none is provided by the class. Since this is a book, you are blessed with unlimited time to think in contrast to the lecture setting, where you typically get on the order of minutes. Some exercises took the author on the order of hours to solve at first, so do not feel disparaged if you find them difficult. At the back of the book, you find hints and solutions for selected exercises. If you fail to solve an exercise, first check if it has a hint, and give it another attempt.

This book can also be used to improve your competitive programming skills. Some parts are unique to competitive programming (in particular Chapter 20 on contest strategy). This knowledge is extracted into *competitive tips*:

Competitive Tip

A competitive tip contains some information specific to competitive programming. These can be safely ignored if you are interested only in the problem solving aspect and not the competitions.

The book often refers to exercises from the [Kodsport.dev](https://kodsport.dev) online judge:

Problem 0.1

Problem Name – problemid

The URL of such a problem is <https://kodsport.dev/problems/problemid>.

The C++ code in this book makes use of some preprocessor directives from a template. Even if you are familiar with C++ (or do not wish to learn it) we still recommend that you read through this template (section 2.14) to better understand the C++ code in the book.

Part I

Preliminaries

1 Algorithms and Problems

The greatest technical invention of the last century was probably the digital general purpose computer. It was the start of the revolution which provided us with the Internet, smartphones, tablets, and the computerization of society.

To harness the power of computers we use *programming*. Programming is the art of developing a solution to a *computational problem*, in the form of a set of instructions that a computer can execute. These instructions are what we call *code*, and the language in which they are written a *programming language*. The abstract method that such code describes is what we call an *algorithm*.

The aim of *algorithmic problem solving* is thus to, given a computational problem, devise an algorithm that solves it. One does not necessarily need to complete the full programming process (i.e. write code that implements the algorithm in a programming language) to enjoy solving algorithmic problems. However, it often provides more insight and trains you at finding simpler algorithms to problems.

In this chapter, we begin our journey into algorithmic problem solving by taking a closer look at these concepts and showing a solution to a common problem.

1.1 Computational Problems

A *computational problem* generally consists of two parts. First, it needs an *input description*, such as “a sequence of integers”, “a text string”, or some other kind of mathematical object. Using this input, we have a goal which we want to accomplish defined by an *output description*. For example, a computational problem might require us to sort a given sequence of integers. This particular problem is called the *Sorting Problem*:

Sorting

Your task is to sort a sequence of integers in ascending order, i.e. from the lowest to the highest.

Input

The input is a sequence of N integers a_0, a_1, \dots, a_{N-1} .

Output

Output a permutation a' of the sequence a , such that $a'_0 \leq a'_1 \leq \dots \leq a'_{N-1}$.

A particular input to a computational problem is called an *instance* of the problem. To the sorting problem, the sequence 3, 6, 1, -1, 2, 2 would be an instance. The correct output for this particular problem would be -1, 1, 2, 2, 3, 6.

Exercise 1.1. If you were given cards with 5 different integers 1 and 1 000 000 written on them, how would you sort them in ascending order? How would your approach change if you had 30 integers? 1000? 1 000 000?

Some variations of this problem format appears later (such as problems without inputs) but in general this is what the problems look like.

Competitive Tip

Problem statements sometimes contain huge amounts of text. Skimming through the input and output sections before any other text in a problem can often give you a quick idea about its topic and difficulty. This helps in determining what problems to solve first when posed with a large number of problems and little time.

Exercise 1.2. What is the input and output descriptions for the following computational problems?

- 1) Compute the greatest common divisor (see Def. 19.5, page 319 if you are not familiar with the concept) of two numbers.
- 2) Find a root (i.e. a zero) of a polynomial.
- 3) Multiply two numbers.

Exercise 1.3. Consider the following problem. I am thinking of an integer between 1 and 100. Your task is to find this number by giving me integers, one at a time. I will tell you whether the given integer is higher, lower or equal to x .

This is an *interactive*, or *online*, computational problem. How would you describe the input and output to it? Why do you think it is called interactive?

1.2 Algorithms

Algorithms are solutions to computational problems. They define methods that use the input to a problem in order to produce the correct output. A computational problem can have many solutions. Efficient algorithms to solve the sorting problem form an entire research area! Let us look at one possible sorting algorithm, called selection sort, as an example.

Selection Sort

We construct the answer, the sorted sequence, iteratively one element at a time, starting with the smallest.

Assume that we have chosen and sorted the K smallest elements of the original sequence. Then, the smallest unchosen element remaining in that sequence must be the $(K + 1)$ 'st smallest element of the original sequence. Thus, by finding the smallest element among those that remain we know what the $(K + 1)$ 'st element of the sorted sequence is. By appending this element to the already sorted K smallest elements we get the sorted $K + 1$ smallest elements of the output.

If we repeat this process N times, the result is the N numbers of the original sequence, but sorted. ■

You can see this algorithm performed on our previous example instance (the sequence 3, 6, 1, -1, 2, 2) in Figures 1.1a-1.1f.

So far, we have been vague about what exactly an algorithm is. Looking at our Selection Sort example, we do not have any particular structure or rigor in the description of our method. There is nothing inherently wrong with describing algorithms this way. It is easy to understand and gives the writer an opportunity to provide context as to why certain actions are performed, making the correctness of the algorithm more obvious. The main downsides of such a description are ambiguity and a lack of detail.

Until an algorithm is described in sufficient detail, it is possible to accidentally abstract away operations we may not know how to perform behind a few English words. As a somewhat contrived example, our plain text description of selection sort includes actions such as “choosing the smallest number of a sequence”. While such an operation may seem very simple to us humans, algorithms are generally constructed with regards to some kind of computer. Unfortunately, computers can not map such English expressions to their code counterparts yet. Instructing a computer to execute an algorithm thus requires us to formulate our

3	6	1	-1	2	2
---	---	---	----	---	---

(a) Originally, we start out with the unsorted sequence (3, 6, 1, -1, 2, 2).

-1	3	6	1	2	2
----	---	---	---	---	---

(b) The smallest element of the sequence is -1, so this is the first element of the sorted sequence.

-1	1	3	6	2	2
----	---	---	---	---	---

(c) We find the next element of the output by removing the -1 and finding the smallest remaining element - in this case 1.

-1	1	2	3	6	2
----	---	---	---	---	---

(d) Here, there is no unique smallest element. We can choose any of the two 2's in this case.

-1	1	2	2	3	6
-1	1	2	2	3	6

(e) The next two elements chosen will be a 2 and a 3.

-1	1	2	2	3	6
----	---	---	---	---	---

(f) Finally, we choose the last remaining element of the input sequence - the 6. This concludes the sorting of our sequence.

Figure 1.1: An example execution of selection sort.

algorithm in steps small enough that even a computer knows how to perform them. In this sense, a computer is rather stupid.

The English language is also ambiguous. We are sloppy with references to “this variable” and “that set”, relying on context to clarify meaning for us. We use confusing terminology and frequently misunderstand each other. Real code does not have this problem. It forces us to be specific with what we mean. However, as all programmers know, we often manage to construct highly specific algorithms that *do the wrong thing* due to our own erroneous thought processes.

We will generally describe our algorithms in a representation called pseudo code (Section 1.4), accompanied by an online exercise to implement the code. Sometimes, we will instead give explicit code that solves a problem. This will be the case whenever an algorithm is very complex, or care must be taken to make the implementation efficient. The goal is that you should get to practice understanding pseudo code, while still ending up with correct implementations

of the algorithms (thus the online exercises).

Exercise 1.4. Do you know any algorithms, for example from school? (Hint: you use many algorithms to solve certain arithmetic and algebraic problems, such as those in Exercise 1.2.)

Exercise 1.5. In Exercise 1.1, you were asked to come up with your own approaches to the sorting problem. Attempt to write them down formally as descriptions of algorithms.

Exercise 1.6. Construct an algorithm that solves the guessing problem in exercise 1.3 using as few questions as possible. How many questions does it use?

Correctness

One subtle, albeit important, point that we glossed over is what it means for an algorithm to actually be *correct*.

There are two common notions of correctness – *partial correctness* and *total correctness*. Partial correctness requires an algorithm to, upon termination, have produced an output that fulfills all the criteria laid out in the output description. Total correctness additionally requires an algorithm to finish within finite time. When we talk about correctness of our algorithms later on, we generally focus on the partial correctness. Termination is instead proved implicitly, as we consider a more granular measure of efficiency (called *time complexity*, in Chapter 5) than just finite termination. This measure implies the termination of the algorithm, completing the proof of total correctness.

Proving that the selection sort algorithm finishes in finite time is quite easy. It performs one iteration of the selection step for each element in the original sequence (which is finite). Furthermore, each such iteration can be performed in finite time by looking at each remaining element of the selection when finding the smallest one. The remaining sequence is a subsequence of the original one and is therefore also finite.

Proving that the algorithm produces the correct output is a bit more difficult to prove formally. The main idea behind a formal proof is contained within our description of the algorithm itself.

While this definition seems clear enough – our algorithm should simply do what the problem asks of it! – we will compromise on both conditions at later points in the book. Generally, we are satisfied with an algorithm terminating in

expected finite time or answering correctly with, say, probability 0.75 for every input. Similarly, we are sometimes happy to find an *approximate* solution to a problem. What this means more concretely will become clear in due time when we study such algorithms.

Competitive Tip

Proving your algorithm correct is sometimes quite difficult. In a competition, a correct algorithm is correct even if you cannot prove it. If you have an idea you *think* is correct it may be worth testing. This is not a strategy without problems though, since it makes distinguishing between an incorrect algorithm and an incorrect implementation even harder.

Exercise 1.7. Prove the correctness of your algorithm to the guessing problem from Exercise 1.6 and your sorting algorithms from Exercise 1.5.

Exercise 1.8. Why would an algorithm that is correct with e.g. probability 0.75 still be very useful to us?

Why is it important that such an algorithm is correct with probability 0.75 on *every* problem instance, instead of always being correct for 75% of all cases?

1.3 Programming Languages

The purpose of *programming languages* is to formulate methods at a level of detail where a computer could perform them. While we in textual descriptions of methods are often satisfied with describing *what* we wish to do, programming languages require considerably more constructive descriptions. Computers are quite basic creatures compared to us humans. They only understand a very limited set of instructions such as adding numbers, multiplying numbers, or moving data around within its memory. The syntax of programming languages often seems a bit arcane at first, but it grows on you with coding experience.

To complicate matters further, programming languages themselves define a spectrum of expressiveness. On the lowest level, programming deals with electrical current in your processor. Current above or below a certain threshold is used to represent the binary digits 0 and 1. Above these circuit-level electronics lies a processor's own programming, often called *microcode*. Using this, a processor implements *machine code*, such as the x86 instruction set. Machine code is often written using a higher-level syntax called *Assembly*. While some code is written in this rather low-level language, we mostly abstract away details of them in high-level languages such as C++ (this book's language of choice).

This knowledge is somewhat useless from a problem solving standpoint, but intimate knowledge of how a computer works is of high importance in software engineering, and is occasionally helpful in programming competitions. Therefore, you should not be surprised about certain remarks relating to these low-level concepts.

These facts also provide some motivation for why we use something called *compilers*. When programming in C++ we can not immediately tell a computer to run our code. As you now know, C++ is code at a higher level than what the processor of a computer can run. A compiler takes care of this problem by translating our C++ code into *machine code* that the processor knows how to handle. It is a program of its own and takes the code files we write as input and produces *executable* files that we can run on the computer. The process and purpose of a compiler is somewhat like what we do ourselves when translating a method from English sentences or our own thoughts into the lower level language of C++.

1.4 Pseudo Code

Somewhere in between describing algorithms in English text and in a programming language we find something called *pseudo code*. As hinted by its name it is not quite real code. The instructions we write are not part of the programming language of any particular computer. The point of pseudo code is to be independent of the computer it is implemented on. Instead, it tries to convey the main points of an algorithm in a detailed manner so that it can easily be translated into any particular programming language. Secondly, we sometimes fall back to the liberties of the English language. At some point, we may decide that “choose the smallest number of a sequence” is clear enough for our audience.

With an explanation of this distinction in hand, let us look at a concrete example of pseudo code. The honor of being an example again falls upon selection sort, now described in pseudo code:

```
1: procedure SELECTIONSORT(sequence A)
2:   Let  $A'$  be an empty sequence
3:   while A is not empty do
4:      $minIndex \leftarrow 0$ 
5:     for every element  $A_i$  in A do
```

```
6:         if  $A_i < A_{minIndex}$  then
7:              $minIndex \leftarrow i$ 
8:         Append  $A_{minIndex}$  to  $A'$ 
9:         Remove  $A_{minIndex}$  from  $A$ 
10:    return the sequence  $A'$ 
```

Pseudo code reads somewhat like our English language variant of the algorithm, except the actions are broken down into much smaller pieces. Most of the constructs of our pseudo code are more or less obvious. The notation $variable \leftarrow value$ is how we denote an *assignment* in pseudo code. For those without programming experience, this means that the variable named *variable* now takes the value *value*. Pseudo code appears when we try to explain some part of a solution in great detail but programming language specific aspects would draw attention away from the algorithm itself.

Competitive Tip

In team competitions where a team only have a single computer, a team will often have solved problems waiting to be coded. Writing pseudo code of the solution to one of these problems while waiting for computer time is an efficient way to parallelize your work. This can be practiced by writing pseudo code on paper even when you are solving problems by yourself.

Exercise 1.9. Write pseudo code for your algorithm to the guessing problem from Exercise 1.6.

1.5 The KS.Dev Online Judge

Most of the exercises in this book exist as problems on the **KS.Dev** web system. You can find it at <https://kodsport.dev>. KS.Dev is a so called **online judge**. It contains a large collection of computational problems, and allows you to submit a program you have written that purports to solve a problem. KS.Dev will then run your program on a large number of predetermined instances of the problem called the problem's **test data**.

Problems on an online judge include some additional information compared to our example problem. Since actual computers only have a finite amount of time and memory, the amount of these resources available to our programs are limited when solving an instance of a problem. This also means that the size of inputs to a problem need to be constrained as well, or else the resource limits for

a given problem would not be obtainable – an arbitrarily large input generally takes arbitrarily large time to process, even for a computer. A more complete version of the sorting problem as given in a competition could look like this:

Sorting

Time: 1s, memory: 1MB

Your task is to sort a sequence of integers in ascending order, i.e. from the lowest to the highest.

Input

The input is a sequence of N integers ($1 \leq N \leq 1000$) a_0, a_1, \dots, a_{N-1} ($|a_i| \leq 10^9$).

Output

Output a permutation a' of the sequence a , such that $a'_0 \leq a'_1 \leq \dots \leq a'_{N-1}$.

If your program exceeds the allowed resource limits (i.e. takes too much time or memory), crashes, or gives an invalid output, KS.Dev will tell you so with a *rejected* judgment. There are many kinds of rejected judgments, such as *Wrong Answer*, *Time Limit Exceeded*, and *Run-time Error*. These mean your program gave an incorrect output, took too much time, and crashed, respectively. Assuming your program passes all the instances, it will be given the *Accepted* judgment.

Note that getting a program accepted by KS.Dev is not the same as having a correct program – it is a necessary but not sufficient criterion for correctness. This is also a fact that can sometimes be exploited during competitions by writing a knowingly incorrect solution that one thinks will pass all test cases that the judges of the competitions designed.

We strongly recommend that you get a (free) account on KS.Dev so that you can follow along with the book's exercises.

Exercise 1.10. Register an account on KS.Dev.

Many other online judges exists, such as:

- Kattis (<https://open.kattis.com>)
- Codeforces (<http://codeforces.com>)
- CSAcademy (<https://csacademy.com>)

- AtCoder (<https://atcoder.jp>)
- TopCoder (<https://topcoder.com>)
- HackerRank (<https://hackerrank.com>)

Chapter Exercises

Exercise 1.11. Pick two sorting algorithms from Wikipedia's list of sorting algorithms: https://en.wikipedia.org/wiki/Category:Sorting_algorithms. Try to understand them and their proof of correctness. Use them by hand to sort the integers 5, 1, 2, 7, 5, 6, 2, 9.

Exercise 1.12. Consider the following problems:

Palindrome

A word is a palindrome if it reads the same forwards and backwards, for example *tacocat*, *madam*, or *abba*. Determine if a word is a palindrome.

Input

The input consists of a single word, containing only lowercase letters a-z.

Output

Output *yes* if the word is a palindrome and *no* otherwise.

Primality

We call an integer $n > 1$ a *prime* if its only positive divisors are 1 and n . Determine if a particular integer is a prime.

Input

The input consists of a single integer $n > 1$.

Output

Output *yes* if the number n was a prime and *no* otherwise.

For each of them,

1. devise an algorithm to solve it,
2. formalize the algorithm and write it down in pseudo code, and
3. prove the correctness of the algorithm.

Chapter Notes

The introductions given in this chapter are very bare, mostly stripped down to what you need to get by when solving algorithmic problems.

Many other books delve deeper into the theoretical study of algorithms than we do, in particular regarding subjects not relevant to algorithmic problem solving. *Introduction to Algorithms* [7] is a rigorous introductory text book on algorithms with both depth and breadth.

For a gentle introduction to the technology that underlies computers, *CODE* [23] is a well-written journey from the basics of bits and bytes all the way up to assembly code and operating systems. It requires no knowledge of programming to read.

2 Programming in C++

In this chapter we learn the basics of the C++ programming language. This language is the most common programming language within the competitive programming community for a few reasons (aside from C++ being a popular language in general). Programs coded in C++ are generally somewhat faster than those written in most other competitive programming languages. There are also many routines in the accompanying standard code libraries that are useful when implementing algorithms.

Of course, no language is without downsides. C++ is a bit difficult to learn as your first programming language to say the least. Its error management is unforgiving, often causing erratic behavior in programs instead of crashing with an error. Programming certain things become quite verbose, compared to many other languages.

After bashing the difficulty of C++, you might ask if it really is the best language in order to get started with algorithmic problem solving. While there certainly are simpler languages we believe that the benefits outweigh the disadvantages in the long term even though it demands more from you as a reader. Either way, it is definitely the language we have the most experience of teaching problem solving with.

When you study this chapter, you will see a lot of example code. **Type this code and run it.** We can not really stress this point enough. Learning programming from scratch – in particular a complicated language such as C++ – is not possible unless you try the concepts yourself. Additionally, we strongly recommend that you do **every** exercise in this chapter, even moreso than in the other chapters.

Finally, know that our treatment of C++ is minimal. We do not explain all the details behind the language, nor do we teach good coding style or general software engineering principles. In fact, we frequently make use of bad coding practices. If you want to delve deeper, you can find more resources in the chapter notes.

2.1 Development Environments

Before we get to the juicy parts of C++ you need to install a compiler for C++ and (optionally) a code editor.

We recommend the editor *Visual Studio Code*. The installation procedure varies depending on what operating system you use. We provide them for Windows, Ubuntu and macOS. If you choose to use some other editor, compiler or operating system you must find out how to perform the corresponding actions (such as compiling and running code) yourself.

Note that instructions like these tend to rot, with applications disappearing from the web, operating systems changing names, and so on. In that case, you are on your own and have to find instructions by yourself.

Windows

Installing a C++ compiler is somewhat complicated in Windows. We recommend installing the Mingw-w64 compiler from <http://www.mingw-w64.org/>.

After installing the compiler, you can download the installer for Visual Studio Code from <https://code.visualstudio.com/>.

Ubuntu

On Ubuntu, or similar Linux-based operating systems, you need to install the GCC C++ compiler, which is the most popular compiled for Linux-based systems. It is called `g++` in most package managers and can be downloaded with the command `sudo apt-get install g++`. After installing the compiler, you can download the installer for Visual Studio Code from <https://code.visualstudio.com/>. Choose the deb installer.

macOS

When using macOS, you first need to install the Clang compiler by installing *Xcode* from the Mac App Store. This is also a code editor, but the compiler is bundled with it.

After installing the compiler, you can download the installer for Visual Studio Code from <https://code.visualstudio.com/>. It is available as a normal macOS package for installation.

Installing the C++ tools

Now that you have installed the compiler and Visual Studio Code, you need to install the C++ plugin for Visual Studio Code. You can do this by opening the program, launching Quick Open (using Ctrl+P), typing `ext install ms-vscode.cpptools`, and pressing Enter. Then, launch Quick Open again, but this time type `ext install formulahendry.code-runner` instead.

The tools need to be configured a bit. Press Ctrl+Shift+P and search for Open Settings. Select Open Settings (JSON) in the list. Here, enter the following configuration:

```
{
  "code-runner.runInTerminal": true,
  "code-runner.saveAllFilesBeforeRun": true,
  "code-runner.executorMap": {
    "cpp": "cd $dir && g++ $fileName -fsanitize=undefined,address -o ↵
          $fileNameWithoutExt -std=c++17 -Wall && $dir$fileNameWithoutExt"
  }
}
```

Note that the line ending with ↵ denotes that the text on the following line should be on the same line.

Then, restart your editor again.

2.2 Hello World!

Now that you have a compiler and editor ready, it is time to learn the basic structure of a C++ program. The classical example of a program when learning a new language is to print the text `Hello World!`. We also solve our first KS.Dev problem in this section.

Exercise 2.1. Throughout this chapter, you will learn many concepts within C++. We recommend that you create a notebook (for example in a file on your computer) where you write down how the different constructs are used when programming to keep as a reference for later.

Start by opening Visual Studio Code and create a new file by going to File ⇒ New File. Save the file as `hello.cpp` by pressing Ctrl+S. Make sure to save it somewhere you can find it.

Now, type the code from Listing 2.1 into your editor.

Listing 2.1 Hello World!

```
1  #include <iostream>
2
3  using namespace std;
4
5  int main() {
6      // Print Hello World!
7      cout << "Hello World!" << endl;
8  }
```

To run the program in Visual Studio Code, you press Ctrl+Alt+N. A tab below your code named `TERMINAL` containing the text `Hello World!` should appear. If no window appears, you probably mistyped the program.

Coincidentally, KS.Dev happens to have a problem whose output description dictates that your program should print the text `Hello World!`. How convenient. This is a great opportunity to get familiar with KS.Dev.

Problem 2.1

Hello World! – hello

When you submit your solution, KS.Dev grades it and give you its judgment. If you typed everything correctly, KS.Dev tells you it got Accepted. Otherwise, you probably got Wrong Answer, meaning your program output the wrong text (and you mistyped the code).

Now that you have managed to solve the problem, it is time to talk a bit about the code you typed.

The first line of the code,

```
#include <iostream>
```

is used to include the `iostream` – *input and output stream* – file from the so-called *standard library* of C++. The standard library is a large collection of ready-to-use algorithms, data structures, and other routines which you can use when coding. For example, there are sorting routines in the C++ standard library, meaning you do not need to implement your own sorting algorithm when coding solutions.

Later on, we will see other useful examples of the standard library and include many more files. The `iostream` file contains routines for reading and writing data to your screen. Your program used code from this file when it printed `Hello World!` upon execution.

On some platforms, there is a special include file called `bits/stdc++.h`. This file includes the entire standard library. You can check if it is available on your platform by including it using

```
#include <bits/stdc++.h>
```

in the beginning of your code. If your program still compiles, you can use this and not include anything else. By using this line you do not have to care about including any other files from the standard library which you wish to use.

The third line,

```
using namespace std;
```

tells the compiler that we wish to use code from the **standard** library. If we did not use it, we would have to specify this every time we used code from the standard library later in our program by prefixing what we use from the library by `std::` (for example `std::cout`).

The fifth line defines our **main function**. When we instruct the computer to run our program the computer starts looking at this point for code to execute. The first line of the main function is thus where the program starts to run with further lines in the function executed sequentially. Later on we learn how to define and use additional functions as a way of structuring our code. Note that the code in a function – its *body* – must be enclosed by curly brackets. Without them, we would not know which lines belonged to the function.

On line 6, we wrote a **comment**

```
// Print Hello World!
```

Comments are explanatory lines which are not executed by the computer. The purpose of a comment is to explain what the code around it does and why. They begin with two slashes `//` and continue until the end of the current line.

It is not until the seventh line that things start happening in the program. We use the standard library utility `cout` to print text to the screen. This is done by writing e.g.:

```
cout << "this is text you want to print. ";  
cout << "you can " << "also print " << "multiple things. ";  
cout << "to print a new line" << endl << "you print endl" << endl;  
cout << "without any quotes" << endl;
```

Lines that do things in C++ are called **statements**. **Note the semi colon** at the end of the line! Semi colons are used to specify the end of a statement, and are mandatory.

Exercise 2.2. Must the main function be named `main`? What happens if you changed `main` to something else and try to run your program?

Exercise 2.3. Play around with `cout` a bit, printing various things. For example, you can print a pretty haiku.

2.3 Variables and Types

The Hello World! program is boring. It only prints text – seldom the only necessary component of an algorithm (aside from the Hello World! problem on KS.Dev). We now move on to a new but hopefully familiar concept.

When we solve mathematical problems, it often proves useful to introduce all kinds of names for known and unknown values. Math problems often deal with classes of N students, ice cream trucks with velocity v_{car} km/h, and candy prices of p_{candy} \$/kg.

This concept naturally translates into C++ but with a twist. In most programming languages, we first need to say what *type* a variable has! We do not bother with this in mathematics. We say “let $x = 5$ ”, and that is that. In C++, we need to be a bit more verbose. We must write that “I want to introduce a variable x now. It is going to be an integer – more specifically, 5”. Once we have decided what kind of value x will be (in this case integer) it will always be an integer. We cannot just go ahead and say “oh, I’ve changed my mind. $x = 2.5$ now!” since 2.5 is of the wrong type (a decimal number rather than an integer).

Listing 2.2 Variables

```
1  #include <iostream>
2  using namespace std;
3
4  int main() {
5      int five = 5;
6      cout << five << endl;
7      int seven = 7;
8      cout << seven << endl;
9      five = seven + 2; // = 7 + 2 = 9
10     cout << five << endl;
11     seven = 0;
12     cout << five << endl; // five is still 9
13     cout << 5 << endl; // we print the integer 5 directly
14 }
```

Another major difference is that variables in C++ are not tied to a single value for the entirety of their lifespans. Instead, we are able to modify the value which our variables have using something called *assignment*. Some languages

does not permit this, preferring their variables to be *immutable*.

In Listing 2.2 we demonstrate how variables are used in C++. Type this program into your editor and run it. What is the output? What did you expect the output to be?

The *first* time we use a variable in C++ we must decide what kind of values it may contain. This is called *declaring* the variable of a certain type. For example the statement

```
int five = 5;
```

declares an integer variable `five` and assigns the value 5 to it. The `int` part is C++ for integer and is what we call a type. After the type, we write the name of the variable – in this case `five`. Finally, we may assign a value to the variable. Note that further use of the variable never include the `int` part. We declare the type of a variable once and only once.

Later on in Listing 2.2 we decide that 5 is a somewhat small value for a variable called `five`. We can change the value of a variable by using the *assignment operator* – the equality sign `=`. The assignment

```
five = seven + 2;
```

states that from now on the variable `five` should take the value given by the expression `seven + 2`. Since (at least for the moment) `seven` has the value 7 the expression evaluates to $7 + 2 = 9$. Thus `five` will actually be 9, explaining the output we get from line 12.

On line 14 we change the value of the variable `seven`. Note that line 15 still prints the value of `five` as 9. Some people find this model of assignment confusing. We first performed the assignment `five = seven + 2;`, but the value of `five` did not change with the value of `seven`. This is mostly an unfortunate consequence of the choice of `=` as operator for assignment. One could think that “once an equality, always an equality” – that the value of `five` should always be the same as the value of `seven + 2`. **This is not the case.** An assignment sets the value of the variable on the left hand side to the value of the expression on the right hand side at a particular moment in time, nothing more.

The snippet also demonstrates how to print the value of a variable on the screen – we `cout` it the same way as with text. This also clarifies why text needs to be enquoted. Without quotes, we can not distinguish between the text string “hi” and the variable `hi`.

Note that it is possible to **declare a variable without assigning a value to it**. When this is done, the variable may receive a *random* value instead. This is useful when you immediately want to assign a value to variable that the user can input (see the next Section 2.4).

Exercise 2.4. What values will the variables *a*, *b*, and *c* have after executing the following code:

```
int a = 4;
int b = 2;
int c = 7;
b = a + c;
c = b - 2;
a = a + a;
b = b * 2;
c = c - c;
```

Here, the operator `-` denotes subtraction and `*` represents multiplication. Once you have arrived at an answer, type this code into the main function of a new program and print the values of the variables. Did you get it right?

Exercise 2.5. What happens when an integer is divided by another integer? Try printing the result of the following divisions: $\frac{5}{3}$, $\frac{15}{5}$, $\frac{2}{2}$, $\frac{7}{2}$, $\frac{-7}{2}$, and $\frac{7}{-2}$.

Exercise 2.6. C++ allows declarations of immutable (constant) variables, using the keyword `const`. For example

```
const int FIVE = 5;
```

What happens if you try to perform an assignment to such a variable?

There are many other types than `int`. We have seen one (although without its correct name), the type for text. You can see some of the most common types in Listing 2.3.

The text data type is called `string`. Values of this type must be enclosed with double quotes. If we want to include an actual quote character in a string, we type `\`.

There exists a data type containing one single letter, the `char`. Such a value is surrounded by single quotes. The `char` value containing the single quote is written `'\'`, similarly to how we included double quotes in strings.

Then comes the `int`, which we discussed earlier. The `long long` type contains integers just like the `int` type. They differ in how large integers they can contain. An `int` can only contain integers between -2^{31} and $2^{31} - 1$ while a `long long` extends this range to -2^{63} to $2^{63} - 1$.

Listing 2.3 Types

```
1  string text = "Johan said: \"heya!\" ";
2  cout << text << endl;
3
4  char letter = '@';
5  cout << letter << endl;
6
7  int number = 7;
8  cout << number << endl;
9
10 long long largeNumber = 888888888888LL;
11 cout << largeNumber << endl;
12
13 double decimalNumber = 513.23;
14 cout << decimalNumber << endl;
15
16 bool thisisfalse = false;
17 bool thisistrue = true;
18 cout << thisistrue << " and " << thisisfalse << endl;
```

Exercise 2.7. Since `\` is used to include a double quote in a string, we can not include backslashes in a string like any other character. For example, how would you output the verbatim string `\`? Find out how to include a literal backslash in a string (for example by searching the web or thinking about how we included the different quote characters).

Exercise 2.8. Write a program that assigns the minimum and maximum values of an `int` to a `int` variable `x`. What happens if you increment or decrement this value using `x = x + 1`; or `x = x - 1`; respectively and print its new value?

Competitive Tip

One of the most common sources for errors in code is trying to store an integer value outside the range of the type. Always make sure your values fit inside the range of an `int` if you use it. Otherwise, use `long longs`!

One of the reasons for why we do not simply use `long long` all the time is that some operations involving `long longs` can be slower using `ints` under certain conditions.

Next comes the `double` type. This type represents decimal numbers. Note that the decimal sign in C++ is a **dot**, not a comma. There is also another similar type called the `float`. The difference between these types are similar to that of

the `int` and `long long`. A `double` can represent “more” decimal numbers than a `float`. This may sound weird considering that there is an infinite number of decimal numbers even between 0 and 1. However, a computer can clearly not represent every decimal number – not even those between 0 and 1. To do this, it would need infinite memory to distinguish between these numbers. Instead, they represent a limited set of numbers – with about 15 significant digits, and about 308 zeroes to the left or right of those digits. Floats have fewer significant digits, and can only represent smaller numbers.

The last of our common types is the `bool` (short for *boolean*). This type can only contain one of two values – it is either `true` or `false`. While this may look useless at a first glance, the importance of the boolean becomes apparent later.

Exercise 2.9. In the same way the integer types had a valid range of values, a `double` cannot represent arbitrarily large values. Find out what the minimum and maximum values a `double` can store is.

C++ has a construct called the *typedef*, or type definition. It allows us to give certain types new names. Since typing `long long` for every large integer variable is very annoying, we could use a type definition to alias it with the much shorter `ll` instead. Such a `typedef` statement looks like this:

```
typedef long long ll;
```

On every line after this statement, we can use `ll` just as if it were a `long long`:

```
ll largeNumber = 888888888888LL;
```

Sometimes we use types with very long names but do not want to shorten them using type definitions. This could be the case when we use many different such types and typedefing them would take unnecessarily long time. We then resort to using the `auto` “type” instead. If a variable is declared as `auto` and assigned a value at the same time its type is inferred from that of the value. This means we could write

```
auto str = 123;
```

instead of

```
int str = 123;
```

2.4 Input and Output

In previous sections we occasionally printed things onto our screen. To spice our code up a bit we are now going to learn how to do the reverse – reading values which we type on our keyboards into a running program! When we run a program we may type things in the window that appears. Pressing the Enter key allows the program to read what we have written so far.

Reading input data is done just as you would expect, almost entirely symmetric to printing output. Instead of `cout` we use `cin`, and instead of `<< variable` we use `>> variable`, i.e.

```
cin >> variable;
```

Type in the program from Listing 2.4 to see how it works.

Listing 2.4 Input

```
1  #include <iostream>
2  using namespace std;
3
4  int main() {
5      string name;
6      cout << "What's your first name?" << endl;
7      cin >> name;
8      int age;
9      cout << "How old are you?" << endl;
10     cin >> age;
11     cout << "Hi, " << name << "!" << endl;
12     cout << "You are " << age << " years old." << endl;
13 }
```

Exercise 2.10. What happens if you type an invalid input, such as your first name instead of your age?

When the program reads input into a string variable it only reads the text until the first whitespace.

We revisit more advanced input and output concepts in Section 3.10 about the standard library. For example, we learn how to read entire lines of text and not only single words.

Problem 2.2*Echo* – echo**Note:** only solve part 1, receiving 1/2 points

2.5 Operators

Earlier we saw examples of what is called *operators*, such as the assignment operator `=`, and the arithmetic operators `+` `-` `*` `/`, which stand for addition, subtraction, multiplication and division. They work almost like they do in mathematics, and allow us to write code such as the one in Listing 2.5.

Exercise 2.11. Type in Listing 2.5 and test it on a few different values. Most importantly, test:

- $b = 0$
- Negative values for a and/or b
- Values where the expected result is outside the valid range of an `int`

As you probably noticed, the division operator of C++ performs so-called *integer division*. This means the answer is rounded to an integer (towards 0). Hence $7 / 3 = 2$, with remainder 1, and $-7 / 3 = -2$.

Exercise 2.12. If division rounds down towards zero, how do you compute $\frac{x}{y}$ rounded to an integer *away from* zero?

The snippet also introduces the *modulo* operator, `%`. It computes the remainder of the first operand when divided by the second. As an example, $7 \% 3 = 1$. Different programming languages have different behaviours regarding modulo operations on negative integers. In particular, the value of a modulo operation can be negative when including negative operands.

In case we want the answer to be a decimal number instead of performing integer division one of the operands must be a `double` (Listing 2.6).

We end this section with some shorthand operators. Check out Listing 2.7 for some examples. Each arithmetic operator has a corresponding combined assignment operator. Such an operator, e.g. `a += 5;`, is equivalent to `a = a + 5;` They act as if the variable on the left hand side is also the left hand side of the corresponding arithmetic operator and assign the result of this computation to said variable. Hence, the above statement increases the variable `a` with 5.

Listing 2.5 Operators

```
1  #include <iostream>
2  using namespace std;
3
4  int main() {
5      int a = 0;
6      int b = 0;
7      cin >> a >> b;
8      cout << "Sum: " << (a + b) << endl;
9      cout << "Difference: " << (a - b) << endl;
10     cout << "Product: " << (a * b) << endl;
11     cout << "Quotient: " << (a / b) << endl;
12     cout << "Remainder: " << (a % b) << endl;
13 }
```

Listing 2.6 Division Operators

```
1  int a = 6;
2  int b = 4;
3  cout << (a / b) << endl;
4
5  double aa = 6.0;
6  double bb = 4.0;
7  cout << (aa / bb) << endl;
```

It turns out that addition and subtraction with 1 is a fairly common operation. So common, in fact, that additional operators were introduced into C++ for this purpose of saving an entire character compared to the highly verbose `+=1` operator. These operators consist of two plus signs or two minus signs. For instance, `a++` increments the variable by 1.

We sometimes use the fact that these expressions also evaluate to a value. Which value this is depends on whether we put the operator before or after the variable name. By putting `++` before the variable, the value of the expression will be the incremented value. If we put it afterwards we get the original value. To get a better understanding of how this works it is best if you type the code in Listing 2.7 in yourself and analyze the results.

We end the discussion on operators by saying something about *operator precedence*, i.e. the order in which operators are evaluated in expressions. In mathematics, there is a well-defined precedence: brackets go first, then

Listing 2.7 Shorthand Operators

```
1  int num = 0;
2  num += 1;
3  cout << num << endl;
4  num *= 2;
5  cout << num << endl;
6  num -= 3;
7  cout << num << endl;
8  cout << num++ << endl;
9  cout << num << endl;
10 cout << ++num << endl;
11 cout << num << endl;
12 cout << num-- << endl;
13 cout << num << endl;
```

exponents, followed by division, multiplication, addition, and subtraction. Furthermore, most operations (exponents being a notable exception) have left-to-right associativity so that $5 - 3 - 1$ equals $((5 - 3) - 1) = 1$ rather than $(5 - (3 - 1)) = 3$. In C++, there are a lot of operators, and knowing precedence rules can easily save you from bugs in your future code.

Exercise 2.13. Research online C++ documentation on operator precedence to determine what the expression

$2 * 4 - 7 * 2 \% 4 / 2$

evaluates to in C++. Run it as a program to see if you got it correct.

Problem 2.3

Two-sum – twosum

Triangle Area – triarea

Bijele – bijele

Digit Swap – digitswap

Pizza Crust – pizzacrust

R2 – r2

2.6 If Statements

In addition to assignment and arithmetic there are a large number of *comparison operators*. These compare two values and evaluate to a bool value depending on the result of the comparison (see Listing 2.8).

Listing 2.8 Comparison Operators

```
1  a == b // check if a equals b
2  a != b // check if a and b are different
3  a > b // check if a is greater than b
4  a < b // check if a is less than b
5  a <= b // check if a is less than or equal to b
6  a >= b // check if a is greater than or equal to b
```

A `bool` can also be *negated* using the `!` operator. So the expression `!false` (which we read as “not false”) has the value `true` and vice versa `!true` evaluates to `false`. The operator works on any boolean expressions, so that if `b` would be a boolean variable with the value `true`, then the expression `!b` evaluates to `false`.

There are two more important boolean operators. The *and* operator `&&` takes two boolean values and evaluates to `true` if and only if both values were `true`. Similarly, the *or* operator `||` evaluates to `true` if and only if at least one of its operands were `true`.

Exercise 2.14. Write a program that reads two integers as input, and prints the result of the different comparison operators from Listing 2.8, e.g

```
cout << (a == b) << endl;
```

Note the parenthesis used due to operator precedence!

A major use of boolean variables is in conjunction with *if* statements (also called *conditional* statements). They come from the necessity of executing certain lines of code *if* (and only if) some condition is true. Let us write a program that takes an integer as input, and tells us whether it is odd or even. We can do this by computing the remainder of the input when divided by 2 (using the modulo operator) and checking if it is 0 (even number), 1 (positive odd number) or, -1 (negative odd number). An implementation of this can be seen in Listing 2.9.

An if statement consists of two parts – a *condition*, given inside brackets after the `if` keyword, followed by a body – some lines of code surrounded by curly brackets. The code inside the body will be executed in case the condition evaluates to `true`.

Our odd or even example contains a certain redundancy. If a number is not even we already know it is odd. Checking this explicitly using the modulo

Listing 2.9 Odd or Even

```
1  int input;
2  cin >> input;
3  if (input % 2 == 0) {
4      cout << input << " is even!" << endl;
5  }
6  if (input % 2 == 1 || input % 2 == -1) {
7      cout << input << " is odd!" << endl;
8  }
```

operator seems to be a bit unnecessary. Indeed, there is a construct that saves us from this verbosity – the *else* statement. It is used after an if statement and contains code that should be run if the condition given to the condition of an if statement is false. We can adopt this to simplify our odd and even program to the one in Listing 2.10.

Listing 2.10 Odd or Even 2

```
1  int input;
2  cin >> input;
3  if (input % 2 == 0) {
4      cout << input << " is even!" << endl;
5  } else {
6      cout << input << " is odd!" << endl;
7  }
```

There is one last if-related construct – the *else if*. Since code is worth a thousand words, we demonstrate how it works in Listing 2.11 by implementing a helper for the children’s game *FizzBuzz*. In *FizzBuzz*, one goes through the natural numbers in increasing order and say them out loud. When the number is divisible by 3 you instead say Fizz. If it is divisible by 5 you say Buzz, and if it is divisible by both you say FizzBuzz.

Exercise 2.15. Run the program in Listing 2.11 with the values 30, 10, 6, 4. Explain the output you get.

Problem 2.4

Expected Earnings – casino

Grading – grading
Three-Sort – threesort
Spavanac – spavanac
Cetvrta – cetvrta

2.7 For Loops

Another rudimentary building block of programs is the *for loop*. A for loop is used to execute a block of code multiple times. The most basic loop repeats code a fixed number of times as in the example from Listing 2.12.

A for loop is built up from four parts. The first three parts are the semi-colon separated statements immediately after the `for` keyword. In the first of these parts you write some expression, such as a variable declaration. In the second part you write an expression that evaluates to a `bool`, such as a comparison between two values. In the third part you write another expression.

The first part will be executed only once – it is the first thing that happens in a loop. In this case, we declare a new variable `i` and set it to 0. The loop will then be repeated until the condition in the second part is false. Our example loop will repeat until `i` is no longer less than `repetitions`. The third part executes after each execution of the loop. Since we use the variable `i` to count how many times the loop has executed, we want to increment this by 1 after each iteration.

Together, these three parts make sure our loop will run exactly `repetitions` times. The final part of the loop is the statements within curly brackets. Just as with the `if` statements, this is called the body of the loop and contains the code that will be executed in each repetition of the loop. A repetition of a loop is in algorithm language more commonly referred to as an *iteration*.

Exercise 2.16. What happens if you enter a negative value as the number of loop repetitions?

Exercise 2.17. Design a loop that instead counts backwards, from `repetitions - 1` to 0.

Problem 2.5

N-Sum – nsum
Building Pyramids – pyramids
Echo – echo

Note: solve both parts now, reciving 2/2 points

Cinema Crowds – cinema

Refridgerator Transports – refridgerator

Within a loop, two useful keywords can be used to modify the loop – `continue` and `break`. Using `continue`; inside a loop exits the current iteration and starts the next one. `break`; on the other hand, exits the loop altogether. For an example, consider Listing 2.13.

Exercise 2.18. What will the following code snippet output?

```
1  for (int i = 0; false; i++) {
2      cout << i << endl;
3  }
4
5  for (int i = 0; i >= -10; --i) {
6      cout << i << endl;
7  }
8
9  for (int i = 0; i <= 10; ++i) {
10     if (i % 2 == 0) continue;
11     if (i == 8) break;
12     cout << i << endl;
13 }
```

Problem 2.6

Cinema Crowds 2 – cinema2

Lamps – lamps

2.8 While Loops

There is a second kind of loop, which is simpler than the `for` loop. It is called a **while loop**, and works like a `for` loop where the initial statement and the update statement are removed, leaving only the condition and the body. It can be used when you want to loop over something until a certain condition is false (Listing 2.14).

The `break`; and `continue`; statements work the same way as the `do` in a `for` loop.

Problem 2.7

$3n+1 - 3nplus1$

Soda Sluper – sodaslurper

Listing 2.11 Else If

```
1  int input;
2  cin >> input;
3  if (input % 15 == 0) {
4      cout << "FizzBuzz" << endl;
5  } else if (input % 5 == 0) {
6      cout << "Buzz" << endl;
7  } else if (input % 3 == 0) {
8      cout << "Fizz" << endl;
9  } else {
10     cout << input << endl;
11 }
```

Listing 2.12 For Loops

```
1  int repetitions = 0;
2  cin >> repetitions;
3  for (int i = 0; i < repetitions; i++) {
4      cout << "This is repetition " << i << endl;
5  }
```

2.9 Functions

In mathematics a *function* is something that takes one or more arguments and computes some value based on them. Common functions include the squaring function $\text{square}(x) = x^2$, the addition function $\text{add}(x, y) = x + y$ or, the minimum function $\text{min}(a, b)$ which evaluates to the smallest of its arguments.

Functions exist in programming as well but work a bit differently. Indeed, we have already seen a function – the `main()` function. We have implemented the example functions in Listing 2.15.

In the same way that a variable declaration starts by proclaiming what data type the variable contains a function declaration states what data type the function evaluates to. Afterwards, we write the name of the function followed by its arguments (which is a comma-separated list of variable declarations). Finally, we give it a body of code wrapped in curly brackets.

All of these functions contain a statement with the `return` keyword, unlike our `main` function. A `return` statement says “stop executing this function, and return the following value!”. Thus, when we call the squaring function by `square(x)`,

Listing 2.13 Break and Continue

```
1  int check = 36;
2
3  for (int divisor = 2; divisor * divisor <= check; ++divisor) {
4      if (check % divisor == 0) {
5          cout << check << " is not prime!" << endl;
6          cout << "It equals " << divisor << " x " << (check / divisor) << endl;
7          break;
8      }
9  }
10
11 for (int divisor = 1; divisor <= check; ++divisor) {
12     if (check % divisor == 0) {
13         continue;
14     }
15     cout << divisor << " does not divide " << check << endl;
16 }
```

Listing 2.14 While

```
1  int num = 9;
2  while (num != 1) {
3      if (num % 2 == 0) {
4          num /= 2;
5      } else {
6          num = 3 * num + 1;
7      }
8      cout << num << endl;
9  }
```

the function will compute the value $x * x$ and make sure that `square(x)` evaluates to just that.

Why have we left a return statement out of the main function? In `main()`, the compiler inserts an implicit `return 0;` statement at the end of the function.

Exercise 2.19. What will the following function calls evaluate to?

```
min(square(10), add(square(9), 23));
```

Exercise 2.20. . We declared all of the new arithmetic functions above our `main` function in the example. Why did we do this? What happens if you move one below the `main` function instead? (Hint: what happens if you try to use a variable before declaring it?)

Listing 2.15 Functions

```
1  #include <iostream>
2
3  using namespace std;
4
5  int square(int x) {
6      return x * x;
7  }
8
9  int min(int x, int y) {
10     if (x < y) {
11         return x;
12     } else {
13         return y;
14     }
15 }
16
17 int add(int x, int y) {
18     return x + y;
19 }
20
21 int main() {
22     int x, y;
23     cin >> x >> y;
24     cout << x << "^2 = " << square(x) << endl;
25     cout << x << " + " << y << " = " << add(x, y) << endl;
26     cout << "min(" << x << ", " << y << ") = " << min(x, y) << endl;
27 }
```

Exercise 2.21. . Research online what a *forward declaration* of a function is, and how it resolves the problem from Exercise 2.20.

Problem 2.8

Arithmetic Functions – arithmeticfunctions

An important caveat to note when calling functions is that the arguments we send along are copied. If we try to change them by assigning values to our arguments, we will not change the original variables in the calling function (see Listing 2.16 for an example).

We can also choose to not return *anything* by using the `void` return type. This may seem useless since nothing ought to happen if we call a function but does not get anything in return. However, there are ways we can affect the program

Listing 2.16 Argument Copying

```
1 void change(int val) {
2     val = 0;
3 }
4
5 int main() {
6     int variable = 100;
7     change(variable);
8     cout << "Variable is " << variable << endl;
9 }
```

without returning.

The first one is by using *global* variables. It turns out that variables may be declared outside of a function. It is then available to every function in your program. Changes to a global variable by one function are also be seen by other functions (try out Listing 2.17 to see them in action).

Listing 2.17 Global Variables

```
1 int currentMoney = 0;
2
3 void deposit(int newMoney) {
4     currentMoney += newMoney;
5 }
6 void withdraw(int withdrawal) {
7     currentMoney -= withdrawal;
8 }
9
10 int main() {
11     cout << "Currently, you have " << currentMoney << " money" << endl;
12     deposit(1000);
13     withdraw(2000);
14     cout << "Oh-oh! Your balance is " << currentMoney << " :(" << endl;
15 }
```

Problem 2.9

Counting Days – countingdays

Secondly, we may actually change the variables given to us as arguments by declaring them as *references*. Such an argument is written by adding a & before

the variable name, for example `int &x`. If we perform assignments to the variable `x` within the function we change the variable used for this argument in the calling function instead. Listing 2.18 contains an example of using references.

Listing 2.18 References

```
1  // Note &val instead of val
2  void change(int &val) {
3      val = 0;
4  }
5
6  int main() {
7      int variable = 100;
8      cout << "Variable is " << variable << endl;
9      change(variable);
10     cout << "Variable is " << variable << endl;
11 }
```

Problem 2.10

Logic Functions – logicfunctions

Exercise 2.22. Why is the function call `change(4)` not valid C++? (Hint: what exactly are we changing when we assign to the reference in `func`?)

2.10 Structures

Algorithms operate on data, usually lots of it. Programming language designers therefore came up with many ways of organizing the data our programs use. One of these constructs is the **structure** (also called a record, and in C++ almost equivalent to something called a class). Structures are a special kind of data type that can contain member variables – variables inside them – and member functions – functions which can operate on member variables.

The basic syntax used to define a structure looks like this:

```
struct Point {
    double x;
    double y;
};
```

This particular structure contains two member variables, `x` and `y`, representing the coordinates of a point in 2D Euclidean space.

Once we have defined a structure we can create **instances** of it. Every instance has its own copy of the member variables of the structure. Structs

essentially encapsulate concepts – like books – while instances of the struct represent individual, particular books (like this one!).

To create an instance of a struct, use the same syntax as with other variables; We can get the value of a member variable of a structure using the syntax `instance.variable`:

```
Point origin; // create an instance of the Point structure

// set the coordinates to (0, 0)
origin.x = 0;
origin.y = 0;

cout << "The origin is (" << origin.x << ", "
      << origin.y << ")." << endl;
```

As you can see structures allow us to group certain kinds of data together in a logical fashion. Later on, this will simplify the coding of certain algorithms and data structures immensely.

There is an alternate way of constructing instances called **constructors**. A constructor looks like a function inside our structure and allows us to pass arguments when we create a new instance of a struct. The constructor receives these arguments to help set up the instance.

Let us add a constructor to our point structure, to more easily create instances:

```
struct Point {
    double x;
    double y;

    Point(double theX, double theY) {
        x = theX;
        y = theY;
    }
};
```

The newly added constructor lets us pass two arguments when constructing the instance to set the coordinates correctly. With it, we avoid the two extra statements to set the member variables.

```
Point p(4, 2.1);
cout << "The point is (" << p.x << ", " << p.y << ")." << endl;
```

Structure values can also be constructed outside of a variable declaration using the syntax

```
Point(1, 2);
```

so that we can reassign a previously declared variable with

```
p = Point(1, 2);
```

We can also define functions inside the structure. These functions work just like any other functions except they can also access the member variables of the instance that the member function is called on. For example, we might want a convenient way to mirror a certain point in the x-axis. This could be accomplished by adding a member function:

```
struct Point {  
    double x;  
    double y;  
  
    Point(double theX, double theY) {  
        x = theX;  
        y = theY;  
    }  
  
    Point mirror() {  
        return Point(x, -y);  
    }  
};
```

To call the member function `mirror()` on the point `p`, we write `p.mirror()`, for example:

```
Point p(1, 2);  
Point mirrored = p.mirror();  
cout << "(" << mirrored.x << ", " << mirrored.y << ")" << endl;
```

In this example we see yet another use of a void function. Such member functions can still modify the member variables of the struct they belong to.

Exercise 2.23. Add a `translate` member function to the point structure. It should take two double values `x` and `y` as arguments, returning a new point which is the instance point translated by (x, y) .

Similarly to the `const` modifier that could be added to a variable declaration, one can also declare a *member function* to be `const`:

```
Point mirror() const {  
    return Point(x, -y);  
}
```

The keyword must be added right before the last brace. Such a function is *unable to modify* any of the member variables. It can not call other member functions that are not declared as `const` either. Generally, you will never have to worry about declaring functions to be `const`.

Exercise 2.24. What happens if we try to change a member variable in a `const` member function?

Finally, C++ has a powerful mechanism called *operator overloading*. It allows us to define how various operators such as `+` should behave if we **apply them to instances of a struct**. For example, we could define what happens when we write

`a + b`

where `a` and `b` are `Points`. The syntax for the binary operators looks like this:

```
Point operator+(Point other) {  
    double newX = x + other.x;  
    double newY = y + other.y;  
    return Point(newX, newY);  
}
```

Try this function out by defining two points and computing their sum.

Exercise 2.25. One can use operator overloading for binary operators where the types are different as well. For example,

```
Point operator*(double m) { ... }
```

would define what happens if you multiply a point by a double. Add such a function to your point, that returns a point with its coordinates scaled by the given double.

Exercise 2.26. Fill in the remaining code to implement this structure:

```
1  struct Quotient {  
2      // .. member variables?  
3      // Construct a new Quotient with the given numerator and denominator  
4      Quotient(int n, int d) { }  
5      // Return a new Quotient, this instance plus the "other" instance  
6      Quotient add(const Quotient &other) const { }  
7      // Return a new Quotient, this instance times the "other" instance  
8      Quotient multiply(const Quotient &other) const { }  
9      // Output the value on the screen in the format n/d  
10     void print() const { }  
11 };
```

2.11 Arrays

In the Sorting Problem from Chapter 1 we often spoke of the data type “sequence of integers”. Until now, none of the data types we have seen in C++ represents this kind of data. We present the *array*. It is a special type of variable, which

can contain a large number of variables of the same type. For example, it could be used to represent the recurring data type “sequence of integers” from the Sorting Problem in Chapter 1. When declaring an array, we specify the type of variable it should contain, its name, and its size using the syntax:

```
type name[size];
```

For example, an integer array of size 10 named `seq` would be declared with

```
int seq[10];
```

This creates 10 integer “variables” which we can refer to using the syntax `seq[index]`, starting from zero (they are zero-indexed). Thus we can use `seq[0]`, `seq[1]`, etc., all the way up to `seq[9]`. The values are called the *elements* of the array.

size = 10

<code>seq[0]</code>	<code>seq[1]</code>	<code>seq[2]</code>	<code>seq[3]</code>	<code>seq[4]</code>	<code>seq[5]</code>	<code>seq[6]</code>	<code>seq[7]</code>	<code>seq[8]</code>	<code>seq[9]</code>
---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------

Figure 2.1: A 10-element array called `seq`.

Be aware that using an index outside the valid range for a particular array (i.e. below 0 or above the size – 1) can cause erratic behavior in the program without crashing it.

If you declare a global array all elements get a default value. For numeric types this is 0, for booleans this is `false`, for strings this is the empty string and so on. If, on the other hand, the array is declared in the body of a function that guarantee does not apply. Instead of being zero-initialized, the elements *can have random values*. For this reason, arrays are mostly declared globally in competitive programming.

You can see an example of arrays in action in Listing 2.19, which computes a few of the the possible scores of a roll in the dice game Yatzee.

Later on (Section 3.1) we transition from using arrays to a much more powerful structure from the standard library which serves the same purpose – the *vector*.

Problem 2.11

Reversal – reverse

N-Back – `nback`

Modulo – modulo

I've Been Everywhere, Man – everywhere

2.12 Lambdas

We will now briefly discuss a somewhat complex language construct – *lambdas*. It is very seldom necessary to solve problems, but we occasionally use it in code throughout the book.

A lambda expression is essentially an *unnamed function* that can be defined within another function and assigned to a variable of the function type:

```
function<int,int(int)> op = [](int a, int b) -> int {  
    return a * b + a + b;  
};  
cout << op(5, op(1, 2)) << endl;
```

Here, we have defined a function that takes two values *a* and *b*, and returns the value *a * b + a + b*. We have assigned the function to the variable *op*, and can invoke it as if it was a regular function with that name.

Generally, definitions look simpler than this – if the function is “simple enough”, we can ignore the *-> int* part, which we use to specify the return value of the lambda. We also tend to use the *auto* type instead of the more convoluted *function<...>* type, as long as **the lambda does not call itself** through the name of the variable to which it is assigned.

Thus, the declaration may also look like this:

```
auto op = [](int a, int b) {  
    return a * b + a + b;  
};
```

What is the point of doing this rather than simply using regular functions? Lambdas can also be given access to *variables of the enclosing function*:

```
int x = 5;  
auto addToX = [&](int y) {  
    x += y;  
};
```

Here, note the added ampersand in *[&]*. This means that all variables defined before the lambda in the function should be accessible within the lambda **as references**.

Exercise 2.27. Use the internet to figure out:

- how to only make a single variable from the enclosing function available in a lambda.

Listing 2.19 Arrays

```
1  #include <iostream>
2
3  using namespace std;
4
5  int rolls[7];
6
7  int main() {
8      cout << "Enter 5 dice rolls between 1 and 6: " << endl;
9      for (int i = 0; i < 5; i++) {
10         int roll;
11         cin >> roll;
12         rolls[roll]++;
13     }
14     cout << "Yatzee scores: " << endl;
15     for (int i = 1; i <= 6; i++) {
16         cout << i << "'s: " << (i * rolls[i]) << endl;
17     }
18 }
```

- how to make variables within the enclosing function available **as copies** rather than as references.
- how lambdas can be passed as arguments to other functions.

2.13 The Preprocessor

C++ has a powerful tool called the *preprocessor*. This utility is able to read and modify your code using certain rules during compilation. The commonly used `#include` is a preprocessor directive that includes a certain file in your code.

Besides file inclusion, we mostly use the `#define` directive. It allows us to replace certain tokens in our code with other ones. The most basic usage is

```
#define TOREPLACE REPLACEWITH
```

which replaces the token `TOREPLACE` in our program with `REPLACEWITH`. The true power of the `define` comes when using `define` directives with parameters. These look similar to functions and allows us to replace certain expressions with another one, additionally inserting certain values into it. We call these *macros*. For example the macro

```
#define rep(i,a,b) for (int i = a; i < b; i++)
```

means that the expression

```
rep(i,0,5) {  
    cout << i << endl;  
}
```

is expanded to

```
for (int i = 0; i < 5; ++i) {  
    cout << i << endl;  
}
```

You can probably get by without ever using macros in your code. The reason we discuss them is because we are going to use them in code in the book so it is a good idea to at least be familiar with their meaning. They are also used in competitive programming in general,

2.14 Template

In competitive programming, one often uses a *template*, with some shorthand typedef's and preprocessor directives. In Listing 2.20, we give an example of the template used in some of the C++ code in this book.

Listing 2.20 Coding Template

```
1  #include <bits/stdc++.h>  
2  using namespace std;  
3  
4  #define rep(i, a, b) for(int i = a; i < (b); ++i)  
5  #define trav(a, x) for(auto& a : x)  
6  #define all(x) x.begin(), x.end()  
7  #define sz(x) (int)(x).size()  
8  typedef long long ll;  
9  typedef pair<int, int> pii;  
10 typedef vector<int> vi;  
11  
12 int main() {  
13 }
```

The `rep(i,a,b)` macro is the one we saw in the previous section – it can be used to write a simple counting loop in a compact way.

The `trav(a, x)` macro is used to iterate through all members of a data structure from the standard library such as the vector – the first topic of Chapter 3.

The `all(x)` macro is used together with certain operations from the standard library – we'll see concrete examples in the next chapter.

The `sz(x)` macro is used get the size of a data structure from the standard library.

Chapter Exercises

Problem 2.12

Cubes – kuber

Islands – oar

Grading – betygsattning

Faroffistanian Personal Numbers – checksum

Mini Golf – minigolf

Booking – booking

Tomatoes – tomat

Will Roger's Phenomena – willrogers

Yatzee – yatzee

Memory – memory

Chapter Notes

C++ was invented by Danish computer scientist Bjarne Stroustrup. Bjarne has also published a book on the language, *The C++ Programming Language*[27], that contains a more in-depth treatment of the language. It is rather accessible to C++ beginners but is better read by someone who have some prior programming experience (in any programming language).

C++ is standardized by the International Organization for Standardization (ISO). These standards are the authoritative source on what C++ is. The final drafts of the standards can be downloaded at the homepage of the Standard C++ Foundation¹.

There are many online references of the language and its standard library. The two we use most are:

- <http://en.cppreference.com/w/>
- <http://www.cplusplus.com/reference/>

¹<https://isocpp.org/>

3 The C++ Standard Library

In this chapter we study parts of the C++ standard library – that is, data structures, algorithms and utilities that are already provided for us without having to code them ourselves.

We start by examining a number of basic *data structures*. Data structures help us organize the data we work with in the hope of making processing both easier and more efficient. Different data structures serve widely different purposes and solve different problems. Whether a data structure fits our needs depends on what operations we wish to perform on the data. We consider neither the efficiency of the various operations in this chapter nor how they are implemented. These concerns are postponed until Chapter 6, when we have the tools to analyze the efficiency of data structures.

The standard library also contains many useful algorithms such as sorting and various mathematical functions. These are discussed after the data structures.

In the end, we take a deeper look at string handling in C++ and some more input/output routines.

3.1 vector

One of the latter things discussed in the C++ chapter was the fixed-size array. As you might remember the array is a special kind of data type that allows us to store multiple values of the same data type inside what appeared to us as a single variable. Arrays are a bit awkward to work with in practice. When passing them as parameters we must also pass along the size of the array. We are also unable to change the size of arrays once declared nor can we easily remove or insert elements, or copy arrays.

The *dynamic array* is a special type of array that can change size (hence the name dynamic). It also supports operations such as removing and inserting elements at any position in the list.

The C++ standard library includes a dynamic array called a *vector*, which is an alternative name for dynamic arrays in some languages. To use it you must include the `vector` file by adding the line

```
#include <vector>
```

among your other includes at the top of your program.

When declaring vectors, they need to know what type of data they should store, just like ordinary arrays. This is done using a somewhat peculiar syntax. To create a vector containing strings named `words` we write

```
vector<string> words;
```

This angled bracket syntax appears again later when using other C++ structures from the standard library.

Once a vector is created elements can be appended to it using the `push_back` member function. The following four statements would add the words `Simon is a fish` as separate elements to the vector:

```
words.push_back("Simon");  
words.push_back("is");  
words.push_back("a");  
words.push_back("fish");
```

To refer to a specific element in a vector you can use the same operator `[]` as for arrays. Thus, `words[i]` refers to the *i*'th value in the vector (starting at 0).

```
cout << words[0] << " " << words[1] << " "; // Prints Simon is  
cout << words[2] << " " << words[3] << " "; // Prints a fish
```

Like arrays, accessing indices outside the valid range of the vector can cause weird behaviour in your program.

We can get the current size of an array using the `size()` member function:

```
cout << "The vector contains " << words.size() << " words" << endl;
```

There is also an `empty()` function that can be used to check if the vector contains no elements. These two functions are part of basically every standard library data structure.

Problem 3.1

Vector Functions – vectorfunctions

You can also create dynamic arrays that already contain a number of elements. This is done by passing an integer argument when first declaring the vector. They are filled with the same default value as (global) arrays are when created:

```
vector<int> vec(5); // creates a vector containing 5 zeroes
```

The value that such an array is filled with can also be set explicitly by using a two-argument constructor; the second argument is the value to fill the array with:

```
vector<int> vec(5, -1); // creates a vector containing 5 -1's
```

Exercise 3.1. What happens when you create vectors of a *struct*?

Try using structures:

- without a constructor,
- with a zero argument constructor,
- with only non-zero argument constructors and
- with both zero argument and non-zero argument constructors.

We can create vectors that also contain other vectors, to make *multidimensional* vectors. For example, we could make a 2-dimensional vector (i.e. a grid of values) in the following way:

```
vector<vector<int>> grid(7, vector<int>(5));
```

Since we filled the vector with 7 vectors of length 5, we get 7×5 grid of integers. The values in the grid are then referred to using `grid[a][b]` where $0 \leq a < 5$ and $0 \leq b < 7$.

Similarly, one can create N -dimensional vectors by creating vectors of vectors of ... and so on.

Problem 3.2

Cinema Seating – cinemaseating

Other occasionally useful functions that a vector support are:

- `pop_back()`: remove the last element of a vector
- `clear()`: remove all elements of a vector
- `front()`: get the first element of a vector
- `back()`: get the first element of a vector
- `assign(n, val)`: replace the contents of the vector with n copies of `val`.

3.2 Iterators

A concept central to the standard library is the *iterator*. An iterator is an object which “points to” an element in some kind of data structure (such as a vector). Essentially, they are a generalization of the role played by an integer representing an index of a vector. The reason we could not simply eliminate their use and use integer indexes directly wherever iterators appear is that some data structures do not support accessing values directly by their index. Not all data structures support iterators either.

The type of an iterator for a data structure of type `t` is `t::iterator`. An iterator of a `vector<string>` thus has the type `vector<string>::iterator`. Most of the time we instead use the `auto` type since this is very long to type.

To get an iterator to the first element of a vector, we use `begin()`:

```
auto first = words.begin();
```

We can get the value that an iterator points at using the `*` operator:

```
cout << "The first word is " << *first << endl;
```

If we have an iterator `it` pointing at the i 'th element of a vector we can get a new iterator pointing to another value in one of two ways. For iterators of a vector, we add or subtract an integer value to the iterator. For example, `it + 4` points to the $(i + 4)$ 'th element of the vector, and `it - 1` is the iterator pointing to the $(i - 1)$ 'st element.

For those structures that do not support access by indexes, the iterators can instead be *moved* forwards and backwards using the `++` and `--` operators, i.e. by writing `it++` and `it--`.

There is a special kind of iterator which points to the first position *after* the last element. We get this iterator by using the function `end()`. It allows us to iterate through a vector in the following way:

```
for (auto it = words.begin(); it != words.end(); it++) {  
    string value = *it;  
    cout << value << endl;  
}
```

In this loop we start by creating an iterator which points to the first element of the vector. Our update condition will repeatedly move the iterator to the next element in the vector. The loop condition ensures that the loop breaks when the iterator first points to the element past the end of the vector.

In modern C++ language versions, there is a shorter construct that is equivalent to this loop:

```
for (auto value : words) {  
    cout << value << endl;  
}
```

In addition to the `begin()` and `end()` pair of iterators, there is also `rbegin()` and `rend()`. They work similarly, except that they are *reverse iterators* - they iterate in the other direction. Thus, `rbegin()` actually points to the last element of the vector, and `rend()` to an imaginary element before the first element of the vector. If we move a reverse iterator in a positive direction, we will actually move it in the opposite direction (i.e. adding 1 to a reverse iterator makes it point to the element *before* it in the vector).

Exercise 3.2. Use the `rbegin()/rend()` iterators to code a loop that iterates through a vector in the reverse order.

Certain operators on a vector require the use of vector iterators. For example, the `insert` and `erase` member functions, used to insert and erase elements at arbitrary positions, take iterators to describe positions. When removing the second element, we write

```
words.erase(words.begin() + 1);
```

The `insert()` function uses an iterator to know at what position an element should be inserted. If it is passed the `begin()` iterator the new element will be inserted at the start of the array. Similarly, as an alternative to `push_back()` we could have written

```
words.insert(words.end(), "food");
```

to insert an element at the end of the vector.

Exercise 3.3. After adding these two lines, what would the loop printing every element of the vector `words` output?

Problem 3.3

Cut in Line – cutinline

3.3 queue

The `queue` structure corresponds to a plain, real-life queue. It supports mainly two operations: appending an element to the back of the queue, and extracting the first element of the queue. The structure is in the `queue` file so it must be included using

```
#include<queue>
```

As with all standard library data structures declaring a queue requires us to provide the data type which we wish to store in it. To create a queue storing ints we would use the statement

```
queue<int> q;
```

We use mainly five functions when dealing with queues:

- `push(x)`: add the element `x` to the **back of the queue**
- `pop()`: remove the element from the **front of the queue**
- `front()`: return the element from the **front of the queue**
- `empty()`: return `true` if and only if the queue is empty
- `size()`: return the number of elements in the queue

Exercise 3.4. There is a similar data structured called a *deque*. The standard library version is named after the abbreviation deque instead. Use one of the C++ references from the C++ chapter notes (Section 2.14) to find out what this structure does and what its member functions are called.

3.4 stack

A stack is a structure very similar to a queue. The difference is that when `push()`ing something to a stack, it is added first to the stack instead of last. To use it you need to include

```
#include<stack>
```

Creating a stack containing e.g. ints is as easy as creating a queue:

```
stack<int> q;
```

We use mainly five functions when dealing with queues:

- `push(x)`: add the element `x` to the **top of the stack**
- `pop()`: remove the element from the **top of the stack**
- `top()`: return the element from the **top of the stack**

- `empty()`: return true if and only if the stack is empty
- `size()`: return the number of elements in the stack

Note the change in terminology. Instead of the first and last elements being called front and back as for the queue they are instead called top and bottom in a stack.

3.5 priority_queue

The queue and stack structures are arguably unnecessary, since they can be emulated using a vector (see Sections 6.2, 6.3). This is not the case for the next structure, the `priority_queue`.

The structure is similar to a queue or a stack, but instead of insertions and extractions happening at one of the endpoints of the structure, the **greatest** element is always returned during the extraction.

The structure is located in the same file as the queue structure, so add

```
#include<queue>
```

to use it.

To initialize a priority queue, use the same syntax as for the other structures:

```
priority_queue<int> pq;
```

This time there is one more way to create the structure that is important to remember. It is not uncommon to prefer the sorting to be done according to some other order than descending. For this reason there is another way of creating a priority queue. One can specify a ***comparison function*** that takes two arguments of the type stored in the queue and returns true if the first one should be considered less than the second. This function can be given as an argument to the type in the following way:

```
bool cmp(int a, int b) {  
    return a > b;  
}
```

```
priority_queue<int, vector<int>, cmp> pq;
```

```
// or equivalently  
priority_queue<int, vector<int>, greater<int>> pq;
```

Note that a priority queue by default returns the greatest element. If we want to make it return the smallest element, the comparison function needs to instead

say that the smallest of the two elements actually is the greatest, somewhat counter-intuitively.

Interactions with the queue is similar to that of the other structures:

- `push(x)`: add the element `x` to the priority queue
- `pop()`: remove the **greatest** element from the priority queue
- `top()`: return the **greatest** element from the priority queue
- `empty()`: return `true` if and only if the priority queue is empty
- `size()`: return the number of elements in the priority queue

Problem 3.4

I Can Guess the Data Structure! – `guessthedatastructure`

Akcija – `akcija`

Cookie Selection – `cookieselection`

Pivot – `pivot`

3.6 set and map

The final data structures to be studied in this chapter are also the most powerful: the `set` and the `map`.

The `set` structure is similar to a mathematical set (Section A.2), in that it contains a collection of unique elements. Unlike the `vector`, particular positions in the structure can not be accessed using the `[]` operator. This may make sets seem worse than vectors. The advantage of sets is twofold. First, we can determine membership of elements in a set much more efficiently compared to when using vectors (in Chapters 5 and 6, what this means will become clear). Secondly, sets are also sorted. This means we can quickly find the smallest and greatest values of the set.

Elements are instead accessed only through iterators, obtained using the `begin()`, `end()` and `find()` member functions. These iterators can be moved using the `++` and `--` iterators, allowing us to navigate through the set in sorted (ascending) order (with `begin()` referring to the smallest element).

Elements are inserted using the `insert` function and removed using the `erase` function. A concrete example usage is found in Listing 3.1

A structure similar to the `set` is the `map`. It is essentially the same as a `set`, except the elements are called *keys* and have associated *values*. When declaring

Listing 3.1 Sets

```
1  set<int> s;
2  s.insert(4);
3  s.insert(7);
4  s.insert(1);
5
6  // find returns an iterator to the element if it exists
7  auto it = s.find(4);
8  // ++ moves the iterator to the next element in order
9  ++it;
10 cout << *it << endl;
11
12 // if nonexistent, find returns end()
13 if (s.find(7) == s.end()) {
14     cout << "7 is not in the set" << endl;
15 }
16
17 // erase removes the specific element
18 s.erase(7);
19
20 if (s.find(7) == s.end()) {
21     cout << "7 is not in the set" << endl;
22 }
23
24 cout << "The smallest element of s is " << *s.begin() << endl;
```

a map two types need to be provided – that of the key and that of the value. To declare a map with string keys and int values you write

```
map<string, int> m;
```

Accessing the value associated with a key x is done using the `[]` operator, for example, `m["Johan"]`; would access the value associated with the "Johan" key.

Problem 3.5

Secure Doors – securedoors

Babelfish – babelfish

3.7 Math

Many algorithmic problems require mathematical functions. In particular you there is a heavy use of square roots and trigonometric functions in geometry problems. These of these functions are be found in the

Listing 3.2 Maps

```
1  map<string, int> age;
2  age["Johan"] = 22;
3  age["Simon"] = 23;
4
5  if (age.find("Aron") == age.end()) {
6      cout << "No record of Aron's age" << endl;
7  }
8
9  cout << "Johan is " << age["Johan"] << " years old" << endl;
10 cout << "Anton is " << age["Anton"] << " years old" << endl;
11
12 age.erase("Johan");
13 cout << "Johan is " << age["Johan"] << " years old" << endl;
14
15 auto last = --age.end();
16 cout << (*last).first << " is " << (*last).second << " years old" << endl;
```

`#include <cmath>`

library.

We list some of the most common such functions here:

- `abs(x)`: computes $|x|$ (x if $x \geq 0$, otherwise $-x$)
- `sqrt(x)`: computes \sqrt{x}
- `pow(x, y)`: computes x^y
- `exp(x)`: computes e^x
- `log(x)`: computes $\ln(x)$
- `cos(x)` / `acos(x)`: computes $\cos(x)$ and $\arccos(x)$ respectively
- `sin(x)` / `asin(x)`: computes $\sin(x)$ and $\arcsin(x)$ respectively
- `tan(x)` / `atan(x)`: computes $\tan(x)$ and $\arctan(x)$ respectively
- `ceil(x)` / `floor(x)`: computes $\lceil x \rceil$ and $\lfloor x \rfloor$ respectively

There are also `min(x, y)` and `max(x, y)` functions which compute the minimum and maximum of the values x and y respectively. These are not in the `cmath` library however. Instead, they are in `algorithm`.

Problem 3.6*Vacuumba* – vacuumba*Half a Cookie* – halfacookie*Ladder* – ladder*A1 Paper* – a1paper

3.8 Algorithms

A majority of the algorithms we regularly use from the standard library operate on sequences. To use algorithms, you need to include

```
#include <algorithm>
```

Sorting

Sorting a sequences is very easy in C++. The function for doing so is named `sort`. It takes two iterators marking the beginning and end of the interval to be sorted and sorts it in-place in ascending order. For example, to sort the first 10 elements of a vector named `v` you would use

```
sort(v.begin(), v.begin() + 10);
```

Note that the right endpoint of the interval is *exclusive* – it is not included in the interval itself. This means that you can provide `v.end()` as the end of the interval if you want to sort the vector until the end.

As with `priority_queues` or `sets`, the sorting algorithm can take a custom comparator if you want to sort according to some other order than that defined by the `<` operator. For example,

```
sort(v.begin(), v.end(), greater<int>());
```

would sort the vector `v` in descending order. You can provide other sorting functions as well. For example, you can sort numbers by their absolute value by passing in the following comparator:

```
bool cmp(int a, int b) {  
    return abs(a) < abs(b);  
}  
sort(v.begin(), v.end(), cmp);
```

What happens if two values have the same absolute value when sorted with the above comparator? With `sort`, this behaviour is not specified: they can be ordered in any way. Occasionally you want that values compared by your comparison function as equal are sorted in the same order as they were given

in the input. This is called a *stable sort*, and is implemented in C++ with the function `stable_sort`.

To check if a vector *is* sorted, the `is_sorted` function can be used. It takes the same arguments as the `sort` function.

Problem 3.7

Shopaholic – `shopaholic`

Busy Schedule – `bussyschedule`

Sort of Sorting – `sortofsorting`

Searching

The most basic search operation is the `find` function. It takes two iterators representing an interval and a value. If one of the elements in the interval equals the value, an iterator to the element is returned. In case of multiple matches the first one is returned. Otherwise, the iterator provided as the end of the interval is returned. The common usage is

```
find(v.begin(), v.end(), 5);
```

which would return an iterator to the first instance of 5 in the vector.

To find out how many times an element appears in a vector, the `count` function takes the same arguments as the `find` function and returns the total number of matches.

If the array is sorted, you can use the much faster *binary search* operations instead. The `binary_search` function takes as argument a **sorted** interval given by two iterators and a value. It returns true if the interval contains the value. The `lower_bound` and `upper_bound` functions takes the same arguments as `binary_search`, but instead returns an iterator to the first element **not less** and **greater** than the specified value, respectively. For more details on how these are implemented, read Section 12.3.

Permutations

In some problems, the solution involves iterating through all *permutations* (Section ??) of a vector. As one of few languages, C++ has a built-in functions for this purpose: `next_permutation`. The function takes two iterators as arguments and rearranges the interval they specify to be the next permutation in lexicographical order. If there is no such permutation, the interval instead becomes sorted and the function returns `false`. This suggests the following

common pattern to iterate through all permutations of a vector *v*:

```
sort(v.begin(), v.end());
do {
    // do something with v
} while (next_permutation(v.begin(), v.end()));
```

This do-while-syntax is similar to the while loop, except the condition is checked *after* each iteration instead of before. It is equivalent to

```
sort(v.begin(), v.end());
while (true) {
    // do something with v
    if (!next_permutation(v.begin(), v.end())) {
        break;
    }
}
```

Problem 3.8

Veci – veci

3.9 Strings

We have already used the `string` type many times before. Until now one of the essential features of a string has been omitted – a string is to a large extent like a vector of `chars`. This is especially true in that you can access the individual characters of a string using the `[]` operator. For a string

```
string thecowsays = "boo";
```

the expression `thecowsays[0]` is the character `'b'`. Furthermore, you can `push_back` new characters to the end of a string.

```
thecowsays.push_back('p');
```

would instead make the string `boop`.

Problem 3.9

Detailed Differences – detaileddifferences

Autori – autori

Skener – skener

Conversions

In some languages, the barrier between strings and e.g. integers is more fuzzy than in C++. In Java, for example, the code `"4" + 2` would append the *character*

'2' to the string "4", yielding the string "42". This is not the case in C++ (what errors do you get if you try to do this?).

Instead, there are other ways to convert between strings and other types. The easiest way is through using the `stringstream` class. A `stringstream` essentially works as a combined `cin` and `cout`. An empty stream is declared by

```
stringstream ss;
```

Values can be written to the stream using the `<<` operator and read from it using the `>>` operator. This can be exploited to convert strings to and from e.g. numeric types like this:

```
stringstream numToString;
numToString << 5;
string val;
numToString >> val; // val is now the string "5"
```

```
stringstream stringToNum;
stringToNum << "5";
int val;
stringToNum >> val; // val is now the integer 5
```

Just as with `cin`, you can use a `stringstream` to determine what type the next word is. If you try to read from a `stringstream` into an `int` but the next word is not an integer, the expression will evaluate to false:

```
stringstream ss;
ss << "notaninteger";
int val;
if (ss >> val) {
    cout << "read an integer!" << endl;
} else {
    cout << "next word was not an integer" << endl;
}
```

Problem 3.10

Filip – filip

Stacking Cups – cups

3.10 Input/Output

Input and output is primarily handled by the `cin` and `cout` objects, as previously witnessed. While they are very easy to use, adjustments are sometimes necessary.

Detecting End of File

The first advanced usage is reading input until we run out of input (often called reading until the *end-of-file*). Normally, input formats are constructed so that

you always know beforehand how many tokens of input you need to read. For example, lists of integers are often either prefixed by the size of the list or terminated by some special sentinel value. For those few times when we need to read input until the end we use the fact that `cin >> x` is an expression that evaluates to `false` if the input reading failed. This is also the case if you try to read an `int` but the next word is not actually an integer. This kind of input loop thus looks something like the following:

```
int num;
while (cin >> num) {
    // do something with num
}
```

Problem 3.11

A Different Problem – different

Statistics – statistics

Input Line by Line

As we stated briefly in the C++ chapter, `cin` only reads a single word when used as input to a string. This is a problem if the input format requires us to read input line by line. The solution to this is the `getline` function, which reads text until the next newline:

```
getline(cin, str);
```

Be warned that if you use `cin` to read a single word that is the last on its line, the final newline is not consumed. That means that for an input such as

word

blah blah

the code

```
string word;
cin >> word;
string line;
getline(cin, line);
```

would produce an *empty* line! After `cin >> word` the newline of the line `word` still remains, meaning that `getline` only reads the (zero) remaining characters until the newline. To avoid this problem, you need to use `cin.ignore()`; to ignore the extra newline before your `getline`.

Once a line has been read we often need to process all the words on the line one by one. For this, we can use the `stringstream`:

```
stringstream line(str);
string word;
while (line >> word) {
    // do something with word
}
```

The stringstream takes an argument that is the string you want to process. After this, it works just like cin does, except reading input from the string instead of the terminal. To use stringstream, add the include

```
#include <sstream>
```

Problem 3.12

Bacon Eggs and Spam – baconeggsandspam

Compound Words – compoundwords

Output Decimal Precision

Another common problem is that outputting decimal values with cout produces numbers with too few decimals. Many problems stipulate that an answer is considered correct if it is within some specified relative or absolute precision of the judges' answer. The default precision of cout is 10^{-6} . If a problem requires higher precision, it must be set manually using e.g.

```
cout << setprecision(10);
```

If the function argument is x , the precision is set to 10^{-x} . This means that the above statement would set the precision of cout to 10^{-10} . This precision is normally the relative precision of the output (i.e. the total number of digits to print). If you want the precision to be absolute (i.e. specify the number of digits after the decimal point) you write

```
cout << fixed;
```

Problem 3.13

A Real Challenge – areal

Chapter Exercises

Problem 3.14

Apaxiaaaaaaaaaaans! – apaxiaaans

Different Distances – differentdistances

Odd Man Out – oddmanout

Timebomb – timebomb

Missing Gnomes – missinggnomes

Chapter Notes

In this chapter, only the parts from the standard library we deemed most important to problem solving were extracted. The standard library is much larger than this, of course. While you will almost always get by using only what we discussed additional knowledge of the library can make you a faster, more effective coder.

For a good overview of the library, [cppreference.com](http://en.cppreference.com)¹ contains lists of the library contents categorized by topic.

¹<http://en.cppreference.com/w/cpp>

4 Implementation Problems

The “simplest” kind of problem we solve is those where the statement of a problem is so detailed that the difficult part is not figuring out the solution, but implementing it in code. This type of problem is mostly given in the form of performing some calculation or simulating some process based on a list of rules stated in the problem.

The Recipe

Swedish Olympiad in Informatics 2011, School Qualifiers (CC BY-SA 3.0)

You have decided to cook some food. The dish you are going to make requires N different ingredients. For every ingredient, you know the amount you have at home, how much you need for the dish, and how much it costs to buy (per unit).

If you do not have a sufficient amount of some ingredient you need to buy the remainder from the store. Your task is to compute the cost of buying the remaining ingredients.

Input

The first line of input is an integer $N \leq 10$, the number of ingredients in the dish.

The next N lines contain the information about the ingredients, one per line. An ingredient is given by three space-separated integers $0 \leq h, n, c \leq 200$ – the amount you have, the amount you need, and the cost per unit for this ingredient.

Output

Output a single integer – the cost for purchasing the remaining ingredients needed to make the dish.

This problem is not particularly hard. For every ingredient we need to calculate the amount which we need to purchase. The only gotcha in the problem is the mistake of computing this as $n - h$. The correct formula is $\max(0, n - h)$, required in case of the luxury problem of having more than we need. We then multiply this number by the ingredient cost and sum the costs up for all the ingredients. A solution would look something like the following.

```
1: procedure RECIPE( $N$ ,  $int[]$   $h$ ,  $int[]$   $n$ ,  $int[]$   $c$ )
2:    $ans \leftarrow 0$ 
3:   for  $i \leftarrow 0$  to  $N - 1$  do
4:      $ans \leftarrow ans + \max(0, n[i] - h[i]) \cdot c[i]$ 
5:   return  $ans$ 
```

Generally, the implementation problems are the easiest type of problems in a contest. They do not require much algorithmic knowledge so more contestants are able to solve them. However, not every implementation problem is easy to code. Just because implementation problems are easy to spot, understand, and formulate a solution to, you should not underestimate the difficulty coding them. Contestants usually fail implementation problems either because the algorithm you are supposed to implement is very complicated with many easy-to-miss details, or because the amount of code is very large. In the latter case, you are more prone to bugs because more lines of code tend to include more bugs.

Let us study a straightforward implementation problem that turned out to be rather difficult to code.

Game Rank

Nordic Collegiate Programming Contest 2016 – Jimmy Mårdell (CC BY-SA 3.0)

The gaming company Sandstorm is developing an online two player game. You have been asked to implement the ranking system. All players have a rank determining their playing strength which gets updated after every game played. There are 25 regular ranks, and an extra rank, “Legend”, above that. The ranks are numbered in decreasing order, 25 being the lowest rank, 1 the second highest rank, and Legend the highest rank.

Each rank has a certain number of “stars” that one needs to gain before advancing to the next rank. If a player wins a game, she gains a star. If before the game the player was on rank 6-25, and this was the third or more consecutive win, she gains an additional bonus star for that win. When she has all the stars for her rank (see list below) and gains another star, she will instead gain one rank and have one star on the new rank.

For instance, if before a winning game the player had all the stars on her current rank, she will after the game have gained one rank and have 1 or 2 stars (depending on whether she got a bonus star) on the new rank. If on the other hand she had all stars except one on a rank, and won a game that also gave her a

bonus star, she would gain one rank and have 1 star on the new rank.

If a player on rank 1-20 loses a game, she loses a star. If a player has zero stars on a rank and loses a star, she will lose a rank and have all stars minus one on the rank below. However, one can never drop below rank 20 (losing a game at rank 20 with no stars will have no effect).

If a player reaches the Legend rank, she will stay legend no matter how many losses she incurs afterwards.

The number of stars on each rank are as follows:

- Rank 25-21: 2 stars
- Rank 20-16: 3 stars
- Rank 15-11: 4 stars
- Rank 10-1: 5 stars

A player starts at rank 25 with no stars. Given the match history of a player, what is her rank at the end of the sequence of matches?

Input

The input consists of a single line describing the sequence of matches. Each character corresponds to one game; ‘w’ represents a win and ‘l’ a loss. The length of the line is between 1 and 10 000 characters (inclusive).

Output

Output a single line containing a rank after having played the given sequence of games; either an integer between 1 and 25 or “Legend”.

A very long problem statement! The first hurdle is finding the energy to read it from start to finish without skipping any details. Not much creativity is needed here – indeed, the algorithm to implement is given in the statement. Despite this, it is not as easy as one would think. Although it was the second most solved problem at the contest where it was used in, it was also the one with the *worst success ratio*. On average, a team needed 3.59 attempts before getting a correct solution, compared to the runner-up problem at 2.92 attempts. None of the top 6 teams in the contest got the problem accepted on their first attempt. Failed attempts cost a lot. Not only in absolute time, but many forms of competition include additional penalties for submitting incorrect solutions.

Implementation problems get much easier when you know your programming

language well and can use it to write good, structured code. Split code into functions, use structures, and give your variables good names and implementation problems become easier to code. A solution to the Game Rank problem which attempts to use this approach is given here:

```
1  #include <bits/stdc++.h>
2
3  using namespace std;
4
5  int curRank = 25, curStars = 0, conseqWins = 0;
6
7  int starsOfRank() {
8      if (curRank >= 21) return 2;
9      if (curRank >= 16) return 3;
10     if (curRank >= 11) return 4;
11     if (curRank >= 1) return 5;
12     assert(false);
13 }
14
15 void addStar() {
16     if (curStars == starsOfRank()) {
17         --curRank;
18         curStars = 0;
19     }
20     ++curStars;
21 }
22
23 void addWin() {
24     int curStarsWon = 1;
25     ++conseqWins;
26     if (conseqWins >= 3 && curRank >= 6) curStarsWon++;
27
28     for (int i = 0; i < curStarsWon; i++) {
29         addStar();
30     }
31 }
32
33 void loseStar() {
34     if (curStars == 0) {
35         if (curRank == 20) return;
36         ++curRank;
37         curStars = starsOfRank();
38     }
39     --curStars;
40 }
41
42 void addLoss() {
43     conseqWins = 0;
44     if (curRank <= 20) loseStar();
```

```

45 }
46
47 int main() {
48     string seq;
49     cin >> seq;
50     for (char res : seq) {
51         if (res == 'W') addWin();
52         else addLoss();
53         if (curRank == 0) break;
54         assert(1 <= curRank && curRank <= 25);
55         assert(0 <= curStars && curStars <= starsOfRank());
56     }
57     if (curRank == 0) cout << "Legend" << endl;
58     else cout << curRank << endl;
59 }

```

Note the use of the `assert()` function. The function takes a single boolean parameter and crashes the program with an assertion failure if the parameter evaluated to `false`. This is helpful when solving problems since it allows us to verify that assumptions we make regarding the internal state of the program indeed holds. In fact, when the above solution was written the assertions in it actually managed to catch some bugs before submitting the problem!

Problem 4.1

Game Rank – gamerrank

Next, we work through a complex implementation problem, starting with a long, hard-to-read solution with a few bugs. Then, we refactor it a few times until it is correct and easy to read.

Mate in One

Introduction to Algorithms at Danderyds Gymnasium

"White to move, mate in one."

When you are looking back in old editions of the *New in Chess* magazine, you find loads of chess puzzles. Unfortunately, you realize that it was way too long since you played chess. Even trivial puzzles such as finding a mate in one now far exceed your ability.

But, perseverance is the key to success. You realize that you can instead use your new-found algorithmic skills to solve the problem by coding a program to find the winning move.

You will be given a chess board, which satisfy:

- No player may castle.
- No player can perform an en passant¹.
- The board is a valid chess position.
- White can mate black in a single, unique move.

Write a program to output the move white should play to mate black.

Input

The board is given as a 8×8 grid of letters. The letter . represent an empty space, the characters pbnrqk represent a white pawn, bishop, knight, rook, queen and king, and the characters PBNRQK represents a black pawn, bishop, knight, rook, queen and king.

Output

Output a move on the form a1b2, where a1 is the square to move a piece from (written as the column, a-h, followed by the row, 1-8) and b2 is the square to move the piece to.

Our first solution attempt clocks in at about 300 lines.

```
1  #include <bits/stdc++.h>
2  using namespace std;
3
4  #define rep(i,a,b) for (int i = (a); i < (b); ++i)
5  #define trav(it, v) for (auto& it : v)
6  #define all(v) (v).begin(), (v).end()
7  typedef pair<int, int> ii;
8  typedef vector<ii> vii;
9  template <class T> int size(T &x) { return x.size(); }
10
11 char board[8][8];
12
13 bool is_empty(int x, int y) {
14     return board[x][y] == '.';
15 }
16
17 bool is_white(int x, int y) {
18     return board[x][y] >= 'A' && board[x][y] <= 'Z';
19 }
```

¹If you are not aware of this special pawn rule, do not worry – knowledge of it is irrelevant with regard to the problem.

```

20
21 bool is_valid(int x, int y) {
22     return x >= 0 && x < 8 && y >= 0 && y < 8;
23 }
24
25 int rook[8][2] = {
26     {1, 2},
27     {1, -2},
28     {-1, 2},
29     {-1, -2},
30
31     {2, 1},
32     {-2, 1},
33     {2, -1},
34     {-2, -1}
35 };
36
37 void display(int x, int y) {
38     printf("%c%d", y + 'a', 7 - x + 1);
39 }
40
41 vii next(int x, int y) {
42     vii res;
43
44     if (board[x][y] == 'P' || board[x][y] == 'p') {
45         // pawn
46
47         int dx = is_white(x, y) ? -1 : 1;
48
49         if (is_valid(x + dx, y) && iz_empty(x + dx, y)) {
50             res.push_back(ii(x + dx, y));
51         }
52
53         if (is_valid(x + dx, y - 1)
54             && is_white(x, y) != is_white(x + dx, y - 1)) {
55             res.push_back(ii(x + dx, y - 1));
56         }
57
58         if (is_valid(x + dx, y + 1)
59             && is_white(x, y) != is_white(x + dx, y + 1)) {
60             res.push_back(ii(x + dx, y + 1));
61         }
62
63     } else if (board[x][y] == 'N' || board[x][y] == 'n') {
64         // knight
65
66         for (int i = 0; i < 8; i++) {
67             int nx = x + rook[i][0],
68                 ny = y + rook[i][1];
69

```

```
70     if (is_valid(nx, ny) && (iz_empty(nx, ny) ||
71         is_white(x, y) != is_white(nx, ny))) {
72         res.push_back(ii(nx, ny));
73     }
74 }
75
76 } else if (board[x][y] == 'B' || board[x][y] == 'b') {
77     // bishop
78
79     for (int dx = -1; dx <= 1; dx++) {
80         for (int dy = -1; dy <= 1; dy++) {
81             if (dx == 0 && dy == 0)
82                 continue;
83
84             if ((dx == 0) != (dy == 0))
85                 continue;
86
87             for (int k = 1; ; k++) {
88                 int nx = x + dx * k,
89                     ny = y + dy * k;
90
91                 if (!is_valid(nx, ny)) {
92                     break;
93                 }
94
95                 if (iz_empty(nx, ny) || is_white(x, y) != is_white(nx, ny)) {
96                     res.push_back(ii(nx, ny));
97                 }
98
99                 if (!iz_empty(nx, ny)) {
100                     break;
101                 }
102             }
103         }
104     }
105
106 } else if (board[x][y] == 'R' || board[x][y] == 'r') {
107     // rook
108
109     for (int dx = -1; dx <= 1; dx++) {
110         for (int dy = -1; dy <= 1; dy++) {
111             if ((dx == 0) == (dy == 0))
112                 continue;
113
114             for (int k = 1; ; k++) {
115                 int nx = x + dx * k,
116                     ny = y + dy * k;
117
118                 if (!is_valid(nx, ny)) {
119                     break;
```

```

120         }
121
122         if (iz_empty(nx, ny) || is_white(x, y) != is_white(nx, ny)) {
123             res.push_back(ii(nx, ny));
124         }
125
126         if (!iz_empty(nx, ny)) {
127             break;
128         }
129     }
130 }
131 }
132
133 } else if (board[x][y] == 'Q' || board[x][y] == 'q') {
134     // queen
135
136     for (int dx = -1; dx <= 1; dx++) {
137         for (int dy = -1; dy <= 1; dy++) {
138             if (dx == 0 && dy == 0)
139                 continue;
140
141             for (int k = 1; ; k++) {
142                 int nx = x + dx * k,
143                     ny = y + dy * k;
144
145                 if (!is_valid(nx, ny)) {
146                     break;
147                 }
148
149                 if (iz_empty(nx, ny) || is_white(x, y) != is_white(nx, ny)) {
150                     res.push_back(ii(nx, ny));
151                 }
152
153                 if (!iz_empty(nx, ny)) {
154                     break;
155                 }
156             }
157         }
158     }
159
160
161 } else if (board[x][y] == 'K' || board[x][y] == 'k') {
162     // king
163
164     for (int dx = -1; dx <= 1; dx++) {
165         for (int dy = -1; dy <= 1; dy++) {
166             if (dx == 0 && dy == 0)
167                 continue;
168
169             int nx = x + dx,

```

```
170         ny = y + dy;
171
172         if (is_valid(nx, ny) && (iz_empty(nx, ny) ||
173             is_white(x, y) != is_white(nx, ny))) {
174             res.push_back(ii(nx, ny));
175         }
176     }
177 }
178 } else {
179     assert(false);
180 }
181
182 return res;
183 }
184
185 bool is_mate() {
186
187     bool can_escape = false;
188
189     char new_board[8][8];
190
191     for (int x = 0; !can_escape && x < 8; x++) {
192         for (int y = 0; !can_escape && y < 8; y++) {
193             if (!iz_empty(x, y) && !is_white(x, y)) {
194
195                 vii moves = next(x, y);
196                 for (int i = 0; i < size(moves); i++) {
197                     for (int j = 0; j < 8; j++)
198                         for (int k = 0; k < 8; k++)
199                             new_board[j][k] = board[j][k];
200
201                 new_board[moves[i].first][moves[i].second] = board[x][y];
202                 new_board[x][y] = '.';
203
204                 swap(new_board, board);
205
206                 bool is_killed = false;
207                 for (int j = 0; !is_killed && j < 8; j++) {
208                     for (int k = 0; !is_killed && k < 8; k++) {
209                         if (!iz_empty(j, k) && is_white(j, k)) {
210                             vii nxts = next(j, k);
211
212                             for (int l = 0; l < size(nxts); l++) {
213                                 if (board[nxts[l].first][nxts[l].second] == 'k') {
214                                     is_killed = true;
215                                     break;
216                                 }
217                             }
218                         }
219                     }
220                 }
221             }
222         }
223     }
224 }
```

```

220         }
221     }
222
223     swap(new_board, board);
224
225     if (!is_killed) {
226         can_escape = true;
227         break;
228     }
229 }
230
231 }
232 }
233 }
234
235 return !can_escape;
236 }
237
238 int main()
239 {
240     for (int i = 0; i < 8; i++) {
241         for (int j = 0; j < 8; j++) {
242             scanf("%c", &board[i][j]);
243         }
244
245         scanf("\n");
246     }
247
248     char new_board[8][8];
249     for (int x = 0; x < 8; x++) {
250         for (int y = 0; y < 8; y++) {
251             if (!is_empty(x, y) && is_white(x, y)) {
252
253                 vii moves = next(x, y);
254
255                 for (int i = 0; i < size(moves); i++) {
256
257                     for (int j = 0; j < 8; j++)
258                         for (int k = 0; k < 8; k++)
259                             new_board[j][k] = board[j][k];
260
261                     new_board[moves[i].first][moves[i].second] = board[x][y];
262                     new_board[x][y] = '.';
263
264                     swap(new_board, board);
265
266
267                     if (board[moves[i].first][moves[i].second] == 'P' &&
268                         moves[i].first == 0) {
269

```

```
270         board[moves[i].first][moves[i].second] = 'Q';
271         if (is_mate()) {
272             printf("%c%d%c%d\n", y + 'a', 7 - x + 1,
273                 moves[i].second + 'a', 7 - moves[i].first + 1);
274             return 0;
275         }
276
277         board[moves[i].first][moves[i].second] = 'N';
278         if (is_mate()) {
279             printf("%c%d%c%d\n", y + 'a', 7 - x + 1,
280                 moves[i].second + 'a', 7 - moves[i].first + 1);
281             return 0;
282         }
283
284     } else {
285         if (is_mate()) {
286             printf("%c%d%c%d\n", y + 'a', 7 - x + 1,
287                 moves[i].second + 'a', 7 - moves[i].first + 1);
288             return 0;
289         }
290     }
291
292     swap(new_board, board);
293 }
294 }
295 }
296 }
297
298 assert(false);
299
300 return 0;
301 }
```

That is a lot of code! Note how there are a few obvious mistakes which makes the code harder to read, such as typo of `iz_empty` instead of `is_empty`, or how the list of moves for the knight is called `rook`. Our final solution reduces this to less than half the size.

Exercise 4.1. Read through the above code carefully and consider if there are better ways to solve the problem. Furthermore, it has a bug – can you find it?

First, let us clean up the move generation a bit. Currently, it is implemented as the function `next`, together with some auxillary data (lines 25-179). It is not particularly abstract, plagued by a lot of code duplication.

The move generation does not need a lot of code. Almost all the moves of the pieces can be described in the same way, as: “pick a direction out of a list D and move at most L steps along this direction, stopping either before exiting

the board or taking your own piece, or when taking another piece.”. For the king and queen, D is all 8 directions one step away, with $L = 1$ for the king and $L = \infty$ for the queen.

Implementing this abstraction is done with little code.

```

1  const vii DIAGONAL = {{-1, 1}, {-1, 1}, {1, -1}, {1, 1}};
2  const vii CROSS = {{0, -1}, {0, 1}, {-1, 0}, {1, 0}};
3  const vii ALL_MOVES = {{-1, 1}, {-1, 1}, {1, -1}, {1, 1},
4    {0, -1}, {0, 1}, {-1, 0}, {1, 0}};
5  const vii KNIGHT = {{-1, -2}, {-1, 2}, {1, -2}, {1, 2},
6    {-2, -1}, {-2, 1}, {2, -1}, {2, 1}};
7  vii directionMoves(const vii& D, int L, int x, int y) {
8    vii moves;
9    trav(dir, D) {
10      rep(i, 1, L+1) {
11        int nx = x + dir.first * i, ny = y + dir.second * i;
12        if (!isValid(nx, ny)) break;
13        if (isEmpty(nx, ny)) moves.emplace_back(nx, ny);
14        else {
15          if (isWhite(x, y) != isWhite(nx, ny)) moves.emplace_back(nx, ny);
16          break;
17        }
18      }
19    }
20    return moves;
21  }

```

A short and sweet abstraction, that will prove very useful. It handles all possible moves, except for pawns. These have a few special cases.

```

1  vii pawnMoves(int x, int y) {
2    vii moves;
3    if (x == 0 || x == 7) {
4      vii queenMoves = directionMoves(ALL_MOVES, 16, x, y);
5      vii knightMoves = directionMoves(KNIGHT, 1, x, y);
6      queenMoves.insert(queenMoves.begin(), all(knightMoves));
7      return queenMoves;
8    }
9    int mv = (isWhite(x, y) ? -1 : 1);
10   if (isValid(x + mv, y) && isEmpty(x + mv, y)) {
11     moves.emplace_back(x + mv, y);
12     bool canMoveTwice = (isWhite(x, y) ? x == 6 : x == 1);
13     if (canMoveTwice && isValid(x + 2 * mv, y) && isEmpty(x + 2 * mv, y)) {
14       moves.emplace_back(x + 2 * mv, y);
15     }
16   }
17   auto take = [&](int nx, int ny) {
18     if (isValid(nx, ny) && !isEmpty(nx, ny)
19       && isWhite(x, y) != isWhite(nx, ny))

```

```
20     moves.emplace_back(nx, ny);
21 };
22 take(x + mv, y - 1);
23 take(x + mv, y + 1);
24 return moves;
25 }
```

This pawn implementation also takes care of promotion, rendering the logic previously implementing this obsolete.

The remainder of the move generation is now implemented as:

```
1  vii next(int x, int y) {
2      vii moves;
3      switch(toupper(board[x][y])) {
4          case 'Q': return directionMoves(ALL_MOVES, 16, x, y);
5          case 'R': return directionMoves(CROSS, 16, x, y);
6          case 'B': return directionMoves(DIAGONAL, 16, x, y);
7          case 'N': return directionMoves(KNIGHT, 1, x, y);
8          case 'K': return directionMoves(ALL_MOVES, 1, x, y);
9          case 'P': return pawnMoves(x, y);
10     }
11     return moves;
12 }
```

These make up a total of about 50 lines – a reduction to a third of how the move generation was implemented before. The trick was to rework all code duplication into a much cleaner abstraction.

We also have a lot of code duplication in the `main` (lines 234-296) and `is_mate` (lines 181-232) functions. Both functions loop over all possible moves, with lots of duplication. First of all, let us further abstract the move generation to not only generate the moves a certain piece can make, but all the moves a player can make. This is done in both functions, so we should be able to extract this logic into only one place:

```
1  vector<pair<ii, ii>> getMoves(bool white) {
2      vector<pair<ii, ii>> allMoves;
3      rep(x,0,8) rep(y,0,8) if (!isEmpty(x, y) && isWhite(x, y) == white) {
4          vii moves = next(x, y);
5          trav(it, moves) allMoves.emplace_back(ii{x, y}, it);
6      }
7      return allMoves;
8  }
```

We also have some duplication in the code *making* the moves. Before extracting this logic, we will change the structure used to represent the board. A

`char[8][8]` is a tedious structure to work with. It is not easily copied or sent as parameter. Instead, we use a `vector<string>`, typedef'd as `Board`:

```
typedef vector<string> Board;
```

We then add a function to make a move, returning a new board:

```
1 Board doMove(pair<ii, ii> mv) {
2     Board newBoard = board;
3     ii from = mv.first, to = mv.second;
4     newBoard[to.first][to.second] = newBoard[from.first][from.second];
5     newBoard[from.first][from.second] = '.';
6     return newBoard;
7 }
```

Hmm... there should be one more thing in common between the `main` and `is_mate` functions. Namely, to check if the current player is in check after a move. However, it seems this is not done in the `main` function – a bug. Since we do need to do this twice, it should probably be its own function:

```
1 bool inCheck(bool white) {
2     trav(mv, getMoves(!white)) {
3         ii to = mv.second;
4         if (!isEmpty(to.first, to.second)
5             && isWhite(to.first, to.second) == white
6             && toupper(board[to.first][to.second]) == 'K') {
7             return true;
8         }
9     }
10    return false;
11 }
```

Now, the long `is_mate` function is much shorter and readable, thanks to our refactoring:

```
1 bool isMate() {
2     if (!inCheck(false)) return false;
3     Board oldBoard = board;
4     trav(mv, getMoves(false)) {
5         board = doMove(mv);
6         if (!inCheck(false)) return false;
7         board = oldBoard;
8     }
9     return true;
10 }
```

A similar transformation is now possible of the `main` function, that loops over all moves white make and checks if black is in mate:

```
1  int main() {
2      rep(i,0,8) {
3          string row;
4          cin >> row;
5          board.push_back(row);
6      }
7      Board oldBoard = board;
8      trav(mv, getMoves(true)) {
9          board = doMove(mv);
10         if (!inCheck(true) && isMate()) {
11             outputSquare(mv.first.first, mv.first.second);
12             outputSquare(mv.second.first, mv.second.second);
13             cout << endl;
14             break;
15         }
16     }
17     return 0;
18 }
```

Now, we have actually rewritten the entire solution. From the 300-line behemoth with gigantic functions, we have refactored the solution into few, short functions with are easy to follow. The rewritten solution is less than half the size, clocking in at less than 140 lines (the author's own solution is 120 lines). Learning to code such structured solutions comes to a large extent from experience. During a competition, we might not spend time thinking about how to structure our solutions, instead focusing on getting it done as soon as possible. However, spending 1-2 minutes thinking about how to best implement a complex solution could pay off not only in faster implementation times (such as halving the size of the program) but also in being less buggy.

Problem 4.2

Mate in One – mateinone

To sum up: implementation problems should not be underestimated in terms of implementation complexity. Work on your coding best practices and spend time practicing coding complex solutions and you will see your implementation performance improve.

Chapter Exercises

Problem 4.3

Flexible Spaces – flexiblesspaces

Permutation Encryption – permutationencryption

Jury Jeopardy – juryjeopardy

Fun House – funhouse

Settlers of Catan – settlers2

Cross – cross

Basic Interpreter – basicinterpreter

Cat Coat Colors – catcoat

Chapter Notes

Many good sources exist to become more proficient at writing readable and simple code. *Clean Code*[17] describes many principles that helps in writing better code. It includes good walk-throughs on refactoring, and shows in a very tangible fashion how coding cleanly also makes coding easier.

Code Complete[18] is a huge tome on improving your programming skills. While much of the content is not particularly relevant to coding algorithmic problems, chapters 5-19 give many suggestions on coding style.

Different languages have different best practices. Some resources on improving your skills in whatever language you code in are:

C++ *Effective C++*[20], *Effective Modern C++*[21], *Effective STL*[19], by Scott Meyers,

Java *Effective Java*[5] by Joshua Bloch,

Python *Effective Python*[25] by Brett Slatkin, *Python Cookbook*[4] by David Beazley and Brian K. Jones.

5 Time Complexity

How do you know if your algorithm is fast enough before you have coded it? In this chapter we examine this question from the perspective of *time complexity*, a common tool of algorithm analysis to determine roughly how fast an algorithm is.

We start our study of complexity by looking at a new sorting algorithm – *insertion sort*. Just like selection sort (studied in Chapter 1), insertion sort works by iteratively sorting a sequence.

5.1 The Complexity of Insertion Sort

The insertion sort algorithm works by ensuring that all of the first i elements of the input sequence are sorted. First for $i = 1$, then for $i = 2$, etc, up to $i = n$, at which point the entire sequence is sorted.

Insertion Sort

Assume we wish to sort the list a_0, a_1, \dots, a_{N-1} of N integers. If we know that the first K elements a_0, \dots, a_{K-1} numbers are sorted, we can make the list a_0, \dots, a_K sorted by taking the element a_K and inserting it into the correct position of the already-sorted prefix a_0, \dots, a_{K-1} .

For example, we know that a list of a single element is always sorted, so we can use that a_0 is sorted as a base case. We can then sort a_0, a_1 by checking whether a_1 should be to the left or to the right of a_0 . In the first case, we swap the two numbers.

Once we have sorted a_0, a_1 , we insert a_2 into the sorted list. If it is larger than a_1 , it is already in the correct place. Otherwise, we swap a_1 and a_2 , and keep going until we either find the correct location, or determine that the number was the smallest one in which case the correct location is in the beginning.

This procedure is then repeated for every remaining element.

In this section we determine how long time insertion sort takes to run. When analyzing an algorithm we do not attempt to compute the actual wall clock time an algorithm takes. Indeed, this would be nearly impossible a priori – modern

computers are complex beasts with often unpredictable behavior. Instead, we try to approximate the *growth* of the running time, as a function of the size of the input.

Competitive Tip

While it is difficult to measure exact wall-clock time of your algorithm just by analyzing the algorithm and code, it is sometimes a good idea to benchmark your solution before you submit it to the judge. This way you trade a few minutes of time (constructing the worst-case input) for avoiding many time limit exceeded verdicts. If you are unsure of your solution and the competition format penalizes you for rejected submissions, this tradeoff can have good value.

When sorting fixed-size integers the size of the input would be the number of elements we are sorting, N . We denote the time the algorithm takes in relation to N as $T(N)$. Since an algorithm often has different behaviours depending on how an instance is constructed, this time is taken to be the *worst-case* time, over every instance of N elements.




Figure 5.1: Insertion Sort sorting the sequence 2, 1, 4, 5, 3, 0.

To properly analyze an algorithm, we need to be more precise about exactly what it does. We give the following pseudo code for insertion sort:

```

1: procedure INSERTIONSORT( $A$ )  ▷ Sorts the sequence  $A$  containing  $N$  elements
2:   for  $i \leftarrow 0$  to  $N - 1$  do
3:      $j \leftarrow i$ 
4:     while  $j > 0$  and  $A[j] < A[j - 1]$  do
5:       Swap  $A[j]$  och  $A[j - 1]$ 
6:        $j \leftarrow j - 1$ 

```

To analyze the running time of the algorithm, we make the assumption that any “sufficiently small” operation takes the same amount of time – exactly 1 (of some undefined unit). We have to be careful in what assumptions we make regarding what a sufficiently small operation means. For example, sorting N numbers is not a small operation, while adding or multiplying two fixed-size numbers is. Multiplication of integers of arbitrary size is not a small operation (see the Karatsuba algorithm, Section 12.4).

In our program every line happens to represent a small operation. However, the two loops may cause some lines to execute more than once. The outer for loop will execute N times. The number of times the inner loop runs depends on how the input looks. We introduce the notation t_i to mean the number of iterations the inner loop runs during the i 'th iteration of the outer loop. These are included in figure 5.1 for every iteration.

Now we can annotate our pseudo code with the number of times each line executes.

```

1: procedure INSERTIONSORT( $A$ )  ▷ Sorts the sequence  $A$  containing  $N$  elements
2:   for  $i \leftarrow 0$  to  $N - 1$  do                                ▷ Runs  $N$  times, cost 1
3:      $j \leftarrow i$                                               ▷ Runs  $N$  times, cost 1
4:     while  $j > 0$  and  $A[j] < A[j - 1]$  do                    ▷ Runs  $\sum_{i=0}^{N-1} t_i$  times, cost 1
5:       Swap  $A[j]$  och  $A[j - 1]$                                 ▷ Runs  $\sum_{i=0}^{N-1} t_i$  times, cost 1
6:        $j \leftarrow j - 1$                                        ▷ Runs  $\sum_{i=0}^{N-1} t_i$  times, cost 1

```

We can now express $T(N)$ as

$$T(N) = N + N + \left(\sum_{i=0}^{N-1} t_i \right) + \left(\sum_{i=0}^{N-1} t_i \right) + \left(\sum_{i=0}^{N-1} t_i \right)$$

$$= 3 \left(\sum_{i=0}^{N-1} t_i \right) + 2N$$

We still have some t_i variables left so we do not truly have a function of N . We can eliminate this by realizing that in the worst case $t_i = i$. This occurs when the list we are sorting is in descending order. Each element must then be moved to the front, requiring i swaps for the i 'th element.

With this substitution we can simplify the expression:

$$\begin{aligned} T(N) &= 3 \left(\sum_{i=0}^{N-1} i \right) + 2N \\ &= 3 \frac{(N-1)N}{2} + 2N \\ &= \frac{3}{2}(N^2 - N) + 2N \\ &= \frac{3}{2}N^2 + \frac{N}{2} \end{aligned}$$

This function grows quadratically with the number of elements of N . Since the approximate growth of the time a function takes is assigned such importance a notation was developed for it.

5.2 Asymptotic Notation

Almost always when we express the running time of an algorithm we use what is called *asymptotic notation*. The notation captures the behavior of a function as its arguments grow. For example, the function $T(N) = \frac{3}{2}N^2 + \frac{N}{2}$ which described the running time of insertion sort, is bounded by $c \cdot N^2$ for large N , for some constant c . We write

$$T(N) = O(N^2)$$

to state this fact.

Similarly, the linear function $2N + 15$ is bounded by $c \cdot N$ for large N , with $c = 3$, so $2N + 15 = O(N)$. The notation only captures upper bounds though (and not the actual rate of growth). We could therefore say that $2N + 15 = O(N^2)$, even though this particular upper bound is very lax. However, N^2 is not bounded by $c \cdot N$ for any constant c when N is large, so $N^2 \neq O(N)$. This also corresponds to intuition – quadratic functions grows much faster than linear functions.

Definition 5.1 — O -notation

Let f and g be non-negative functions from $\mathbb{R}_{\geq 0}$ to $\mathbb{R}_{\geq 0}$. If there exists positive constants n_0 and c such that $f(n) \leq cg(n)$ whenever $n \geq n_0$, we say that $f(n) = O(g(n))$.




Figure 5.2: All linear functions are eventually outgrown by N^2 , so $an + b = O(n^2)$.

Intuitively, the notation means that $f(n)$ grows *slower than or as fast as* $g(n)$, within a constant factor. Any quadratic function $an^2 + bn + c = O(n^2)$. Similarly, any linear function $an + b = O(n^2)$ as well. This definition implies that for two functions f and g which are always within a constant factor of each other, we have that both $f(n) = O(g(n))$ and $g(n) = O(f(n))$.

We can use this definition to prove that the running time of insertion sort is $O(N^2)$, even in the worst case.

Example 5.1 Prove that $\frac{3}{2}N^2 + \frac{N}{2} = O(N^2)$.

Proof. When $N \geq 1$ we have $N^2 \geq N$ (by multiplying both sides with N). This means that

$$\frac{3}{2}N^2 + \frac{N}{2} \leq \frac{3}{2}N^2 + \frac{N^2}{2} = \frac{4}{2}N^2 = 2N^2$$

for $N \geq 1$. Using the constants $c = 2$ and $n_0 = 1$ we fulfill the condition from the definition. \square

For constants k we say that $k = O(1)$. This is a slight abuse of notation, since neither k nor 1 are functions, but it is a well-established abuse. If you prefer, you can instead assume we are talking about functions $k(N) = k$ and $1(N) = 1$.

Competitive Tip

The following table describes approximately what complexity you need to solve a problem of size n if your algorithm has a certain complexity when the time limit is about 1 second.

Complexity	n
$O(\log n)$	$2^{(10^7)}$
$O(\sqrt{n})$	10^{14}
$O(n)$	10^7
$O(n \log n)$	10^6
$O(n\sqrt{n})$	10^5
$O(n^2)$	$5 \cdot 10^3$
$O(n^2 \log n)$	$2 \cdot 10^3$
$O(n^3)$	300
$O(2^n)$	24
$O(n2^n)$	20
$O(n^2 2^n)$	17
$O(n!)$	11

Table 5.1: Approximations of needed time complexities

Note that this is in no way a general rule – while complexity does not bother about constant factors, wall clock time does!

Complexity analysis can also be used to determine *lower bounds* of the time an algorithm takes. To reason about lower bounds we use Ω -notation. It is similar to O -notation except it describes the reverse relation.

Definition 5.2 – Ω -notation

Let f and g be non-negative functions from $\mathbb{R}_{\geq 0}$ to $\mathbb{R}_{\geq 0}$. If there exists positive constants n_0 and c such that $cg(n) \leq f(n)$ for every $n \geq n_0$, we say that $f(n) = \Omega(g(n))$.

We know that the complexity of insertion sort has an upper bound of $O(N^2)$ in the worst-case, but does it have a lower bound? It actually has the same lower

bound as upper bound, i.e. $T(N) = \Omega(N^2)$.

Example 5.2 Prove that $\frac{3}{2}N^2 + \frac{N}{2} = \Omega(N^2)$.

Proof. When $N \geq 1$, we have $\frac{3}{2}N^2 + \frac{N}{2} \geq N^2$. Using the constants $c = 1$ and $n_0 = 1$ we fulfill the condition from the definition. \square

In this case, both the lower and the upper bound of the worst-case running time of insertion sort coincided (asymptotically). We have another notation for when this is the case:

Definition 5.3 — Θ -notation

If $f(n) = O(g(n))$ and $f(n) = \Omega(g(n))$, we say that $f(n) = \Theta(g(n))$.

Thus, the worst-case running time for insertion sort is $\Theta(n^2)$.

There are many ways of computing the time complexity of an algorithm. The most common case is when a program has K nested loops, each of which performs $O(M)$ iterations. The complexity of these loops are then $O(M^K \cdot f(N))$ if the inner-most operation takes $O(f(N))$ time. In Chapter 12, you also see some ways of computing the time complexity of a particular type of recursive solution, called Divide and Conquer algorithms.

Exercise 5.1. Find a lower and an upper bound that coincide for the best-case running time for insertion sort.

Exercise 5.2. Give a $\Theta(n)$ algorithm and a $\Theta(1)$ algorithm to compute the sum of the n first integers.

Exercise 5.3. Prove, using the definition, that $10n^2 + 7n - 5 + \log^2 n = O(n^2)$. What constants c, n_0 did you get?

Exercise 5.4. Prove that $f(n) + g(n) = \Theta(\max\{f(n), g(n)\})$ for non-negative functions f and g .

Exercise 5.5. Determine whether, with proof:

- 1) Is $2^{n+1} = O(2^n)$?
- 2) Is $2^{2n} = O(2^n)$?

Exercise 5.6. Prove that $(n + a)^b = \Theta(n^b)$ for positive constants a, b .

Amortized Complexity

Consider the following algorithm:

```
1: procedure COUNT_OCCURANCES( $A, v$ )  $\triangleright$  Count the number of occurrences of the  
   value  $v$  in the sequence  $A$   
2:    $N \leftarrow \text{length of } A$   
3:    $i \leftarrow 0$   
4:    $\text{ans} \leftarrow 0$   
5:   while  $i \neq N$  do  
6:     while  $i < N$  and  $A[i] \neq v$  do  
7:        $i \leftarrow i + 1$   
8:     if  $i \neq N$  then  
9:        $\text{ans} \leftarrow \text{ans} + 1$   
10:     $i \leftarrow i + 1$ 
```

It computes the number of occurrences of a certain value in a sequence (in a somewhat cumbersome way). We repeatedly scan through the sequence A to find the next occurrence of v . Whenever we find one, we increase the answer by 1 and resume our scanning. This procedure is continued until we have scanned through the entire sequence.

What is the time complexity of this procedure? How do we handle the fact that neither the outer nor the inner loop of the algorithm repeats a predictable number of iterations each time? For example, the outer loop would iterate N times if every element equals v , but only once if it does not contain the value at all. Similarly, the number of iterations of the inner loop depends on the *position* of the elements in A which equals v .

We find help in a technique called *amortized complexity*. The concept of amortization is best known from loans. It represents the concept of periodically paying of some kind of debt until it has been paid off in full.

When analyzing algorithms, the “debt” is the running time the algorithm has. We try to prove that the algorithm has a given running time by looking at many, possibly uneven parts that together over the long run sums up to the total running time, similarly to how a loan can be paid of by amortization.

A quick analysis of the counting algorithm gives us a good guess that the algorithm runs in time $\Theta(N)$. It should be clear that it is the inner loop at lines 6-7 that dominates the running time of the algorithm. The question is how we

compute the number of times it executes even though we do not know how many times it executes nor how many iterations each execution takes? We try the amortization trick by looking at how many iterations it performs over all those executions, no matter how many they are. Assume that the loop is run k times (**including the final time when the condition first is false**) and each run iterates b_i times ($1 \leq i \leq k$). We claim that

$$b_1 + b_2 + \cdots + b_k = \Theta(N)$$

Our reasoning is as follows. There are two ways the variable i can increase. It can either be increased inside the loop at line 7, or at line 10. If the loop executes N times in total, it will certainly complete and never be executed again since the loop at line 5 completes too. This gives us $\sum_{i=1}^k b_i = O(N)$.

On the other hand, we get one iteration for every time i is increased. If i is increased on line 7, it was done within a loop iteration. If i is increased on line 9, we instead count the final check if the loop just before it once. Each addition of i seems happen together with an iteration of the loop, so $\sum_{i=1}^k b_i = \Omega(N)$. Together, these two results prove our claim.

This particular application of amortized complexity is called the *aggregate method*.

Exercise 5.7. Consider the following method of implementing adding 1 to an integer represented by its binary digit.

```
1: procedure BINARYINCREMENT( $D$ )
2:    $i \leftarrow 0$ 
3:   while  $D[i] = 1$  do                                 $\triangleright$  Add 1 to the  $i$ 'th digit
4:      $D[i] = 0$        $\triangleright$  We add 1 to a 1 digit, resulting in a 0 digit plus a carry
5:      $i \leftarrow i + 1$ 
6:    $D[i] = 1$        $\triangleright$  We add 1 to a digit not resulting in a carry
```

The algorithm is the binary version of the normal addition algorithm where the two addends are written above each other and each resulting digit is computed one at a time, possibly with a carry digit.

What is the amortized complexity of this procedure over 2^n calls, if D starts out as 0?

5.3 NP-complete problems

Of particular importance in computer science are the problems that can be solved by algorithms running in polynomial time (often considered to be the “tractable” problems). For some problems we do not yet know if there is an algorithm whose time complexity is bounded by a polynomial. One particular class of these are the ***NP-complete problems***. They have the property that they are all reducible to one another, in the sense that a polynomial-time algorithm to any one of them yields a polynomial-time algorithm to all the others. Many of these NP-complete problems (or problems reducible to such a problem, a property called ***NP-hardness***) appear in algorithmic problem solving. It is good to know that they exist and that it is unlikely that you can find a polynomial-time solution. During the course of this book, you will occasionally see such problems with their NP-completeness mentioned.

5.4 Other Types of Complexities

There are several other types of complexities aside from the time complexity. For example, the ***memory complexity*** of an algorithm measures the amount of memory it uses. We use the same asymptotic notation when analyzing memory complexity. In most modern programming competitions, the allowed memory usage is high enough for the memory complexity not be a problem – if you get memory limit problems you also tend to have time limit problems. However, it is still of interest in computer science (and thus algorithmic problem solving) and computer engineering in general.

Another common type of complexity is the ***query complexity***. In some problems (like the Guessing Problem from chapter 1), we are given access to some kind of external procedure (called an *oracle*) that computes some value given parameters that we provide. A procedure call of this kind is called a *query*. The number of queries that an algorithm makes to the oracle is called its query complexity. Problems where the algorithm is allowed access to an oracle often bound the number of queries the algorithm may make. In these problems the query complexity of the algorithm is of interest.

5.5 The Importance of Constant Factors

In this chapter, we have essentially told you not to worry about the constant factors that the Θ notation hides from you. While this is true when you are

solving problems in theory, only attempting to get a good asymptotic time complexity, constant factors can unfortunately be of large importance when implementing the problems subject to time limits.

Speeding up a program that has the correct time complexity but still gets *time limit exceeded* when submitted to an online judge is half art and half engineering. The first trick is usually to generate the worst-case test instance. It is often enough to create a test case where the input matches the input limits of the problem, but sometimes your program behaves differently depending on how the input looks. In these cases, more complex reasoning may be required.

Once the worst-case instance has been generated, what remains is to improve your code until you have gained a satisfactory decrease in how much time our program uses. When doing this, focus should be placed on those segments of your code that takes the longest wall-clock time. Decreasing time usage by 10% in a segment that takes 1 second is clearly a larger win than decreasing time usage by 100% in a segment that takes 0.01 seconds.

There are many tricks to improving your constant factors, such as:

- using symmetry to perform less calculations
- precomputing oft-repeated expressions, especially involving trigonometric functions
- passing very large data structures by reference instead of by copying
- avoiding to repeatedly allocate large amounts of memory
- using SIMD (single instruction, multiple data) instructions

5.6 Additional Exercises

Exercise 5.8. Prove that if $a(x) = O(b(x))$ and $b(x) = O(c(x))$, then $a(x) = O(c(x))$, i.e. asymptotic notation is a *transitive* property. This means functions can be ordered by their asymptotic growth.

Exercise 5.9. Order the following functions by their asymptotic growth *with proof*!

- x
- \sqrt{x}

- x^2
- 2^x
- e^x
- $x!$
- $\log x$
- $\frac{1}{x}$
- $x \log x$
- x^3

Exercise 5.10. 1) Prove that if $a(x) = O(b(x))$ and $c(x) = O(d(x))$, then $a(x) + c(x) = O(b(x) + d(x))$.

2) Prove that if $a(x) = O(b(x))$ and $c(x) = O(d(x))$, then $a(x) \cdot c(x) = O(b(x) \cdot d(x))$.

5.7 Chapter Notes

Advanced algorithm analysis sometimes use complicated discrete mathematics, including number theoretical (as in Chapter 19) or combinatorial (as in Chapter 17) facts. *Concrete Mathematics* [15] by Donald Knuth, et al, does a thorough job on both accounts.

An Introduction to the Analysis of Algorithms [11] by Sedgewick and Flajolet has a more explicit focus on the analysis of algorithms, mainly discussing combinatorial analysis.

The study of various kinds of complexities constitute a research area called *computational complexity theory*. *Computational Complexity* [22] by Papadimitriou is a classical introduction to computational complexity, although *Computational Complexity: A Modern Approach* [3] by Arora and Barak is a more modern textbook, with recent results that the book by Papadimitriou lack.

While complexity theory is mainly concerned about the limits of specific computational models on problems that *can* be solved within those models, what can not be done by computers is also interesting. This is somewhat out of scope for an algorithmic problem solving book (since we are interested in those problems which can be solved), but is still of general interest. A book

on e.g. automata theory (such as *Introduction to Automata Theory, Languages, and Computation* [28] by Ullman et al) can be a good compromise, mixing both some foundations of the theory of computation with topics more applicable to algorithms (such as automata and languages).

6 Data Structures

Solutions to algorithmic problems consists of two constructs – algorithms and *data structures*. Data structures are used to organize the data that the algorithms operate on. For example, the array is such a data structure.

Many data structures has been developed to handle particular common operations we need to perform on data quickly. In this chapter we discuss some of the basic data structures used in programming. We have chosen an approach that is perhaps more theoretic than many other problem solving texts when it comes to the basic structures. In particular, we have chosen to explicitly discuss not only the data structures themselves and their complexities, but also their implementations. We do this mainly because we believe that their implementations show useful algorithmic techniques. While you may feel that you can simply skip this chapter if you are familiar with all the data structures, we advise that you still read through the sections for the structures for which you lack confidence in their implementations.

6.1 Dynamic Arrays

The most basic data structure is the *fixed-size array*. It consists of a fixed size block of memory and can be viewed as a sequence of N variables of the same type T . It supports the operations:

- $T\ arr[] = \text{new } T[size]$: creating a new array, with a given size.

Complexity: $\Theta(1)$ ¹

- $\text{delete } arr$: deleting an existing array.

Complexity: $\Theta(1)$

- $arr[index]$: accessing the value in a certain access a value.

Complexity: $\Theta(1)$

¹This complexity is debatable, and highly dependent on what computational model one uses. In practice, this is roughly “constant time” in most memory management libraries used in C++. In all Java and Python implementations we tried, it is instead linear in the size.

In Chapter 2 we saw how to create fixed-size arrays where we knew the size beforehand. In C++ we can create fixed-size arrays using an expression as size instead. This is done using the syntax above, for example:

```
int size = 5;
int arr[] = new int[size];
arr[2] = 5;
cout << arr[2] << endl;
delete arr;
```

Exercise 6.1. What happens if you try to create an array with a negative size?

The fixed-size array can be used to implement a more useful data structure, the *dynamic array*. This is an array that can change size when needed. For example, we may want to repeatedly insert values in the array without knowing the total number of values beforehand. This is a very common requirement in programming problems. In particular, we want to support two additional operations in addition to the operations supported by the fixed-size array.

- `insert(pos, val)`: inserting a value in the array at a given position.

Amortized complexity: $\Theta(\text{size} - \text{pos})$

Worst case complexity: $\Theta(\text{size})$

- `remove(pos)`: erase a position in an array.

Complexity: $\Theta(\text{size} - \text{pos})$

The complexities we list above are a result of the usual implementation of the dynamic array. A key consequence of these complexities is that addition and removal of elements to the end of the dynamic array takes $\Theta(1)$ amortized time.

A dynamic array can be implemented in numerous ways, and underlies the implementation of essentially every other data structure that we use. A naive implementation of a dynamic array is using a fixed-size array to store the data and creating a new one with the correct size whenever we add or remove an element, copying the elements from the old array to the new array. The complexity for this approach is linear in the size of the array for every operation that changes the size since we need to copy $\Theta(\text{size})$ elements during every update. We need to do better than this.

To achieve the targeted complexity, we can instead modify this naive approach slightly by not creating a new array *every* time we have to change the size of the dynamic array. Whenever we need to increase the size of the dynamic array, we

create a fixed-size array that is *larger* than we actually need it to be. For example, if we create a fixed-size array with n more elements than our dynamic array needs to store, we will not have to increase the size of the backing fixed-size array until we have added n more elements to the dynamic array. This means that a dynamic array does not only have a size, the number of elements we currently store in it, but also a capacity, the number of elements we *could* store in it. See Figure 6.1 for a concrete example of what happens when we add elements to a dynamic array that is both within its capacity and when we exceed it.




Figure 6.1: The resizing of an array when it overflows its capacity.

To implement a dynamic array in C++, we could use a structure storing as members the capacity, size and backing fixed-size array of the dynamic array. An example of how such a structure could look can be found in Listing 6.1.

Listing 6.1 The dynamic array structure

```
1  struct DynamicArray {
2      int capacity;
3      int size;
4      int backing[];
5
6      DynamicArray() {
7          capacity = 10;
8          size = 0;
9          backing = new int[10];
10     }
11 };
```

We are almost ready to add and remove elements to our array now. First, we

need to handle the case where insertion of a new element would result in the size of the dynamic array would exceed its capacity, that is when $size = capacity$. Our previous suggestion was to allocate a new, bigger one, but just how big? If we always add, say, 10 new elements to the capacity, we have to perform the copying of the old elements with every 10th addition. This still results in additions to the end of the array taking linear time on average. There is a neat trick that avoids this problem; creating the new backing array with **double** the current capacity.

This ensures that the complexity of all copying needed for an array up to some certain capacity have an amortized complexity of $\Theta(cap)$. Assume that we have just increased the capacity of our array to cap , which required us to copy $\frac{cap}{2}$ elements. Then, the previous increase will have happened at around capacity $\frac{cap}{2}$ and took time $\frac{cap}{4}$. The one before that occurred at capacity $\frac{cap}{4}$ and so on.

We can sum up all of this copying:

$$\frac{cap}{2} + \frac{cap}{4} + \dots \leq cap$$

using the formula for the sum of a geometric series.

Since each copy is assumed to take $\Theta(1)$ time, the total time to create this array was $\Theta(cap)$. As $\frac{cap}{2} \leq size \leq cap$, this is also $\Theta(size)$, meaning that adding $size$ elements to the end of the dynamic array takes amortized $\Theta(size)$ time.

When implementing this in code, we use a function that takes as argument the capacity we require the dynamic array to have and ensures that the backing array have at least this size, possibly by creating a new one double in size until it is sufficiently large. Example code for this can be found in Listing 6.2.

With this method in hand, insertion and removal of elements is actually pretty simple. Whenever we remove an element, we simply need to move the elements coming after it in the dynamic array forward one step. See Figure 6.2 for an illustration of an element being removed.

When adding an element, we reverse this process by moving the elements coming after the position we wish to insert a new element at one step towards the back. An example of this is shown in Figure 6.3.

Exercise 6.2. Implement insertion and removal of elements in a dynamic array.


Dynamic arrays are called vectors in C++ (Section 3.1). They have the same complexities as the one described at the beginning of this section.

Listing 6.2 Ensuring that a dynamic array have sufficient capacity

```

1 void ensureCapacity(int need) {
2     while (capacity < need) {
3         int newBacking[] = new int[2 * capacity];
4         for (int i = 0; i < size; i++)
5             newBacking[i] = backing[i];
6         delete backing;
7         backing = newBacking;
8     }
9 }

```

**Figure 6.2:** The removal of the element 2 at index 2.

Exercise 6.3. How can *any* element be removed in $\Theta(1)$ if we ignore the ordering of values in the array?

6.2 Stacks

The **stack** is a data structure that contains an ordered lists of values and supports the following operations:

- **push(val):** inserting a value at the top of the stack.
Amortized complexity: $\Theta(1)$
- **pop():** remove the value at the top of the stack.
Complexity: $\Theta(1)$
- **top():** get the value at the top of the stack.
Complexity: $\Theta(1)$

The structure is easily implemented with the above time complexities using




Figure 6.3: The insertion of the element 2 into index 2.

a dynamic vector. After all, the vector supports exactly the same operations that a stack requires. In C++, the stack is called a `stack` (Section 3.4).

Exercise 6.4. Implement a stack using a dynamic vector.

6.3 Queues

The *queue* is, like the vector and the stack, an ordered list of values. Instead of removing and getting values from the end like the stack, it gets the value from the *front*. The supported operations are thus:

- `push(val)`: inserting a value at the end of the queue.
Amortized complexity: $\Theta(1)$
- `pop()`: remove the value at the front of the queue.
Amortized complexity: $\Theta(1)$
- `front()`: get the value at the front of the queue.
Complexity: $\Theta(1)$

As previously seen, C++ has an implementation of the queue called `queue` (Section 3.3).

Implementing a queue can also be done using a vector. After all, the operations and complexities are nearly the same; only removing the value of the front is wrong. To fix this, one can simply hold a *pointer* to what the front of the queue is in the vector. Removing the front element is then equivalent to moving the pointer forward one step. To see this how this would work in practice, see an example push and pop operation in Figure 6.4.

Exercise 6.5. Implement a queue using a vector.




Figure 6.4: Pushing and popping elements in a queue

Exercise 6.6. A queue can also be implemented using two stacks. How?

In a similar manner, a stack can also be implemented using two queues. How?

Exercise 6.7. This naive suggestion of a queue implementation suffers a slight problem. After pushing and popping k elements, the backing dynamic array has size at least k even though none of its elements are in use, thus occupying memory unnecessarily. Devise a strategy that ensures the backing dynamic array never uses more than cn elements (where n is the current size of the queue) for some constant c .

6.4 Priority Queues

Now, let us look at our first more complicated data structure. The *priority queue* is an unordered bag of elements, from which we can get and remove the largest one quickly. It supports the operations

- $\text{push}(val)$: inserting a value into the heap.
Complexity: $O(\log n)$
- $\text{pop}()$: remove the largest value in the heap.
Complexity: $O(\log n)$
- $\text{getMax}()$: get the largest value in the heap.
Complexity: $\Theta(1)$

This is implemented as `priority_queue` in C++ (Section 3.5), although one often instead use a `set` which not only supports the same operations with the same complexities, but also supports erasing elements.

The backing implementation of the priority queue structure we study is called a *heap*. The heap itself will be implemented using another data structure called a *binary tree*, which we need to describe first.

Binary Trees

A *binary tree* is a rooted tree where every vertex have either 0, 1 or 2 children. In Figure 6.5a, you can see an example of a binary tree.




Figure 6.5: Examples of binary trees.

We call a binary tree *complete* if every level of the tree is completely filled, except possibly the bottom one. If the bottom level is not filled, all the vertices need to be as far as possible to the left. In a complete binary tree, we can order every vertex as we do in Figure 6.5b, i.e. from the top down, left to right at each layer.

The beauty of this numbering is that we can use it to store a binary tree in a vector. Since each vertex is given a number, we can map the number of each vertex into a position in a vector. The n vertices of a complete binary tree then occupy all the indices $[1, n]$. An important property of this numbering is that it is easy to compute the number of the parent, left child and right child of a vertex. If a vertex has number i , the parent simply have number $\lfloor \frac{i}{2} \rfloor$, the left child has number $2i$ and the right child has number $2i + 1$.

Exercise 6.8. Prove that the above properties of the numbering of a complete binary tree hold.

A heap is implemented as a complete binary tree, which we in turn implement

by a backing vector in the manner described. In the implementation, we use the following convenience functions:

- 1: **function** PARENT(i) **return** $i/2$
- 2: **function** LEFT(i) **return** $2i$
- 3: **function** RIGHT(i) **return** $2i + 1$

Note: if you use a vector to represent a complete binary tree in this manner it needs to have the size $n + 1$ where n is the number of vertices, since the tree numbering is 1-indexed and the vector is 0-indexed!

Heaps

A *heap* is a special kind of complete binary tree. More specifically, it should always satisfy the following property: *a vertex always have a higher value than its immediate children*. Note that this condition acts transitively, which means that a vertex also has a higher value than its grand-children, and their children and so on. In particular, a consequence of this property is that the root of the tree is always be the largest value in the heap. As it happens, this property is exactly what we need to quickly get the maximum value of the heap. You can see an example of a heap in Figure 6.6




Figure 6.6: A heap of the elements 1, 4, 4, 5, 7, 10.

We thus start our description of a heap somewhat backwards, with the function needed to get the largest element:

- 1: **procedure** GET-MAX($tree$) **return** $tree[1]$

The complicated operations on the heap is to add and remove elements ensuring that the heap keeps satisfying this property. We start with the simplest one, adding a new element. Since we represent the heap using a vector, adding a new element to a heap can be done by appending the element to the vector. In this manner, we ensure that the underlying binary tree is still complete. However,

it may be that the value we added is now larger than its parent. If this is the case, we can fix the violation of our heap property by swapping the value with its parent. This does not guarantee that the value still is not larger than its parent. In fact, if the newly added element is largest in the heap, it would have to be repeatedly swapped up to the top! This procedure, of moving the newly added element up in the tree until it is no longer larger than its parent (or it becomes the root) is called *bubbling up*:

```

1: procedure BUBBLE-UP(idx, tree)
2:   while idx > 1 do
3:     if tree[idx] > tree[Parent(idx)] then
4:       Swap tree[idx] and tree[Parent(idx)]
5:       idx ← Parent(idx)
6:     else
7:       break

```

Pushing a value now reduces to appending it to the tree and bubbling it up. You can see this procedure in action in Figure 6.7.

```

1: procedure PUSH(x, tree)
2:   tree.push_back(x)
3:   Bubble-Up(tree.size() - 1, tree)

```




Figure 6.7: Adding a new value and bubbling it up.

Removing a value is slightly harder. First of, the tree will no longer be a binary tree – it is missing its root! To rectify this, we can take the last element of the tree and put it as root instead. This keeps the binary tree complete, but may cause it to violate the heap property since our new root may be smaller than either or both of its children.

The solution to this problem is similar as to that of adding an element. Instead of bubbling up, we bubble it down by repeatedly swapping it with one of its children until it no longer is greater than any of its children. The only question mark is which of its children we should bubble down to, in case the element is smaller than both of its children. The answer is clearly the largest of the two children. If we take the smaller of the two children, we will again violate the heap property. Just as with pushing, popping a value is done by a combination of removing the value and fixing the heap to satisfy the heap property again.

```
1: procedure REMOVE-MAX( $x$ ,  $tree$ )
2:    $tree[1] \leftarrow tree[tree.size() - 1]$ 
3:    $tree.pop\_back()$ 
4:   Bubble-Down(1,  $tree$ )
5: procedure BUBBLE-DOWN( $idx$ ,  $tree$ )
6:   while true do
7:      $largest \leftarrow idx$ 
8:     if Left( $idx$ ) <  $tree.size()$  and  $tree[Left(idx)] > tree[largest]$  then
9:        $largest \leftarrow Left(idx)$ 
10:    if Right( $idx$ ) <  $tree.size()$  and  $tree[Right(idx)] > tree[largest]$  then
11:       $largest \leftarrow Right(idx)$ 
12:    if  $largest = idx$  then
13:      break
14:    else
15:      Swap  $tree[idx]$  and  $tree[largest]$ 
16:       $idx \leftarrow largest$ 
```

A final piece of our analysis is missing. It is not yet proven that the time complexity of adding and removing elements are indeed $O(\log n)$. To do this, we first need to state a basic fact of complete binary trees: their height is at most $\log_2 n$. This is easily proven by contradiction. Assume that the height of the tree

is at least $\log_2 n + 1$. We claim that any such tree must have strictly more than n vertices. Since all but the last layers of the tree must be complete, it must have at least $1 + 2 + \dots + 2^{\log_2 n} = 2^{\log_2 n + 1} - 1$ vertices. But $2^{\log_2 n + 1} - 1 = 2n - 1 > n$ for positive n – the tree has more than n vertices. This means that a tree with n vertices cannot have more than height $\log_2 n$.

The next piece of the puzzle is analyzing just how many iterations the loops in the bubble up and bubble down procedures can perform. In the bubble up procedure, we keep an index to a vertex that, for every iteration, moves up in the tree. This can only happen as many times as there are levels in the tree. Similarly, the bubble down procedure tracks a vertex that moves down in the tree for every iteration. Again, this is bounded by the number of levels in the tree. We are forced to conclude that since the complexity of each iteration is $\Theta(1)$ as they only perform simple operations, the complexities of the procedures as a whole are $O(\log n)$.

Problem 6.1

Binary Heap – heap

Exercise 6.9. Prove that adding an element using Push never violates the heap property.

Exercise 6.10. To construct a heap with n elements by repeatedly adding one at a time takes $O(n \log n)$ time, since the add function takes $O(\log n)$ time in the worst case.

One can also construct it in $\Theta(n)$ time in the following way: arbitrarily construct a complete binary tree with all the n elements, and then call Bubble-Down on each of the elements in reverse order $n, n - 1, \dots, 2, 1$.

Prove that this correctly constructs a heap, and that it takes $\Theta(n)$ time.

6.5 Bitsets

We move on to the simplest data structure of the chapter. The *bitset* can be viewed as a specialization of a static array in the case where the values stored are booleans, i.e. supporting only the operations of setting and getting values in the array.

The idea behind it is simple. Booleans are essentially values 0 (false) or 1 (true), i.e. equivalent to a binary digit. What else consists of binary digits – integers. In the chapter on programming, we noted that the memory of a

computer is just a long sequence of binary digits. Any interpretation of what these digits actually mean is up to us. For example, your typical 64-bit integers could just as well represent an array of 64 booleans, indexed from 0 to 63. That is the *bit* part.

An array of booleans of size N can also be interpreted as a subset of the integers $\{0, 1, \dots, N - 1\}$. If the i 'th value of the array is true, we say that i is in the subset, otherwise it is not. This is the *set* part. An example of this equivalence between integers and subsets is visualized in Figure 6.8.

$$90 = \begin{array}{c} 7 \quad 6 \quad 5 \quad 4 \quad 3 \quad 2 \quad 1 \quad 0 \\ \boxed{\begin{array}{|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 0 & 1 & 1 & 0 & 1 & 0 \\ \hline \end{array}} = \{1, 3, 4, 6\}$$

Figure 6.8: The equivalence between 90 and set of elements $\{1, 3, 4, 6\}$

Interpreting N -bit integers as subsets of $\{0, 1, \dots, N - 1\}$ is surprisingly useful. It allows us to easily pass subsets as parameters to functions, and even use them as indexes into a vector rather than using a map!

To operate on a bitset, we use *bitwise operators* in C++. To construct the representation of the set $\{i\}$, the expression $1 \ll i$ is used. This operator is called the **left-shift operator**. It works by taking the binary representation of the left operand and moves it i steps to the left. Since the binary representation of 1 is one, the expression results in a binary number with an 1 in its i 'th place, which is the correct representation of the set.

To take the union of two bitsets x and y , we use the bitwise or operator: $x \mid y$. This operator takes the two integers and gives us a new one, with a 1 as the i 'th digit if either x or y had a 1 as their i 'th digit.

Similarly, the bitwise and operator $x \& y$ computes the intersection of the two sets.

This allows us to check for set membership of i in the bitset x using the expression $x \& (1 \ll i)$ which is 0 if i was not a member of the set, and $1 \ll i$ otherwise.

One can also compute the symmetric difference of two sets with the bitwise exclusive or operator, $x \wedge y$. An element is a member of the bitset if it was a member of exactly one of x and y . Thus, one can toggle the presence of an element i in a bitset using $x \wedge (1 \ll i)$.

Finally, the bitwise negation operator \bar{x} computes the complement of a set.

To compute the size of a bitset, the most common compilers support the macro `__builtin_popcount(x)` which returns the number of 1 digits in x .

There are several neat tricks involving bitsets. Some worth mentioning are:

- computing the representation of $\{0, 1, \dots, N - 1\}$ with the expression $(1 \ll N) - 1$,
- removing the lowest-numbered element of the set with $x \& (x - 1)$,
- retrieve the lowest-numbered element using $x \& -x$, and
- iterating through all (non-empty) subsets of a bitset x using the loop

```
for (int sub = x; sub != 0; sub = (sub - 1) & x)
```

Exercise 6.11. Given a bitset, use bitwise operators to compute the next higher bitset with the same number of elements.

6.6 Hash Tables

In Section 3.6, we looked at a data structure called `map`, which stored a mapping from a set of *keys* to their corresponding *values*. The underlying implementation of this data structure in STL is called a *self-balancing tree*, something we explore further in Chapter 22.1. If one is willing to forego having the structure being sorted by keys, a *hash table* can be used instead. Hash tables are an easier implementation of the `map`, and can, depending on read and write patterns, be faster than the self-balancing tree implementation. It exists in STL as well, called `unordered_map`.

The operations the hash table supports are:

- `set(x, y)`: set the value of key x to y .
Complexity: expected $\Theta(1)$
- `get(x)`: return the value of key x .
Complexity: expected $\Theta(1)$
- `erase(x)`: remove the key x .
Complexity: expected $\Theta(1)$
- `contains(x)`: returns whether the table contains the key x or not.
Complexity: expected $\Theta(1)$

The main idea behind supporting these operations quickly is essentially the same as that of the dynamic array. Assume that the set of all possible keys, the *universe*, were the integers $0, 1, \dots, N$ for some fixed N . If so, we could easily implement the above operations by storing the values in a dynamic array of this size. What do we do if this is not the case?

We apply the concept of *hashing*. Imagine your universe consists of all possible integers that fits in an `int`. Clearly, we could not use a dynamic array of that size (2^{32}) to store the table – at least not in competitions. Instead, the goal of hashing is to shrink this huge universe into a small one, that we could store in an array. For example, we could take the last K digits of the key for some small K (i.e take $\text{key} \bmod 10^K$) and use it as the index into an array for the value of that key. There are only 10^K such keys, which can fit in an array for a small K .

Given such a mapping, we can store the key and value as a pair on the corresponding index in the array. This is illustrated for $K = 1$ in Figure 6.9.




Figure 6.9: An example of where the hash table would store the values of a few keys when $K = 1$.

This kind of transformation, that takes an arbitrarily large value (any integer) and maps it into a set of constant size, is called a *hash function*².

Unfortunately, we are bitten by one of the fundamental limitations of mathematics – if a function maps a set to a smaller one, there must be at least two values which map to the same value. Our table can thus be subject to a *collision*, two keys that map to the same index in the array. Resolving this situation is actually straightforward. Instead of using the table to store a single key-value pair, we can store a dynamic array of key-value pairs³.

This complicates an implementation only slightly. To retrieve a value from the hash table, we look up the correct index of the array backing the table and check all key-value pairs stored in the array at that index. If we find a pair with

²You might have heard about a version of this often used in cryptography, the *cryptographic hash function*, which aims to provide stronger guarantees than we care about

³Traditionally, most introductions to the hash table decides to instead use a data structure called a *linked list* to pairs that collide. We elect not to, since it's mostly a question about real-world performance, and expect you to mostly use the implementation from STL if you use C++.




Figure 6.10: When several keys with the same hash value are stored, we save all pairs in a sub-array.

the correct key, we return its associated value.

Setting the value of a key becomes a two-step process. First, we need to check the correct sub-array to see if it already contains a key-value pair with the same key. If it does, we update the value of the pair to the new value. If it does not, we insert the new key and value as a pair in that list.

We are not done yet. If the observant reader attempts to reason about the time complexity of the above solution, they would notice a problem. When all the keys inserted into the table have the same hash value – in the case of our simple function, the same last K digits – they map into the same position in the backing array. Searching through the sub-array to determine if it already contains the key would then become a linear time operation as our hash table degrades into a single array! Similarly, if we pick a K that is much smaller than the number of keys, we will have to scan through very large arrays. For example, when $K = 1$ and we have 10^6 keys, every operation would need to scan through on average 10^5 values. This leaves us with two issues to resolve – how do we choose a reasonable hash function, and how big must our hash table be?

As for the second question, the goal is essentially to always keep the size as a multiple of the number of keys. When the number of keys grows too large for the table, the size is increased and the hash of all keys are recomputed with regards to the new table size. Typically, one would double the size of the hash table when it grows too dense, similar to how dynamic arrays were implemented. This makes sure inserting keys keep the time complexity amortized constant over resizing. Furthermore, if one starts with a table size that is a power of two, doubling it keeps the size a power of two. In the remainder of the section, we will assume it is equal to 2^N at the time of hashing.

The first question depends on the context. For the sake of programming problems where the hash table keys are 32-bit integers, one can usually use

something taking the upper N bits of Ax as the hash, where A is a large constant odd constant, i.e:

```
(A * x) >> (64 - N)
```

Using an odd constants makes sure A is relatively prime to 2^N , which gives us more randomness. Otherwise, Ax in binary would just have a few extra zeroes at the end.

Exercise 6.12. What happens if one takes the lower N bits of the product Ax as the hash instead?

Exercise 6.13. The above hash function is easily broken when x can be a 64-bit integer – how?

When we now move on to trying to compute the complexity of a hash table, we assume that a randomly chosen key has the same probability of being mapped to any of the possible hash values. By the resizing trick above, we can also assume that the table size is always within a constant of the number of keys.

Assume that K operations are performed on a table of some size M . The complexity of the i 'th operation is exactly that of the length of the sub-array to which the key involved in the operation would map. The expected complexity of the operation is then equal to the expected length of the sub-array. Let a_j be 0 if the j 'th operation inserted a key into this sub-array, and 1 otherwise (for $1 \leq j < i$). The expected length of the sub-array is then $\mathbb{E}[\sum_{j=1}^i a_j] = \sum_{j=1}^i \mathbb{E}[a_j]$ by the linearity of expectation. By the assumption that keys map randomly into hash values, $\mathbb{E}[a_j] \leq \frac{1}{M}$ so that the sum above is bounded by $\frac{j}{M} \leq \frac{K}{M}$. Since $\frac{K}{M} \leq c$ for some constant c by the dynamic resizing, the expected length is also bounded by a constant, meaning the complexity is as well.

Note that this analysis says nothing about the *worst-case* complexity of an operation (which can be linear if all keys map to the same hash value) or the expected length of the *longest* sub-array (which is $\frac{\log K}{\log \log K}$).

Universal Hashing

Certain competition forms include a stage where contestants may challenge the solutions of others for correctness, by providing a test case they believe the solution would fail. In this case, the hash function above is not good enough. Another contestant can easily generate values of x that all map to the same hash value, by generating a large number of values and evaluating your hash function

on them, picking a large set of collisions from them. This also applies when using `unordered_map` from STL⁴. To resolve this, one picks a hash function from a family of functions at random at every invocation of your program, a concept called **universal hashing**. In practice, the randomness tends come from reading the current time at a sufficiently granular level to be hard to predict.

The hash function we will look at is again Ax (as an unsigned 64-bit integer), but this time A is a random (odd) 64-bit integer. We claim that taking the top K bits of Ax makes a good hash function from 64-bit integers to K -bit integers, i.e.

$$h(x) = \left\lfloor \frac{Ax}{2^{64-K}} \right\rfloor$$

Theorem 6.1

For any two fixed 64-bit integers x, y , their hashes $h(x)$ and $h(y)$ are equal with probability $\frac{2}{2^K}$ over all choices of the hash function parameter A .

Proof. This proof uses some basic number theoretic facts – if you are not familiar with modular inverses, you might need to work through Chapter 19.

Assume that $h(x) = h(y)$. This means that the top K bits of Ax and Ay are equal. Thus, the top K bits of $Ax - Ay = A(x - y)$ must either be all zeroes (if $Ax \geq Ay$) or all ones (if $Ax < Ay$, causing a carry bit from the top K bits).

Now, we introduce the following variables. Let z be the odd part of $x - y$ such that $x - y = z2^i$ for some i . Also, let B be the top 63 bits of A so that $A = 2B + 1$. Since A is a uniformly random odd 64-bit integer, B is uniformly random. We can now perform the rewrite $A(x - y) = (2B + 1)(x - y) = (2B + 1)z2^i = Bz2^{i+1} + z2^i$. Since B is uniformly random over 2^{63} and z is odd, $Bz \bmod 2^{63}$ is uniformly random over 2^{63} (this follows from the fact that z as an odd number is relatively prime with 2^{63}). Thus, the integer $A(x - y) = Bz2^{i+1} + z2^i$ is uniformly random in its top $63 - i$ bits and contains only zeroes in its lower i bits.

Note that $Ax = Ay + A(x - y)$. Since $A(x - y)$ has zeroes in the lower i bits and a 1 in the i 'th bit, the i 'th bit of Ay will change when adding $A(x - y)$ to it, so that it will differ from the i 'th bit in Ax . By assumption the top K bits of Ax and Ay , are equal, which thus forces $i \leq 64 - K$.

⁴Using `map` is fine however, since it is not backed by a hash table.

Since the top $63 - i \geq 63 - (64 - K) = K$ are thus uniformly random, we get that they are all ones or all zeroes with probability $\frac{2}{2^K}$. □

Exercise 6.14. Is it a problem that the hash has pairwise collisions with probability $\frac{2}{2^K}$ rather than $\frac{1}{2^K}$ with regard to hash table complexity?

Chapter Exercises

Exercise 6.15. Assume that you want to implement shrinking of a dynamic array (or a hash table) where many elements were deleted so that the capacity is unnecessarily large. This will be implemented by calling a particular function after any removal, to see if the array should be shrunk. What is the problem with the following implementation?

```

1: procedure SHRINKVECTOR(V)
2:   while  $2 \cdot V.capacity > V.size$  do
3:     arr  $\leftarrow$  new T[ $V.capacity/2$ ]
4:     copy the elements of V.backing to arr
5:     V.backing  $\leftarrow$  arr
6:     V.capacity  $\leftarrow V.capacity/2$ 

```

Chapter Notes

For a more rigorous treatment of the basic data structures, we again refer to *Introduction to Algorithms* [7]. In particular it goes through other techniques regarding hash tables more thoroughly, something we skipped since it is the hashing technique and general knowledge of the structure we deemed important here – an efficient implementation is something your language standard library will provide.

If you want to dive deeper into proper implementations of the algorithms in C++, *Data Structures and Algorithm Analysis in C++*[29] covers what we brought up in this chapter and a bit more.

7 Recursion

This chapter introduces the first proper algorithmic technique of the book, that of *recursion*. The first four chapters of the next part – brute force, greedy algorithms, dynamic programming and divide and conquer – are all based on this concept. Recursion is perhaps the first truly creatively tricky (rather than technically difficult) technique faced by the fresh programmer, so we have chosen to dedicate an entire chapter for a primer on the topic.

The remainder of this book, and computer science as a whole, strongly depends on a solid understanding of recursion. You are therefore urged to read it more carefully than the previous chapters. Even better; once you have read it, read it again.

7.1 Recursive Definitions

The first example of recursion most people become acquainted with is the definition of a particular mathematical sequence, the *Fibonacci numbers*. The infinite sequence starts with the numbers 0, 1, 1, 2, 3, 5, 8, 13, Except for the first two, each number is computed by taking the sum of the two previous ones. A formal mathematical definition of the i 'th Fibonacci number F_i can look like this:

$$F_i = \begin{cases} 0 & \text{if } i = 0 \\ 1 & \text{if } i = 1 \\ F_{i-1} + F_{i-2} & \text{for } i \geq 2. \end{cases} \quad (7.1)$$

By the definition, we would have e.g. $F_6 = F_5 + F_4 = 5 + 3 = 8$, which holds.

Exercise 7.1. Use the definition to compute the 15 first Fibonacci numbers.

This is a so-called *recursive definition*, meaning that it refers back to itself – the definition of a Fibonacci number depends on the definition of (earlier) Fibonacci numbers. A program to directly implement the recursion looks very similar to the mathematical definition Eq. 7.1.

Computing Fibonacci Numbers

```

1  int F(int n) {
2      if (n == 0) return 0;
3      if (n == 1) return 1;
4      return F(n - 1) + F(n - 2);
5  }

```

Note that this function, just like the recursive definition, computes its result $F(n)$ by calling itself to compute the (smaller) Fibonacci numbers $F(n - 1)$ and $F(n - 2)$. A knee-jerk reaction might be that such a function could never finish. After all, in order to compute a single Fibonacci number, the function calls itself, not just one, but two times! The solution is one of the key ideas of recursion, namely that there is some base case where the self-referential – recursive – computation eventually bottoms out, so that the definition does not refer back to itself forever and ever. In the case of Fibonacci, once you try computing F_0 or F_1 , the definition gives us the values immediately without having to apply the recursive case. One can visualize the computation as in Figure 7.1.




Figure 7.1: A visualization of the computation of $F(5)$

Other possible applications of the recursive principle would be computing a^n where n is non-negative integer. Since a^n is defined as the product of a n times, we can base a recursion around first computing a^{n-1} and then multiplying it with a :

$$a^n = \begin{cases} 1 & \text{if } n = 0 \\ a \cdot a^{n-1} & \text{for } n \geq 1. \end{cases} \quad (7.2)$$

The implementation is similarly straightforward.

Recursive Exponentiation

```
1 int power(int a, int n) {  
2     if (n == 0) return 1;  
3     return a * power(a, n - 1);  
4 }
```

Even though recursive definitions are the simplest examples of recursion, they often come up in practice.

One can also solve problems traditionally programmed using loops by formulating them recursively¹. Consider the problem of summing an array of integers. Normally, you'd use loops for such a task. However, there is nothing preventing you from using the following recursive definition. Let $A = (a_0, \dots, a_{n-1})$ be an array of integers. If $S(k)$ is the sum of the first k elements of A , we have that

$$S(k) = \begin{cases} 0 & \text{if } k = 0 \\ a_{k-1} + S(k-1) & \text{otherwise.} \end{cases} \quad (7.3)$$

To compute the sum of the entire array, we would call $S(n)$. Even though we now have to deal with a vector, the implementation is similar:

```
1 // Invoked with sum(A, A.size())  
2 int sum(const vi& A, int k) {  
3     if (k == 0) return 0;  
4     return A[k - 1] + sum(A, k - 1);  
5 }
```

Equation 7.3 is a recursive definition too: it reduces the problem of summing the entire array A to summing a smaller part of A . **This is the essence of recursion** – it was the common factor in all three examples. A recursive definition is meant to express the solution to a problem in terms of other instances of the same problem. The goal is that the new instances should be smaller than the first, in order to make progress on the problem.

Exercise 7.2. Write recursive functions to compute:

- the product of all integers in an array,
- the largest element of an array, and
- the greatest sum of two consecutive elements of an array.

¹In fact, some programming languages do not have loops. Instead, you must formulate them recursively.

7.2 The Time Complexity of Recursive Functions

How fast is a recursive program? It is not as easy to compute as most of the algorithms we have seen so far. The work is distributed over some number of recursive calls, and we need to sum all of it up.

In all problems of this chapter, the time complexity of each recursive call is the same. This is true of all of our examples so far – they only performed constant-time work plus some recursive calls. Computing the complexity then boils down to two factors: the number of function calls in total, and the time complexity of a single function call (excluding the recursive calls). Summing all the work is as simple as taking the product of these two things.

In the example of summing a vector of size m , there are a total of $m+1 = \Theta(m)$ function calls. One call is made for each element, and one final call for the base case of the recursion. A single call performs only constant-time operations, so it has complexity $\Theta(1)$. This means the time complexity is $\Theta(m) \cdot \Theta(1) = \Theta(m)$.

The number of function calls made in total can be considerably harder to compute, such as for the Fibonacci recursion.

Exercise 7.3. Write a program that uses the recursive function to compute Fibonacci numbers. Try computing all the Fibonacci numbers starting from F_{30} and upwards until the execution takes over 30 seconds. Take note on how long your program takes. What complexity does the function seem to have?

From the above exercise, one thing should be clear – the running time is not a linear function. In fact, it turns out to be exponential!

A simple lower bound is $2^{\frac{n}{2}}$ function calls, which we can prove by induction. Let $T(n)$ be the time taken to compute F_n . If $T(n) \geq 2^{\frac{n}{2}}$ for all n up to some $n' - 1$ and $n = 1$, then

$$\begin{aligned} T(n') &\geq T(n' - 2) + T(n' - 1) \\ &= 2^{\frac{n'-2}{2}} + 2^{\frac{n'-1}{2}} \\ &\geq 2^{\frac{n'-2}{2}} + 2^{\frac{n'-2}{2}} \\ &= 2^{1+\frac{n'-2}{2}} \\ &= 2^{\frac{n'}{2}} \end{aligned}$$

using the fact that $T(n) = T(n - 2) + T(n - 1) + \Theta(1)$, so the statement holds for $n = n'$ too. By induction, it holds for all $n \geq 1$.

This lower bound is quite lax though – we can do better. Assume that $T(n) \geq x^n$ for some real value x . As $T(n) \geq T(n-2) + T(n-1) = x^{n-1} + x^{n-2}$ (within a constant term), we get the inequality

$$x^n \geq x^{n-1} + x^{n-2}$$

which, after dividing by x^{n-2} results in $x^2 - x - 1 > 0$. Solving for $x = 0$, we find that the inequality holds when x is greater than the so-called the *golden ratio*, $\phi = \frac{1+\sqrt{5}}{2} \approx 1.618$.

Exercise 7.4. Use the same inductive technique as before to prove that $T(n) = \Omega(1.61^n)$ and $T(n) = O(1.62^n)$.

7.3 Choice

While all recursion is based on reducing a problem instance to a smaller instance of the same problem, there are many different conceptual ways to do this. This time, we look at problems involving choices of different kinds.

Stairs

Tasha the Kitty loves playing with the stairs at home while her caretakers are at work. Her favorite game involves jumping up to the top of the stairs by repeatedly skipping either 1 or 2 stairs at a time. She doesn't like jumping on the exact same sequence of stairs during two different climbs.




Figure 7.2: One way Tasha could climb a staircase of 6 stairs.

If the staircase has $1n \leq 20$ steps (including the top), in how many different ways can she climb the stairs?

Solution. With such a small n , computing this efficiently is not the main issue; computing it at all is. The trick lies in formulating Tasha's jumping up the stairs

as a sequence of choices. After Tasha has jumped k steps, she has two choices – should her next jump be up a single stair to $k + 1$, or two stairs to $k + 2$? When dealt such a problem, always ask yourself: what was Tasha’s last choice, just before she climbed up to the top of the stairs? Consider these two options in Figure 7.3.




Figure 7.3: The two stairs leading to the top.

If there are a total of n stairs and Tasha’s last jump was a single step, then she came from step $n - 1$. Similarly, if she took two steps, she came from step $n - 2$. These two options are exhaustive – there is no other way she could have come to step n . They are also exclusive – we assumed that this was Tasha’s last jump, so there is no overlap between these possibilities. This means that the number of ways Tasha can get to the n ’th step must be equal to the number of ways she could get to the $(n - 1)$ ’st step, plus the number of ways she could get to the $(n - 2)$ ’nd step.

A recursive procedure based on this insight is then straightforward:

```

1: procedure STAIRS( $n$ )
2:   if  $n = 0$  then
3:     return 1
4:   if  $n = 1$  then
5:     return 1
6:   return STAIRS( $n - 1$ ) + STAIRS( $n - 2$ )

```

Note the bases cases we added, for the case of an empty staircase or a single stair. The time complexity of the solution is the same as that for Fibonacci, since the recursion is the same. \square

We solve another problem from an early Swedish high school qualifier in the same way.

The Plank

Swedish Olympiad in Informatics, School Qualifiers 2001

You want to construct a long plank using smaller wooden pieces. There are three kinds of pieces of lengths 1, 2 and 3 meters respectively, which you have an unlimited number each of. You can glue together several of the smaller pieces to create a longer plank.




Figure 7.4: There are 7 ways to glue a 4 meter plank.

If the plank should have length n ($1 \leq n \leq 24$) meters, in how many different ways can you glue pieces together to get a plank of the right length?

Solution. The idea here is the same as in the Stairs problem. To compute the size of the set of all possible planks, we find a recursive definition that reduces the problem into counting the number of ways one can build some smaller planks. For any given plank of length n , the rightmost piece of the plank has size either 1, 2 or 3. This means that the number of ways in which we can construct the plank is equal to the number of ways in which planks of sizes $n - 1$, $n - 2$ and $n - 3$ can be computed. While this isn't easier to compute directly, we can apply the same reduction recursively to these smaller planks, ending up with a very similar solution:

```

1: procedure PLANKWAYS( $n$ )
2:   if  $n < 0$  then
3:     return 0
4:   if  $n = 0$  then
5:     return 1

```

6: **return** PlankWays($n - 1$) + PlankWays($n - 2$) + PlankWays($n - 3$)

Again, we had to add a few base cases to give the recursion somewhere to stop. The two base cases we picked here may be slightly less intuitive. We say that there is a single way to construct a plank of length 0, and no ways to construct negative-length planks. \square

Exercise 7.5. The PlankWays algorithm has a time complexity lower bound of $\Omega(1.83^n)$ and $O(1.84^n)$. Prove this.

Problem 7.1

The Plank – plankan

Note: solve the first subtask for 50 points.

While these two problems in particular are much alike, many other recursive problems also follow this template:

- formulate the problem as a sequence of choices,
- look at what the last choice was, and
- find out if “backtracking” along that choice reduces the problem to smaller instances of the same problem.

Now that we have warmed up, we are going to look at a slightly harder recursive problem, where it is less obvious to figure out how to reduce the problem to a smaller one.

Dominoes

In how many ways can a $2 \times n$ ($1 \leq n \leq 20$) grid be tiled by n dominoes, i.e. bricks of size 1×2 or 2×1 such that no dominoes overlap?




Figure 7.5: An example tiling of a 2×7 grid.

Solution. Looking at the example tiling in figure 7.5 might help us. Let us denote the number of tilings of a $2 \times n$ grid with $S(n)$. In general, a recursion would somehow reduce the problem of computing $S(n)$ to computing smaller values of this function. By considering the rightmost domino of the example, a partial solution idea should form. If the rightmost tile is placed vertically, the remaining grid has size $2 \times (n - 1)$, so there are $S(n - 1)$ such tilings. If it is not placed vertically, the two rightmost squares must instead be occupied by two horizontal tiles. In this case, the remaining grid would have size $2 \times (n - 2)$, meaning there would be $S(n - 2)$ ways to complete the remainder of the tiling (see Figure 7.6).




Figure 7.6: The two resulting subproblems after covering the rightmost column.

Since these are the only two options, the total number of tilings must be $S(n) = S(n - 1) + S(n - 2)$, and thus we get our recursive solution. Here too we got the same recursion as the one for Fibonacci, with the same time complexity. \square

Exercise 7.6. 1) Write a recursion to compute the number of strings of length n consisting of only letters A and B, with no two A's next to each other.

2) Write a recursion to compute the number of subsets of $\{1, 2, \dots, n\}$ of length n , where at least one of $i, i + 1$ must be in the subset for all $1 \leq i < n$.

In the next chapter on brute force, we will revisit this way of thinking in a new light as we use recursion to solve optimization problems rather than simply counting ways.

7.4 Multidimensional Recursion

So far, every recursive solution we produced were about a single sequence of numbers – the input was an integer n , and we computed the n 'th value of the sequence by a recursive definition.

This is far from the only area where recursion can be applied. Problem instances comes in all forms and shapes. Recursing on more advanced problems can sometimes give us several recursive sequences that not only refer to themselves, but smaller values of the other sequences.

Varied Amusements

Marika and Lisa loves going to amusement parks. This time, they have their eyes set on a park with lots of exciting rides of three different types: tilt-a-whirls, roller coasters and drop towers. There are $1 \leq a \leq 10$ different tilt-a-whirls, $1 \leq b \leq 10$ roller coasters and $1 \leq c \leq 10$ drop towers. They want to ride $1 \leq n \leq 10$ different rides in sequence, but never two rides of the same type in a row. In how many ways can they choose such sequences of n ?

Solution. On the surface, the problem is a prime candidate for the choice-strategy. There are n choices – what ride to go on each time. However, once we choose the last ride the girls took, we are faced with a problem. If we chose a roller coaster, the first $n - 1$ rides may not end with a roller coaster. This is not a smaller instance of the same problem, where the last ride could be anyone we wanted. Instead, depending on the type of ride we choose as the last one, we get three different problems: How many sequences of $n - 1$ rides are there that does not end with A) a tilt-a-whirl? B) a roller coaster? or C) a drop tower?

What happens if we apply the same strategy to these three *new problems*? Well, in the problem where we have to choose an $n - 1$ ride sequence that does not end with a tilt-a-whirl, there are two options for the last ride. If it was a roller coaster, we have to choose the remaining $n - 2$ rides such that they do not end with a roller coaster. If it was a drop tower, the remaining $n - 2$ rides may not end with a drop tower. Either way, both cases reduce to a smaller problem of the other two types!

By introducing the three new problems $A(n)$, $B(n)$ and $C(n)$, defined as the number of ride sequences of length n not ending in a tilt-a-whirl, roller coaster or drop tower respectively, we can produce recursive definitions that refer only

to these recursions:

$$A(n) = b \cdot B(n-1) + c \cdot C(n-1)$$

$$B(n) = a \cdot A(n-1) + c \cdot C(n-1)$$

$$C(n) = a \cdot A(n-1) + b \cdot B(n-1)$$

with the only required base case of $A(0) = B(0) = C(0) = 1$. The answer then becomes $a \cdot A(n-1) + b \cdot B(n-1) + c \cdot C(n-1)$.

When implementing the solution in C++, don't forget results from Exercises 2.20-2.21 to resolve the circular dependencies between functions calling each other. \square

Exercise 7.7. Determine, with proof, the time complexity of the solution to Varied Amusements.

Problem 7.2

Varied Amusements – variedamusements

Note: solve the first subtask for 1 point.

7.5 Recursion vs. Iteration

After looking through the problems we solved in this chapter, you might wonder if recursion really is needed. When computing the Fibonacci numbers by hand, we tend to not use a method nearly as complicated as the recursive function. Instead, we write them down one by one after each other, taking the sum of the last two ones to compute the next. This approach could be implemented iteratively in code, looking something like this:

```
1  int F(int n) {
2      int secondLast = 0; // stores F_(i-2)
3      int last = 1; // stores F_(i-1)
4      for (int i = 2; i <= n; i++) {
5          int current = last + secondLast; // Compute F_i = F_(i-2) + F_(i-1)
6          // Since i will be increased by 1, the old F_(i-1) becomes
7          // F_(i-2), and the old F_i becomes F_(i-1)
8          secondLast = last;
9          last = current;
10     }
11     return last;
12 }
```

Algorithmically, recursion – in the sense of functions calling themselves – is not needed. As a programming construct it doesn't bring any additional

computational powers. Recursive functions can even be simulated with a single loop and a stack, by storing the current chain of recursive calls in the stack and processing them one at a time in the loop. This is actually what your computer does behind the scenes².

The reason behind the strong focus in recursion is another, namely that it is an incredibly powerful mode of thinking. To us, it is unclear how one would find a natural solution to the Dominoes problem without going in with a recursive mindset and looking for that reduction to a smaller instance. That being said, once a recursive formulation has been deduced, an iterative implementation can many times be simpler or faster. For example, try to compute F_{46} using both the recursive and iterative approach. You'll see that one of the two versions finish, and one do not.

Problem 7.3

Varied Amusements – variedamusements

Note: solve the first two subtasks for 2 points.

The Plank – plankan

Note: solve both subtasks for 100 points.

Chapter Exercises

Exercise 7.8. There are n lines drawn in the plane, no three lines intersecting in the same point. What is the number of connected regions they split the plane into?

Problem 7.4

$3 \times n$ Dominoes – 3xndominoes

Note: solve the first subtask for 2 points.

3 -close Sets – 3close

Note: solve the first two subtasks for 2 point.

Even A's, Odd B's – evenaoddb

Note: solve the first two subtasks for 2 point.

²At a low level, modern processors are essentially a single execution loop with a stack for function-related memory.

Chapter Notes

Recursion as a problem solving technique is a common one both in mathematics and algorithmics. In mathematics, there tends to be a larger focus in finding closed forms for the recursions though, while we are happy with any kind of efficient computation. There is a rich combinatorial theory behind finding such closed forms. As previously mentioned, *Concrete Mathematics*[?] is one of the mathematics books that really excel on teaching these techniques.

8 Graph Theory

We finish this foundational part with an introduction to **graph theory**, the study of mathematical objects known as graphs. As a mathematical area of study, it dates back to the early 1700s, when Euler first studied the famous *Seven Bridges of Königsberg* problem. It is one of the most well-studied areas in algorithmic problem solving, being one of only two topics (together with data structures) to make an appearance in all three parts of this book. In almost every programming contests you can find a problem relating to graphs.

8.1 Graphs

A graph is an abstract way of representing various types of relations, such as roads between cities, friendships between people, networks links between computers and so on. Graphs are essentially a set of objects where certain pairs of objects are connected. Formally, graphs are defined in the following way.

Definition 8.1 A **simple graph** $G = (V, E)$ consists of a pair V of **vertices**, and a set E of **edges**. An edge consists of a pair of vertices $\{u, v\}$ called the **endpoints** of the edge.

A graph lends itself naturally to a graphical representation, where vertices are represented by points in the plane with lines drawn between the two vertices of an edge. For example, the graph given by $V = \{1, 2, 3, 4, 5\}$ and $E = \{\{1, 2\}, \{2, 3\}, \{3, 4\}, \{2, 4\}\}$ can be drawn as in Figure 8.1.




Figure 8.1: An example graph with 6 vertices and 4 edges.

Exercise 8.1. Draw the graphical representation of the graph with vertices

$\{a, b, c, d\}$ and edges $\{\{a, b\}, \{b, c\}, \{c, d\}, \{a, d\}, \{b, d\}\}$.

Exercise 8.2. The graph on n vertices containing *all* possible edges is called the *complete graph*, or K_n . How many edges does K_n have?

Trip Planning

Lars is planning to do a backpacking tour by train throughout N cities in Europe. He has a list of the M train lines that go back and forth between pairs of these cities. He wants to visit the cities in the order $1, 2, \dots, N$, finally returning back to his home in city 1.

Since Lars have limited vacation days, he only has time to take exactly N direct trains during his trip. Can you determine if this is possible, and tell Lars which trains to take?

Input

The first line contains the number of cities $N \leq 10^6$ that Lars wants to visit, and $M \leq 10^6$, the number of direct trains.

The next M lines each contain two integers $1 \leq a \neq b \leq N$, indicating that there is a train line traveling between cities a and b . No two train lines will have the same two integers.

Output

If Lars cannot perform his trip taking only N trains, output `no trip`. Otherwise, output the numbers of the N train lines that lines Lars should take (in order of travel), where the train lines are numbered from 1 to M in the order they appear in the input.

Solution. The problem essentially asks if there are direct train lines between cities $(1, 2), (2, 3), \dots, (N - 1, N), (N, 1)$. If there are, we want to find the indices of these lines in the input. This is a typical problem that can be modelled as a graph. In those terms, we have a graph on N vertices with its M edges given in a list. We are asked if the graph contains a certain list of edges.

A possible solution would be to keep a vector of the indices of these particular edges while we read the list of edges. If we find the edge $\{k, k + 1\}$ for a given k , we can store the index of the edge in the k 'th position in the vector. Only if we managed to find every edge should we reply with their indexes. Otherwise, we would output `no trip`. \square

Problem 8.1*Trip Planning – tripplanning*

Of particular importance is the set of vertices that are connected to a given vertex by edges.

Definition 8.2 For a graph $G = (V, E)$, if $u, v \in E$, u and v are called **neighbours**.

The set of neighbours of a vertex v is denoted $N(v)$, called the **neighbourhood** of v .

The size of the neighbourhood of a vertex v is called the **degree** of v . It is denoted $\deg(v)$.

The degrees of the vertices in a graph fulfill many useful properties.

Theorem 8.1

The sum of degrees of a graph $G = (V, E)$ is even. Specifically,

$$\sum_{v \in V} \deg(v) = 2|E|.$$

Proof. We prove the statement by induction. If the graph contains no edges, the degree of each vertex is clearly 0, and $|E| = 0$, so the equality holds.

Now, assume that it holds for all graphs of k edges. If we add a single edge $\{u, v\}$, to this graph, we add one more vertex to the neighbourhoods of u and v . This increases $\deg(u)$ and $\deg(v)$ by 1, increasing the left hand side by 2. It also increases $|E|$ by 1, so the right hand side is increased by 2 as well, thus the identity holds for all graphs on $k + 1$ edges too.

By the principle of induction, the statement must hold for all graphs. \square

Example 8.1 In a graph G , the degrees of the vertices are 3, 5, 4, 4, 4, 6, 6 respectively. Prove that there is a sequence of neighbouring vertices starting and ending with the vertices of degree 3 and 5.

Solution. Let S be the set of vertices x for which there is a sequence of neighbouring vertices starting at the vertex of degree 3 and ending at x . Let H be a new graph with S as vertices, containing as edges only those that

have vertices of S as endpoints. Note that S must contain every neighbour of every vertex x in S – if there is a sequence of neighbouring vertices from the degree 3 vertex to x , such a sequence to the neighbours of x can be constructed by appending the neighbour it to the sequence of vertices ending at x . This means that the degree of all vertices in H have the same degree as they had in G .

Now, the sum of the degrees of all vertices in S is even by Theorem 8.1. Since S contains a vertex of degree 3, it must contain another vertex of odd degree (or the sum would be odd). The only such vertex is the one of degree 5. By the definition of S , this means the requested sequence of neighbouring vertices between the two vertices exist. \square

Exercise 8.3. Prove that in a simple graph of at least 2 vertices, there must exist 2 vertices of the same degree.

Problem 8.2

Given a graph, print all the vertices with the highest degrees.

While the simple graph is able to represent many kinds of relations, we sometimes need variations to capture all the information we are interested in. Consider a graph representing roads between cities, where the vertices represent the cities and the edges correspond to the roads between them. In this case, we might be interested in also capturing the lengths of the roads between all the cities as well: a situation depicted in Figure 8.2.




Figure 8.2: A road network on 5 cities.

We can modify our definition of an edge to include such a number, giving us a new type of graph.

Definition 8.3 A **weighted graph** is a graph (V, E) together with a **weight function** w associating each edge $e \in E$ with a real-valued **weight** $w(e)$.

Weighted graphs often appear when there is a natural measure of an edge, such as the distance between two cities, the latency between two computers, or the cost of a train ticket between two airports.

Problem 8.3

Compute the shortest triangle in a weighted graph, where $E \leq 100$.

Finally, not all relations we model are symmetric in the way indicated by graphs. In many situations, we would instead prefer if an edge could have a certain direction, going *from* a vertex *to* another vertex. For example, when modelling all the car roads in a city, certain roads may be one-way, a nuance the simple graph would miss. We fix this by making edges ordered pairs rather than sets:

Definition 8.4 A **directed graph** is a graph (V, E) where E consists of **directed edges**, i.e. ordered pairs $e = (u, v)$ of vertices. The edge e is called an **out-edge** of u and **in-edge** of v .

When representing directed graphs graphically, edges will be arrows, with the arrowhead pointing from u to v (Figure 8.3).




Figure 8.3: The graph given by $V = \{1, 2, 3, 4\}$ and $E = \{(1, 2), (3, 1), (4, 2), (4, 1)\}$.

Problem 8.4

Determine if there exists a directed triangle in a graph, where $E \leq 100$.

8.2 Representing Graphs

When dealing with graphs in algorithms, there are three common data structures used to represent them: *adjacency matrices*, *adjacency lists* and *adjacency maps*. Occasionally we also represent graphs *implicitly* – for example by a function that takes a vertex v and provides us with all edges it's an endpoint of.

This latter representation is common when dealing with searches in the graph corresponding to the positions in a game such as chess.

In the following sections, we present the representation of the directed, unweighted graph in Figure 8.4.




Figure 8.4: An example graph with 6 vertices and 5 edges.

Adjacency Matrices

An adjacency matrix represents a graph $G = (V, E)$ with $|V|$ as a 2D $|V| \times |V|$ matrix in the following way:

Definition 8.5 If $G = (V, E, w)$ is a directed, weighted graph and $V = \{v_1, \dots, v_n\}$, the graph's **adjacency matrix** the $|V| \times |V|$ -matrix $A = (a_{i,j})$ where $a_{i,j} = w(v_i v_j)$.

For undirected graphs, we set $a_{i,j} = a_{j,i} = w(\{v_i, v_j\})$.

For unweighted graphs, we set $a_{i,j} = 1$.

This representation uses $\Theta(|V|^2)$ memory, and takes $O(1)$ time for adding, modifying and removing edges. To iterate through the neighbours of a vertex, you need $\Theta(|V|)$ time, independent of the number of neighbours of the vertex itself.

Adjacency matrices are best to use when $|V|^2 \approx |E|$, i.e. when the graph is *dense*.

The adjacency matrix for the directed, unweighted graph in Figure 8.4 is:

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

Adjacency Lists

Another way to represent graphs is by storing lists of neighbours for every vertex. This approach is called **adjacency lists**. This only requires $\Theta(|E| + |V|)$ memory, which is better when your graph have few edges, i.e. it is *sparse*. If you use a vector to represent each list of neighbours, you also get $\Theta(1)$ addition and removal (if you know the index of the edge and ignore their order) of edges, but it takes $O(|V|)$ time to determine if an edge exists. On the upside, iterating through the neighbours of a vertex takes time proportional to the number of neighbours instead of the number of vertices in the graph. This means that iterating through all the neighbours of all vertices takes time $\Theta(|E| + |V|)$ instead of $\Theta(|V|^2)$ as for the adjacency matrix. For large, sparse graphs this is clearly better!

When representing weighted graphs, the list usually stores the edges as pairs of (*neighbour, weight*) instead. For undirected graphs, both endpoints of an edge contains the other in their adjacency lists.

This representation is common in many graph search algorithms to be studied in Chapter 14.

Adjacency Maps

An *adjacency map* combines the adjacency matrix with the adjacency list to get the benefits of both the matrix ($\Theta(1)$ time to check if an edge exists) and the lists (low memory usage and fast neighbourhood iteration). Instead of using lists of neighbours for each vertex, we can use a hash table for each vertex.

This has the same time and memory complexities as the adjacency lists, but it also allows for checking if an edge is present in $\Theta(1)$ time. The downsides are that hash tables have a higher constant factor than the adjacency list, and that you lose the ordering you have of your neighbours (if this is important). The adjacency map also inherits another sometimes important property from the matrix: you can *remove* arbitrary edges in $\Theta(1)$ time!

This representation is mostly used when one is dynamically modifying a graph.

Exercise 8.4. Given a graph, which representation or representations are suitable if

- a) $|V| = 1000$ and $|E| = 499500$
- b) $|V| = 10000$ and $|E| = 20000$

c) $|V| = 1000$ and $|E| = 10000$

Exercise 8.5. Implement programs that take unweighted, directed graphs as input and:

- output the degree of each vertex,
- can quickly determine whether a certain edge exists or not, and
- finds a triangle – i.e. a cycle of length 3 – in a graph.

8.3 Breadth-First Search

A *path* is a sequence of *distinct* vertices $p_0, p_1, \dots, p_{l-1}, p_l$ such that $\{p_i, p_{i+1}\} \in E$ ($i = 0, 1, \dots, l-1$). This means any two vertices on a path must be connected by an edge. We say that this path has **length** l , since it consists of l edges. In Figure ??, the sequence 3, 1, 4, 2 is a path of length 3.

One of the most common basic graph algorithms is the **breadth-first search**. It is used to find the distances from a certain vertex in an unweighted graph.

Single-Source Shortest Path, Unweighted Edges

Given an unweighted graph $G = (V, E)$ and a source vertex s , compute the shortest distances $d(s, v)$ for all $v \in V$.

For simplicity, we first consider the problem on a grid graph, where the unit squares constitute vertices, and vertices which share an edge are connected. Additionally, some squares are blocked (and don't have a corresponding vertex). An example can be seen in Figure 14.1.




Figure 8.5: An example grid graph, with source marked s .

Let us solve this problem inductively. First of all, what vertices have distance 0? Clearly, this is only the source vertex s itself. This seems like a reasonable

base case, since the problem is about shortest paths from s . Then, what vertices have distance 1? These are exactly those with a path consisting of a single edge from s , meaning they are the neighbors of s (marked in Figure 14.2).




Figure 8.6: The square with distance 1 from the source.

If a vertex v has distance 2, it must be a neighbor of a vertex u with distance 1 (except for the starting vertex). This is also a sufficient condition, since we can construct a path of length 2 simply by extending the path of any neighbor of distance 1 with the edge (u, v) .





Figure 8.7: The squares with distance 2, 3 and 4.

In fact, this reasoning generalizes to any particular distance, i.e., that all the vertices that have exactly distance k are those that have a neighbor of distance $k - 1$ but no neighbor to a vertex with a smaller distance. Using this, we can construct an algorithm to solve the problem. Initially, we set the distance of s to 0. Then, for every $dist = 1, 2, \dots$, we mark all vertices that have a neighbor with distance $dist - 1$ as having $dist$. This algorithm is called the **breadth-first search**.

Exercise 8.6. Use the BFS algorithm to compute the distance to every square in the following grid:



A simple implementation of this would be to iteratively construct the lists of the vertices which have distance 0, 1, \dots , and so on.

```

1: procedure BREADTHFIRSTSEARCH(vertices  $V$ , vertex  $s$ )
2:    $distances \leftarrow \text{new int}[|V|]$ 
3:   fill  $distances$  with  $\infty$ 
4:    $curDist \leftarrow 0$ 
5:    $curVertices \leftarrow \text{new vector}$ 
6:    $curVertices.add(s)$ 
7:    $distances[s] \leftarrow curDist$ 
8:   while  $curDistVertices \neq \emptyset$  do
9:      $nextVertices \leftarrow \text{new vector}$ 
10:    for  $from \in curVertices$  do
11:      for  $v \in from.neighbours$  do
12:        if  $distances[v] = \infty$  then
13:           $nextVertices.add(v)$ 
14:           $distances[v] \leftarrow curDist + 1$ 
15:     $curDist \leftarrow curDist + 1$ 
16:     $curVertices = nextVertices$ 
17:  return  $distances$ 
    
```

Each vertex is added to $nextVertices$ at most once, since it is only pushed if $distances[v] = \infty$ whereupon it is immediately set to something else. We then iterate through every neighbor of all these vertices. In total, the number of all neighbours is $2E$, so the algorithm in total uses $\Theta(V + E)$ time.

Usually, the outer loop are often coded in another way. Instead of maintaining two separate vectors, we can merge them into a single queue:

```

1: while  $curVertices \neq \emptyset$  do
    
```



```

2:  from ← curVertices.front()
3:  curVertices.pop()
4:  for from ∈ curVertices do
5:      for v ∈ from.neighbours do
6:          if distances[v] = ∞ then
7:              curVertices.add(v)
8:              distances[v] = distances[from] + 1

```

The order of iteration is equivalent to the original order.

Exercise 8.7. Prove that the shorter way of coding the BFS loop (Algorithm ??) is equivalent to the longer version (Algorithm ??).

Exercise 8.8. Implement BFS problem

In many problems the task is to find a shortest path between some pair of vertices where the graph is given *implicitly*.

8-puzzle

In the *8-puzzle*, 8 tiles are arranged in a 3×3 grid, with one square left empty. A move in the puzzle consists of sliding a tile into the empty square. The goal of the puzzle is to perform some moves to reach the target configuration. The target configuration has the empty square in the bottom right corner, with the numbers in order 1, 2, 3, 4, 5, 6, 7, 8 on the three lines.

← 8	6	
7	1	4
2	5	3

8		6
7	1	4
2	5	3

1	2	3
4	5	6
7	8	

Figure 8.8: An example 8-puzzle, with a valid move. The rightmost puzzle shows the target configuration.

Given a puzzle, determine how many moves are required to solve it, or if it cannot be solved.

This is a typical BFS problem, characterized by a starting state (the initial puzzle), some transitions (the moves we can make), and the task of finding a short sequence of transitions to some goal state. We can model this kind of problem using a graph. The vertices represent the possible arrangements of

the tiles in the grid, and an edge connects two states if they differ by a single move. A sequence of moves from the starting state to the target configuration then represents a path in this graph. The minimum number of moves required is the same as the distance between those vertices in the graph, meaning we can use a BFS.

In such a problem, most of the code usually deals with the representation of a state as a vertex, and generating the edges that a certain vertex is adjacent to. When an implicit graph is given, we generally do not compute the entire graph explicitly. Instead, we use the states from the problems as-is, and generate the edges of a vertex only when it is being visited in the breadth-first search. In the 8-puzzle, we can represent each state as a 3×3 2D-vector. The difficult part is generating all the states that we can reach from a certain state.

Generating 8-puzzle Moves

```
1  typedef vector<vi> Puzzle;
2
3  vector<Puzzle> edges(const Puzzle& v) {
4      int emptyRow, emptyCol;
5      rep(row,0,3)
6          rep(col,0,3)
7              if (v[row][col] == 0) {
8                  emptyRow = row;
9                  emptyCol = col;
10             }
11     vector<Puzzle> possibleMoves;
12     auto makeMove = [&](int rowMove, int colMove) {
13         int newRow = emptyRow + rowMove;
14         int newCol = emptyCol + colMove;
15         if (newRow >= 0 && newCol >= 0 && newRow < 3 && newCol < 3) {
16             Puzzle newPuzzle = v;
17             swap(newPuzzle[emptyRow][emptyCol], newPuzzle[newRow][newCol]);
18             possibleMoves.push_back(newPuzzle);
19         }
20     };
21     makeMove(-1, 0);
22     makeMove(1, 0);
23     makeMove(0, -1);
24     makeMove(0, 1);
25     return possibleMoves;
26 }
```

With the edge generation in hand, the rest of the solution is a normal BFS, slightly modified to account for the fact that our vertices are no longer numbered $0, \dots, V - 1$. We can solve this by using e.g. maps instead.

8-puzzle BFS

```

1  int puzzle(const Puzzle& S, const Puzzle& target) {
2      map<Puzzle, int> distances;
3      distances[S] = 0;
4      queue<Puzzle> q;
5      q.push(S);
6      while (!q.empty()) {
7          const Puzzle& cur = q.front(); q.pop();
8          int dist = distances[cur];
9          if (cur == target) return dist;
10         for (const Puzzle& move : edges(cur)) {
11             if (distances.find(move) != distances.end()) continue;
12             distances[move] = dist + 1;
13             q.push(move);
14         }
15     }
16     return -1;
17 }

```

Besides this kind of search problems that can be solved using a BFS, some problems require modifications of a BFS, or use the distances generated only as an intermediary result.

Shortest Cycle

Compute the length of the shortest simple cycle in a graph.

Problem 8.5

Button Bashing – buttonbashing

8.4 Depth-First Search

If a path additionally satisfy that $\{p_0, p_l\} \in E$, we may append this edge to make the path cyclical. This is called a *cycle*. Similarly, a walk with starts and ends at the same vertex is called a **closed walk**. If a trail starts and ends at the same vertex, we call it a **closed trail**.

A graph where any pair of vertices have a path between then is called a **connected** graph. The (maximal) subsets of a graph which *are* connected form the **connected components** of the graph. In Figure ??, the graph consists of two components, $\{1, 2, 3, 4\}$ and $\{5\}$.

The depth-first search is an analogue to the breadth-first search that visits vertices in another order. Similarly to how the BFS grows the set of visited vertices using a wide frontier around the source vertex, the depth-first search

proceeds its search by, at every step, trying to plunge deeper into the graph. This order is called the depth-first order. More precisely, the search starts at some source vertex s . Then, any neighbor of s is chosen to be the next vertex v . Before visiting any other neighbor of s , we first visit any of the neighbours of v , and so on.

Implementing the depth-first search is usually done with a recursive function, using a vector *seen* to keep track of visited vertices:

```
1: procedure DEPTH-FIRST SEARCH(vertex  $at$ , adjacency list  $G$ )
2:   if  $seen[at]$  then
3:     return
4:    $seen[at] = \text{true}$ 
5:   for  $neighbour \in G[at]$  do
6:      $dfs(neighbour, G)$ 
```

In languages with limited stack space, it is possible to implement the DFS iteratively using a stack instead, keeping the vertices which are currently open in it.

Due to the simplicity of coding the DFS compared to a BFS, it is usually the algorithm of choice in problems where we want to visit all the vertices.

Coast Length

KTH Challenge 2011 – Ulf Lundström

The residents of Sotesholm value their coast highly and therefore want to maximize its total length. For them to be able to make an informed decision on their position in the issue of global warming, you have to help them find out whether their coastal line will shrink or expand if the sea level rises. From height maps they have figured out what parts of their islands will be covered by water, under the different scenarios described in the latest IPCC report on climate change, but they need your help to calculate the length of the coastal lines.




Figure 8.9: Gray squares are land and white squares are water. The thick black line is the sea coast.

You will be given a map of Sotesholm as an $N \times M$ grid. Each square in the grid has a side length of 1 km and is either water or land. Your goal is to compute the total length of sea coast of all islands. Sea coast is all borders between land and sea, and sea is any water connected to an edge of the map only through water. Two squares are connected if they share an edge. You may assume that the map is surrounded by sea. Lakes and islands in lakes are not contributing to the sea coast.

Solution. We can consider the grid as a graph, where all the water squares are vertices, and two squares have an edge between them if they share an edge. If we surround the entire grid by an water tiles (a useful trick to avoid special cases in this kind of grid problems), the sea consists exactly of those vertices that are connected to these surrounding water tiles. This means we need to compute the vertices which lie in the same connected component as the sea – a typical DFS task¹. After computing this component, we can determine the coast length by looking at all the squares which belong to the sea. If such a square share an edge with a land tile, that edge contributes 1 km to the coast length.

```

1  const vpi moves = {pii(-1, 0), pii(1, 0), pii(0, -1), pii(0, 1)};
2
3  int coastLength(const vector<vector<bool>>& G) {
4      int H = sz(G) + 4;
5      W = sz(G[0]) + 4;
6      vector<vector<bool>> G2(H, vector<bool>(W, true));
7      rep(i,0,sz(G)) rep(j,0,sz(G[i])) G2[i+2][j+2] = G[i][j];
8      vector<vector<bool>> sea(H, vector<bool>(W));
9

```

¹This particular application of DFS, i.e. computing a connected area in a 2D grid, is called a *flood fill*.

```
10  function<void(int, int)> floodFill = [&](int row, int col) {
11      if (row < 0 || row >= H || col < 0 || col >= W) return;
12      if (sea[row][col]) return;
13      sea[row][col] = true;
14      trav(move, moves) floodFill(row + move.first, col + move.second);
15  };
16  dfs(0, 0);
17
18  int coast = 0;
19  rep(i, 1, sz(G)+1) rep(j, 1, sz(G[0])+1) {
20      if (sea[i][j]) continue;
21      trav(move, moves) if (!sea[i + move.first][j + move.second]) coast++;
22  }
23  return coast;
24 }
```

□

Problem 8.6*Mårten's DFS* – martensdfs

8.5 Trees

A **tree** is a special kind of graph – a **connected** graph which does not contain any cycle. The graph in Figure ?? is not a tree, since it contains the cycle 1, 2, 4, 1. The graph in Figure 8.10 on the other hand, contains no cycle.

Figure 8.10: The tree given by $V = \{1, 2, 3, 4\}$ and $E = \{\{1, 2\}, \{3, 1\}, \{4, 1\}\}$.

Exercise 8.9. Prove that a tree of n vertices have exactly $n - 1$ edges.

Chapter Exercises

Given a graph, determine how many edges must be added to make it regular.

Chapter Notes

Part II

Basics

9 Brute Force

Many problems are solved by testing a large number of possibilities. For example, chess engines work by testing countless variations of moves and choosing the ones resulting in the “best” positions. This approach is called **brute force**. Brute force algorithms exploit that computers are fast, resulting in you having to be less smart. Just as with chess engines, brute force solutions might still require some ingenuity. A brute force problem might have a simple algorithm which requires a computer to evaluate 2^{40} options, while some deeper analysis might be able to reduce this to 2^{20} . This would be a huge reduction in running time. Different approaches to brute force may be the key factor in reaching the latter case instead of the former. In this chapter, we look at four different techniques used to solve brute force problems, ranging from the simplicity of just evaluating every single option to an advanced memory-time tradeoff called meet-in-the-middle.

9.1 Optimization Problems

In an **optimization problem**, we have some **solution set** S and a **value function** f . The goal is to find an $x \in S$ which maximize $f(x)$, i.e., optimizing the function.

Optimization problems constitute a large class of the problems that we solve in algorithmic problem solving, such as the Max Clique problem and the Buying Books problems we examine in this chapter. One of the most famous optimization problems is the NP-complete *Travelling Salesman Problem*. The problem asks for the shortest cycle that visits all vertices of a weighted graph. The practical applications of this problem are many. A logistics company that must perform a number of deliveries probably want to minimize the distance traveled when visiting all the points of delivery. When planning your backpacking vacation, you may prefer to minimize the cost of travelling between all your destinations. In this problem, the solution set S would consist of all cycles in the graph that visit all the vertices, with $f(x)$ being the sum of all edges in the cycle x .

The brute force technique essentially consists of evaluating $f(x)$ for a large number (sometimes even all) of $x \in S$. For large S , this is slow.

The focus of this chapter and the chapters on Greedy Algorithms (Chapter 10) and Dynamic Programming (Chapter 11) is to develop techniques that exploit particular structures of optimization problems to avoid evaluating the entire set S .

9.2 Generate and Test

Our first brute force method is the *generate and test* method. This particular brute force strategy consists of *generating* solutions – naively constructing candidate solutions to a problem – and then *testing* them – evaluating the value function of them and at the same time removing invalid solutions. It is applicable whenever the number of candidate solutions is quite small.

Max Clique

In a graph, a subset of the vertices form a *clique* if each pair of vertices is connected by an edge.




Figure 9.1: A graph with a clique of size 4 highlighted.

Given a graph on V vertices and E edges, determine the size of the largest clique.

This problem is one of the so-called NP-complete problems we mentioned in Chapter 5. Thus, a polynomial-time solution is currently out of reach. We solve the problem for $V \leq 15$.

Is a generate and test approach suitable? To know, we must first define what our candidate solutions are. In this problem, only one object comes naturally; subsets of vertices. For every such subset candidate, we must test

whether it is a clique. If it is, its size must be computed and the largest clique chosen.

In the Max Clique problem, there are only 2^V subsets of vertices (and candidate solutions) – a quite small number. Given such a set, we can verify whether it is a clique in $O(V^2)$ time by checking if every pair of vertices in the candidate set has an edge between them. To perform this check in $\Theta(1)$ time, we keep a 2D vector *adj* such that *adj*[*i*][*j*] is true if and only if vertices *i* and *j* are adjacent to each other. This gives us a total complexity of $\Theta(2^V \cdot V^2)$ in the worst case. According to our table of approximate allowed input sizes for various complexities (p. 88), this should be fast enough for $V = 15$.

Max Clique

```

1  int V, E;
2  cin >> V >> E;
3  vector<vector<bool>> adj(V, vector<bool>(V));
4  rep(i,0,E) {
5      int a, b;
6      cin >> a >> b;
7      adj[a][b] = adj[b][a] = true;
8  }
9  rep(i,0,V) adj[i][i] = true;
10
11 int ans = 0;
12 rep(subset,0,1<<V) {
13     bool isClique = true;
14     rep(i,0,V) {
15         // Skip if the subset does not contain i
16         if ((subset & (1 << i)) == 0) continue;
17         rep(j,0,V) {
18             // Skip if the subset does not contain j
19             if ((subset & (1 << j)) == 0) continue;
20             if (!adj[i][j]) {
21                 // Subset contained both i, j and are not neighbors.
22                 isClique = false;
23             }
24         }
25     }
26     if (isClique) {
27         ans = max(ans, __builtin_popcount(subset));
28     }
29 }
30 cout << ans << endl;

```

Note the nifty use of integers interpreted as bitsets to easily iterate over every possible subset of an V -element set, a common technique in generate and test solutions based on subsets.

Problem 9.1*Max Clique* – maxclique

This kind of brute force problem is often easy to spot. There will be a very small input limit on the parameter you are to brute force over. The solution will often be subsets of some larger base set (such as the vertices of a graph), or combinations of several small sets.

Exercise 9.1. Algorithm ?? can be made to run in $\Theta(2^N \cdot N)$ by using bitsets to represent the neighbourhoods. Figure out how to do this.

Problem 9.2*4 thought* – 4thought*Lifting Walls* – walls

Let us look at another example of this technique, where the answer is not just a subset.

The Clock

Swedish Olympiad in Informatics 2004, School Qualifiers (CC BY-SA 3.0)

When someone asks you what time it is, most people respond “a quarter past five”, “15:29” or something similar. If you want to make things a bit harder, you can answer with the angle from the minute hand to the hour hand, since this uniquely determines the time. However, most people are not used to this way of specifying the time, so it would be nice to have a program which translates this to a more common format.




Figure 9.2: The angle between the hands at time :30.

We assume that our clock have no seconds hand, and only display the time at whole minutes (i.e., both hands only move forward once a minute). The angle is determined by starting at the hour hand and measuring the number of degrees clockwise to the minute hand. To avoid decimals, this angle is specified in tenths of a degree.

Input

The first and only line of input contains a single integer $0 \leq A < 3600$, the angle specified in tenths of a degree.

Output

Output the time in the format `hh:mm` between `00:00` and `11:59`.

It is difficult to come up with a formula that gives the correct times as a function of the angles between the hands on a clock. Instead, we can turn the problem around. If we know what the time is, can we compute the angle between the two hands of the clock?

Assume that the time is currently h hours and m minutes. The minutes hand is then at angle $\frac{360}{60}m = 6m$ degrees clockwise from straight up. Similarly, the hour hand moves $\frac{360}{12}h = 30h$ degrees clockwise after h whole hours, with an extra $\frac{360}{12} \frac{1}{60}m = 0.5m$ degrees added due to the minute. While computing the current time directly from the angle is difficult, computing the angle from the current time is easy.

The brute force solution is to test the $60 \cdot 12 = 720$ different times, and pick

the one which matched the given angle:

```
1: procedure CLOCK(A)
2:   for h  $\leftarrow$  0 to 11 do
3:     for m  $\leftarrow$  0 to 59 do
4:       hourAng  $\leftarrow$  300h + 5m     $\triangleright$  Angles are 10'ths of degrees to avoid
       storing half of degrees.
5:       minuteAng  $\leftarrow$  60m
6:       angBetween  $\leftarrow$  (minuteAng - hourAng + 3600) mod 3600
7:       if angBetween = A then
8:         return h:m
```

Exercise 9.2. Can there be two times that produce the same angle? If yes, produce such an example. If no, prove that there are no two such times.

Competitive Tip

Competitions sometimes pose problems which are solvable quite fast, but where a brute force algorithm suffice. Code the simplest correct solution that is fast enough, even if you see a faster one.

Problem 9.3

The Clock – theclock

All about that base – allaboutthatbase

Perket – perket

9.3 Backtracking

Backtracking is a variation of the generate and test method. It can be much faster than a generate and test solution, but it is not always applicable and is sometimes be more difficult to code (in particular when the solutions are subsets).

Consider our solution to the Max Clique problem. In our solution, we generated all the candidate solutions (i.e., subsets of vertex) by using bitsets. When solving problems where the candidate solutions are other objects than subsets, or the number of subsets is too large to iterate through, we need to construct the solutions in another way. Generally, we do this recursively.

When generating subsets, we would go through every element one at a time, deciding whether to include it or not in a recursive fashion. Backtracking extends




Figure 9.3: Recursively backtracking over subsets.

generate and test to not only testing all the candidate solutions, but also these partial candidates. Once a partial candidate is identified as being infeasible – for the clique example, once we include two non-neighbouring vertices – the backtracking can stop early. If there are much fewer valid partial candidates than total candidates checked by a generate and test approach, this saves time.

In Figure 9.3 an example of this approach is demonstrated. The example illustrates the beginning of a backtracking recursion that generates subsets. Subsets are recursively created by either including or excluding every element, one at a time. One subset, $\{1, 2\}$ was identified as not being permissible – for example, because the vertices they represented in a clique problem was not neighbours. Thus, no further backtracking was performed.

As a concrete example, consider a variant of the clique problem, where we are interested in computing the *number of cliques* with at most 6 vertices. In this case, we can solve the problem for larger instances than what generating all possible subsets and testing them allows us to. The trick is that while there are 2^V subsets of V vertices, there are only $O(V^7)$ subsets containing up to 6 vertices (see Chapter 17 on why this is). Thus, if we can restrict ourselves to never generating any other kind of subset, the solution ought to be much faster. This generation is easily implemented using backtracking:

6-cliques, Recursive Variant

```
1  // Recursively construct cliques of at most 6 vertices. We keep track
2  // of what vertices are included so far in the bitset included, and
3  // have yet to make a decision of whether vertices [0, 1, ..., at]
4  // should be included or not.
5  int clique6(int at, bitset<40> included,
6              const vector<bitset<40>>& neighbours) {
7      if (at < 0) return 1;
8      int answer = 0;
9      // Case 1: Not including the current vertex.
10     answer += clique6(at - 1, included, neighbours);
11     // Case 2: All the previously included vertices are neighbours,
12     // and we have only chosen at most 5 vertices previously.
13     if ((included & neighbours[at]) == included
14         && included.count() <= 5) {
15         included.set(at);
16         answer += clique6(at - 1, included, neighbours);
17     }
18     return answer;
19 }
20
21 int main() {
22     int V, E;
23     cin >> V >> E;
24     // Bitsets containing all the neighbours of all vertices
25     vector<bitset<40>> neighbours(V);
26     rep(i,0,E) {
27         int A, B;
28         cin >> A >> B;
29         neighbours[A].set(B);
30         neighbours[B].set(A);
31     }
32     cout << clique6(V - 1, bitset<40>(), neighbours) << endl;
33 }
```

How fast is this solution? Analysis is slightly tricky, but we can give an upper bound that is small enough. If we have about 40 vertices, there are about $7.6 \cdot 10^5$ subsets with at most 5 vertices. Since no subset requires more than 40 recursive calls to construct it (one for each vertex in the set), constructing these subsets require no more than about $3.0 \cdot 10^7$ recursive calls. Each subset of 6 elements is constructed from one of these subsets and results in only a single additional recursive call, so in total the function is invoked at most $6.0 \cdot 10^7$. The function performs only a few constant-time operations, so this should be fine. Compare this with a non-backtracking generate and test solution, which would need to construct about $1.1 \cdot 10^{12}$ subsets instead – clearly too much.

Backtracking in principle works if we can:

- construct candidate solutions recursively,
- quickly determine whether a partial candidate solution can possibly be completed to a candidate solution, and
- the number of valid partial solutions is sufficiently small.

Problem 9.4*6-cliques* – 6clique*Class Picture* – classpicture*Boggle* – boggle*Geppetto* – geppetto*Map Colouring* – mapcolouring*Picking Apples* – apples

As a general principle, backtracking seems simple enough. Some backtracking solutions require a bit more ingenuity, as in the next problem.

Basin City Surveillance

Nordic Collegiate Programming Contest 2014

Authors: Pål G. Drange and Markus S. Dregi (CC BY-SA 3.0)

BASIN CITY is known for her incredibly high crime rates. The police see no option but to tighten security. They want to install traffic drones at different intersections to observe who's running on a red light. If a car runs a red light, the drone will chase and stop the car to give the driver an appropriate ticket. The drones are quite stupid, however, and a drone will stop before it comes to the next intersection as it might otherwise lose its way home, its home being the traffic light to which it is assigned. The drones are not able to detect the presence of other drones, so the police's R&D department found out that if a drone was placed at some intersection, then it was best not to put any drones at any of the neighbouring intersections. As is usual in many cities, there are no intersections in Basin City with more than four other neighbouring intersections.

The drones are government funded, so the police force would like to buy as many drones as they are allowed to. Being the programmer-go-to for the Basin City Police Department, they ask you to decide, for a given number of drones, whether it is feasible to position exactly this number of drones.

Input

The first line contains an integer k ($0 \leq k \leq 15$), giving the number of drones to position. Then follows one line with $1 \leq n \leq 100\,000$, the total number of intersections in Basin City. Finally follow n lines describing consecutive intersections. The i 'th line describes the i 'th intersection in the following format: The line starts with one integer d ($0 \leq d \leq 4$) describing the number of intersections neighbouring the i 'th one. Then follow d integers denoting the indices of these neighbouring intersections. They will be all distinct and different from i . The intersections are numbered from 1 to n .

Output

If it is possible to position k drones such that no two neighbouring intersections have been assigned a drone, output a single line containing `possible`. Otherwise, output a single line containing `impossible`.

At a first glance, it is not even obvious whether the problem is a brute force problem, or if some smarter principle should be applied. After all, 100 000 intersections is a huge number of intersections! We can make the problem a bit more reasonable with our first insight. If we have a large number of intersections, and every intersection is adjacent to very few other intersection, it is probably very easy to place the drones at appropriate intersections. To formalize this insight, consider what happens when we place a drone at an intersection.




Figure 9.4: The intersections affected by placing a drone at an intersection.

By placing a drone at the intersection marked in black in Figure 9.4, at most five intersections are affected – the intersection we placed the drone at, along with its neighbouring intersections marked in gray. If we would remove these

five intersections, we would be left with a new city where we need to place $k - 1$ drones. This simple fact – which is the basis of a recursive solution to the problem – tells us that if we have $N \geq 5k - 4$ intersections, we immediately know the answer is *possible*. The -4 terms comes from the fact that when placing the final drone, we no longer care about removing its neighbourhood, since no further placements will take place.

Therefore, we can assume that the number of intersections is less than $5 \cdot 15 - 4 = 71$, i.e., $n \leq 70$. This certainly makes the problem seem much more tractable. Now, let us start developing solutions to the problem.

First of all, we can attempt to use the same algorithm as we used for the Max Clique problem. We could recursively construct the set of our k drones by, for each intersection, trying to either place a drone there or not. If placing a drone at an intersection, we would forbid placing drones at any neighbouring intersection.

Unfortunately, this means that we test every intersection when placing a certain drone somewhere. This would give us a complexity of $O(n^k)$. More specifically, the execution time $T(n, k)$ would satisfy $T(n, k) \approx T(n - 1, k) + T(n - 1, k - 1)$, which implies $T(n, k) \approx \binom{n}{k} = \Omega(n^k)$ (see Section ?? for more details). For $n = 70, k = 15$, this is too high. The values of n and k do suggest that an exponential complexity is in order, just not of this kind. Instead, something similar to $O(c^k)$ where c is a small constant would be a better fit. One way of achieving such a complexity would be to limit the number of intersections we must test to place a drone at before trying one that definitely works. If we could manage to test only c such intersections, we would get a complexity of $O(c^k)$.

Competitive Tip

In this problem, we tried to use the size of the parameters n and k together with the time limit to guide the kind of solution we need to design. While this works most of the time, note that this can sometimes be severely misleading – as this problem was before we realized that having 100 000 intersections was red herring.

The trick, yet again, comes from Figure 9.4. Assume that we choose to include the black intersection in our solution, but still can not construct a solution. The only reason this case can happen is (aside from bad previous choices) that no optimal solution includes this intersection. What could possibly stop this intersection from being included in an optimal solution? It must be because one of its gray neighbours is included in every optimal solution. If this was not the case, then we could just pick an optimal solution where none of the gray intersections were included and instead include the black vertex. Fortunately for us, this gives us just what we need to improve our algorithm – either a given intersection, or one of its neighbours, must be included in any optimal solution.

We have accomplished our goal of reducing the number of intersections to test for each drone to a mere 5, which will give us a complexity of about $O(5^k)$ (possibly with an additional polynomial factor in n depending on implementation). This is still too much, unless, as the jury, noted, some “clever heuristics” are applied. Fortunately for us, applying a common principle will speed things up dramatically (even giving us a better time complexity).

Our next trick is to assume that the graph we are working with is connected. In many problems the connected components of the graph are all independent of each other. This is also the case in this problem. Clearly placing drones in one component does not affect how we place drones in any other component, so we can solve them separately by computing the maximal number of drones we can place in every such set, until we have processed enough sets to place k drones.

How does the connectedness help us? Consider what happens when we have placed our first drone on the black intersection as in Figure 9.4. By removing it and the gray neighbours, the white intersections must now have at most 3 neighbours instead. Recursing on one of the white intersections would then leave us with only 4 choices, placing a drone on either the white vertex or one of its (at most 3) neighbours. In fact, we can extend this reasoning to show that there must always be an intersection with at most 3 neighbours! Proving this is a straightforward proof by contradiction. Assume that we in a connected graph have placed at least one drone, but all remaining intersections have four

neighbours. Then, none of these intersections can be neighbours with any vertex we have removed so far. This means that the set of intersections removed and the set of intersections remaining are actually disconnected, to the contrary of our assumption of connectedness. Taking this insight to its conclusion, we achieve a complexity of $O(4^k)$ by always branching on the intersection with the fewest neighbours.

While such an algorithm is significantly faster than $O(5^k)$, further improvements are possible. Again, let us consider under what circumstances a certain intersection is excluded from any optimal solution. We have already concluded that if this is the case, then one of its neighbours must be included in any optimal solution. Can it ever be the case that only *one* of its neighbours are included in an optimal solution, as in Figure 9.5 where we chose to place a drone on the black vertex but none of the white intersections?




Figure 9.5: Placing a drone at a single neighbour of an intersection.

This is actually never the case. We can always move the drone from the black intersection to the white intersection it has as neighbour, since none of its other white intersections contain a drone. Now, we are basically done; for any intersection, there will either be an optimal solution including it, or (at least) two of its neighbours. Since an intersection has at most 4 neighbours, it has at most 6 pairs of neighbours. This means our recursion will take time $T(k) = T(k-1) + 6T(k-2)$ in the worst case. This recurrence has the solution 3^k , since $3^{k-1} + 6 \cdot 3^{k-2} = 3^{k-1} + 2 \cdot 3^{k-1} = 3 \cdot 3^{k-1} = 3^k$. A final improvement would be to combine this insight with the independence of the connected subsets of intersections. The second term of the time recurrence would then be a 3

instead of a 6 (as 3 neighbours make 3 pairs). Solving this recurrence would give us the complexity $O(2.31^k)$ instead.

Problem 9.5

Basin City – basincity

The general version of this problem (without the bounded degree) is called **Independent Set**, and is also one of the NP-complete problems.

So, what is the take-away regarding backtracking? Start by finding a way to construct candidate solutions iteratively. Then, try to integrate the process of testing the validity of a complete solution with the iterative construction, in the hope of significantly reducing the number of candidate solutions which need evaluating. Finally, we might need to use some additional insights, such as what to branch on (which can be something complicated like the neighborhood of a vertex), deciding whether to backtrack or not (i.e., improving the testing part) or reducing the number of branches necessary (speeding up the generation part).

Problem 9.6

Domino – domino

Fruit Baskets – fruitbaskets

Infiltration – infiltration

Vase Collection – vase

9.4 Fixing Parameters

The parameter fixing technique is also similar to the generate and test method, but is not used to test the solution set itself. Instead, you perform brute force to *fix* some parameter in the problem by trying possible values they can assume. Hopefully, fixing the correct parameters allows you to solve the remaining problem easily. The intuition is that while any one choice of parameter may be wrong, testing every choice allows us to assume that we at some point used the correct parameter.

The path to a solution usually starts in the other end. Let us illustrate the technique with a problem.

Buying Books

Swedish Olympiad in Informatics 2010, Finals

You are going to buy N books, and are currently checking the different M internet book stores for prices. Each book is sold by at least one book store, and can vary in prices between the different stores. Furthermore, each book store incur a postage fee if you order from it. Postage may vary between the various book stores, but it is always the same for a book store no matter how many books you decide to order. You may order any number of books from any number of the book stores. Compute the smallest amount of money you need to pay for all the books.

Input

The first line contains two integers $1 \leq N \leq 100$ – the number of books, and $1 \leq M \leq 15$ – the number of book stores.

Then, M descriptions of the book stores follow. The description of the i 'th store starts with a line containing two integers $0 \leq P_i \leq 1\,000$ (the postage for this book store), and $1 \leq L_i \leq N$ (the number of books this store sells). The next L_i lines contains the books sold by the store. The j 'th book is described by two integers $0 \leq B_{i,j} < N$ – the (zero-indexed) number of a book being sold here, and $1 \leq C_{i,j} \leq 1000$ – the price of this book at the current book store.

Output

Output a single integer – the smallest amount of money you need to pay for the books.

If we performed naive generate and test on this problem, we would probably get something like 15^{100} solutions, by testing every book stores for every book. This is infeasible. So, why can we do better than this? There must be some hidden structure in the problem that makes testing all those possibilities unnecessary. To find this structure, we will analyze a candidate solution as given by the naive generate and test method, i.e. an assignment from each book to a book store where we should purchase it from.

For the sake of example, let's assume that in this candidate solution, we purchased books from the book stores 1, 4 and 5. If we then purchased a book from store 4, but it was actually cheaper from store 1, we should have picked it from there instead. Thus there seems to be quite a bit of redundancy in this set of candidate solutions – a strong hint that we might have found some crucial insight. We *could* decide to use this fact to turn our generate and test into a

backtracking algorithm, by pruning away any solution where we at some point decide to purchase a book from a book store that contains a book that we already purchased at a cheaper price. Unfortunately, this is easily defeated by giving most of the books equal prices at most of the book stores.

Instead, let us use the insight differently. Our observation hints that making a choice for every book is not especially good, since choices elsewhere affects the optimality of an individual book choice greatly. Digging deeper, we find that we do not really have any choice in where we buy the books. During the course of the suggested backtracking, what matters is never what particular store we purchased a book from. Rather, it is only of interest what stores we have decided to use so far – this will uniquely determine where we purchase every book from, since we are forced to buy the book from the cheapest store that we buy books from.

At this point, we are basically done. We have reduced the amount of information we need to know in order to easily find a solution “which book stores will we purchase from?”. This parameter has much fewer possibilities – only 2^{15} . By testing each possible choice and *fixing* it, we can immediately find the best candidate solution using the “cheapest store for each book” rule. Since we test every option for this parameter in the full set of candidate solutions, we must also test the one that is given by the optimal solution. This results in the following pseudo code:

```
1: procedure BUYINGBOOKS(books  $N$ , stores  $M$ , costs  $C$ , postages  $P$ )
2:    $answer \leftarrow \infty$ 
3:   for every  $S \subseteq [M]$  do
4:      $cost \leftarrow 0$ 
5:     for every  $s \in S$  do
6:        $cost \leftarrow cost + P_s$ 
7:       for every  $b \in [N]$  do
8:          $cost \leftarrow cost + \min_{i \in S} C_{i,b}$ 
9:        $answer \leftarrow \min(answer, cost)$ 
10:  return  $answer$ 
```

Alternatively, we could have come to the same insight by asking ourselves “can we brute force over the number of book stores?”. Whenever you have a parameter about this size in a problem, this is a question worth asking. However,

the parameter to brute force over is not always this explicit, as in the following problem, which asks us to find all integer solutions to an equation in a certain interval.

Problem 9.7

Buying Books – buyingbooks

Integer Equation

Codeforces Round #262, Problem B

Given integers a ($|a| \leq 10\,000$), b ($1 \leq b \leq 5$), and c ($|c| \leq 10\,000$), determine the integers x ($1 \leq x \leq 10^9$) satisfying

$$x = a \cdot s(x)^b + c,$$

where $s(x)$ is the digit sum of x .

In this problem, the only explicit object we have is x . Unfortunately, 10^9 is a tad too many possibilities. If we study the equation a bit closer, we see that $s(x)$ also varies as a function of x . This is helpful, since $s(x)$ has far fewer options than x . In fact, a number bounded by 10^9 has at most 9 digits, meaning it has a maximum digit sum of $9 \cdot 9 = 81$. Thus, we can solve the problem by looping over all the possible values of $s(x)$. This uniquely fixes our right hand side, which equals x . Given x , we verify that $s(x)$ has the correct value. Whenever we have such a function, i.e., one with a large domain (like x in the problem) but a small image (like $s(x)$ in the problem), this technique can be used by brute forcing over the image instead. This is actually what we did in the book store problem as well. Our function was then from a given candidate solution to the book stores used. Since the image of this function (all the possible subsets of book stores) was small, the problem could be attacked by brute forcing over the image of the function rather than the domain.

Problem 9.8

Shopping Plan – shoppingplan

9.5 Meet in the Middle

The *meet in the middle* technique is a special case of the parameter fixing technique. The general theme will be to fix half of the parameter space and

build some fast data structure such that when testing the other half of the parameter space, we can quickly find the best parameters for the first half. It is a space-time tradeoff, in the sense that we improve the time usage (testing half of the parameter space much faster), by paying with increased memory usage (to save the pre-computed structures).

Subset Sum

Given a set of integers S , is there some subset $A \subseteq S$ with a sum equal to T ?

Input

The first line contains an integer N , the size of S , and T . The next line contains N integers s_1, s_2, \dots, s_N , separated by spaces – the elements of S . It is guaranteed that $s_i \neq s_j$ for $i \neq j$.

Output

Output possible if such a subset exists, and impossible otherwise.

In this problem, a simple generate and test solution would have N parameters to brute force over. For each element of S , we either choose to include it in A or not – a total of two choices for each parameter. This naive attempt at solving the problem (which amounts to computing the sum of every subset) gives us a complexity of $O(2^N)$. While sufficient for e.g. $N = 20$, we can make an improvement that makes the problem tractable even for $N = 40$.

To figure out if a meet in the middle solution is applicable, the two halves of the parameter space must to a large extent be independent. Individual choices in one half should have little effect on the other. This is for example not the case in the max clique problem. Deciding what vertices to include from, say, the lower-numbered half will put very complicated constraints on what vertices we could pick from the other half. Such a situation should discourage you from attempting to meet in the middle.

In the subset sum problem, our parameters are to a large extent independent. When fixing the $\frac{N}{2}$ first parameters, which may mean we include elements with a sum of U , a single constraint is placed on the remaining $\frac{N}{2}$ parameters; they must sum to $T - U$, to together make the correct sum.

Thus, if we could quickly answer the question “can we choose the latter half of the integers such that they have a given sum?” we could solve the problem by fixing the first half of the parameters. Individually, this constraint takes $O(2^{\frac{N}{2}})$ time to check if we use brute force. However, we can compute the answer for

all such questions in one go by computing the sum of every subset of the latter half of elements in $\Theta(N2^{\frac{N}{2}})$. The resulting sums can be inserted into a hash set, letting us determine if a sum can be formed by those elements in $\Theta(1)$ instead. Here is the space-time tradeoff – we sacrifice an exponential amount of memory (a set of $2^{\frac{N}{2}}$ elements) to win an exponential speedup of $2^{\frac{N}{2}}$. The complexity is thus $\Theta(N2^{\frac{N}{2}})$, dominated by the precomputation. In pseudo code, a meet in the middle solution looks something like this:

```

1: procedure SUBSETSUM(set  $S$ , target  $T$ )
2:    $N \leftarrow |S|$ 
3:    $left \leftarrow \frac{N}{2}$ 
4:    $right \leftarrow N - left$ 
5:    $Lset \leftarrow$  the  $left$  first elements of  $S$ 
6:    $Rset \leftarrow S \setminus Lset$ 
7:    $Lsums \leftarrow$  new set
8:   for each  $L \subseteq Lset$  do
9:      $Lsums.insert(\sum_{l \in L} l)$ 
10:  for each  $R \subseteq Rset$  do
11:     $sum \leftarrow \sum_{r \in R} r$ 
12:    if  $Lsums.contains(T - sum)$  then
13:      output true
14:    return
15:  output false

```

Problem 9.9

Maximum Loot – maxloot

We will end the chapter by solving a brute force problem combining two techniques.

Limited Correspondence

Greg Hamerly

Emil, a Polish mathematician, sent a simple puzzle by post to his British friend, Alan. Alan sent a reply saying he didn't have an infinite amount of time he could spend on such non-essential things. Emil modified his puzzle (making it a bit

more restricted) and sent it back to Alan. Alan then solved the puzzle.

Here is the original puzzle Emil sent: given a sequence of pairs of strings $(a_1, b_1), (a_2, b_2), \dots, (a_k, b_k)$, find a non-empty sequence s_1, s_2, \dots, s_m such that the following is true:

$$a_{s_1} a_{s_2} \dots a_{s_m} = b_{s_1} b_{s_2} \dots b_{s_m}$$

where $a_{s_1} a_{s_2} \dots$ indicates string concatenation. The modified puzzle that Emil sent added the following restriction: for all $i \neq j$, $s_i \neq s_j$.

You don't have enough time to solve Emil's original puzzle. Can you solve the modified version?

Input

The input starts with a line containing an integer $1 \leq k \leq 11$, followed by k lines. Each of the k lines contains two lowercase alphabetic strings which represent a pair of strings. Each individual string will be non-empty and at most 100 characters long.

Output

Output the sequence found (if it is possible to form one) or IMPOSSIBLE (if it is not possible to solve the problem). If it is possible but there are multiple sequences, you should prefer the shortest one (in terms of the number of characters output). If there are multiple shortest sequences, choose the one that is lexicographically first.

The original problem as posed by Emil¹ is called the Post correspondence problem and is an *undecidable* problem, i.e. there is no algorithm in the familiar sense that can solve the problem in finite time.

Alan's added restriction to the problem, that if $i \neq j$, then $s_i \neq s_j$, means that each pair of strings may be used in the sequence at most once. This clearly allows us to solve the problem in finite time – we could simply test all $k!$ permutations of the pairs, and walk through them to see if the two strings formed by the respective strings of all the pairs match. This would require around $k! \cdot k \cdot 100$ operations, which is about $4 \cdot 10^{10}$ for the maximum $k = 11$, which is clearly too slow.

The big difference in this problem compared to the one where we could perform meet in the middle is that our selection of strings is highly *order dependent*. We can't arbitrarily split up our word pairs into a "first half" and a

¹Emil Post, an American mathematician

“second half” and attempt to combine them. After all, the correct solution might involve constructing a string where words from the two halves are intertwined. Of course, we have previously seen how to deal with such a problem. If an arbitrary choice is not good enough, we try *all* the choices. We simply fix the parameter that is the *subset of word pairs constituting the first half*. There are only 462 ways in which one can pick the first 5 word pairs out of a maximum eleven – a small price to pay.

Fixing parameters thus allows us to split up the words in two halves. This begs the question – if we attempt all possible permutations of the words in the first half, what constraint do they put on the second half? Let’s check. Assume that a given permutation of the (currently fixed) first half of the pairs give a concatenated string of a ’s is the string S , and without loss of generality shorter than the concatenation of the b ’s. First of all, it should be clear that S must be a substring of the concatenation of the b ’s – otherwise, concatenating the strings of the second half can’t make them equal. Thus, we assume that the concatenation of the b ’s is ST .

Symmetrically, this tells us that if the concatenation of all b ’s in the second half is the string U , the concatenation of all the a ’s must be TU , to together make the strings STU . Thus, the question is – can we order the words of the second part so that they create strings of the form U and TU for any U ?

This is precisely the kind of simple question that makes meet in the middle possible. To answer it quickly, we check all the 6 words in the second half. If they are of the form U and XU , for some X , we store it in a hash map from X to the lexicographically smallest U that we found. With this map in hand, we can determine if there is a way to complete the strings formed by the first half in constant time.

In total, the cost is somewhere around about $462 \cdot (5! \cdot 5 + 6! \cdot 6) \approx 2 \cdot 10^6$ hash map operations, and $462 \cdot (5! \cdot 5 + 6! \cdot 6) \cdot 600 \approx 1.3 \cdot 10^9$ individual character operations depending on implementation. While the latter seems like a lot, they are very fast – solving the worst-case on the author’s computer 5 times takes less than a second.

Problem 9.10*Closest Sums* – closestsums*Celebrity Split* – celebritiesplit*Circuit Counting* – countcircuits

Indoorienteering – indoorienteering

Key to Knowledge – keytoknowledge

Knights in Fen – knightsfen

Rubrik's Revenge in ... 2D!? 3D? – rubriksrevenge

9.6 Chapter Notes

Brute force problems come in many varieties. While the basis of most solutions often involve some kind of backtracking with pruning, finding the right ones can often need a lot of creativity.

A deeper dive into brute force together can be found in *Exact Exponential Algorithms*[1]. The book also discusses several techniques that are common within algorithm problem solving.

A concept that often comes up in research and occasionally in programming competitions, is that of *fixed parameter tractability*. A problem that might (as of now) only have exponential time algorithms, such as the max clique problem, may have a polynomial solution when fixing some parameter, such as the size of the maximum clique. Indeed, we saw that the problem of finding the largest clique of size at most 6 admitted a polynomial solution in the number of vertices. *Parameterized Algorithms*[8] provide a comprehensive toolbox of such problems.

Another neat technique with elements of brute force is *color-coding*[2] to solve certain graph theoretical problems.

10 Greedy Algorithms

In this chapter, we are going to look at another standard technique to solve some categories of search and optimization problems faster than naive brute force, by exploiting properties of *local optimality*.

10.1 Change-making Problem

Let us start by looking at a well-known algorithmic problem, the *change-making problem*.

Change-making Problem, Denominations 1, 2, 5

Given an infinite number of coins of denominations 1, 2, 5, determine the smallest number of coins needed to sum up to T ($1 \leq T \leq 10^9$).

We can use the tools from the previous chapter on brute force to formulate a plain backtracking solution. Using a recursive function that takes in T , we can attempt to add a single coin of each type, and keep searching:

```
1: procedure MAKECHANGE(integer  $T$ )
2:   if  $T = 0$  then
3:     return 0
4:    $answer \leftarrow 1 + \text{MakeChange}(T - 1)$ 
5:   if  $T \geq 2$  then
6:      $answer \leftarrow \min(answer, 1 + \text{MakeChange}(T - 2))$ 
7:   if  $T \geq 5$  then
8:      $answer \leftarrow \min(answer, 1 + \text{MakeChange}(T - 5))$ 
9:   return  $answer$ 
```

However, this solution is exponential in time. It branches three times per call, and each call decreases T by at most 5, so the recursion will have a depth of at least $\frac{T}{5}$. It thus has a lower bound of $\Omega(3^{\frac{T}{5}})$.

We can phrase this problem using the kind of graph we previously discussed. Let the graph have vertices labeled $0, 1, 2, \dots, T$, representing the amount of

money we wish to sum up to. For a vertex labeled x , we add edges to vertices labeled $x - 1$, $x - 2$, $x - 5$ (if those vertices exist), weighted 1. Traversing such an edge represents adding a coin of denomination 1, 2 or 5. Then, the Change-making Problem can be phrased as computing the shortest path from the vertex T to 0. The corresponding graph for $T = 5$ can be seen in Figure 10.1.




Figure 10.1: The Change-making Problem, formulated as finding the shortest path in a DAG, for $K = 5$.

So, how does this graph formulation help us? Solving the problem on the graph as before using simple recursion would be very slow (with an exponential complexity, even). In Chapter 11 on Dynamic Programming, we will see how to solve such problems in polynomial time. For now, we will settle with solving problems exhibiting yet another property besides having optimal substructure – that of local optimality.

Exercise 10.1. Compute the shortest path from each vertex in Figure 10.1 to T using the optimal substructure property.

10.2 Optimal Substructure

Most optimization problems we study can be formulated as having to make sequential choices. For example, we managed to formulate the max clique problem as a sequence of V choices – for each vertex, should we include it or not?

A useful formulation of this process is the following: Given a weighted, directed acyclic graph (DAG) on N vertices, what is the “best” path from a vertex S to another vertex T ? This graph is almost never be given explicitly. Instead, it hides behind the problem as a set of states (the vertices). At each state, we are to make some choice that takes us to another state (traversing an edge). In the case of backtracking, you may recognize elements of this formulation in

Figure 9.3 on page 155, since we essentially explore a big kind of graph in such algorithms.

If the path consists of edges e_1, e_2, \dots, e_k , the function we are to maximize will be of the form

$$G(e_1, e_2, \dots, e_k) = g(e_1, g(e_2, g(\dots, g(e_k, 0))))$$

where g is a function from $E \times R$ to R . We will denote $B(v)$ as the maximum value of $G(e_1, e_2, \dots, e_k)$ over all paths from v to T . Often, g will be the negative sum of the edge weights, meaning we look for the shortest path from S to T . If $g(e, x)$ is increasing in x , we say that problems exhibiting this property has *optimal substructure*.

One way to solve the problem would be to evaluate G for every path in the graph. Unfortunately, the number of paths in a general graph can be huge (growing exponentially with the number of vertices). However, the optimal substructure property allows us to make a simplification. Assume that S has neighbors v_1, v_2, \dots, v_m . Then by the definition of B and g , $B(s) = \max(g(\{s, v_i\}, B(v_i)))$. Thus, we can solve the problem by first solving **the same problem** on all vertices v_i instead.

10.3 Locally Optimal Choices

Greedy algorithms solve this kind of problem by making what is called a *locally optimal choice*. We generally construct our shortest path iteratively, one edge at a time. We start out at the first vertex, S , and need to choose an edge to traverse. A greedy algorithm chooses the edge which locally looks like a good choice, without any particular thought about future edges.

For example, consider how you would solve the Change-making Problem yourself. You would probably attempt to add the 5 coin as many times as possible, until you need to add less than 5. This makes sense locally, since 5 is the largest amount we may charge at a time. When we need to add less than 5, we would probably instead add coins worth 2, until we need to add either 0 or 1. In the latter case, we would add a final 1 coin.

Intuitively, this makes sense. Adding the highest amount every time *ought* to be the best possible. For this problem, this is actually true. However, if the denominations were different, this could fail (Exercise 10.2).

Exercise 10.2. Prove that the greedy choice may fail if the coins have denominations 1, 6 and 7.

Assume that the optimal solution uses the 1, 2 and 5 coins a , b and c times respectively. We then have that either:

$a = 1, b \leq 1$ If $b \geq 2$, we could exchange one 1 coin and two 2 coins for one 5 coin.

$a = 0, b \leq 2$ If $b \geq 3$, we could exchange three 2 coins one 1 coin and one 5 coin.

If $a \geq 2$, we could instead add a single 2 coin instead of two 1 coins.

This means that the possibilities for a and b are few:

- $a = 0, b = 0$: value 0
- $a = 1, b = 0$: value 1
- $a = 0, b = 1$: value 2
- $a = 1, b = 1$: value 3
- $a = 0, b = 2$: value 4

Now, assume that we use exactly c coins of value 5, so that $T = 5c + r$. We know that $0 \leq r < 5$. Otherwise, we would be summing up to an amount larger than 4 without using any 5 coins, but this is impossible based on the above list. This means c must be as large as possible, i.e. it is optimal to add as many 5 coins as possible – which the greedy choice will. Then, only the cases 0, 1, 2, 3 and 4 remain. Looking at the list of those cases, we see that their optimal solutions correspond to how the greedy algorithm works. Thus, the greedy algorithm will always give the optimal answer.

Competitive Tip

If you have the choice between a greedy algorithm and another algorithm (such as one based on brute force or dynamic programming), use the other algorithm unless you are certain the greedy choice works.

Proving the correctness of a locally optimal choice is sometimes very cumbersome. In the remainder of the chapter, we are going to look at a few standard problems that are solvable using a greedy algorithm. Take note of the kind of arguments we are going to use – there are a few common types of proofs which are often used in proofs of greedy algorithms.

10.4 Scheduling

Scheduling problems are a class of problems which deals with constructing large subsets of non-overlapping intervals, from some given set of intervals.

The classical Scheduling Problem is the following.

Scheduling Problem

Given is a set of half-open (being open on the right) intervals S . Determine the largest subset $A \subseteq S$ of non-overlapping intervals.

Input

The input contains the set of intervals S , where $|S|$.

Output

The output should contain the subset A .




Figure 10.2: An instance of the scheduling problem, with the optimal solution at the bottom.

We will construct the solution iteratively, adding one interval at a time. When looking for greedy choices to perform, extremal cases are often the first ones you should consider. Hopefully, one of these extremal cases can be proved

to be included in an optimal solution. For intervals, some extremal cases would be:

- a shortest interval,
- a longest interval,
- an interval overlapping as few other intervals as possible,
- an interval with the leftmost left endpoint (and symmetrically, the rightmost right endpoint),
- an interval with the leftmost right endpoint (and symmetrically, the rightmost left endpoint).

As it turns out, we can always select an interval satisfying the fifth case. In the example instance in Figure 10.2, this results in four intervals. First, the interval with the leftmost right endpoint is the interval $[1, 2]$. If we include this in the subset A , intervals $[0, 3]$ and $[1, 6]$ must be removed since they overlap $[1, 2]$. Then, the interval $[3, 4]$ would be the one with the leftmost right endpoint of the remaining intervals. This interval overlaps no other interval, so it should obviously be included. Next, we would choose $[4, 6]$ (overlapping with $[4, 7]$), and finally $[7, 8]$. Thus, the answer would be $A = \{[1, 2], [3, 4], [4, 6], [7, 8]\}$.

```

1: procedure SCHEDULING(set  $S$ )
2:    $ans \leftarrow$  new set
3:   Sort  $S$  by right endpoint
4:    $highest \leftarrow \infty$ 
5:   for each interval  $[l, r) \in S$  do
6:     if  $l \geq highest$  then
7:        $ans.insert([l, r))$ 
8:        $highest \leftarrow r$ 
9:    $Lsums.insert(\sum_{l \in L} l)$ 
10:  return  $ans$ 

```

We can prove that this strategy is optimal using a *swapping argument*, one of the main greedy proof techniques. In a swapping argument, we attempt to prove that given any solution, we can always modify it in such a way that our greedy choice is no worse. This is what we did in the Change-making Problem,

where we argued that an optimal solution had to conform to a small set of possibilities, or else we could swap some set of coins for another (such as two coins worth 1 for a single coin worth 2).

Assume that the optimal solution does not contain the interval $[l, r)$, an interval with the leftmost right endpoint. The interval $[l', r')$ that has the leftmost right endpoint *of the intervals in the solution*, must have $r' > r$. In particular, this means any other interval $[a, b)$ in the solution must have $a \geq r'$ (or else the intervals would overlap). However, this means that the interval $[l, r)$ does not overlap any other interval either, since $a \geq r' > r$ so that $a \geq r$. Then, swapping the interval $[l, r)$ for $[l', r')$ still constitute a valid solution, of the same size. This proves that we could have included the interval $[l, r)$ in the optimal solution. Note that the argument in no way say that the interval $[l, r)$ *must* be in an optimal solution. It is possible for a scheduling problem to have many distinct solutions. For example, in the example in Figure 10.2, we might just as well switch $[4, 6)$ for $[4, 7)$ and still get an optimal solution.

This solution can be implemented in $\Theta(|S| \log |S|)$ time. In the Scheduling algorithm, this is accomplished by sort performing a sort (in $\Theta(|S| \log |S|)$), followed by a loop for each interval, where each iteration takes $\Theta(1)$ time.

Exercise 10.3. For each of the first four strategies, find a set of intervals where they fail to find an optimal solution.

We can extend the problem to choosing exactly K disjoint subsets of non-overlapping intervals instead, maximizing the sum of their sizes. The solution is similar to the original problem, in that we always wish to include the interval with the leftmost right endpoint in one of the subsets if possible. The question is then, what subset? Intuitively, we wish to choose a subset there the addition of our new interval causes as little damage as possible. This subset is the one with the rightmost right endpoint that we can place the interval in. Proving this is similar to the argument we used when deciding what interval to choose, and is a good exercise to practice the swapping argument.

Exercise 10.4. Prove that choosing to place the interval in the subset with the rightmost right endpoint is optimal.

Problem 10.1

Entertainment Box – entertainmentbox

Disastrous Downtime – downtime

10.5 Huffman Coding

When storing binary data with a non-negligible degree of regularity, it is possible to encode this data using a representation that uses fewer bits than the data itself. For example, we could store the string “01010101010101010101010101” as “01, 14 times”. One such encoding scheme uses a lossless type of coding called *Huffman Coding*.

Definition 10.1 — Prefix-free code

Given a set of symbols S , a binary *prefix-free code* of S is a function $P : S \rightarrow \{0, 1\}^*$ such that $P(x)$ is a prefix of $P(y)$ only if $x = y$.

A prefix-free code is useful since it allows us to encode a string of symbols using a sequence of bits, such that the original string of symbols can be uniquely determined from the bit sequence. This can be done greedily. If the binary encoding is $B = b_1b_2 \dots b_n$, we find any symbol c such that $P(c)$ is some prefix $b_1b_2 \dots b_k$ of the encoding. We know only one such prefix can exist. Otherwise, if the encodings $P(c)$ and $P(d)$ of two symbols c and d would both be prefix of B , the shorter of them would also be a prefix of the other. Then, P would not be a prefix-free code. We thus choose the unique symbol whose encoding is a prefix of B as the first symbol that was encoded, remove the corresponding symbols and continue until B is empty.

Example 10.1 Let $S = \{a, b, c, d\}$. Then, the code defined by $P(a) = 0$, $P(b) = 1$, $P(c) = 10$, $P(d) = 11$ is not prefix-free, since $P(b) = 1$ is a prefix of $P(c) = 10$.

On the other hand, the code defined by $P(a) = 0$, $P(b) = 11$, $P(c) = 100$, $P(d) = 101$ is prefix-free. If 0110101100 is an encoding of a string of symbols, the code given by P gives us that the first symbol must be a (0), the second b (11), the third a (0), followed by d (101) and c (100).

There exists a straightforward fixed-length construction of a prefix-free code. If $2^{n-1} < |S| \leq 2^n$, we can choose $|S|$ strings of exactly n bits as encodings for the symbols. For compression purposes, we instead want a *short* encoding.

The *Huffman code* is a special kind of prefix-free code that takes into account the frequency of every symbol. For example, if the string we are to encode consist to 99% of the letter a , we probably want a to have a very short encoding. On the other hand, a symbol which only appear with a frequency of 1% does not contribute much to the length of the encoding.

Problem 10.2*Whether Report* – whether**Chapter Notes**

Determining whether coins of denominations D can even be used to construct an amount T is an NP-complete problem in the general case[16]. It possible to determine what cases can be solved using the greedy algorithm described in polynomial time though[6]. Such a set of denominations is called a *canonical coin system*.

Introduction to Algorithms[7] also treats the scheduling problem in its chapter in greedy algorithms. It also brings up the connection between greedy problems and a concept known as *matroids*, which is well worth studying.

11 Dynamic Programming

This chapter will study a technique called *dynamic programming* (often abbreviated DP). In one sense, it is simply a technique to solve the general case of the best path in a directed acyclic graph problem (Section 10.2) in cases where the graph does not admit locally optimal choices, in time approximately equal to the number of edges in the graph. For graphs which are essentially trees with a unique path to each vertex, dynamic programming is no better than brute force. In more interconnected graphs, where many paths lead to the same vertex, the power of dynamic programming shines through. It can also be seen as a way to speed up recursive functions (called *memoization*), which will be our first application.

First, we will see a familiar example – the Change-making problem, with a different set of denominations. Then, we will discuss a little bit of theory, and finally round off with a few concrete examples and standard problems.

11.1 Best Path in a DAG

We promised a generalization that can find the best path in a DAG that exhibit optimal substructure even when locally optimal choices does not lead to an optimal solution. In fact, we already have such a problem – the Change-making Problem from Section 10.2, but with certain other sets of denominations. Exercise 10.2 even asked you to prove that the case with coins worth 1, 6 and 7 could not be solved in the same greedy fashion.

So, how can we adapt our solution to this case? The secret lies in the graph formulation of the problem which we constructed (Figure 10.1). For the greedy solution, we essentially performed a recursive search on this graph, except we always knew which edge to pick. When we do not know, the solution ought to be obvious – let us test *all* the edges.

```
1: procedure CHANGEMAKING(denominations  $D$ , target  $T$ )
2:   if  $T = 0$  then
3:     return 0
```

```

4:  ans ← ∞
5:  for denomination d ∈ {1, 6, 7} do
6:    if T ≥ d then
7:      ans = min(ans, 1 + ChangeMaking(D, T − d)
8:  return ans

```

This solution as written is actually exponential in T (it is $\Omega(3^{\frac{T}{7}})$). The recursion tree for the case $T = 10$ can be seen in Figure 11.1.




Figure 11.1: The recursion tree for the Change-making problem with $T = 10$.

The key behind the optimal substructure property is that the answer for any particular call in this graph with the same parameter c is the same, independently of the previous calls in the recursion. Right now, we perform calls with the same parameters multiple times. Instead, we can save the result of a call the first time we perform it:

```

1: memo = new int[T + 1]
2: set memo[i] = −1 for all i
3: procedure CHANGEMAKING(denominations D, target T)
4:   if T = 0 then
5:     return 0
6:   if memo[T] ≠ −1 then

```

```

7:   return memo[T]
8:   ans  $\leftarrow \infty$ 
9:   for denomination d  $\in D$  do
10:    if T  $\geq d$  then
11:      ans = min(ans, 1 + ChangeMaking(D, T - d))
12:    memo[T]  $\leftarrow$  ans
13:  return ans

```

This new algorithm is actually linear in T instead of exponential. The call graph now looks very different (Figure 11.2), since all calls with the same parameter will be merged (as such calls are only evaluated once).




Figure 11.2: The recursion tree for the Change-making problem with $T = 10$, with duplicate calls merged.

Note the similarity between this graph and our previous DAG formulation of the Change-making problem (Figure 10.1).

11.2 Dynamic Programming

With these examples in hand, we are ready to give a more concrete characterization of dynamic programming. In principle, it can be seen as solving the kind of “sequence of choices” problems that we used brute force to solve, but where different choices can result in the same situation. For example, in the Change-making Problem, after adding two coins worth 1 and 6, we might just as well have added a 7 coin instead. After we have performed a sequence of

choices, *how* we got to the resulting state is no longer relevant – only where we can go from there. Basically, we throw away the information (what exact coins we used) that is no longer needed. This view of dynamic programming problems as having a “forgetful” property, that the exact choices we have made do not affect the future, is useful in most dynamic programming problems.

Another, more naive view, is that dynamic programming solutions are simple recursions, where we happen to solve the same recursive subproblem a large number of times. In this view, a DP solution is basically nothing more than a recursive solution – find the correct base cases, a fast enough recursion, and memoize the results.

More pragmatically, DP consists of two important parts – the states of the DP, and the computation of a state. Both of these parts are equally important. Fewer states generally mean less computations to make, and a better complexity per state gives a better complexity overall.

Bottom-Up Computation

When applied to a dynamic programming problem, memoization is sometimes called *top-down dynamic programming* instead. The name is inspired from the way we compute the solution to our problem by starting at the largest piece at the top of the recursion tree, and recursively breaking it down to smaller and smaller pieces.

There is an alternative way of implementing a dynamic programming solution, which (not particularly surprisingly) is called *bottom-up dynamic programming*. This method instead constructs the solutions to our sub-problems in the other order, starting with the base case and iteratively computing solutions to larger sub-problems.

For example, we might just as well compute the solution to the Change-making problem the following way:

```
1: procedure CHANGEMAKING(denominations  $D$ , target  $T$ )
2:    $ans = \text{new int}[T + 1]$ 
3:   set  $ans[i] = \infty$  for all  $i$ 
4:   set  $ans[0] = 0$ 
5:   for  $i = 1$  to  $T$  do
6:     for denomination  $d \in D$  do
7:       if  $i \geq d$  then
```

```
8:         ans[i] = min(ans[i], 1 + ans[i - d])
9:     return ans[T]
```

How do you choose between bottom-up and top-down? Mostly, it comes down to personal choice. A dynamic programming solution will almost always be fast enough no matter if you code it recursively or iteratively. There are some performance concerns, both ways. A recursive solution is affected by the overhead of recursive function calls. This problem is not as bad in C++ as in many other languages, but it is still noticeable. When you notice that the number of states in your DP solution is running a bit high, you might want to consider coding it iteratively. Top-down DP, on the other hand, has the upside that only the states reachable from the starting state will be visited. In some DP solutions, the number of unreachable states which are still in some sense “valid” enough to be computed bottom-up is significant enough that excluding them weighs up for the function call overhead. In extreme cases, it might turn out that an entire parameter is uniquely given by other parameters (such as the Ferry Loading problem in Section 11.3). While we probably would notice when this is the case, the top-down DP saves us when we do not.

Order of Computation and Memory Usage

For top-down DP, the memory usage is often quite clear and unavoidable. If a DP solution has N states, it will have an $\Omega(N)$ memory usage. For a bottom-up solution, the situation is quite different.

Firstly, let us consider one of the downsides of bottom-up DP. When coding a top-down DP, you do not need to bother with the structure of the graph you are solving your problem on. For a bottom-up DP, you need to ensure that whenever you solve a subproblem, you have already solved its subproblems too. This requires you to define an *order of computation*, such that if the subproblem a is used in solving subproblem b , a is computed before b .

In most cases, such an order is quite easy to find. Most parameters can simply be increasing or decreasing, using a nested loop for each parameter. When a DP is done over intervals, the order of computation is often over increasing or decreasing length of the interval. DP over trees usually require a post-order traversal of the tree. In terms of the recurring graph representation we often use for DP problems, the order of computation must be a topological ordering of the graph.

While this order of computation business may seem to be nothing but a

nuisance that bottom-up users have to deal with, it is related to one of the perks of bottom-up computation. If the order of computation is chosen in a clever way, we need not save every state during our computation. Consider e.g. the Change-making Problem again, which had the following recursion:

$$\text{change}(n) = \begin{cases} 0 & \text{if } n = 0 \\ \min(\text{change}(n-1), \text{change}(n-6), \text{change}(n-7)) & \text{if } n > 0 \end{cases}$$

It should be clear that using the order of computation $0, 1, 2, 3, \dots$, once we have computed e.g. $\text{change}(k)$, the subproblems $\text{change}(k-7)$, $\text{change}(k-8)$, ... etc. are never used again.

Thus, we only need to save the value of 7 subproblems at a time. This $\Theta(1)$ memory usage is pretty neat compared to the $\Theta(K)$ usage needed to compute $\text{change}(K)$ otherwise.

Competitive Tip

Generally, memory limits are very generous nowadays, somewhat diminishing the art of optimizing memory in DP solutions. It can still be a good exercise to think about improving the memory complexity of the solutions we will look at, for the few cases where these limits are still relevant.

11.3 Multidimensional DP

Now, we are going to look at a DP problem where our state consists of more than one variable. The example will demonstrate the importance of carefully choosing your DP parameters.

Ferry Loading

Swedish Olympiad in Informatics 2013, Online Qualifiers

A ferry is to be loaded with cars of different lengths, with a long line of cars currently queued up for a place. The ferry consists of four lanes, each of the same length. When the next car in the line enters the ferry, it picks one of the lanes and parks behind the last car in that line. There must be safety margin of 1 meter between any two parked cars.

Given the length of the ferry and the length of the cars in the queue, compute the maximal amount of cars that can park if they choose the lanes optimally.




Figure 11.3: An optimal placement on a ferry of length 5 meters, of the cars with lengths 2, 1, 2, 5, 1, 1, 2, 1, 1, 2 meters. Only the first 8 cars could fit on the ferry.

Input

The first line contains the number of cars $0 \leq N \leq 200$ and the length of the ferry $1 \leq L \leq 60$. The second line contains N integers, the length of the cars $1 \leq a_i \leq L$.

Output

Output a single integer – the maximal number of cars that can be loaded on the ferry.

The ferry problem looks like a classical DP problem. It consists of a large number of similar choices. Each car has 4 choices – one of the lanes. If a car of length m chooses a lane, the remaining length of the chosen lane is reduced by $m + 1$ (due to the safety margin). After the first c cars have parked on the ferry, the only thing that has changed are the lengths of the ferry. As a simplification, we increase the initial length of the ferry by 1, to accommodate an imaginary safety margin for the last car in a lane in case it is completely filled.

This suggests a DP solution with nL^4 states, each state representing the number of cars so far placed and the lengths of the four lanes:

Ferry Loading

```

1  int dp[201][62][62][62][62] = {-1};
2
3  int ferry(int c, vi used, const vi& A) {
4      if (c == sz(A)) return 0;
5      int& ans = dp[c][left[0]][left[1]][left[2]][left[3]];
6      if (ans != -1) return ans;
7      rep(i,0,4) {
8          if (used[i] + A[i] + 1 > L + 1) continue;
9          used[i] += A[i] + 1;
10         ans = max(ans, ferry(c + 1, left, A) + 1);

```

```
11     used[i] -= A[i] + 1;  
12 }  
13 return ans;  
14 }
```

Unfortunately, memoizing this procedure would not be sufficient. The size of our memoization array is $200 \cdot 60^4 \approx 2.6 \cdot 10^9$, which needs many gigabytes of memory.

The trick in improving this is basically the same as the fundamental principle of DP. In DP, we reduce a set of choices to a smaller set of information, which represent the effects of those choices. This removes information which turned out to be redundant. In our case, we do not care about what lanes cars chose, only their remaining lengths. Our suggested solution still has some lingering redundancy though.

In Figure 11.3 from the problem statement, we have an example assignment of the cars 2, 1, 2, 5, 1, 1, 2, 1. These must use a total of $3+2+3+6+2+2+3+2 = 23$ meters of space on the ferry. We define $U(c)$ to be the total usage (i.e., lengths plus the safety margin) of the first c cars. Note that $U(c)$ is a strictly increasing function in c , meaning it is in bijection to $[n]$. Let $u_1(c), u_2(c), u_3(c), u_4(c)$ be the usage of the four lanes individually in some given assignment. Then, we have that $U(c) = u_1(c) + u_2(c) + u_3(c) + u_4(c)$. The four terms on the right are four parameters in our memoization. The left term is not, but it has a bijection with c , which is a parameter in the memoization. Thus, we actually have a redundancy in our parameters. We can eliminate the parameter c , since it is uniquely defined given the values $u_1(c), u_2(c), u_3(c), u_4(c)$. This simplification leaves us with $60^4 \approx 13\,000\,000$ states, which is well within reason.

11.4 Subset DP

Another common theme of DP is subsets, where the state represents a subset of something. The subset DP is used in many different ways. Sometimes (as in Subsection 11.6.4), the problem itself is about sets and subsets. Another common usage is to reduce a solution which requires us to test permutations of something into instead constructing permutations iteratively, using DP to remember only what elements so far used in the permutation, and not the exact assignment.

Amusement Park

Swedish Olympiad in Informatics 2012, Online Qualifiers

Lisa has just arrived at an amusement park, and wants to visit each of the N attractions exactly once. For each attraction, there are two identical facilities at different locations in the park. Given the locations of all the facilities, determine which facility Lisa should choose for each attraction, in order to minimize the total distance she must walk. Originally, Lisa is at the entrance at coordinates $(0, 0)$. Lisa must return to the entrance once she has visited every attraction.

Input

The first line contains the integer $1 \leq N \leq 15$, the number of attractions Lisa wants to visit. Then, N lines follow. The i 'th of these lines contains four integers $-10^6 \leq x_1, y_1, x_2, y_2 \leq 10^6$. These are the coordinates (x_1, y_1) and (x_2, y_2) for the two facilities of the i 'th attraction.

Output

First, output the smallest distance Lisa must walk. Then, output N lines, one for each attraction. The i 'th line should contain two numbers a and f – the i 'th attraction Lisa visited (a number between 1 and N), and the facility she visited (1 or 2).

Consider a partial walk, where we have visited a set S of attractions and currently stand at coordinates (x, y) . Then, any choice up to this point is irrelevant for the remainder of the problem, which suggests that these parameter S, x, y is a good DP state. Note that (x, y) only have at most 31 possibilities – two for each attraction, plus the entrance at $(0, 0)$. Since we have at most 15 attractions, the set S of visited attractions has 2^{15} possibilities. This gives us $31 \cdot 2^{15} \approx 10^6$ states. Each state can be computed in $\Theta(N)$ time, by choosing what attraction to visit next. All in all, we get a complexity of $\Theta(N^2 2^N)$. When coding DP over subsets, we generally use bitsets to represent the subset, since these map very cleanly to integers (and therefore indices into a vector):

Amusement Park

```
1 double best(int at, int visited) {
2     // 2N is the number given to the entrance point
3     if (visited == (1<<N) - 1) return dist(at, 2*N);
4     double ans = inf;
5     rep(i,0,N) {
6         if (visited & (1<<N)) continue;
7         rep(j,0,2) {
8             // 2i + j is the number given to the j'th facility
9             // of the i'th attraction
10            int nat = 2 * i + j;
11            ans = min(ans, dist(at + nat) + best(nat, visited | (1<<i)));
12        }
13    }
14    return ans;
15 }
```

11.5 Digit DP

Digit DP are a class of problems where we count numbers with certain properties that contain a large number of digits, up to a certain limit. These properties are characterized by having the classical properties of DP problems, i.e. being easily computable if we would construct the numbers digit-by-digit by remembering very little information about what those numbers actually were.

Palindrome-Free Numbers

Baltic Olympiad in Informatics 2013 – Antti Laaksonen

A string is a palindrome if it remains the same when it is read backwards. A number is palindrome-free if it does not contain a palindrome with a length greater than 1 as a substring. For example, the number 16276 is palindrome-free whereas the number 17276 is not because it contains the palindrome 727. The number 10102 is not valid either, since it has 010 as a substring (even though 010 is not a number itself).

Your task is to calculate the total number of palindrome-free numbers in a given range.

Input

The input contains two numbers $0 \leq a \leq b \leq 10^{18}$.

Output

Your output should contain one integer: the total number of palindrome-free numbers in the range $a, a + 1, \dots, b - 1, b$ (including a and b).

First, a common simplification when solving counting problems on intervals. Instead of computing the answer for the range $a, a + 1, \dots, b - 1, b$, we will solve the problem for the intervals $[0, a)$ and $[0, b + 1)$. The answer is then the answer for the second interval with the answer for the first interval removed. Our lower limit will then be 0 rather than a , which simplifies the solution.

Next up, we need an essential observation to turn the problem into a standard application of digit DP. Palindromes as general objects are very unwieldy in our situation. Any kind of iterative construction of numbers would have to bother with digits far back in the number since any of them could be the edge of a palindrome. Fortunately, it turns out that any palindrome must contain a rather short palindromic subsequence, namely one of length 2 (for even-length palindromes), or length 3 (for odd-length palindromes). This means that when constructing the answer recursively, we only need to care about the last two digits. When adding a digit to a partially constructed number, it may not be equal to either of the last two digits.

Before arriving at the general solution, we will solve the problem when the upper limit was 999...999 – the sequence consisting of n nines. In this case, a simple recursive function will do the trick:

Palindrome-Free Numbers

```
1 ll sol(int at, int len, int b1, int b2) {
2     if (at == len) return 1; // we have successfully constructed a number
3     ll ans = 0;
4     rep(d,0,10) {
5         // this digit would create a palindrome
6         if (d == b2 || d == b1) continue;
7         // let -1 represent a leading 0, to avoid the palindrome check
8         bool leadingZero = b2 == -1 && d == 0;
9         ans += sol(at + 1, len, b2, leadingZero ? -1 : d);
10    }
11    return ans;
12 }
13
14 // we start with an empty number with leading zeroes
15 sol(0, n, true, -1, -1);
```

We fix the length of all numbers to have length n , by giving shorter numbers leading zeroes. Since leading zeroes in a number are not subject to the palindrome restriction, they must be treated differently. In our case, they are given the special digit -1 instead, resulting in 11 possible “digits”. Once this function is memoized, it will have $n \cdot 2 \cdot 11 \cdot 11$ different states, with each state using a loop iterating only 10 times. Thus, it uses on the order of $1000n$ operations. In our problem, the upper limit has at most 19 digits. Thus, the solution only requires about 20 000 operations.

Once a solution has been formulated for this simple upper limit, extending it to a general upper limit is quite natural. First, we will save the upper limit as a sequence of digits L . Then, we need to differentiate between two cases in our recursive function. The partially constructed number is either equal to the corresponding partial upper limit, or it is less than the limit. In the first case, we are still constrained by the upper limit – the next digit of our number can not exceed the next digit of the upper limit. In the other case, the the upper limit is no longer relevant. If a prefix of our number is strictly lower than the prefix of the upper limit, our number can never exceed the upper limit.

This gives us our final solution:

Palindrome-Free Numbers, General Case

```

1 vector<int> L;
2
3 ll sol(int at, int len, bool limEq, int b1, int b2) {
4     if (at == len) return 1;
5     ll ans = 0;
6     // we may not exceed the limit for this digit
7     // if equal to the prefix of the limit
8     rep(d, 0, (limEq ? L[at] + 1 : 10)) {
9         if (d == b2 || d == b1) continue;
10        // the next step will be equal to the prefix if it was now,
11        // and our added digit was exactly the limit
12        bool limEqNew = limEq && d == L[at];
13        bool leadingZero = b2 == -1 && d == 0;
14        ans += sol(at + 1, len, limEqNew, b2, leadingZero ? -1 : d);
15    }
16    return ans;
17 }
18
19 // initially, the number is equal to the
20 // prefix limit (both being the empty number)
21 sol(0, n, true, true, -1, -1);

```

11.6 Standard Problems

Many problems are variations of known DP problems, or have them as parts of their solutions. This section will walk you through a few of them.

Knapsack

The knapsack problem is one of the most common standard DP problem. The problem itself has countless variations. We will look at the “original” knapsack problem, with constraints making it suitable for a dynamic programming approach.

Knapsack

Given is a knapsack with an integer capacity C , and n different objects, each with an integer weight and value. Your task is to select a subset of the items with maximal value, such that the sum of their weights does not exceed the capacity of the knapsack.

Input

The integer C giving the capacity of the knapsack, and an integer n , giving the number of objects. This is followed by the n objects, given by their value v_i and weight w_i .

Output

Output the indices of the chosen items.

We are now going to attempt to formulate an $O(nC)$ solution. As is often the case when the solution is a subset of something in DP solutions, we solve the problem by looking at the subset as a sequence of choices – to either include an item in the answer or not. In this particular problem, our DP state is rather minimalistic. Indeed, after including a few items, we are left only with the remaining items and a smaller knapsack to solve the problem for.

Letting $K(c, i)$ be the maximum value using at most weight c and the i first items, we get the recursion

$$K(c, i) = \max \begin{cases} K(c, i - 1) \\ K(c - w_i, i - 1) + v_i & \text{if } w_i \leq c \end{cases}$$

Translating this recursion into a bottom-up solution gives a rather compact algorithm:

```

1: procedure KNAPSACK(capacity  $C$ , items  $n$ , values  $V$ , weights  $W$ )
2:    $best \leftarrow \text{new int}[n + 1][C + 1]$ 
3:   fill  $best$  with  $-\infty$ 
4:    $best[0][C] = 0$ 
5:   for  $i$  from 0 to  $n - 1$  do
6:     for  $j$  from 0 to  $C$  do
7:       if  $j \geq W[i]$  then
8:          $best[i + 1][j] \leftarrow \max(best[i][j], best[i][j - W[i]] + V[i])$ 
9:   return  $best$ 

```

However, this only helps us compute the answer. The problem asks us to explicitly construct the subset. This step, i.e., tracing what choices we made to arrive at an optimal solution is called *backtracking*.

For this particular problem, the backtracking is relatively simple. One usually proceeds by starting at the optimal state, and then consider all transitions that lead to this state. Among these, the “best” one is picked. In our case, the transitions correspond to either choosing the current item, or not choosing it.

Both lead to two other states which are simple to compute. In the first case, the state we arrived from must have the same value and capacity, while in the second case the value should differ by $V[i]$ and the weight by $W[i]$:

Make sure to study the implementation closely; this kind of reconstruction is a bit tricky to get right, but most reconstructions look something like it.

```

1: procedure KNAPSACKCONSTRUCT(capacity  $C$ , items  $n$ , values  $V$ , weights  $W$ )
2:    $best \leftarrow Knapsack(C, n, V, W)$ 
3:    $bestCap \leftarrow C$ 
4:   for  $i$  from  $C$  to 0 do
5:     if  $best[N][i] > best[N][bestCap]$  then
6:        $bestCap \leftarrow i$ 
7:   for  $i$  from  $N$  to 1 do
8:     if  $W[i] \leq bestCap$  then
9:        $newVal \leftarrow best[i - 1][bestCap - W[i]]$ 
10:      if  $newVal = best[i][bestCap] + V[i]$  then
11:         $ans.add(i)$ 
12:         $bestCap \leftarrow bestCap - W[i]$ 
13:   output  $ans$ 

```

Problem 11.1

Knapsack – knapsack

Walrus Weights – walrusweights

A common knapsack variation is a variation of *coin change* problem. In this problem, we are instead allowed to use each item an unlimited number of times, and ask ourselves how many different ways we can choose items in order to fill the knapsack. We formulate our DP solution in a similar way to ordinary knapsack. As subproblem, we count the number of ways to fill a knapsack of capacity c using items $i, i + 1, \dots$. In such a situation, we have two choices. We can either keep using item i once, leaving us with the same set of items but capacity $c - w_i$, or decide that we have finished using item i . In the latter case, our knapsack will have the same capacity, but we will instead use items $i + 1, i + 2, \dots$.

In this case, the recursion is slightly altered:

$$K(c, i) = K(c, i + 1) + K(c - w_i, i)$$

Longest Common Subsequence

Longest Common Subsequence

A sequence a_1, a_2, \dots, a_n has c_1, c_2, \dots, c_k as a *subsequence* if there exists indices $p_1 < p_2 < \dots < p_k$ such that $a_{p_i} = c_i$. For example, the sequence $\langle 1, 1, 5, 3, 3, 7, 5 \rangle$, has $\langle 1, 3, 3, 5 \rangle$ as one of its sub-sequences.

Given two sequences $A = \langle a_1, a_2, \dots, a_n \rangle$ and $B = \langle b_1, b_2, \dots, b_m \rangle$, find the longest sequence c_1, \dots, c_k that is a subsequence of both A and B .

When dealing with DP problems on pairs of sequences, a natural subproblem is to solve the problem for all *prefixes* of A and B . For the subsequence problem, some reasoning about what a common subsequence is leads to a recurrence expressed in this way. Basically, we can do a case analysis on the last letter of the strings A and B . If the last letter of A is not part of a longest increasing subsequence, we can simply ignore it, and solve the problem on the two strings where the last letter of A is removed. A similar case is applicable when the last letter of B is not part of a longest increasing subsequence. A single case remains – when both the last letter of A and the last letter of B are part of a longest increasing subsequence. In this case, we argue that these two letters must correspond to the same letter c_i in the common subsequence. In particular, they must correspond to the final character of the subsequence (by the definition of a subsequence). Thus, whenever the two final letters are equal, we may have the case that they are the last letter of the subsequence, and that the remainder of the subsequence is the longest common subsequence of A and B with the final letter removed.

This yields a simple recursive formulation, which takes $\Theta(|A||B|)$ to evaluate (since each state takes $\Theta(1)$ to evaluate).

$$\text{lcs}(A, B, n, m) = \max \begin{cases} 0 & \text{if } n = 0 \text{ or } m = 0 \\ \text{lcs}(A, B, n - 1, m) & \text{if } n > 0 \\ \text{lcs}(A, B, n, m - 1) & \text{if } m > 0 \\ \text{lcs}(A, B, n - 1, m - 1) + 1 & \text{if } a_n = b_m \end{cases}$$

TODO

Problem 11.2

Longest Common Subsequence – longcommonsubseq

Longest Increasing Subsequence

Problem 11.3

Longest Increasing Subsequence – longincsubseq

Set Cover

Set Cover

You are given a family of subsets S_1, S_2, \dots, S_k of some larger set S of size n . Find a minimum number of subsets $S_{a_1}, S_{a_2}, \dots, S_{a_l}$ such that

$$\bigcup_{i=1}^l S_{a_i} = S$$

i.e., *cover* the set S by taking the union of as few of the subsets S_i as possible.

For small k and large n , we can solve the problem in $\Theta(n2^k)$, by simply testing each of the 2^k covers. In the case where we have a small n but k can be large, this becomes intractable. Instead, let us apply the principle of dynamic programming. In a brute force approach, we would perform k choices. For each subset, we would try including it or excluding it. After deciding which of the first m subsets to include, what information is relevant? Well, if we consider what the goal of the problem is – covering S – it would make sense to record what elements have been included so far. This little trick leaves us with a DP of $\Theta(k2^n)$ states, one for each subset of S we might have reached, plus counting how many of the subsets we have tried to use so far. Computing a state takes $\Theta(n)$ time, by computing the union of the current cover with the set we might potentially add. The recursion thus looks like:

$$\text{cover}(C, k) = \begin{cases} 0 & \text{if } C = S \\ \min(\text{cover}(C, k+1), \text{cover}(C \cup S_k, k+1)) & \text{else} \end{cases}$$

This is a fairly standard DP solution. The interesting case occurs when n is small, but k is *really* large, say, $k = \Theta(2^n)$. In this case, our previous complexity $\Theta(nk2^n)$ turns into $\Theta(n4^n)$. Such a complexity is unacceptable for anything but very small n . To avoid this, we must rethink our DP a bit.

The second term of the recursive case of $\text{cover}(C, k)$, i.e. $\text{cover}(C \cup S_k, k+1)$, actually degenerates to $\text{cover}(C, k+1)$ if $S_k \subseteq C$. When k is large, this means many states are essentially useless. In fact, at most n of our k choices will

actually result in us adding something, since we can only add a new element at most n times.

We have been in a similar situation before, when solving the backtracking problem *Basin City Surveillance* in Section 9.3. We were plagued with having many choices at each state, where a large number of them would fail. Our solution was to limit our choices to a set where we *knew* an optimal solution would be found.

Applying the same change to our set cover solution, we should instead do DP over our current cover, and only try including sets which are not subsets of the current cover. So, does this help? How many subsets are there, for a given cover C , which are not its subsets? If the size of C is m , there are 2^m subsets of C , meaning $2^n - 2^m$ subsets can add a new element to our cover.

To find out how much time this needs, we will use two facts. First of all, there are $\binom{n}{m}$ subsets of size m of a size n set. Secondly, the sum $\sum_{m=0}^n \binom{n}{m} 2^m = 3^n$. If you are not familiar with this notation or this fact, you probably want to take a look at Section ?? on binomial coefficients.

So, summing over all possible extending subsets for each possible partial C , we get:

$$\sum_{m=0}^n \binom{n}{m} (2^n - 2^m) = 2^n \cdot 2^n - 3^n = 4^n - 3^n$$

Closer, but no cigar. Intuitively, we still have a large number of redundant choices. If our cover contains, say, $n - 1$ elements, there are 2^{n-1} sets which can extend our cover, but they all extend it in the same way. This sounds wasteful, and avoiding it probably the key to getting an asymptotic speedup.

It seems that we are missing some key function which, given a set A , can respond to the question: “is there some subset S_i , that could extend our cover with some subset $A \subseteq S$?”. If we had such a function, computing all possible extensions of a cover of size m would instead take time 2^{n-m} – the number of possible extensions to the cover. Last time we managed to extend a cover in time $2^n - 2^m$, but this is exponentially better!

In fact, if we do our summing this time, we get:

$$\begin{aligned}\sum_{m=0}^n \binom{n}{m} 2^{n-m} &= \sum_{m=0}^n \binom{n}{n-m} 2^{n-m} \\ &= \sum_{m=0}^n \binom{n}{m} 2^m \\ &= 3^n\end{aligned}$$

It turns out our exponential speedup in extending a cover translated into an exponential speedup of the entire DP.

We are not done yet – this entire algorithm depended on the assumption of our magical “can we extend a cover with a subset A ?” function. Sometimes, this function may be quick to compute. For example, if $S = \{1, 2, \dots, n\}$ and the family S_i consists of all sets whose sum is less than n , an extension is possible if and only if *its* sum is also less than n . In the general case, our S_i are not this nice. Naively, one might think that in the general case, an answer to this query would take $\Theta(nk)$ time to compute, by checking if A is a subset of each of our k sets. Yet again, the same clever trick comes to the rescue.

If we have a set S_i of size m available for use in our cover, just how many possible extensions could this subset provide? Well, S_i itself only have 2^m subsets. Thus, if we for each S_i mark for each of its subsets that this is a possible extension to a cover, this precomputation only takes 3^n time (by the same sum as above).

Since both steps are $O(3^n)$, this is also our final complexity.

Problem 11.4

Square Fields (Easy) – squarefieldseasy

Square Fields (Hard) – squarefieldshard

Chapter Notes

12 Divide and Conquer

A recursive algorithm solves a problem by reducing it to smaller subproblems, hoping that their solutions can be used to solve the larger problem. So far, the subproblems we have considered have been “almost the same” as the problem at hand. We have usually recursed on a series of choices, where each recursive step made one choice, as in the change-making problem. In particular, our subproblems often overlapped – solving two different subproblems required solving a common, third subproblem. In this chapter, we will take another approach altogether, by splitting our instance into large, disjoint (or almost disjoint parts) parts – dividing it – and combining their solutions – conquering it.

12.1 Inductive Constructions

Inductive construction problems comprise a large class of divide and conquer problems. The goal is often to construct something, such as a tiling of a grid, cycles in a graph and so on. In these cases, divide and conquer algorithms aim to reduce the construction of the whole object to instead constructing smaller parts which can be combined into the final answer. Such constructions are often by-products of mathematical induction proofs of the existence of such a construction. In the following example problems, it is not initially clear that the object we are asked to construct even exists.

Grid Tiling

In a square grid of side length 2^n , one unit square is blocked (represented by coloring it black). Your task is to cover the remaining $4^n - 1$ squares with *triominos*, *L*-shaped tiles consisting of three squares in the following fashion. The triominos can be rotated by any multiple of 90 deg (Figure 12.1).



Figure 12.1: The four rotations of a triomino.

The triominoes may not overlap each other, nor cover anything outside the grid. A valid tiling for $n = 2$ would be




Figure 12.2: A possible tiling for $n = 2$.

Input

The input consists of three integers $1 \leq n \leq 8$, $0 \leq x < 2^n$ and $0 \leq y < 2^n$. The black square has coordinates (x, y) .

Output

Output the positions and rotations of any valid tiling of the grid.

When tiling a $2^n \times 2^n$ grid, it is not immediately clear how the divide and conquer principle can be used. To be applicable, we must be able to reduce the problem into smaller instances of the same problem and combine them. The peculiar side length 2^n does hint about a possible solution. Aside from the property that $2^n \cdot 2^n - 1$ is evenly divisible by 3 (a necessary condition for a tiling to be possible), it also gives us a natural way of splitting an instance, namely into its 4 quadrants.




Figure 12.3: Splitting the $n = 3$ case into its four quadrants.

Each of these have the size $2^{n-1} \times 2^{n-1}$, which is also of the form we require of grids in the problem. The crux lies in that these four new grids do not comply with the input specification of the problem. While smaller and disjoint, three of them contain no black square, a requirement of the input. Indeed, a grid of this size without any black squares can not be tiled using triominos.

The solution lies in the trivial solution to the $n = 1$ case, where we can easily reduce the problem to four instances of the $n = 0$ case:



Figure 12.4: A solution to the $n = 1$ case.

In the solution, we use a single triomino which blocked a single square of each of the four quadrants. This gives us four trivial subproblems of size 1×1 , where each grid has one blocked square. We can actually place such a triomino in every grid by placing it in the center (the only place where a triomino may cover three quadrants at once).




Figure 12.5: Placing a triomino in the corners of the quadrants without a black square.

After this transformation, we can now apply the divide and conquer principle. We split the grid into its four quadrants, each of which now contain one black square. This allows us to recursively solve four new subproblems. At some point, this recursion will finally reach the base case of a 1×1 square, which must already be filled.

```

1: procedure TILE( $N, (B_x, B_y), (T_x, T_y)$ )
2:   if  $N = 0$  then
3:     return
4:    $mid \leftarrow 2^{N-1}$ 
5:    $blocked \leftarrow \{(0, 0), (mid - 1, 0), (mid - 1, mid - 1), (0, mid - 1)\}$ 
6:   if  $B_x \geq mid$  and  $B_y \geq mid$  then
7:      $blockedQuad \leftarrow TOP\_RIGHT$ 
8:   if  $B_x < mid$  and  $B_y \geq mid$  then
9:      $blockedQuad \leftarrow TOP\_LEFT$ 
10:  if  $B_x < mid$  and  $B_y < mid$  then
11:     $blockedQuad \leftarrow BOTTOM\_LEFT$ 
12:  if  $B_x \geq mid$  and  $B_y < mid$  then
13:     $blockedQuad \leftarrow BOTTOM\_RIGHT$ 
14:   $place(T_x + mid, T_y + mid, blockedQuad)$ 
15:   $tile(N - 1, blocked[0], T_x + mid, T_y + mid)$ 
16:   $tile(N - 1, blocked[1], T_x, T_y + mid)$ 
17:   $tile(N - 1, blocked[2], T_x, T_y)$ 
18:   $tile(N - 1, blocked[3], T_x + mid, T_y)$ 

```

The time complexity of the algorithm can be computed easily if we use the fact that each call to `tile` only takes $\Theta(1)$ time except for the four recursive calls. Furthermore, each call places exactly one tile on the board. Since there are $\frac{4^n - 1}{3}$ tiles to be placed, the time complexity must be $\Theta(4^n)$. This is asymptotically optimal, since this is also the size of the output.

Exercise 12.1. It is possible to tile such a grid with triominos colored red, blue and green such that no two triominos sharing an edge have the same color. Prove this fact, and give an algorithm to generate such a coloring.

Divisible Subset

Let $n = 2^k$. Given a set A of $2n - 1$ integers, find a subset S of size exactly n such that

$$\sum_{x \in S} x$$

is a multiple of n .

Input

The input contains an integer $1 \leq n \leq 2^{15}$ that is a power of two, followed by the $2n - 1$ elements of A .

Output

Output the n elements of S .

When given a problem, it is often a good idea to solve a few small cases by hand. This applies especially to this kind of construction problems, where constructions for small inputs often shows some pattern or insight into how to solve larger instances. The case $n = 1$ is not particularly meaningful, since it is trivially true (any integer is a multiple of 1). When $n = 2$, we get an insight which might not seem particularly interesting, but is key to the problem. We are given $2 \cdot 2 - 1 = 3$ numbers, and seek two numbers whose sum is even. Given three numbers, it must have either two numbers which both are even, or two odd numbers. Both of these cases yield a pair with an even sum.

It turns out that this construction generalizes to larger instances. Generally, it is easier to do the “divide” part of a divide and conquer solution first, but in this problem we will do it the other way around. The recursion will follow quite naturally after we attempt to find a way in combining solutions to the smaller instance to a larger one.

We will lay the the ground work for a reduction of the case $2n$ to n . First, assume that we could solve the problem for a given n . The larger instance then contains $2(2n - 1) = 4n - 1$ numbers, of which we seek $2n$ numbers whose sum is a multiple of $2n$. This situation is essentially the same as for the case $n = 2$, except everything is scaled up by n . Can we scale our solution up as well?

If we have three **sets of n numbers** whose respective sums are all **multiples of n** , we can find **two sets of n numbers** whose total sum is **divisible by $2n$** . This construction essentially use the same argument as for $n = 2$. If the three subsets have sums an, bn, cn and we wish to find two whose sum is a multiple of $2n$, this is the same as finding two numbers of a, b, c whose sum is a multiple of 2. This is possible, according to the case $n = 2$.

A beautiful generalization indeed, but we still have some remnants of wishful thinking we need to take care of. The construction assumes that, given $4n - 1$ numbers, we can find three sets of n numbers whose sum are divisible by n . We have now come to the recursive aspect of the problem. By assumption, we could solve the problem for n . This means we can pick any $2n - 1$ of our $4n - 1$ numbers to get our first subset. The subset uses up n of our $4n - 1$ numbers,

leaving us with only $3n - 1$ numbers. We keep going, and pick any $2n - 1$ of these numbers and recursively get a second subset. After this, $2n - 1$ numbers are left, exactly how many we need to construct our third subset.

The division of the problem was thus into *four* parts. Three subsets of n numbers, and one set of $n - 1$ which we throw away. Coming up with such a division essentially required us to solve the combination part first with the generalizing of the case $n = 2$:

```
1 void fillWith(hashset<int>& toFill, hashset<int>& from, int size) {
2     while (sz(from) < size) {
3         from.insert(toFill.begin());
4         toFill.erase(toFill.begin());
5     }
6 }
7
8 hashset<int> divisibleSubset(int k, hashset<int> A) {
9     if (k == 1) return nums;
10
11     hashset<int> part;
12     // Find three subsets of size k/2 with sums divisible by k/2
13     fillWith(part, A, k - 1);
14     hashset<int> pa = divisibleSubset(k / 2, A);
15     A.erase(all(pa));
16
17     fillWith(part, A, k - 1);
18     hashset<int> pb = divisibleSubset(k / 2, A);
19     A.erase(all(pb));
20
21     fillWith(part, A, k - 1);
22     hashset<int> pc = divisibleSubset(k / 2, A);
23     A.erase(all(pc));
24
25     // Choose two who sum to k/2
26     int as = accumulate(all(pa), 0);
27     int bs = accumulate(all(pa), 0);
28     int cs = accumulate(all(pa), 0);
29
30     hashset<int> ans;
31     if ((as + bs) % k == 0) {
32         ans.insert(all(pa));
33         ans.insert(all(pb));
34     } else if ((as + cs) % k == 0) {
35         ans.insert(all(pa));
36         ans.insert(all(pc));
37     } else {
38         ans.insert(all(pb));
39         ans.insert(all(pc));
40     }
```

```

41  return ans;
42  }

```

This complexity is somewhat more difficult to analyze. Now, each call to `divisibleSubset` takes linear time to n , and makes 3 recursive calls with $n/2$. Thus, the complexity obeys the recurrence $T(n) = 3T(n/2) + \Theta(n)$. By the master theorem, this has complexity $\Theta(n^{\log_2 3}) \approx \Theta(n^{1.6})$.

Exercise 12.2. What happens if we, when solving the problem for some k , construct $k - 1$ pairs of integers whose sum are even, throw away the remaining element, and scale the problem down by 2 instead? What is the complexity then?

Exercise 12.3. The problem can be solved using a similar divide and conquer algorithm for *any* k , not just those which are powers of two¹. In this case, those k which are prime numbers can be treated as base cases. How is this done for composite k ? What is the complexity?

Exercise 12.4. The knight piece in chess can move in 8 possible ways (moving 2 steps in any one direction, and 1 step in one of the two perpendicular directions). A closed tour exists for an 8×8 grid.




Figure 12.6: A closed tour on an 8×8 grid.

Give an algorithm to construct a tour for any $2^n \times 2^n$ grid with $n \geq 3$.

Exercise 12.5. An n -bit *Gray code* is a sequence of all 2^n bit strings of length n , such that two adjacent bit strings differ in only one position. The first and

¹This result is known as the Erdős–Ginzburg–Ziv theorem

last strings of the sequence are considered adjacent. Possible Gray codes for the first few n are

$$n = 1: 0 \ 1$$

$$n = 2: 00 \ 01 \ 11 \ 10$$

$$n = 3: 000 \ 010 \ 110 \ 100 \ 101 \ 111 \ 011 \ 001$$

Give an algorithm to construct an n -bit Gray code for any n .

Problem 12.1

Bell Ringing – bells

Hamiltonian Hypercube – hypercube

12.2 Merge Sort

Merge sort is a sorting algorithm which uses divide and conquer. It is rather straightforward, and works by recursively sorting smaller and smaller parts of the array. When sorting an array by *dividing it into parts* and combining their solution, there is an obvious candidate for how to perform this partitioning. Namely, splitting the array into two halves and sorting them. When splitting an array in half repeatedly, we will eventually reach a rather simple base case. An array containing a single element *is already sorted*, so it is trivially solved. If we do so recursively, we get the recursion tree in Figure 12.7. Coding this recursive split is easy.




Figure 12.7: The recursion tree given when performing a recursive split of the array [5, 1, 6, 3, 7, 2, 0, 4].

When we have sorted the two halves, we need to combine them to get a sorted version of the entire array. The procedure to do this is based on a simple insight. If an array A is partitioned into two smaller arrays P_1 and P_2 , the smallest value

of A must be either the smallest value of P_1 or the smallest value of P_2 . This insight gives rise to a simple iterative procedure, where we repeatedly compare the smallest values of P_1 and P_2 , extract the smaller one of them, and append it to our sorted array.

```
1 void sortVector(vector<int>& a, int l, int r){
2     int m = (l + r)/2;
3     sortVector(a, l, m); // sort left half
4     sortVector(a, m, r); // sort right half
5
6     //combine the answers
7     vector<int> answer;
8     int x = l;
9     int y = m;
10    while(x < m || y < r){
11        if(x == m || (y < r && a[y] < a[x])){
12            answer.push_back(a[y]);
13            y++;
14        } else {
15            answer.push_back(a[x]);
16            x++;
17        }
18    }
19    for (int i = l; i < r; ++i) {
20        a[i] = answer[i - l];
21    }
22 }
```

To compute the complexity, consider the recursion tree in Figure 12.7. We make one call with 8 elements, two calls with 4 elements, and so on. Further, the combining procedure takes $\Theta(l)$ time for a call with l elements. In the general case of $n = 2^k$ elements, this means merge sort takes time

$$\sum_{i=0}^k 2^i \cdot \Theta(2^{k-i}) = \Theta(k2^k)$$

Since $k = \log_2 n$, this means the complexity is $\Theta(n \log n)$.

Exercise 12.6. Our complexity analysis assumed that the length of the array is a power of 2. The complexity is the same in the general case. Prove this fact.

Exercise 12.7. Given an array A of size n , we call the pair $i < j$ an *inversion* of A if $A[i] > A[j]$. Adapt the merge sort algorithm to count the number of inversions of an array in $\Theta(n \log n)$.

12.3 Binary Search

The *binary search* is a common component of other solution. Given a number L and a **non-decreasing** function $f : \mathbb{R} \rightarrow \mathbb{R}$, we wish to find the greatest x such that $f(x) \leq L$. Additionally, we need two numbers lo and hi , such that $f(lo) \leq L < f(hi)$.

It is not difficult to see how we solve this problem. Consider the number $mid = \frac{lo+hi}{2}$. If $f(mid) \leq L$, then we know that the answer must lie somewhere in the interval $[mid, hi)$. On the other hand, $L < f(mid)$ gives us a better upper bound on the answer, which must be contained in $[lo, mid)$. Computing $f(mid)$ allowed us to halve the interval in which the answer can be.




Figure 12.8: Two iterations of binary search.

We can repeat this step until we get close to x .

```

1  const double precision = 1e-7;
2
3  double binarySearch(double lo, double hi, double lim) {
4      while (hi - lo > precision) {
5          double mid = (lo + hi) / 2;
6          if (lim < f(mid)) hi = mid;
7          else lo = mid;
8      }
9      return lo;
10 }
```

Notice that we are actually computing an approximation of x within some given precision (10^{-7} in the example implementation), but this is often useful enough.

Competitive Tip

Remember that the double-precision floating point type only have a precision of about 10^{15} . If the limits in your binary search are on the order of 10^x , this means that using a binary search precision of something smaller than 10^{x-15} may cause an infinite loop. This happens because the difference between lo and the next possible double is actually larger than your precision.

As an example, the following parameters cause our binary search with precision 10^{-7} to fail.

```

1 double f(double x) {
2     return 0;
3 }
4
5 double lo = 1e12;
6 double hi = nextafter(lo, 1e100);
7 binarySearch(lo, hi, 0);

```

An alternative when dealing with limit precision is to perform binary search a fixed number of iterations:

```

1 double binarySearch(double lo, double hi, double lim) {
2     rep(i,0,60) {
3         double mid = (lo + hi) / 2;
4         if (lim < f(mid)) hi = mid;
5         else lo = mid;
6     }
7     return lo;
8 }

```

The complexity of binary search depends on how good of an approximation we want. Originally, the interval we are searching in has length $hi - lo$. After halving the interval c times, it has size $\frac{hi-lo}{2^c}$. If we binary search until our interval has some size p , this means we must choose c such that

$$\begin{aligned}\frac{hi - lo}{2^c} &\leq p \\ \frac{hi - lo}{p 2^c} &\leq 1 \\ \frac{hi - lo}{p} &\leq 2^c \\ \log_2 \frac{hi - lo}{p} &\leq c\end{aligned}$$

For example, if we have an interval of size 10^9 which we wish to binary search down to 10^{-7} , this would require $\log_2 10^{16} = 54$ iterations of binary search.

Now, let us study some applications of binary search.

Optimization Problems

Cutting Hot Dogs

The finals of the Swedish Olympiad in Informatics is soon to be arranged. During the competition the participants must have a steady supply of hot dogs to get enough energy to solve the problems. Being foodies, the organizing committee couldn't just *buy* ready-to-cook hot dogs – they have to make them on their own!

The organizing committee has prepared N long rods of hot dog. These rods have length a_1, a_2, \dots, a_N centimeters. It is now time to cut up these rods into actual hot dogs. A rod can be cut into any number k of hot dogs using $k - 1$ cuts. Note that parts of one rod cannot be combined with parts of another rod, due to the difference in taste. It is allowed to leave leftovers from a rod that will not become a hot dog.

In total, the contest has M participants. Each contestant should receive a single hot dog, and all of their hot dogs should be of the same length. What is the maximal hot dog length L the committee can use to cut M hot dogs of length L ?

Input

The first line contains two integers $1 \leq N \leq 10\,000$ and $1 \leq M \leq 10^9$.

The next line contains N real numbers $0 < a_1, a_2, \dots, \leq 10^6$, the length of the hot dog rods in centimeters.

Output

Output a single real number – the maximal hot dog length possible in centimeters. Any answer with a relative or absolute error of 10^{-6} is acceptable.

Solution. This problem is not only an optimization problem, but a *monotone* one. The monotonicity lies in that while we only ask for a certain maximal hot dog length, all lengths below it would also work (in the sense of being able to have M hot dogs cut of this length), while all lengths above the maximal length produce less than M hot dogs. Monotone optimization problems makes it possible to remove the optimization aspect by inverting the problem. Instead of asking ourselves what the maximum length is, we can instead ask how many

hot dogs $f(x)$ can be constructed from a given length x . After this inversion, the problem is now on the form which binary search solves: we wish to find the greatest x such that $f(x) = M$ (replacing \leq with $=$ is equivalent in the cases where we know that $f(x)$ assume the value we are looking for). We know that this length is at most $\max_i a_i \leq 10^6$, which gives us the interval $(0, 10^6]$ to search in.

What remains is to actually compute the function $f(x)$. In our case, this can be done by considering just a single rod. If we want to construct hot dogs of length x , we can get at most $\lfloor \frac{a_i}{x} \rfloor$ hot dogs from a rod of length a_i . Summing this for every rod gives us our solution.

```
1: procedure COUNTRODS(lengths  $A$ , minimum length  $M$ )
2:    $dogs \leftarrow 0$ 
3:   for each  $l \in A$  do
4:      $dogs \leftarrow dogs + \lfloor l/M \rfloor$ 
5:   return  $dogs$ 
6: procedure HOTDOGS(lengths  $A$ , participants  $M$ )
7:    $L \leftarrow 0, H \leftarrow 10^6$ 
8:   while  $H - L \leftarrow 10^{-7}$  do
9:      $mid \leftarrow (L + H)/2$ 
10:    if COUNTRODS( $A, mid$ )  $< M$  then
11:       $H \leftarrow mid$ 
12:    else
13:       $L \leftarrow mid$ 
14:   output  $L$ 
```

□

The key to our problem was that the number of hot dogs constructible with a length x was monotonically decreasing with x . It allowed us to perform binary search on the answer, a powerful technique which is a component of many optimization problems. In general, it is often easier to determine *if* an answer is acceptable, rather than computing a maximal answer.

Searching in a Sorted Array

The classical application of binary search is to find the position of an element x in a sorted array A of length n . Applying binary search to this is straightforward.

At first, we know nothing about location of the element – its position could be anyone of $[0, n)$. So, we consider the middle element, $mid = \lfloor \frac{n}{2} \rfloor$, and compare $A[mid]$ to x . Since A is sorted, this leaves us with three cases:

- $A[mid] = x$ – and we are done
- $x < A[mid]$ – since the array is sorted, any occurrence of x must be to the left of mid
- $A[mid] < x$ – by the same reasoning, x can only lie to the right of mid .

The last two cases both halve the size of the sub-array which x could be inside. Thus, after doing this halving $\log_2 n$ times, we have either found x or can conclude that it is not present in the array.

```
1: procedure SEARCH(array  $A$ , target  $x$ )
2:    $L \leftarrow 0, H \leftarrow |A|$ 
3:   while  $H - L > 0$  do
4:      $mid \leftarrow \lfloor (L + H)/2 \rfloor$ 
5:     if  $A[mid] = x$  then
6:       return  $mid$ 
7:     else if  $x < A[mid]$  then
8:        $H = mid$ 
9:     else
10:       $L = mid + 1$ 
11:   return  $-1$ 
```

Competitive Tip

When binary searching over discrete domains, care must be taken. Many bugs have been caused by improper binary searches².

The most common class of bugs is related to the endpoints of your interval (i.e. whether they are inclusive or exclusive). Be explicit regarding this, and take care that each part of your binary search (termination condition, midpoint selection, endpoint updates) use the same interval endpoints.

Exercise 12.8. — Kattis Problems *Ballot Boxes* – ballotboxes

²In fact, for many years the binary search in the standard Java run-time had a bug: http://bugs.java.com/bugdatabase/view_bug.do?bug_id=6412541

Generalized Binary Search

Binary search can also be used to find *all* points where a monotone function changes value (or equivalently, all the intervals on which a monotone function is constant). Often, this is used in problems on large sequences (often with $n = 100\,000$ elements), which can be solved by iterating through all contiguous sub-sequences in $\Theta(n^2)$ time.

Or Max

Petrozavodsk Winter Training Camp 2015

Given is an array A of integers. Let

$$B(i, k) = A[i] \mid A[i + 1] \mid \dots \mid A[i + k - 1]$$

i.e. the bitwise or of the k consecutive numbers starting with the i 'th,

$$M(i, k) = \max\{A[i], A[i + 1], \dots, A[i + k - 1]\}$$

i.e. the maximum of the k consecutive numbers starting with the i 'th, and

$$S(i, k) = B(i, k) + M(i, k)$$

For each $1 \leq k \leq n$, find the maximum of $S(i, k)$.

Input

The first line contains the length $1 \leq n \leq 10^5$ of A .

The next and last line contains the n values of A ($0 \leq A[i] < 2^{16}$), separated by spaces.

Output

Output n integers, the maximum values of $S(i, k)$ for $k = 1, 2, \dots, n$.

As an example, consider the array in Figure 12.9. The best answer for $k = 1$ would be $S(0, 1)$, with both maximal element and bitwise or 5, totaling 10. For $k = 2$, we have $S(6, 2) = 7 + 4 = 11$.

This problem can easily be solved in $\Theta(n^2)$, by computing every $S(i, k)$ iteratively. We can compute all the $B(i, k)$ and $M(i, k)$ using the recursions

$$B(i, k) := \begin{cases} 0 & \text{if } k = 0 \\ B(i, k - 1) \mid A[i + k - 1] & \text{if } k > 0 \end{cases}$$

$i = 0$	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$	$i = 6$	$i = 7$	$i = 8$	$i = 9$
5	1	4	2	2	0	4	3	1	2

101	001	100	010	010	000	100	011	001	010
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

Figure 12.9: Example array, with the numbers additionally written in binary.

$$M(i, k) = \begin{cases} 0 & \text{if } k = 0 \\ \max\{M(i, k - 1), A[i + k - 1]\} & \text{if } k > 0 \end{cases}$$

by looping over k , once we fix an i . With $n = 100\,000$, this approach is too slow.

The difficulty of the problem lies in $S(i, k)$ consisting of two basically unrelated parts – the maximal element and the bitwise or of a segment. When maximizing sums of unrelated quantities that put constraints on each other, brute force often seems like a good idea. This is basically what we did in the Buying Books problem (Section 9.4), where we minimized the sum of two parts (postage and book costs) which constrained each other (buying a book forced us to pay postage to its store) by brute forcing over one of the parts (the set of stores to buy from). Since the bitwise or is much more complicated than the maximal element – it is decided by an entire interval rather than a single element – we are probably better off doing brute force over the maximal element. Our brute force will consist of fixing which element is our maximal element, by assuming that $A[m]$ is the maximal element.

With this simplification in hand, only the bitwise or remains. We could now solve the problem by looping over all the left endpoints of the interval and all the right endpoints of the interval. At a first glance, this seems to actually worsen the complexity. Indeed, this takes quadratic time for each m (on average), resulting in a cubic complexity.

This is where we use our new technique. It turns out that, once we fix m , there are only a few possible values for the bitwise or of the intervals containing the m 'th element. Any such interval $A[l], A[l + 1], \dots, A[m - 1], A[m], A[m + 1], \dots, A[r - 1], A[r]$ can be split into two parts: one to the left, $A[l], A[l + 1], \dots, A[i - 1], A[i]$, and one to the right, $A[i], A[i + 1], \dots, A[r - 1], A[r]$. The bitwise or of either of these two parts is actually a monotone function (in their length), and can only assume at most 16 different values!

Studying Figure 12.10 gives a hint about why. The first row shows the binary

$i = 0$	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$	$i = 6$	$i = 7$	$i = 8$	$i = 9$
101	001	100	010	010	000	100	011	001	010
111	111	110	110	110	100	100	111	111	111
7	7	6	6	6	4	4	7	7	7

Figure 12.10: The bitwise or of the left and right parts, with an endpoint in $m = 6$

values of the array, with $m = 6$ (our presumed maximal element) marked. The second row shows the binary values of the bitwise or of the interval $[i, m]$ or $[m, i]$ (depending on whether m is the right or left endpoint). The third line shows the decimal values of the second row.

For example, when extending the interval $[2, 6]$ (with bitwise or 110) to the left, the new bitwise or will be 110|001. This is the only way the bitwise or can change – when the new value includes bits which so far have not been set. Obviously, this can only happen at most 16 times, since the values in A are bounded by 2^{16} .

For a given m , this gives us a partition of all the elements, by the bitwise or of the interval $[m, i]$. In Figure 12.10, the left elements will be partitioned into $[0, 1]$, $[2, 4]$, $[5, 6]$. The right elements will be partitioned into $[6, 6]$, $[7, 9]$. These partitions are everything we need to compute the final.

For example, if we pick the left endpoint from the part $[2, 4]$ and the right endpoint from the part $[7, 9]$, we would get a bitwise or that is $6 \mid 7 = 7$, of a length between 4 and 8, together with the 4 as the presumed maximal element. For each maximal element, we get at most $16 \cdot 16$ such choices, totaling less than $256N$ such choices. From these, we can compute the final answer using a simple sweep line algorithm.

12.4 Karatsuba's algorithm

Karatsuba's algorithm was developed by Russian mathematician Anatoly Karatsuba and published in the early 1960's. It is one of the earliest examples of a divide and conquer algorithm, and is used to quickly multiply large numbers. While multiplying small numbers (i.e. those that fit in the integer types of your favorite programming language) is considered to be a $\Theta(1)$ operation, this is

not the case for arbitrarily large integers. We will look at Karatsuba as a way of multiplying polynomials, but this can easily be extended to multiplying integers.

Polynomial Multiplication

Given two n -degree polynomials (where n can be large) $p(x) = \sum_{i=0}^n x^i a_i$ and $q(x) = \sum_{i=0}^n x^i b_i$ compute their product

$$(pq)(x) = \sum_{i=0}^{2n} x^i \left(\sum_{j=0}^i a_j b_{i-j} \right)$$

The naive multiplication algorithm evaluates this using $\Theta(n^2)$ multiplications (e.g. by two nested loops).

It turns out we can do this faster, using a recursive transformation. If we split the numbers p and q into their upper and lower $k = \frac{n}{2}$ coefficients (if n is odd, we pad the polynomials with a leading zero), so that $p(x) = p_l(x)x^k + p_r(x)$ and $q(x) = q_l(x)x^k + q_r(x)$, their product is equal to

$$\begin{aligned} (pq)(x) &= (p_l(x)x^k + p_r(x))(q_l(x)x^k + q_r(x)) \\ &= p_l(x)q_l(x)x^n + (p_l(x)q_r(x) + p_r(x)q_l(x))x^k + p_r(x)q_r(x) \end{aligned}$$

This formula requires multiplying 4 pairs of k -degree polynomials instead, which we can recursively compute, resulting in the time complexity recurrence $T(n) = 4T(\frac{n}{2}) + \Theta(n)$. Using the master theorem gives us the solution $T(n) = \Theta(n^2)$, which is no faster than the naive multiplication.

However, we can compute $p_l(x)q_r(x) + p_r(x)q_l(x)$ using only one multiplication instead of two. Both of these terms are part of the expansion of pq , which is only one multiplication. That particular multiplication is on n -degree polynomials, but it is not difficult to see how we can reduce it to a single k -degree multiplication. We simply throw away the multiplicative factors that makes $p_l x^k + p_r$ and $q_l x^k + q_r$ an n -degree polynomials:

$$(p_l(x) + p_r(x))(q_l(x) + q_r(x)) = p_l(x)q_l(x) + p_l(x)q_r(x) + p_r(x)q_l(x) + p_r(x)q_r(x) \blacksquare$$

so that

$$p_l(x)q_r(x) + p_r(x)q_l(x) = (p_l(x) + p_r(x))(q_l(x) + q_r(x)) - p_l(x)q_l(x) - p_r(x)q_r(x) \blacksquare$$

This means we only need to compute three k -degree multiplications: $(p_l(x) + p_r(x))(q_l(x) + q_r(x))$, $p_l(x)q_l(x)$, $p_r(x)q_r(x)$. Our time complexity recurrence

is then reduced to $T(n) = 3T(\frac{n}{2}) + O(n)$, which by the master theorem is $\Theta(n^{\log_2 3}) \approx \Theta(n^{1.585})$.

Exercise 12.9. *Polynomial Multiplication 2 – polymul2*

12.5 Chapter Notes

13 Data Structures

This chapter extends Chapter 6 further by showing the two most common advanced data structures that make an appearance in algorithmic problem solving. We do not take the approach of simply presenting the data structure and the problem that it solves. Instead, we take the same approach as we would for any example problem or algorithm in this book, by gradually improving an initial naive solution using additional insights. This is particularly valuable in this case, since certain problems may require variations of these structures in which only certain of the optimizations we show are applicable. Some of the techniques which we show during this journey can also occasionally be applied to other problems, so make sure to digest not only the final results, but every individual step on the journey to it.

13.1 Disjoint Sets

In the *Connectivity Problem*, we want to determine whether two vertices in the graph are in the same connected component. This problem can be solved using a Depth-First Search (Section 14.2). For now, we will instead focus on an extension of this problem called the *Dynamic Connectivity Problem, Additions only*¹, where we may also alter the graph by adding edges.

Dynamic Connectivity, Additions Only

Given a graph G which initially consists of V disconnected vertices, you will receive Q queries of two types:

1. take two vertices v and w , and add an edge between them
2. determine whether vertices v and w are currently in the same component

This problem can easily be solved in $O(Q^2)$, by after each query performing a DFS to partition the graph into its connected components. Since the graph has at most Q edges, a DFS takes $O(Q)$ time.

¹The Dynamic Connectivity problem where edges may also be removed can be solved by a data structure called a *Link-Cut tree*, which is not discussed in this book.

We can improve this by explicitly computing the connected components at each step. Originally, we have V components of a single vertex each. When adding an edge between two vertices that are in different components, these components will now merge. Note that it is irrelevant which vertices we are adding an edge between – it is only important what components the vertices belong to. When merging two connected components to one, we iterate through all the vertices of one component and add them to the other. Since we make at most V joins of different components, and the components involved contain at most V vertices, these queries only take $O(V^2)$ time. Determining whether two vertices are in the same component can then be done in $O(V)$, leaving us with a total complexity of $O(V(V + Q))$. We can speed this up further by using a simple look-up table $comp[v]$ for the vertices, which stores some identifier for the component a vertex is in. We can then respond to a connectivity query in $\Theta(1)$ by comparing the value of $comp[v]$ and $comp[w]$. Our complexity is then $O(V^2 + Q)$ instead.

Finally, we will use a common trick to improve the complexity to $O(V \log V + Q)$ instead. Whenever we merge two components of size a and b , we can do this in $O(\min(a, b))$ instead of $O(a + b)$ by merging the smaller component into the larger component. Any individual vertex can be part of the smaller component at most $O(\log V)$ times. Since the total size of a component is always at least twice the size of the smaller component, this means that if a vertex is merged as part of the smaller component k times the new component size must be at least 2^k . This cannot exceed V , meaning $k \leq \log_2 V$. If we sum this up for every vertex, we arrive at the $O(V \log V + Q)$ complexity.

Disjoint Set

```
1  struct DisjointSets {
2
3      vector<vector<int>>> components;
4      vector<int> comp;
5      DisjointSets(int elements) : components(elements), comp(elements) {
6          iota(all(comp), 0);
7          for (int i = 0; i < elements; ++i) components[i].push_back(i);
8      }
9
10     void unionSets(int a, int b) {
11         a = comp[a]; b = comp[b];
12         if (a == b) return;
13         if (components[a].size() < components[b].size()) swap(a, b);
14         for (int it : components[b]) {
```

```
15     comp[it] = a;
16     components[a].push_back(it);
17 }
18 }
19
20 };
```

A somewhat faster² version of this structure instead performs the merges of components lazily. When merging two components, we do not update $comp[v]$ for every vertex in the smaller component. Instead, if a and b are the representative vertices of the smaller and the larger component, we only merge a by setting $comp[a] = b$. However, we still need to perform the merges. Whenever we try to find the component a vertex lies in, we perform all the merges we have stored so far.

Improved Disjoint Set

```
1  struct DisjointSets {
2
3      vector<int> comp;
4      DisjointSets(int elements) : comp(elements, -1) {}
5
6      void unionSets(int a, int b) {
7          a = repr(a); b = repr(b);
8          if (a == b) return;
9          if (-comp[a] < -comp[b]) swap(a, b);
10         comp[a] += comp[b];
11         comp[b] = a;
12     }
13
14     bool is_repr(int x) { return comp[x] < 0; }
15
16     int repr(int x) {
17         if (is_repr(x)) return x;
18         while (!is_repr(comp[x])) {
19             comp[x] = comp[comp[x]];
20         }
21         return comp[x];
22     }
23
24 };
```

However, it turns out we can sometimes perform many merges at once. Consider the case where we have k merges lazily stored for some vertex v . Then, we can perform *all* the merges of v , $comp[v]$, $comp[comp[v]]$, \dots at the same time since they all have the same representative: the representative of v .

²With regards to actual time, not asymptotically.

Performing Many Merges

```
1 int repr(int x) {  
2     if (comp[x] < 0) return x;  
3     int par = comp[x];  
4     comp[x] = repr(par);  
5     return comp[x];  
6 }
```

Intuitively, this should be faster. After all, we would be performing performing at least $k + (k - 1) + (k - 2) + \dots = O(k^2)$ merges in $O(k)$ time, only for those vertices. If there were more vertices part of the components merged, this number grows even more.

It turns out that this change improves the complexity asymptotically.

13.2 Range Queries

Problems often ask us to compute some expression based on some interval of an array. The expressions are usually easy to compute in $\Theta(len)$ where len is the length of the interval. We will now study some techniques that trade much faster query responses for a bit of memory and precomputation time.

Prefix Precomputation**Interval Sum**

Given a sequence of integers a_0, a_1, \dots, a_{N-1} , you will be given Q queries of the form $[L, R)$. For each query, compute $S(L, R) = a_L + a_{L+1} + \dots + a_{R-1}$.

Computing the sums naively would require $\Theta(N)$ worst-case time per query if the intervals are large, for a total complexity of $\Theta(NQ)$. If $Q = \Omega(N)$ we can improve this to $\Theta(N^2 + Q)$ by *precomputing* all the answers. To do this in quadratic time, we use the recurrence

$$S(L, R) = \begin{cases} 0 & \text{if } L = R \\ S(L, R - 1) + a_{R-1} & \text{otherwise} \end{cases}$$

Using this recurrence we can compute the sequence $S(L, L), S(L, L+1), S(L, L+2), \dots, S(L, N)$ in average $\Theta(N)$ time for every L . This gives us the $\Theta(N^2 + Q)$ complexity.

If the function we are computing has an inverse, we can speed this precomputation up a bit. Assume that we have computed the values $P(R) = a_0 + a_1 + \dots + a_{R-1}$,

i.e. the prefix sums of a_i . Since this function is invertible (with inverse $-P(R)$), we can compute $S(L, R) = P(R) - P(L)$. Basically, the interval $[L, R)$ consists of the prefix $[0, R)$ with the prefix $[0, L)$ removed. As addition is invertible, we could simply remove the latter prefix $P(L)$ from the prefix $P(R)$ using subtraction. Indeed, expanding this expression shows us that

$$\begin{aligned} P(R) - P(L) &= (a_0 + a_1 + \cdots + a_{R-1}) - (a_0 + a_1 + \cdots + a_{L-1}) \\ &= a_L + \cdots + a_{R-1} = S(L, R) \end{aligned}$$

This leads to the following algorithm:

```

1: procedure PREFIXES(sequence  $A$ )
2:    $P \leftarrow \text{new int}[|A| + 1]$ 
3:   for  $i = 0$  to  $|A| - 1$  do
4:      $P[i + 1] \leftarrow P[i] + A[i]$ 
5:   return  $P$ 

```

for the precomputation, and then

```

1: procedure INTERVALSUM(interval  $[L, R)$ , prefix table  $P$ )
2:   return  $P[R] - P[L]$ 

```

to query for values.

This same technique works for any invertible operation.

Exercise 13.1. The above technique does not work straight-off for non-commutative operations. How can it be adapted to this case?

Sparse Tables

The case where a function does not have an inverse is a bit more difficult.

Interval Minimum

Given a sequence of integers a_0, a_1, \dots, a_{N-1} , you will be given Q queries of the form $[L, R)$. For each query, compute the value

$$M(L, R) = \min(a_L, a_{L+1}, \dots, a_{R-1})$$

Generally, you cannot compute the minimum of an interval based only on a constant number of prefix minimums of a sequence. We need to modify our approach. If we consider the naive approach, where we simply answer the queries by computing it explicitly, by looping over all the $R - L$ numbers in the interval, this is $\Theta(len)$. A simple idea will improve the time used to answer queries by a factor 2 compared to this. If we precompute the minimum of every *pair* of adjacent elements, we cut down the number of elements we need to check in half. We can take it one step further, by using this information to precompute the minimum of all subarrays of four elements, by taking the minimum of two pairs. By repeating this procedure for very power of two, we will end up with a table $m[l][i]$ containing the minimum of the interval $[l, l + 2^i)$, computable in $\Theta(N \log N)$.

Sparse Table

```

1 vector<vi> ST(const vi& A) {
2     vector<vi> ST(__builtin_popcount(sz(A)), vi(sz(A)));
3     ST[0] = A;
4     rep(len, 1, ST.size()) {
5         rep(i, 0, n - (1 << len) + 1) {
6             ST[len][i] = max(ST[len - 1][i], ST[len - 1][i + 1 << (len - 1)]);
7         }
8     }
9     return ST;
10 }
```

Given this, we can compute the minimum of an entire interval in logarithmic time. Consider the binary expansion of the length $len = 2^{k_1} + 2^{k_2} + \dots + 2^{k_l}$. This consists of at most $\log_2 len$ terms. However, this means that the intervals

$$\begin{aligned}
 &[L, L + 2^{k_1}) \\
 &[L + 2^{k_1}, L + 2^{k_1} + 2^{k_2}) \\
 &\dots \\
 &[L + 2^{k_1} + \dots + 2^{k_{l-1}}, L + len)
 \end{aligned}$$

together cover $[L, L + len)$. Thus we can compute the minimum of $[L, L + len)$ as the minimum of $\log_2 len$ intervals.

Sparse Table Querying

```

1  int rangeMinimum(const vector<vi>& table, int L, int R) {
2      int len = R - L;
3      int ans = std::numeric_limits<int>::max();
4      for (int i = sz(table) - 1; i >= 0; --i) {
5          if (len & (1 << i)) {
6              ans = min(ans, table[i][L]);
7              L += 1 << i;
8          }
9      }
10     return ans;
11 }
```

This is $\Theta((N + Q) \log N)$ time, since the preprocessing uses $\Theta(N \log N)$ time and each query requires $\Theta(\log Q)$ time. This structure is called a *Sparse Table*, or sometimes just the *Range Minimum Query* data structure.

We can improve the query time to $\Theta(1)$ by using that the min operation is idempotent, meaning that $\min(a, a) = a$. Whenever this is the case (and the operation at hand is commutative), we can use just two intervals to cover the entire interval. If 2^k is the largest power of two that is at most $R - L$, then

$$[L, L + 2^k)$$

$$[R - 2^k, R)$$

covers the entire interval.

```

1  int rangeMinimum(const vector<vi>& table, int L, int R) {
2      int maxLen = 31 - __builtin_clz(R - L);
3      return min(table[maxLen][L], table[maxLen][R - (1 << maxLen)]);
4  }
```

While most functions either have inverses (so that we can use the prefix precomputation) or has idempotent (so that we can use the $\Theta(1)$ sparse table), some functions do not. In such cases (for example matrix multiplication), we must use the logarithmic querying of the sparse table.

Segment Trees

The most interesting range queries occur on *dynamic* sequences, where values can change.

Dynamic Interval Sum

Given a sequence of integers a_0, a_1, \dots, a_{N-1} , you will be given Q queries. The queries are of two types:

1. Given an interval $[L, R)$, compute $S(L, R) = a_L + a_{L+1} + \dots + a_{R-1}$.
2. Given an index i and an integer v , set $a_i := v$.

To solve the dynamic interval problem, we will use a similar approach as the general sparse table. Using a sparse table as-is for the dynamic version, we would need to update $\Theta(N)$ intervals, meaning the complexity would be $\Theta(\log N)$ for interval queries and $\Theta(N)$ for updates. It turns out the sparse table as we formulated it contains an unnecessary redundancy.

If we accept using $2 \log N$ intervals to cover each query instead of $\log N$, we can reduce memory usage (and precomputation time!) to $\Theta(N)$ instead of $\Theta(\log N)$. We will use the same decomposition as in merge sort (Section 12.2). In Figure 13.1, you can see this decomposition, with an example of how a certain interval can be covered. In this context, the decomposition is called a *segment tree*.




Figure 13.1: The $2N - 1$ intervals to precompute.

Usually, this construction is represented as a flat, 1-indexed array of length $2^{\lceil \log_2 N \rceil}$. The extraneous are set to some sentinel value that does not affect queries (i.e. 0 in the case of sum queries). From this point, we assume N to be a power of two, with the array padded by these sentinel values.

- 1: **procedure** MAKE_TREE(sequence A)
- 2: $tree \leftarrow \text{new int}[2|N|]$


```

3:   for  $i = |N|$  to  $2|N| - 1$  do
4:      $tree[i] \leftarrow A[i - |N|]$ 
5:   for  $i = |N| - 1$  to  $1$  do
6:      $tree[i] \leftarrow tree[2 \cdot i] + tree[2 \cdot i + 1]$ 
7:   return  $P$ 

```

In the construction, we label each interval $1, 2, 3, \dots$ in order, meaning the entire interval will have index 1, the two halves indices 2, 3 and so on. This means that the two halves of the interval numbered i will have indices $2i$ and $2i + 1$, which explains the precomputation loop.

We can compute the sum of each of these intervals in $\Theta(1)$, assuming the sum of all the smaller intervals have already been computed, since each interval is composed by exactly two smaller intervals (except the length 1 leaves). The height of this tree is logarithmic in N .

Note that any particular element of the array is included in $\log N$ intervals – one for each size. This means that updating an element requires only $\log N$ intervals to be updated, which means the update time is $\Theta(\log N)$ instead of $\Theta(N)$ which was the case for sparse tables.

```

1: procedure UPDATETREE( $tree\ T$ , index  $i$ , value  $v$ )
2:    $index \leftarrow i + N$ 
3:    $tree[index] \leftarrow v$ 
4:   while  $index \neq 0$  do
5:      $index \leftarrow index/2$ 
6:      $tree[index] \leftarrow tree[2 \cdot index] + tree[2 \cdot index + 1]$ 

```

It is more difficult to construct an appropriate cover if the interval we are to compute the sum of. A recursive rule can be used. We start at the interval $[0, N)$. One of three cases must now apply:

- We are querying the entire interval $[0, N)$
- We are querying an interval that lies in either $[0, \frac{N}{2})$ or $[\frac{N}{2}, N)$
- We are querying an interval that lies in both $[0, \frac{N}{2})$ or $[\frac{N}{2}, N)$

In the first case, we are done (and respond with the sum of the current interval). In the second case, we perform a recursive call on the half of the

interval that the query lies in. In the third case, we make the same recursive construction for both the left and the right interval.

Since there is a possibility we perform two recursive calls, we might think that the worst-case complexity of this query would be $\Theta(N)$ time. However, the calls that the third call results in will have a very specific form – they will always have one endpoint in common with the interval in the tree. In this case, the only time the recursion will branch is to one interval that is entirely contained in the query, and one that is not. The first call will not make any further calls. All in all, this means that there will be at most two branches of logarithmic height, so that queries are $O(\log N)$.

```
1: procedure QUERYTREE(tree  $T$ , index  $i$ , query  $[L, R]$ , tree interval  $[L', R']$ )
2:   if  $R \leq L'$  or  $L > R'$  then
3:     return 0
4:   if  $L = L'$  and  $R = R'$  then
5:     return  $T[i]$ 
6:    $M = (L' + R')/2$ 
7:    $lsum = \text{QueryTree}(T, 2i, [L, \min(R, M)], [L', M])$ 
8:    $rsum = \text{QueryTree}(T, 2i + 1, [\max(L, M), R], [M, R])$ 
9:   return  $lsum + rsum$ 
```

13.3 Chapter Notes

14 Graph Algorithms

Graph theory is probably the richest of all algorithmic areas. You are almost guaranteed to see at least one graph problem in any given contest, so it is important to be well versed in the common algorithms that operate on graphs. The most important graph algorithms are used to find shortest paths from some vertex. It is these algorithms that we study first.

14.1 Breadth-First Search

One of the most common basic graph algorithms is the *breadth-first search*. It is used to find the distances from a certain vertex in an unweighted graph.

Single-Source Shortest Path, Unweighted Edges

Given an unweighted graph $G = (V, E)$ and a source vertex s , compute the shortest distances $d(s, v)$ for all $v \in V$.

For simplicity, we first consider the problem on a grid graph, where the unit squares constitute vertices, and vertices which share an edge are connected. Additionally, some squares are blocked (and don't have a corresponding vertex). An example can be seen in Figure 14.1.




Figure 14.1: An example grid graph, with source marked s .

Let us solve this problem inductively. First of all, what vertices have distance 0? Clearly, this is only the source vertex s itself. This seems like a reasonable base case, since the problem is about shortest paths from s . Then, what vertices

have distance 1? These are exactly those with a path consisting of a single edge from s , meaning they are the neighbors of s (marked in Figure 14.2).




Figure 14.2: The square with distance 1 from the source.

If a vertex v has distance 2, it must be a neighbor of a vertex u with distance 1 (except for the starting vertex). This is also a sufficient condition, since we can construct a path of length 2 simply by extending the path of any neighbor of distance 1 with the edge (u, v) .





Figure 14.3: The squares with distance 2, 3 and 4.

In fact, this reasoning generalizes to any particular distance, i.e., that all the vertices that have exactly distance k are those that have a neighbor of distance $k - 1$ but no neighbor to a vertex with a smaller distance. Using this, we can construct an algorithm to solve the problem. Initially, we set the distance of s to 0. Then, for every $dist = 1, 2, \dots$, we mark all vertices that have a neighbor with distance $dist - 1$ as having $dist$. This algorithm is called the *breadth-first search*.

Exercise 14.1. Use the BFS algorithm to compute the distance to every square in the following grid:



A simple implementation of this would be to iteratively construct the lists of the vertices which have distance 0, 1, \dots , and so on.

```

1: procedure BREADTHFIRSTSEARCH(vertices  $V$ , vertex  $s$ )
2:    $distances \leftarrow \text{new int}[|V|]$ 
3:   fill  $distances$  with  $\infty$ 
4:    $curDist \leftarrow 0$ 
5:    $curVertices \leftarrow \text{new vector}$ 
6:    $curVertices.add(s)$ 
7:    $distances[s] \leftarrow curDist$ 
8:   while  $curDistVertices \neq \emptyset$  do
9:      $nextVertices \leftarrow \text{new vector}$ 
10:    for  $from \in curVertices$  do
11:      for  $v \in from.neighbours$  do
12:        if  $distances[v] = \infty$  then
13:           $nextVertices.add(v)$ 
14:           $distances[v] \leftarrow curDist + 1$ 
15:     $curDist \leftarrow curDist + 1$ 
16:     $curVertices = nextVertices$ 
17:  return  $distances$ 

```

Each vertex is added to $nextVertices$ at most once, since it is only pushed if $distances[v] = \infty$ whereupon it is immediately set to something else. We then iterate through every neighbor of all these vertices. In total, the number of all neighbours is $2E$, so the algorithm in total uses $\Theta(V + E)$ time.

Usually, the outer loop are often coded in another way. Instead of maintaining two separate vectors, we can merge them into a single queue:

```

1: while  $curVertices \neq \emptyset$  do

```

```

2:  from ← curVertices.front()
3:  curVertices.pop()
4:  for from ∈ curVertices do
5:      for v ∈ from.neighbours do
6:          if distances[v] = ∞ then
7:              curVertices.add(v)
8:              distances[v] = distances[from] + 1

```

The order of iteration is equivalent to the original order.

Exercise 14.2. Prove that the shorter way of coding the BFS loop (Algorithm ??) is equivalent to the longer version (Algorithm ??).

Exercise 14.3. Implement BFS problem

In many problems the task is to find a shortest path between some pair of vertices where the graph is given *implicitly*.

8-puzzle

In the *8-puzzle*, 8 tiles are arranged in a 3×3 grid, with one square left empty. A move in the puzzle consists of sliding a tile into the empty square. The goal of the puzzle is to perform some moves to reach the target configuration. The target configuration has the empty square in the bottom right corner, with the numbers in order 1, 2, 3, 4, 5, 6, 7, 8 on the three lines.

← 8	6	
7	1	4
2	5	3

8		6
7	1	4
2	5	3

1	2	3
4	5	6
7	8	

Figure 14.4: An example 8-puzzle, with a valid move. The rightmost puzzle shows the target configuration.

Given a puzzle, determine how many moves are required to solve it, or if it cannot be solved.

This is a typical BFS problem, characterized by a starting state (the initial puzzle), some transitions (the moves we can make), and the task of finding a short sequence of transitions to some goal state. We can model this kind of problem using a graph. The vertices represent the possible arrangements of

the tiles in the grid, and an edge connects two states if they differ by a single move. A sequence of moves from the starting state to the target configuration then represents a path in this graph. The minimum number of moves required is the same as the distance between those vertices in the graph, meaning we can use a BFS.

In such a problem, most of the code usually deals with the representation of a state as a vertex, and generating the edges that a certain vertex is adjacent to. When an implicit graph is given, we generally do not compute the entire graph explicitly. Instead, we use the states from the problems as-is, and generate the edges of a vertex only when it is being visited in the breadth-first search. In the 8-puzzle, we can represent each state as a 3×3 2D-vector. The difficult part is generating all the states that we can reach from a certain state.

Generating 8-puzzle Moves

```

1  typedef vector<vi> Puzzle;
2
3  vector<Puzzle> edges(const Puzzle& v) {
4      int emptyRow, emptyCol;
5      rep(row,0,3)
6          rep(col,0,3)
7              if (v[row][col] == 0) {
8                  emptyRow = row;
9                  emptyCol = col;
10             }
11     vector<Puzzle> possibleMoves;
12     auto makeMove = [&](int rowMove, int colMove) {
13         int newRow = emptyRow + rowMove;
14         int newCol = emptyCol + colMove;
15         if (newRow >= 0 && newCol >= 0 && newRow < 3 && newCol < 3) {
16             Puzzle newPuzzle = v;
17             swap(newPuzzle[emptyRow][emptyCol], newPuzzle[newRow][newCol]);
18             possibleMoves.push_back(newPuzzle);
19         }
20     };
21     makeMove(-1, 0);
22     makeMove(1, 0);
23     makeMove(0, -1);
24     makeMove(0, 1);
25     return possibleMoves;
26 }
```

With the edge generation in hand, the rest of the solution is a normal BFS, slightly modified to account for the fact that our vertices are no longer numbered $0, \dots, V - 1$. We can solve this by using e.g. maps instead.

8-puzzle BFS

```
1  int puzzle(const Puzzle& S, const Puzzle& target) {
2      map<Puzzle, int> distances;
3      distances[S] = 0;
4      queue<Puzzle> q;
5      q.push(S);
6      while (!q.empty()) {
7          const Puzzle& cur = q.front(); q.pop();
8          int dist = distances[cur];
9          if (cur == target) return dist;
10         for (const Puzzle& move : edges(cur)) {
11             if (distances.find(move) != distances.end()) continue;
12             distances[move] = dist + 1;
13             q.push(move);
14         }
15     }
16     return -1;
17 }
```

Besides this kind of search problems that can be solved using a BFS, some problems require modifications of a BFS, or use the distances generated only as an intermediary result.

Shortest Cycle

Compute the length of the shortest simple cycle in a graph.

Problem 14.1

Button Bashing – buttonbashing

14.2 Depth-First Search

The depth-first search is an analogue to the breadth-first search that visits vertices in another order. Similarly to how the BFS grows the set of visited vertices using a wide frontier around the source vertex, the depth-first search proceeds its search by, at every step, trying to plunge deeper into the graph. This order is called the depth-first order. More precisely, the search starts at some source vertex s . Then, any neighbor of s is chosen to be the next vertex v . Before visiting any other neighbor of s , we first visit any of the neighbours of v , and so on.

Implementing the depth-first search is usually done with a recursive function, using a vector *seen* to keep track of visited vertices:


```
1: procedure DEPTH-FIRST SEARCH(vertex at, adjacency list G)
2:   if seen[at] then
3:     return
4:   seen[at] = true
5:   for neighbour  $\in G[at]$  do
6:     dfs(neighbour, G)
```

In languages with limited stack space, it is possible to implement the DFS iteratively using a stack instead, keeping the vertices which are currently open in it.

Due to the simplicity of coding the DFS compared to a BFS, it is usually the algorithm of choice in problems where we want to visit all the vertices.

Coast Length

KTH Challenge 2011 – Ulf Lundström

The residents of Soteholm value their coast highly and therefore want to maximize its total length. For them to be able to make an informed decision on their position in the issue of global warming, you have to help them find out whether their coastal line will shrink or expand if the sea level rises. From height maps they have figured out what parts of their islands will be covered by water, under the different scenarios described in the latest IPCC report on climate change, but they need your help to calculate the length of the coastal lines.




Figure 14.5: Gray squares are land and white squares are water. The thick black line is the sea coast.

You will be given a map of Soteholm as an $N \times M$ grid. Each square in

the grid has a side length of 1 km and is either water or land. Your goal is to compute the total length of sea coast of all islands. Sea coast is all borders between land and sea, and sea is any water connected to an edge of the map only through water. Two squares are connected if they share an edge. You may assume that the map is surrounded by sea. Lakes and islands in lakes are not contributing to the sea coast.

Solution. We can consider the grid as a graph, where all the water squares are vertices, and two squares have an edge between them if they share an edge. If we surround the entire grid by an water tiles (a useful trick to avoid special cases in this kind of grid problems), the sea consists exactly of those vertices that are connected to these surrounding water tiles. This means we need to compute the vertices which lie in the same connected component as the sea – a typical DFS task¹. After computing this component, we can determine the coast length by looking at all the squares which belong to the sea. If such a square share an edge with a land tile, that edge contributes 1 km to the coast length.

```

1  const vpi moves = {pii(-1, 0), pii(1, 0), pii(0, -1), pii(0, 1)};
2
3  int coastLength(const vector<vector<bool>>& G) {
4      int H = sz(G) + 4;
5      W = sz(G[0]) + 4;
6      vector<vector<bool>> G2(H, vector<bool>(W, true));
7      rep(i,0,sz(G)) rep(j,0,sz(G[i])) G2[i+2][j+2] = G[i][j];
8      vector<vector<bool>> sea(H, vector<bool>(W));
9
10     function<void(int, int)> floodFill = [&](int row, int col) {
11         if (row < 0 || row >= H || col < 0 || col >= W) return;
12         if (sea[row][col]) return;
13         sea[row][col] = true;
14         trav(move, moves) floodFill(row + move.first, col + move.second);
15     };
16     dfs(0, 0);
17
18     int coast = 0;
19     rep(i,1,sz(G)+1) rep(j,1,sz(G[0])+1) {
20         if (sea[i][j]) continue;
21         trav(move, moves) if (!sea[i + move.first][j + move.second]) coast++;
22     }
23     return coast;
24 }

```

□

¹This particular application of DFS, i.e. computing a connected area in a 2D grid, is called a *flood fill*.

Problem 14.2*Mårten's DFS* – martensdfs**Cutvertices and Bridges****Strongly Connected Components****14.3 Weighted Shortest Path**

The theory of computing shortest paths in the case of weighted graphs is a bit more rich than the unweighted case. Which algorithm to use depends on three factors:

- The number of vertices.
- Whether edge weights are non-negative or not.
- If we seek shortest paths only from a single vertex or between all pairs of vertices.

There are mainly three algorithms used: *Dijkstra's Algorithm*, the *Bellman-Ford* algorithm, and the *Floyd-Warshall* algorithm.

Dijkstra's Algorithm

Dijkstra's Algorithm can be seen as an extension of the breadth-first search that works for weighted graphs as well.

Single-Source Shortest Path, non-negative weights

Given a weighted graph $G = (V, E)$ where all weights are non-negative and a source vertex s , compute the shortest distances $d(s, v)$ for all $v \in V(G)$.

It has a similar inductive approach, where we iteratively compute the shortest distances to all vertices in ordered by distance. The difference lies in that we do not immediately know when we have found the shortest path to a vertex. For example, the shortest path from a neighbour to the source vertex may now use several other vertices in its shortest path to s (Figure TODO).

However, can this be the case for *every* vertex adjacent to s ? In particular, can the neighbour to s with the smallest edge weight W to s have a distance smaller than w ? This is never the case in Figure TODO, and actually holds in general. Assume that this is not the case. In Figure TODO, this hypothetical scenario is shown. Any such path must still pass through some neighbour u of s .

By assumption, the weight of the edge (s, u) must be larger than W (which was the minimal weight of the edges adjacent to s). This reasoning at least allows us to find the shortest distance to one other vertex.

Bellman-Ford

Single-Source Shortest Path, non-negative weights

Given a weighted graph $G = (V, E)$ where all weights are non-negative and a source vertex s , compute the shortest distances $d(s, v)$ for all $v \in V(G)$.

When edges can have negative weights, the idea behind Dijkstra's algorithm no longer works. It is very much possible that a negative weight edge somewhere else in the graph could be used to construct a shorter path to a vertex which was already marked as completed. However, the concept of relaxing an edge is still very much applicable and allows us to construct a slower, inductive solution.

Initially, we know the shortest distances to each vertex, assuming we are allowed to traverse 0 edges from the source. Assuming that we know these values when allowed to traverse up to k edges, can we find the shortest paths that traverse up to $k + 1$ edges? This way of thinking is similar to how to solved the BFS problem. Using a BFS, we computed the vertices at distance $d + 1$ by taking the neighbours of the vertices at distance d . Similarly, we can find the shortest path to a vertex v that traverse up to $k + 1$ edges, by attempting to extend the shortest paths using k edges from the neighbours of v . Letting $D(k, v)$ be the shorter distance to v by traversing up to k edges, we arrive at the following recursion:

$$D(k, v) = \min \begin{cases} 0 & \text{if } v = s \\ D(k - 1, v) & \text{if } k > 0 \\ \min_{e=(u,v) \in E} D(k - 1, u) + W(e) & \text{if } k > 0 \end{cases}$$

The implementation is straightforward:

```

1: procedure BELLMANFORD(vertices  $V$ , edges  $E$ , vertex  $s$ )
2:    $D \leftarrow$  new int[ $|V|$ ][ $|V| + 1$ ]
3:   fill  $D$  with  $\infty$ 
4:    $D[0][s] \leftarrow 0$ 
5:   for  $k = 1$  to  $|V|$  do
6:      $D[k] = D[k - 1]$ 

```

```
7:   for  $e = (u, v) \in E$  do
8:        $D[k][v] = \min(D[k][v], D[k-1][u] + W(e))$ 
9:   return  $D$ 
```

All in all, the states $D(v, i)$ for a particular i takes time $\Theta(|E|)$ to evaluate. To compute the distance $d(s, v)$, we still need to know what the maximum possible k needed to arrive at this shortest path could be. It turns out that this could potentially be infinite, in the case where the graph contains a negative-weight cycle. Such a cycle can be exploited to construct arbitrarily short paths.

However, if no such cycle exists, $k = |V|$ will be sufficient. If a shortest path uses more than $|V|$ edges, it must contain a cycle. If this cycle is not of negative weight, we may simply remove it to obtain a path of at most the same length. Thus, the algorithm takes $O\Theta(|V||E|)$.

Exercise 14.4. Bellman-Ford can be adapted to instead use only $\Theta(V)$ memory, by only keeping a current known shortest path and repeatedly relaxing every edge. Sketch out the pseudo code for such an approach, and prove its correctness.

Exercise 14.5. We may terminate Bellman-Ford earlier without loss of correctness, in case $D[k] = D[k-1]$. How can this fact be used to determine whether the graph in question contains a negative-weight cycle?

Floyd-Warshall

All-Pairs Shortest Paths

Given a weighted graph $G = (V, E)$, compute the shortest distance $d(u, v)$ for **every** pair of vertices u, v .

The Floyd-Warshall algorithm is a remarkably short method to solve the all-pairs shortest paths problem. It basically consists of three nested loops containing a single statement, one of the shortest highly useful algorithms there is:

```
1: procedure FLOYD-WARSHALL(distance matrix  $D$ )
2:   for  $k = 0$  to  $|V| - 1$  do
3:       for  $i = 0$  to  $|V| - 1$  do
4:           for  $j = 0$  to  $|V| - 1$  do
5:                $D[i][j] = \min(D[i][j], D[i][k] + D[k][j])$ 
6:   return  $D$ 
```

Initially, the distance matrix D contains the distances of all the edges in E , so that $D[i][j]$ is the weight of the edge (i, j) if such an edge exists, ∞ if there is no edge between i and j or 0 if $i = j$. Note that if multiple edges exist between i and j , $D[i][j]$ must be given the minimum weight of them all. Additionally, if there is a self-loop (i.e. an edge from i to i itself) of negative weight, $D[i][i]$ must be set to this value.

To see why this approach works, we can use the following invariant proven by induction. After the k 'th iteration of the loop, $D[i][j]$ will be at most the minimum distance between i and j that uses vertices $0, 1, \dots, k - 1$. Assume that this is true for a particular k . After the next iteration, there are two cases for $D[i][j]$. Either there is no shorter path using vertex k than those using only vertices $0, 1, \dots, k - 1$. In this case, $D[i][j]$ will fulfill the condition by the induction assumption. If there is a shorter path between i and j if we use the vertex k , this must have length $D[i][k] + D[k][j]$, since $D[i][k]$ and $D[k][j]$ both contain the shortest paths between i and k , and k and j using vertices $0, 1, \dots, k - 1$. Since we set $D[i][j] = \min(D[i][j], D[i][k] + D[k][j])$ in the inner loop, we will surely find this path too in this iteration. Thus, the statement is true after the $k + 1$ 'th iteration too. By induction, it is true for $k = |V|$, meaning $D[i][j]$ contains at most the minimum distance between i and j using any vertex in the graph.

14.4 Minimum Spanning Tree

Using the depth-first search, we were able to find a subtree of a connected graph which spanned the entire set of vertices. A particularly important such spanning tree is the minimum spanning tree.

Minimum Spanning Tree

We say that the weight of a spanning tree is the sum of the weights of its edges. A *minimum spanning tree* is a spanning tree whose weight is minimal.




Figure 14.6: A graph with a corresponding minimum spanning tree.

Given a weighted graph, find a minimum spanning tree.

Mainly two algorithms are used to solve the Minimum Spanning Tree (MST) problem: either **Kruskal's Algorithm** which is based on the Union-Find data structure, or **Prim's Algorithm** which is an extension of Dijkstra's Algorithm. We will demonstrate Kruskal's, since it is by far the most common of the two.

Kruskal's algorithm is based on a greedy, incremental approach that reminds us of the Scheduling problem (Section 10.4). In the Scheduling problem, we tried to find any interval that we could prove to be part of an optimal solution, by considering various extremal cases. Can we do the same thing when finding a minimum spanning tree?

First of all, if we can always find such an edge, we are essentially done. Given an edge (a, b) that we know is part of a minimum spanning tree, we can contract the two vertices a and b to a single vertex ab , with all edges adjacent to the two vertices. Any edges that go between contracted vertices are ignored. An example of this process in action be be seen in Figure 14.7. Note how the contraction of an edge reduces the problem to finding an MST in a smaller graph.

A natural extremal case to consider is the edge with the minimum weight. After all, we are trying to minimize the edge sum. Our proof is similar in structure to the Scheduling problem as well by using a swapping argument.

Assume that a minimum-weight edge $\{a, b\}$ with weight w is not part of any minimum spanning tree. Then, consider the spanning tree with this edge appended. The graph will then contain exactly one cycle. In a cycle, any edge can be removed while maintaining connectivity. This means that if any edge $\{c, d\}$ on this cycle have a weight w' larger weight than w , we can erase it. We

will thus have replaced the edge $\{c, d\}$ by $\{a, d\}$, while changing the weight of the tree by $w - w' < 0$, reducing the sum of weights. Thus, the tree was improved by using the minimum-weight edge, proving that it could have been part of the tree.

Exercise 14.6. What happens if all edges on the cycle that appears have weight w ? Is this a problem for the proof?

When implementing the algorithm, the contraction of the edge added to the minimum spanning tree is generally not performed explicitly. Instead, a disjoint set data structure is used to keep track of which subsets of vertices have been contracted. Then, all the original edges are iterated through in increasing order of weight. An edge is added to the spanning tree if and only if the two endpoints of the edge are not already connected (as in Figure 14.7).

```
1: procedure MINIMUMSPANNINGTREE(vertices  $V$ , edges  $E$ )
2:   sort  $E$  by increasing weight
3:    $uf \leftarrow$  new DisjointSet( $V$ )
4:    $mst \leftarrow$  new Graph
5:   for each edge  $\{a, b\} \in E$  do
6:     if not  $uf.sameSet(a, b)$  then
7:        $mst.append(a, b)$ 
8:        $uf.join(a, b)$ 
9:   return  $mst$ 
```

The complexity of this algorithm is dominated by the sorting (which is $O(E \log V)$), since operations on the disjoint set structure is $O(\log V)$.

14.5 Chapter Notes




Figure 14.7: Incrementally constructing a minimum spanning tree by merging

15 Maximum Flows

This chapter studies so called *flow networks*, and algorithms we use to solve the so-called *maximum flow* and *minimum cut* problems on such networks. Flow problems are common algorithmic problems, particularly in ICPC competitions (while they are out-of-scope for IOI contests). They are often hidden behind statements which seem unrelated to graphs and flows, especially the *minimum cut* problem.

Finally, we will end with a specialization of maximum flow on the case of bipartite graphs (called *bipartite matching*).

15.1 Flow Networks

Informally, a flow network is a directed graph that models any kind of network where paths have a fixed capacity, or throughput. For example, in a road network, each road might have a limited throughput, proportional to the number of lanes on the road. A computer network may have different speeds along different connections due to e.g. the type of material. These natural models are often use when describing a problem that is related to flows. A more formal definition is the following.

Definition 15.1 — Flow Network

A **flow network** is a special kind of **directed** graph (V, E, c) , where each edge e is given a non-negative **capacity** $c(e)$. Two vertices are designated the **source** and the **sink**, which we will often abbreviate to S and T .

In Figure 15.1, you can see an example of a flow network.

In such a network, we can assign another value to each edge, that models the current throughput (which generally does not need to match the capacity). These values are what we call flows.

Definition 15.2 — Flow

A **flow** is a function $f : E \rightarrow \mathbb{R}_{\geq 0}$, associated with a particular flow network (V, E, c) . We call a flow f **admissible** if:




Figure 15.1: An example flow network.

- $0 \leq f(e) \leq c(e)$ – the flow does not exceed the capacity
- For every $v \in V \setminus \{S, T\}$, $\sum_{e \in \text{in}(v)} f(e) = \sum_{e \in \text{out}(v)} f(e)$ – flow is conserved for each vertex, possibly except the source and sink.

The *size* of a flow is defined to be the value

$$\sum_{v \in \text{out}(S)} f(v) - \sum_{v \in \text{in}(S)} f(v)$$

In a computer network, the flows could e.g. represent the current rate of transfer through each connection.

Exercise 15.1. Prove that the size of a given flow also equals

$$\sum_{v \in \text{in}(T)} f(v) - \sum_{v \in \text{out}(T)} f(v)$$

i.e. the excess flow out from S must be equal to the excess flow in to T .

In Figure 15.2, flows have been added to the network from Figure 15.1.

Given such a flow, we are generally interested in determining the flow of the largest size. This is what we call the *maximum flow* problem. The problem is not only interesting on its own. Many problems which we study might initially seem unrelated to maximum flow, but will turn out to be reducible to finding a maximum flow.




Figure 15.2: An example flow network, where each edge has an assigned flow. The size of the flow is 8.

Maximum Flow

Given a flow network (V, E, c, S, T) , construct a maximum flow from S to T .

Input

A flow network.

Output

Output the maximal size of a flow, and the flow assigned to each edge in one such flow.

Exercise 15.2. The flow of the network in Figure 15.2 is not maximal – there is a flow of size 9. Find such a flow.

Before we study problems and applications of maximum flow, we will first discuss algorithms to solve the problem. We can actually solve the problem greedily, using a rather difficult insight, that is hard to prove but essentially gives us the algorithm we will use. It is probably one of the more complex standard algorithms that is in common use.

15.2 Edmonds-Karp

There are plenty of algorithms which solve the maximum flow problem. Most of these are too complicated to be implemented to be practical. We are going to study two very similar classical algorithms that compute a maximum flow. We will start with proving the correctness of the *Ford-Fulkerson* algorithm. Afterwards, a modification known as *Edmonds-Karp* will be analyzed (and found to have a better worst-case complexity).

Augmenting Paths

For each edge, we define a *residual flow* $r(e)$ on the edge, to be $c(e) - f(e)$. The residual flow represents the additional amount of flow we may push along an edge.

In Ford-Fulkerson, we associate every edge e with an additional *back edge* $b(e)$ which points in the reverse order. Each back edge is originally given a flow and capacity 0. If e has a certain flow f , we assign the flow of the back-edge $b(e)$ to be $-f$ (i.e. $f(b(e)) = -f(e)$). Since the back-edge $b(e)$ of e has capacity 0, their residual capacity is $r(b(e)) = c(b(e)) - f(b(e)) = 0 - (-f(e)) = f(e)$.

Intuitively, the residual flow represents the amount of flow we can add to a certain edge. Having a back-edge thus represents “undoing” flows we have added to a normal edge, since increasing the flow along a back-edge will decrease the flow of its associated edge.




Figure 15.3: The residual flows from the network in Figure 15.2.

The basis of the Ford-Fulkerson family of algorithms is the *augmenting path*. An augmenting path is a path from S to T in the network consisting of edges e_1, e_2, \dots, e_l , such that $r(e_i) > 0$, i.e. every edge along the path has a residual flow. Letting m be the minimum residual flow among all edges on the path, we can increase the flow of every such edge with m .

In Figure 15.3, the path S, c, d, b, T is an augmenting path, with minimum residual flow 1. This means we can increase the flow by 1 in the network, by:

- Increasing the flow from S to c by 1
- Increasing the flow from c to d by 1

- Decreasing the flow from b to d by 1 (since (d, b) is a back-edge, augmenting the flow along this edge represents removing flow from the original edge)
- Increasing the flow from d to T

The algorithm for augmenting a flow using an augmenting path is simple:

```
1: procedure AUGMENT(path  $P$ )
2:    $inc \leftarrow \infty$ 
3:   for  $e \in P$  do
4:      $inc \leftarrow \min(inc, c(e) - f(e))$ 
5:   for  $e \in P$  do
6:      $f(e) \leftarrow f(e) + inc$ 
7:      $f(b(e)) \leftarrow f(b(e)) - inc$ 
8:   return  $inc$ 
```

Performing this kind of augmentation on an admissible flow will keep the flow admissible. A path must have either zero or two edges adjacent to any vertex (aside from the source and sink). One of these will be an incoming edge, and one an outgoing edge. Increasing the flow of these edges by the same amount conserves the equality of flows between in-edges and out-edges, meaning the flow is still admissible.

This means that a flow can be maximum only if it contains no augmenting paths. It turns out this is also a necessary condition, i.e. a flow is maximum if it contains no augmenting path. Thus, we can solve the maximum flow problem by repeatedly finding augmenting paths, until no more exists.

Finding Augmenting Paths

The most basic algorithms based on augmenting paths is the *Ford-Fulkerson* algorithm. It uses a simple DFS to find the augmenting paths:

```
1: procedure AUGMENTINGPATH(flow network  $(V, E, c, f, S, T)$ )
2:    $bool[] \text{ seen} \leftarrow \text{new } bool[|V|]$ 
3:    $Stack \text{ stack} \leftarrow \text{new } Stack$ 
4:    $found \leftarrow DFS(S, T, f, c, \text{seen}, \text{stack})$ 
5:   if  $found$  then
6:     return  $stack$ 
```

```

7:   return Nil
8: procedure DFS(vertex at, sink T, flow f, capacity c, path p)
9:   p.push(at)
10:  if at = T then
11:    return true
12:  for every out-edge e = (at, v) from at do
13:    if  $f(e) < c(e)$  then
14:      if DFS(v, T, f, c, p) then
15:        return true
16:  p.pop()
17:  return false

```

For integer flows, where the maximum flow has size m Ford-Fulkerson may require up to $O(Em)$ time. In the worst case, a DFS takes $\Theta(E)$ time to find a path from S to T , and one augmenting path may contribute only a single unit of flow. For non-integral flows, there are instances where Ford-Fulkerson may not even terminate (nor converge to the maximum flow).

An improvement to this approach is simply to use a BFS instead. This is what is called the *Edmonds-Karp* algorithm. The BFS looks similar to the Ford-Fulkerson DFS, and is modified in the same way (i.e. only traversing those edges where the flow $f(e)$ is smaller than the capacity $c(e)$). The resulting complexity is instead $O(VE^2)$ (which is tight in the worst case).

15.3 Applications of Flows

We will now study a number of problems which are reducible to finding a maximum flow in a network. Some of these problems are themselves considered to be standard problems.

Maximum-Flow with Vertex Capacities

In a flow network, each vertex v additionally have a limit C_v on the amount of flow that can go through it, i.e.

$$\sum_{e \in \text{in}(v)} f(e) \leq C_v$$

Find the maximum flow subject to this additional constraint.

This is nearly the standard maximum flow problem, with the addition of vertex capacities. We are still going to use the normal algorithms for maximum flow. Instead, we will make some minor modifications to the network. The additional constraint given is similar to the constraint placed on an edge. An edge has a certain amount of flow passing through it, implying that the same amount must enter and exit the edge. For this reason, it seems like a reasonable approach to reduce the vertex capacity constraint to an ordinary edge capacity, by forcing all the flow that passes through a vertex v with capacity C_v through a particular edge.

If we partition all the edges adjacent to v into incoming and outgoing edges, it becomes clear how to do this. We can split up v into two vertices v_{in} and v_{out} , where all the incoming edges to v are now incoming edges to v_{in} and the outgoing edges instead become outgoing edges from v_{out} . If we then add an edge of infinite capacity from v_{in} to v_{out} , we claim that the maximum flow of the network does not change. All the flow that passes through this vertex must now pass through this edge between v_{in} and v_{out} . This construction thus accomplish our goal of forcing the vertex flow through a particular edge. We can now enforce the vertex capacity by changing the capacity of this edge to C_v .

Maximum Bipartite Matching

Given a bipartite graph, a **bipartite matching** is a subset of edges in the graph, such that no two edges share an endpoint. Determine the matching containing the maximum number of edges.

The maximum bipartite matching problem is probably the most common reduction to maximum flow in use. Some standard problems additionally reduce to bipartite matching, making maximum flow even more important. Although there are others ways of solving maximum bipartite matching than a reduction to flow, this is what how we are going to solve it.

How can we find such a reduction? In general, we try to find some kind of graph structure in the problem, and model what it “means” for an edge to have flow pushed through it. In the bipartite matching problem, we are already given a graph. We also have a target we wish to maximize – the size of the matching – and an action that is already associated with edges – including it in the matching. It does not seem unreasonable that this is how we wish to model the flow, i.e. that we want to construct a network based on this graph where pushing flow along one of the edges means that we include the edge in the matching. No two

selected edges may share an endpoint, which brings only a minor complication. After all, this condition is equivalent to each of the vertices in the graph having a vertex capacity of 1. We already know how to enforce vertex capacities from the previous problem, where we split each such vertex into two, one for in-edges and one for out-edges. Then, we added an edge between them with the required capacity. After performing this modification on the given graph, we are still missing one important part of a flow network. The network does not yet have a source and sink. Since we want flow to go along the edges, from one of the parts to another part of the graph, we should place the source at one side of the graph and the sink at the other, connecting the source to all vertices on one side and all the vertices on the other side to the sink.

Minimum Path Cover

In a directed, acyclic graph, find a minimum set of vertex-disjoint paths that includes every vertex.

This is a difficult problem to derive a flow reduction to. It is reduced to bipartite matching in a rather unnatural way. First of all, a common technique must be used to get introduce a bipartite structure into the graph. For each vertex, we split it into two vertices, one in-vertex and one out-vertex. Note that this graph still have the same minimum path covers as the original graph.

Now, consider any path cover of this new graph, where we ignore the added edges. Each vertex is then adjacent to at most a single edge, since paths are vertex-disjoint. Additionally, the number of paths are equal to the number of in-edges that does not lie on any path in the cover (since these vertices are the origins of the paths). Thus, we wish to select a maximum subset of the original edges. Since the subgraph containing only these edges is now bipartite, the problem reduces to bipartite matching.

Exercise 15.3. The minimum path cover reduction can be modified slightly to find a *minimum cycle cover* in a directed graph instead. Construct such a reduction.

15.4 Chapter Notes

The Edmonds-Karp algorithm was originally published in 1970 by Yefim Dinitz. The paper by Edmonds and Karp

16 Strings

In computing, much of the information we process is text. Therefore, it should not come as a surprise that many common algorithms and problems focus concerns text strings. In this chapter, we will study some of the common string algorithms and data structures.

16.1 Tries

The *trie* (also called a *prefix tree*) is the most common string-related data structure. It represents a set of words as a rooted tree, where every prefix of every word is a vertex, with children from a prefix P to all strings Pc which are also prefixes of a word. If two words have a common prefix, the prefix only appears once as a vertex. The root of the tree is the empty prefix. The trie is very useful when we want to associate some information with prefixes of strings and quickly get the information from neighboring strings.

The most basic operation of the trie is the insertion of strings, which may be implemented as follows.

Trie

```
1 struct Trie {
2     map<char, Trie> children;
3     bool isWord = false;
4
5     void insert(const string& s, int pos) {
6         if (pos != sz(s)) children[s[pos]].insert(s, pos + 1);
7         else isWord = true;
8     }
9
10 };
```

We mark those vertices which corresponds to the inserted word using a boolean flag `isWord`. Many problems essentially can be solved by very simple usage of a trie, such as the following IOI problem.

Type Printer

International Olympiad in Informatics 2008

You need to print N words on a movable type printer. Movable type printers are those old printers that require you to place small metal pieces (each containing a letter) in order to form words. A piece of paper is then pressed against them to print the word. The printer you have allows you to do any of the following operations:

- Add a letter to the end of the word currently in the printer.
- Remove the last letter from the end of the word currently in the printer. You are only allowed to do this if there is at least one letter currently in the printer.
- Print the word currently in the printer.

Initially, the printer is empty; it contains no metal pieces with letters. At the end of printing, you are allowed to leave some letters in the printer. Also, you are allowed to print the words in any order you like. As every operation requires time, you want to minimize the total number of operations.

Your task is to output a sequence of operations that prints all the words using the minimum number of operations needed.

Input

The first line contains the number of words $1 \leq N \leq 25\,000$. The next N lines contain the words to be printed, one per line. Each word is at most 20 letters long and consist only of lower case letters a-z. All words will be distinct

Output

Output a sequence of operations that prints all the words. The operations should be given in order, one per line, starting with the first. Adding a letter c is represented by outputting c on a line. Removing the last letter of the current word is represented by a $-$. Printing the current word is done by outputting P .

Let us start by solving a variation of the problem, where we are not allowed to leave letters in the printer at the end. First of all, are there actions that never make sense? For example, what sequences of letters will ever appear in the type writer during an optimal sequence of operations? Clearly we never wish to input a sequence that is not a prefix of a word we wish to type. For example, if we

input *abcdef* and this is not a prefix of any word, we must at some point erase the last letter *f*, without having printed any words. But then we can erase the entire sequence of operations between inputting the *f* and erasing the *f*, without changing what words we print.

On the other hand, every prefix of a word we wish to print must at some point appear on the type writer. Otherwise, we would not be able to reach the word we wish to print. Therefore, the partial words to ever appear on the type printer are exactly the prefixes of the words we wish to print – strongly hinting at a trie-based solution.

If we build the trie of all words we wish to print, it contains as vertices exactly those strings which will appear as partial words on the printer. Furthermore, the additions and removals of letters form a sequence of vertices that are connected by edges in this trie. We can move either from a prefix *P* to a prefix *Pc*, or from a prefix *Pc* to a prefix *P*, which are exactly the edges of a trie. The goal is then to construct the shortest possible tour starting at the root of the trie and passing through all the vertices of the trie.

Since a trie is a tree, any such trail must pass through every edge of the trie at least twice. If we only passed through an edge once, we can never get back to the root since every edge disconnects the root from the endpoint of the edge further away from the root. It is actually possible to construct a trail which passes through every edge exactly twice (which is not particularly difficult if you attempt this task by hand). As it happens, the depth-first search of a tree passes through an edge exactly twice – once when first traversing the edge to an unvisited vertex, and once when backtracking.

The problem is subtly different once we are allowed to leave some letters in the printer at the end. Clearly, the only difference between an optimal sequence when letters may remain and an optimal sequence when we must leave the printer empty is that we are allowed to skip some trailing removal operations. If the last word we print is *S*, the difference will be exactly $|S|$ “-” operations. An optimal solution will therefore print the longest word last, in order to “win” as many “-” operations as possible. We would like this last word to be the longest word of *all the ones we print* if possible. In fact, we can order our DFS such that this is possible. First of all, our DFS starts from the root and the longest word is $s_1 s_2 \dots s_n$. When selecting which order the DFS should visit the children of the root in, we can select the child s_1 last. Thus, all words starting with the letter s_1 will be printed last. When visiting s_1 , we use the same trick and visit the child

$s_1 s_2$ last of the children of s_1 , and so on. This guarantees S to be the last word to be printed.

Note that the solution requires no additional data to be stored in the trie – the only modification to our basic trie is the DFS.

Typewriter

```
1  struct Trie {
2      ...
3
4      void dfs(int depth, const string& longest) {
5          trav(it, children)
6          if (it->first != longest[depth])
7              dfs2(depth, longest, it->first);
8          dfs2(depth, longest, longest[depth]);
9      }
10
11     void dfs2(int depth, const string& longest, char output) {
12         cout << output << endl;
13         if (isWord) cout << "P" << endl;
14         children[output]->dfs(depth + 1, longest);
15         if (longest[depth] != output) {
16             cout << "." << endl;
17         }
18     }
19 };
```

Generally, the uses of tries are not this simple, where we only need to construct the trie and fetch the answer through a simple traversal. We often need to augment tries with additional information about the prefixes we insert. This is when tries start to become really powerful. The next problem requires only a small augmentation of a trie, to solve a problem which looks complex.

Rareville

In Rareville, everyone must have a distinct name. When a new-born baby is to be given a name, its parents must first visit *NAME*, the Naming Authority under the Ministry of Epithets, to have its name approved. The authority has a long list of all names assigned to the citizens of Rareville. When deciding whether to approve a name or not, a case officer uses the following procedure. They start at the first name in the list, and read the first letter of it. If this letter matches the first letter of the proposed name, they proceed to read the next letter in the word. This is repeated for every letter of the name *in the list*. After reading a letter from the word, the case officer can sometime determine that this could not

possibly be the same name as the proposed one. This happens if either

- the next letter in the proposed name did not match the name in the list
- there was no next letter in the proposed name
- there was no next letter in the name in the list

When this happens, the case officer starts over with the next name in the list, until exhausting all names in the list. For each letter the case officer reads (or attempts to read) from a name in the list, one second pass.

Currently, there are N people in line waiting to apply for a name. Can you determine how long time the decision process will take for each person?

Input

The first line contains integers $1 \leq D \leq 200\,000$ and $1 \leq N \leq 200\,000$, the size of the dictionary and the number of people waiting in line. The next D lines contains one lowercase name each, the contents of the dictionary. The next N lines contains one lowercase name each, the names the people in line wish to apply with. The total size of the lists is at most 10^6 letters.

Output

For each of the N names, output the time (in seconds) the case officer needs to decide on the application.

The problem clearly relates to prefixes in some way. Given a dictionary word A and an application for a name B , the case officer needs to read letters from A corresponding to the longest common prefix of A and B , plus 1. Hence, our solution will probably be to consider all the prefixes of each proposed name, which is exactly what tries are good at.

Instead of thinking about this process one name at a time, we use a common trie technique and look at the transpose of this problem, i.e. for every i , how many names C_i have a longest common prefix of length *at least* i when handling the application for a name S ? This way, we have transformed the problem from being about D individual processes to $|S|$ smaller problems which treat the dictionary as unified group of strings. Then, we will have to read $C_0 + C_1 + \dots + C_{|S|}$ letters.

Now, the solution should be clear. We augment the trie vertex for a particular prefix p with the number of strings P_p in the list that start with this prefix. Initially, an empty trie has $P_p = 0$ for every p . Whenever we insert a new word $W = w_1 w_2 \dots$ in the trie, we need to increment $P_{w_1}, P_{w_1 w_2}, \dots$, to keep all the P_p

correct, since we have added a new string which have those prefixes. Then, we have that $C_i = P_{s_1 s_2 \dots s_i}$, so that we can compute all the numbers P_i by following the word S in the trie. The construction of the trie is linear in the number of characters we insert, and responding to a query is linear in the length of the proposed name.

Rareville

```
1 struct Trie {
2     map<char, Trie> children;
3     int P = 0;
4
5     void insert(const string& s, int pos) {
6         P++;
7         if (pos != sz(s)) children[s[pos]].insert(s, pos + 1);
8     }
9
10    int query(const string& s, int pos) {
11        int ans = P;
12        if (pos != sz(s)) {
13            auto it = children.find(s[pos]);
14            if (it != children.end) ans += it->second.query(s, pos + 1);
15        }
16        return ans;
17    }
18 };
```

16.2 String Matching

A common problem on strings – both in problem solving and real life – is that of searching. Not only do we need to check whether e.g. a set of strings contain some particular string, but also if one string contain another one as a substring. This operation is ubiquitous; operating systems allow us to search the contents of our files, and our text editors, web browsers and email clients all support substring searching in documents. It should come as no surprise that string matching is part of many string problems.

String Matching

Find all occurrences of the pattern P as a substring in the string W .

We can solve this problem naively in $O(|W| \cdot |P|)$. If we assume that an occurrence of P starts at position i in W , we can compare the substring $W[i \dots i + |P| - 1]$ to P in $O(|P|)$ time by looping through both strings, one character at a time:

```

1: procedure STRINGMATCHING(pattern  $P$ , string  $W$ )
2:    $answer \leftarrow$  new vector
3:   for outer:  $i$  from 0 to  $|W| - |P|$  do
4:     for  $j$  from 0 to  $|P| - 1$  do
5:       if  $P[j] \neq W[i + j]$  then
6:         start next iteration of outer
7:        $answer.append(i)$ 
8:   return  $answer$ 

```

Intuitively, we should be able to do better. With the naive matching, our problem is basically that we can perform long stretches of partial matches for every position. Searching for the string $a^{\frac{n}{2}}$ in the string a^n takes $O(n^2)$ time, since each of the $\frac{n}{2}$ positions where the pattern can appear requires us to look ahead for $\frac{n}{2}$ characters to realize we made a match. On the other hand, if we manage to find a long partial match of length l starting at i , we *know* what the next l letters of W are – they are the l first letters of P . With some cleverness, we should be able to exploit this fact, hopefully avoiding the need to scan them again when we attempt to find a match starting at $i + 1$.

For example, assume we have $P = \text{bananarama}$. Then, if we have performed a partial match of *banana* at some position i in W but the next character is a mismatch (i.e., it is not an r), we know that no match can begin at the next 5 characters. Since we have matched *banana* at i , we have that $W[i + 1 \dots i + 5] = \text{anana}$, which does not contain a b .

As a more interesting example, take $P = \text{abbaabborre}$. This pattern has the property that the partial match of *abbaabb* actually contains as a prefix of P itself as a suffix, namely *abb*. This means that if at some position i get this partial match but the next character is a mismatch, we can not immediately skip the next 6 characters. It is possible that the entire string could have been *abbaabbaabborre*. Then, an actual match (starting at the fifth character) overlaps our partial match.

It seems that if we find a partial match of length 7 (i.e. *abbaabb*), we can only skip the first 4 characters of the partial match.

For every possible partial match of the pattern P , how many characters are we able to skip if we fail a k -length partial match? If we could precompute such a table, we should be able to perform matching in linear time, since we would only have to investigate every character of W once. Assume the next possible match is l letters forward. Then the new partial match must consist of the last $k - l$ letters of the partial match, i.e. $P[l \dots k - 1]$. But a partial match is just a prefix of P , so we must have $P[l \dots k - 1] = P[0 \dots l - 1]$. In other words, for every given k , we must find the longest suffix of $P[0 \dots K - 1]$ that is also a prefix of P (besides $P[0 \dots k - 1]$ itself, of course).

We can compute these suffixes rather easily in $O(n^2)$. For each possible position for the next possible match l , we perform a string matching to find all occurrences of prefixes of P within P :

```

1: procedure LONGESTSUFFIXES(pattern  $P$ )
2:    $T \leftarrow$  new int[ $|P| + 1$ ]
3:   for  $l$  from 1 to  $|P| - 1$  do
4:      $matchLen \leftarrow 0$ 
5:     while  $l + matchLen \leq |W|$  do
6:       if  $P[l] \neq P[matchLen]$  then
7:         break
8:        $matchLen \leftarrow matchLen + 1$ 
9:        $T[l + matchLen] = matchLen$ 
10:  return  $T$ 

```

A string such as $P = \textit{bananarama}$, where no partial match could possibly contain a new potential match, this table would simply be:

P	b	a	n	a	n	a	r	a	m	a
T	0	0	0	0	0	0	0	0	0	0

When $P = \textit{abbaabborre}$, the table instead becomes:

P	a	b	b	a	a	b	b	o	r	r	e
T	0	0	0	1	1	2	3	0	0	0	0

With this precomputation, we can now perform matching in linear time. The matching is similar to the naive matching, except we can now use this

precomputed table to determine whether there is a new possible match somewhere within the partial match.

```
1: procedure STRINGMATCHING(pattern  $P$ , text  $W$ )
2:    $matches \leftarrow$  new vector
3:    $T \leftarrow$  LongestSuffixes( $P$ )
4:    $pos \leftarrow 0, match \leftarrow 0$ 
5:   while  $pos + match < |W|$  do
6:     if  $match < |P|$  and  $W[pos + match] = P[match]$  then
7:        $match \leftarrow match + 1$ 
8:     else if  $match = 0$  then
9:        $pos \leftarrow pos + 1$ 
10:    else
11:       $pos \leftarrow pos + match - T[match]$ 
12:       $match \leftarrow T[match]$ 
13:    if  $match = |P|$  then
14:       $matches.append(match)$ 
15:  return  $matches$ 
```

In each iteration of the loop, we see that either $match$ is increased by one, or $match$ is decreased by $match - T[match]$ and pos is increased by the same amount. Since $match$ is bounded by P and pos is bounded by $|W|$, this can happen at most $|W| + |P|$ times. Each iteration takes constant time, meaning our matching is $\Theta(|W| + |P|)$ time.

While this is certainly better than the naive string matching, it is not particularly helpful when $|P| = \Theta(|W|)$ since we need an $O(|P|)$ preprocessing. The solution lies in how we computed the table of suffix matches, or rather, the fact that it is entirely based on string matching itself. We just learned how to use this table to perform string matching in linear time. Maybe we can use this table to extend itself and get the precomputation down to $O(|P|)$? After all, we are looking for occurrences of prefixes of P in P itself, which is exactly what string matching does. If we modify the string matching algorithm for this purpose, we get what we need:

```
1: procedure LONGESTSUFFIXES(pattern  $P$ )
2:    $T \leftarrow$  new int[ $|P| + 1$ ]
```

```
3:  pos ← 1, match ← 0
4:  while pos + match < |P| do
5:      if P[pos + match] = P[match] then
6:          T[pos + match] ← match + 1
7:          match ← match + 1
8:      else if match = 0 then
9:          pos ← pos + 1
10:     else
11:         pos ← pos + match - T[match]
12:         match ← T[match]
13:     if match = |P| then
14:         matches.append(match)
15: return T
```

This string matching algorithm is called the **Knuth-Morris-Pratt** (KMP) algorithm.

Using the same analysis as for the improved string matching, this precomputation is instead $\Theta(|P|)$. The resulting string matching then takes $\Theta(|P| + |W|)$.

Competitive Tip

Most programming languages have functions to find occurrences of a certain string in another. However, they mostly use the naive $O(|W||P|)$ procedure. Be aware of this and code your own string matching if you need it to perform in linear time.

Clock Pictures

Nordic Collegiate Programming Contest 2014

You have two pictures of an unusual kind of clock. The clock has n hands, each having the same length and no kind of marking whatsoever. Also, the numbers on the clock are so faded that you can't even tell anymore what direction is up in the picture. So the only thing that you see on the pictures, are n shades of the n hands, and nothing else.

You'd like to know if both images might have been taken at exactly the same time of the day, possibly with the camera rotated at different angles.

Given the description of the two images, determine whether it is possible that these two pictures could be showing the same clock displaying the same

time.

Input

The first line contains a single integer n ($2 \leq n \leq 200000$), the number of hands on the clock.

Each of the next two lines contains n integers a_i ($0 \leq a_i \leq 360000$), representing the angles of the hands of the clock on one of the images, in thousandths of a degree. The first line represents the position of the hands on the first image, whereas the second line corresponds to the second image. The number a_i denotes the angle between the recorded position of some hand and the upward direction in the image, measured clockwise. Angles of the same clock are distinct and are not given in any specific order.

Output

Output one line containing one word: `possible` if the clocks could be showing the same time, `impossible` otherwise.

16.3 Chapter Notes

rabin karp paper

KMP

hashing

17 Combinatorics

Combinatorics deals with various discrete structures, such as graphs and permutations. In this chapter, we will mainly study the branch of combinatorics known as *enumerative combinatorics* – the art of counting. We will count the number of ways to choose K different candies from N different candies, the number of distinct seating arrangements around a circular table, the sum of sizes of all subsets of a set and many more objects. Many combinatorial counting problems are based on a few standard techniques which we will learn in this chapter.

17.1 The Addition and Multiplication Principles

The ***addition principle*** states that, given a finite collection of **disjoint** sets S_1, S_2, \dots, S_n , we can compute the size of the union of all sets by simply adding up the sizes of our sets, i.e.

$$|S_1 \cup S_2 \cup \dots \cup S_n| = |S_1| + |S_2| + \dots + |S_n|$$

Example 17.1 Assume we have 5 different types of chocolate bars (the set C), 3 different types of bubble gum (the set G), and 4 different types of lollipops (the set L). These form three disjoint sets, meaning we can compute the total number of snacks by summing up the number of snacks of the different types. Thus, we have $|C| + |G| + |L| = 5 + 3 + 4 = 12$ different snacks.

Later on, we will see a generalization of the addition principle that handles cases where our sets are not disjoint.

The ***multiplication principle***, on the other hand, states that the size of the Cartesian product $S_1 \times S_2 \times \dots \times S_n$ equals the product of the individual sizes of these sets, i.e.

$$|S_1 \times S_2 \times \dots \times S_n| = |S_1| \cdot |S_2| \cdot \dots \cdot |S_n|$$

Example 17.2 Assume that we have the same sets of candies C , G and L as in Example 17.1. We want to compose an entire dinner out of snacks, by choosing one chocolate bar, one bubble gum and a lollipop. The multiplication principles tells us that, modeling a snack dinner as a tuple $(c, g, l) \in C \times G \times L$, we can form our dinner in $5 \cdot 3 \cdot 4 = 60$ ways.

The addition principle is often useful when we solve counting problems by case analysis.

Example 17.3 How many four letter words consisting of the letters a , b , c and d contain exactly two letters a ?

There are six possible ways to place the two letters a :

$aa_ _$
 $a_a_$
 $a_ _a$
 $_aa_$
 $_a_a$
 $_ _aa$

For each of these ways, there are four ways of choosing the other two letters (bb , bc , cb , cc). Thus, there are $4 + 4 + 4 + 4 + 4 + 4 = 6 \cdot 4 = 24$ such words.

Let us now apply these basic principle sto solve the following problem:

Kitchen Combinatorics

Northwestern Europe Regional Contest 2015 – Per Austrin

The world-renowned Swedish Chef is planning a gourmet three-course dinner for some muppets: a starter course, a main course, and a dessert. His famous Swedish cook-book offers a wide variety of choices for each of these three courses, though some of them do not go well together (for instance, you of course cannot serve chocolate moose and sooted shreemp at the same dinner).

Each potential dish has a list of ingredients. Each ingredient is in turn available from a few different brands. Each brand is of course unique in its own special way, so using a particular brand of an ingredient will always result in a

completely different dinner experience than using another brand of the same ingredient.

Some common ingredients such as *pølårber* may appear in two of the three chosen dishes, or in all three of them. When an ingredient is used in more than one of the three selected dishes, Swedish Chef will use the same brand of the ingredient in all of them.

While waiting for the meecaroo, Swedish Chef starts wondering: how many different dinner experiences are there that he could make, by different choices of dishes and brands for the ingredients?

Input

The input consists of:

- five integers r, s, m, d, n , where $1 \leq r \leq 1\,000$ is the number of different ingredients that exist, $1 \leq s, m, d \leq 25$ are the number of available starter dishes, main dishes, and desserts, respectively, and $0 \leq n \leq 2\,000$ is the number of pairs of dishes that do not go well together.
- r integers b_1, \dots, b_r , where $1 \leq b_i \leq 100$ is the number of different brands of ingredient i .
- $s + m + d$ dishes – the s starter dishes, then the m main dishes, then the d desserts. Each dish starts with an integer $1 \leq k \leq 20$ denoting the number of ingredients of the dish, and is followed by k distinct integers i_1, \dots, i_k , where for each $1 \leq j \leq k$, $1 \leq i_j \leq r$ is an ingredient.
- n pairs of incompatible dishes.

Output

If the number of different dinner experiences Swedish Chef can make is at most 10^{18} , then output that number. Otherwise, output “too many”.

The solution is a similar addition-multiplication principle combo as used in Example 17.3. First off, we can simplify the problem considerably by brute forcing over the coarsest component of a dinner experience, namely the courses included. Since there are at most 25 dishes of every type, we need to check up to $25^3 = 15\,625$ choices of dishes. By the addition principle, we can compute the number of dinner experiences for each such three-course dinner, and then sum them up to get the answer. Some pairs of dishes do not go well together. At this stage in the process we exclude any triple of dishes that include such a pair.

We can perform this check in $\Theta(1)$ time if we save the incompatible dishes in 2D boolean vectors, so that e.g. `badStarterMain[i][j]` determines if starter i is incompatible with main dish j .

For a given dinner course consisting of starter a , main dish b and dessert c , only the **set** of ingredients of three dishes matters since the chef will use the same brand for an ingredient even if it is part of two dishes. The next step is thus to compute this set by taking the union of ingredients for the three included dishes. This step takes $\Theta(k_a + k_b + k_c)$. Once this set is computed, the only remaining task is to choose a brand for each ingredient. Assigning brands is an ordinary application of the multiplication principle, where we multiply the number of brands available for each ingredient together.

17.2 Permutations

A **permutation** of a set S is an ordering of all the elements in the set. For example, the set $\{1, 2, 3\}$ has 6 permutations:

123	132
213	231
312	321

Our first “real” combinatorial problem will be to count the number of permutations of an n -element set S . When counting permutations, we use the multiplication principle. We will show a procedure that can be used to construct permutations one element at a time. Assume that the permutation is the sequence $\langle a_1, a_2, \dots, a_n \rangle$. The first element of the permutation, a_1 , can be assigned any of the n elements of S . Once this assignment has been made, we have $n - 1$ elements we can choose to be a_2 (any element of S except a_1). In general, when we are to select the $(i + 1)$ ’th value a_{i+1} of the permutation, i elements have already been included in the permutation, leaving $n - i$ options for a_{i+1} . Using this argument for all n elements of the sequence, we can construct a permutation in $n \cdot (n - 1) \cdots 2 \cdot 1$ ways (by the multiplication principle).

This number is so useful that it has its own name and notation. ■

Definition 17.1 — Factorial

The **factorial** of n , where n is a non-negative integer, denoted $n!$, is defined

as the product of the first n positive integers, i.e.

$$n! = 1 \cdot 2 \cdot \dots \cdot n = \prod_{i=1}^n i$$

For $n = 0$, we use the convention that the empty product is 1.

This sequence of numbers thus begin 1, 1, 2, 6, 24, 120, 720, 40 320, 362 880, 3 628 800, 39 916 800 for $n = 0, 1, 2, \dots, 11$. It is good to know the magnitudes of these numbers, since they are frequent in time complexities when doing brute force over permutations. Asymptotically, they grow as $n^{\Theta(n)}$. More precisely, the well-used *Stirling's formula*¹ gives the approximation

$$n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \left(1 + O\left(\frac{1}{n}\right)\right)$$

Exercise 17.1. In how many ways can 8 persons be seated around a round table, if we consider cyclic rotations of a seating to be different? What if we consider cyclic rotations to be equivalent?

Problem 17.1

n'th permutation – nthpermutation

Name That Permutation – namethatpermutation

Permutations as Bijections

The word *permutation* has roots in Latin, meaning “to change completely”. We are now going to look at permutations in a very different light, which gives some justification to the etymology of the word.

Given a set such as $[5]$, we can fix some ordering of its elements such as $\langle 1, 2, 3, 4, 5 \rangle$. A permutation $\pi = \langle 1, 3, 4, 5, 2 \rangle$ of this set can then be seen as a movement of these elements. Of course, this same movement can be applied to any other 5-element set with a fixed permutation, such as $\langle a, b, c, d, e \rangle$ being transformed to $\langle a, c, d, e, b \rangle$. This suggests that we can consider permutation as a “rule” which describes how to move – *permute* – the elements.

Such a movement rule can also be described as a function $\pi : [n] \rightarrow [n]$, where $\pi(i)$ describes what element should be placed at position i . Thus, the

¹Named after James Stirling (who has other important combinatorial objects named after him too), but stated already by his contemporary Abraham de Moivre.

permutation $\langle 1, 3, 4, 5, 2 \rangle$ would have $\pi(1) = 1$, $\pi(2) = 3$, $\pi(3) = 4$, $\pi(4) = 5$, $\pi(5) = 2$.

i	1	2	3	4	5
	↓	↓	↓	↓	↓
$\pi(i)$	1	3	4	5	2

Since each element is mapped to a different element, the function induced by a permutation is actually a bijection. By interpreting permutations as function, all the theory from functions apply to permutations too.

We call $\langle 1, 2, 3, 4, \dots, n \rangle$ the *identity permutation*, since the function given by the identity permutation is actually the identity function. As a function, we can also consider the composition of two permutations. Given two permutations, α and β , their composition $\alpha\beta$ is also a permutation, given by $\alpha\beta(k) = \alpha(\beta(k))$. If we let $\sigma = \langle 5, 4, 3, 2, 1 \rangle$ the composition with $\pi = \langle 1, 3, 4, 5, 2 \rangle$ from above would then be

i	1	2	3	4	5
	↓	↓	↓	↓	↓
$\pi(i)$	1	3	4	5	2
	↓	↓	↓	↓	↓
$\sigma\pi(i)$	5	3	2	1	4

This is called *multiplying* permutations, i.e. $\sigma\pi$ is the product of σ and π . If we multiply a permutation π by itself n times, we call the resulting product π^n .

An important property regarding the multiplication of permutations follows from their functional properties, namely their associativity. We have that the permutation $(\alpha\beta)\gamma = \alpha(\beta\gamma)$, so we will take the liberty of dropping the parentheses and writing $\alpha\beta\gamma$.

Problem 17.2

Permutation Product – permutation product

Permutations also have inverses, which are just the inverses of their functions. The permutation $\pi = \langle 1, 3, 4, 5, 2 \rangle$ which we looked at in the beginning thus have the inverse given by

$$\pi^{-1}(1) = 1 \quad \pi^{-1}(3) = 2 \quad \pi^{-1}(4) = 3 \quad \pi^{-1}(5) = 4 \quad \pi^{-1}(2) = 5$$

written in permutation notation as $\langle 1, 5, 2, 3, 4 \rangle$. Since this is the functional inverse, we expect $\pi^{-1}\pi = \text{id}$.

i	1	2	3	4	5
	↓	↓	↓	↓	↓
$\pi(i)$	1	3	4	5	2
	↓	↓	↓	↓	↓
$\pi^{-1}\pi(i)$	1	2	3	4	5

Problem 17.3

Permutation Inverse – permutationinverse

A related concept is that of the *cycle decomposition* of a permutation. If we start with an element i and repeatedly apply a permutation on this element (i.e. take $i, \pi(i), \pi(\pi(i)), \dots$) we will at some point find that $\pi^k(i) = i$, at which point we will start repeating ourselves.

i	1	2	3	4	5
	↓	↓	↓	↓	↓
$\pi(i)$	2	1	4	5	3
	↓	↓	↓	↓	↓
$\pi^2(i)$	1	2	5	3	4
	↓	↓	↓	↓	↓
$\pi^3(i)$	2	1	3	4	5
	↓	↓	↓	↓	↓
$\pi^4(i)$	1	2	4	5	3
	↓	↓	↓	↓	↓
$\pi^5(i)$	2	1	5	3	4
	↓	↓	↓	↓	↓
$\pi^6(i)$	1	2	3	4	5

We call the k distinct numbers of this sequence the *cycle* of i . For π , we have two cycles: $(1, 2)$ and $(3, 4, 5)$. Note how $\pi(1) = 2$ and $\pi(2) = 1$ for the first cycle, and $\pi(3) = 4, \pi(4) = 5, \pi(5) = 3$. It gives us an alternative way of writing it, namely as the concatenation of its cycles: $(1, 2)(3, 4, 5)$.

To compute the cycle decomposition of a permutation π , we repeatedly pick any element of the permutation which is currently not a part of a cycle, and compute the cycle it is in using the method described above. Since we will consider every element exactly once, this procedure is $\Theta(n)$ for n -element permutations.

Problem 17.4

Cycle Decomposition – cycledecomposition

Given a permutation π , we define its **order**, denoted $\text{ord } \pi$, as the size of the set $\{\pi, \pi^2, \pi^3, \dots\}$. For all permutations except for the identity permutation, this is the smallest integer $k > 0$ such that π^k is the identity permutation. In our example, we have that $\text{ord } \pi = 6$, since π^6 was the first power of π that was equal to the identity permutation. How can we quickly compute the order of π ?

The maximum possible order of a permutation happens to grow rather quickly (it is $e^{(1+o(1))\sqrt{n \log n}}$ in the number of elements n). Thus, trying to compute the order by computing π^k for every k until π^k is the identity permutation is too slow. Instead, we can use the cycle decomposition. If a permutation has a cycle (c_1, c_2, \dots, c_l) , we know that

$$\pi^l(c_1) = c_1, \pi^l(c_2) = c_2, \dots, \pi^l(c_l) = c_l$$

by the definition of the cycle composition. Additionally, this means that $(\pi^l)^k(c_1) = (\pi^{lk})(c_1) = c_1$. Hence, any power of π that is a multiple of l will act as the identity permutation *on this particular cycle*.

This fact gives us an upper bound on the order of π . If its cycle decomposition has cycles of length l_1, l_2, \dots, l_m , the smallest positive number that is the multiple of every l_i is $\text{lcm}(l_1, l_2, \dots, l_m)$. The permutation $\pi = \langle 2, 1, 4, 5, 3 \rangle$ had two cycles, one of length 2 and 3. Its order was $\text{lcm}(2, 3) = 2 \cdot 3 = 6$. This is also a lower bound on the order, a fact that uses the following fact which is left as an exercise:

Exercise 17.2. Prove that if π has a cycle of length l , we must have $l \mid \text{ord } \pi$.

Problem 17.5

Order of a Permutation – permutationorder

Dance Reconstruction

Nordic Collegiate Programming Contest 2013 – Lukáš Poláček

Marek loves dancing, got really excited when he heard about the coming wedding of his best friend Miroslav. For a whole month he worked on a special dance for the wedding. The dance was performed by N people and there were N marks on the floor. There was an arrow from each mark to another mark and every mark had exactly one incoming arrow. The arrow could be also pointing back to

the same mark.

At the wedding, every person first picked a mark on the floor and no 2 persons picked the same one. Every 10 seconds, there was a loud signal when all dancers had to move along the arrow on the floor to another mark. If an arrow was pointing back to the same mark, the person at the mark just stayed there and maybe did some improvised dance moves on the spot.

Another wedding is now coming up a year later, and Marek would like to do a similar dance. He found two photos from exactly when the dance started and when it ended. Marek also remembers that the signal was triggered K times during the time the song was played, so people moved K times along the arrows.

Given the two photos, can you help Marek reconstruct the arrows on the floor? On the two photos it can be seen for every person to which position he or she moved. Marek numbered the people in the first photo from 1 to N and then wrote the number of the person whose place they took in the second photo.

Marek's time is running out, so he is interested in any placement of arrows that could produce the two photos.

Input

Two integers $2 \leq N \leq 10\,000$ and $1 \leq K \leq 10^9$. Then, N integers $1 \leq a_1, \dots, a_N \leq N$, denoting that dancer number i ended up at the place of dancer number a_i . Every number between 1 and N appears exactly once in the sequence a_i .

Output

If it is impossible to find a placement of arrows such that the dance performed K times would produce the two photos, print "Impossible". Otherwise print N numbers on a line, the i 'th number denoting to which person the arrow leads from person number i .

The problem can be rephrased in terms of permutations. First of all, the dance corresponds to some permutation π of the dancers, given by where the arrows pointed. This is the permutation we seek in the problem. We are given the permutation a , so we seek a permutation π such that $\pi^K = a$.

When given permutation problems of this kind, we should probably attack it using cycle decompositions in some way. Since the cycles of π are all independent of each other under multiplication, it is a good guess that the decomposition can simplify the problem. The important question is then how a cycle of π is affected when taking powers. For example, a cycle of 10 elements

in π would decompose into two cycles of length 5 in π^2 , and five cycles of length 2 in π^5 . The general case involves the divisors of l and K :

Exercise 17.3. Prove that a cycle of length l in a permutation π decomposes into $\gcd(l, K)$ cycles of length $\frac{l}{\gcd(l, K)}$ in π^K .

This suggests our first simplification of the problem: to consider all cycles of π^K partitioned by their lengths. By Exercise 17.3, cycles of different lengths are completely unrelated in the cycle decomposition of π^K .

The result also gives us a way to “reverse” the decomposition that happens to the cycles of π . Given $\frac{l}{m}$ cycles of length m in π^K , we can combine them into a l -cycle in π in the case where $m \cdot \gcd(l, K) = l$. By looping over every possible cycle length l (from 1 to N), we can then find all possible ways to combine cycles of π^K into larger cycles of π . This step takes $\Theta(N \log(N + K))$ due to the GCD computation.

Given all the ways to combine cycles, a knapsack problem remains for each cycle length of π^K . If we have a cycles of length l in π^K , we want to partition them into sets of certain sizes (given by previous computation). This step takes $\Theta(a \cdot c)$ ways, if there are c ways to combine a -length cycles.

Once it has been decided what cycles are to be combined, only the act of computing a combination of them remains. This is not difficult on a conceptual level, but is a good practice to do on your own (the solution to Exercise 17.3 basically outlines the reverse procedure).

17.3 Ordered Subsets

A variation of the permutation counting problem is to count the number of ordered sequences containing *exactly* k distinct elements, from a set of n . We can compute this by first consider the permutations of the entire set of n elements, and then group together those whose k first elements are the same. Taking the

set $\{a, b, c, d\}$ as an example, it has the permutations:

abcd	bacd	cabd	dabc
abdc	badc	cadb	dacb
acbd	bcad	cbad	dbac
acdb	bcda	cbda	dbca
adbc	bdac	cdab	dcab
adcb	bdca	cdba	dcba

Once we have chosen the first k elements of a permutation, there are $(n - k)!$ ways to order the remaining $n - k$ elements. Thus, we must have divided our $n!$ permutations into one group for each ordered k -length sequence, with each group containing $(n - k)!$ elements. To get the correct total, this means there must be $\frac{n!}{(n-k)!}$ such groups – and k -length sequences.

We call these objects *ordered k -subsets* of an n -element set, and denote the number of such ordered sets by

$$P(n, k) = \frac{n!}{(n - k)!}$$

Note that this number can also be written as $n \cdot (n - 1) \cdots (n - k + 1)$, which hints at an alternative way of computing these numbers. We can perform the ordering and choosing of elements at the same time. The first element of our sequence can be any of the n elements of the set. The next element any but the first, leaving us with $n - 1$ choices, and so on. The difference to the permutation is that we stop after choosing the k 'th element, which we can do in $(n - k + 1)$ ways.

17.4 Binomial Coefficients

Finally, we are going to do away with the “ordered” part of the ordered k -subsets, and count the number of subsets of size k of an n -element size. This number is called the *binomial coefficient*, and is probably the most important combinatorial number there is.

To compute the number of k -subsets of a set of size n , we start with all the $P(n, k)$ *ordered* subsets. Any particular unordered k -subset can be ordered in

exactly $k!$ different ways. Hence, there must be $\frac{P(n,k)}{k!}$ unordered subsets, by the same grouping argument we used when determining $P(n, k)$ itself.

For example, consider again the ordered 2-subsets of the set $\{a, b, c, d\}$, of which there are 12.

ab	ba	ca	da
ac	bc	cb	db
ad	bd	cd	dc

The subset $\{a, b\}$ can be ordered in $2!$ ways - the ordered subsets ab and ba . Since each unordered subset is responsible for the same number of ordered subsets, we get the number of unordered subsets by dividing 12 with $2!$, giving us the 6 different 2-subsets of $\{a, b, c, d\}$.

ab		
ac	bc	
ad	bd	cd

Definition 17.2 — Binomial Coefficient

The number of k -subsets of an n -set is called the *binomial coefficient*

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

This is generally read as “ n choose k ”.

Note that

$$\binom{n}{k} = \frac{(n-k+1) \cdot (n-k+2) \cdots (n-1) \cdot n}{1 \cdot 2 \cdots (k-1) \cdot k}$$

They are thus the product of k numbers, divided by another k numbers. With this fact in mind, it does not seem unreasonable that they should be computable

in $O(k)$ time. Naively, one might try to compute them by first multiplying the k numbers in the nominator, then the k numbers in the denominator, and finally divide them.

Unfortunately, both of these numbers grow quickly. Indeed, already at 21! we have outgrown a 64-bit integer. Instead, we will compute the binomial coefficient by alternating multiplications and divisions. We will start with storing $1 = \frac{1}{1}$. Then, we multiply with $n - r + 1$ and divide with 1, leaving us with $\frac{n-r+1}{1}$. In the next step we multiply with $n - r + 2$ and divide with 2, having computed $\frac{(n-r+1) \cdot (n-r+2)}{1 \cdot 2}$. After doing this r times, we will be left with our binomial coefficient.

There is one big question mark from performing this procedure - why must our intermediate result always be integer? This must be true if our procedure is correct, or we will at some point perform an inexact integer division, leaving us with an incorrect intermediate quotient. If we study the partial results more closely, we see that they are binomial coefficients themselves, namely $\binom{n-r+1}{1}$, $\binom{n-r+2}{2}$, \dots , $\binom{n-1}{r-1}$, $\binom{n}{r}$. Certainly, these numbers must be integers. As we just showed, the binomial coefficients count things, and counting things tend to result in integers.

As a bonus, we discovered another useful identity in computing binomial coefficients:

$$\binom{n}{r} = \frac{n}{r} \binom{n-1}{r-1}$$

Exercise 17.4. Prove this identity combinatorially, by first multiplying both sides with r . (Hint: both sides count the number of ways to do the same two-choice process, but in different order.)

We have one more useful trick up our sleeves. Currently, if we want to compute e.g. $\binom{10^9}{1}$, we have to perform $10^9 - 1$ operations. To avoid this, we exploit a symmetry of the binomial coefficient. Assume we are working with subsets of some n -element set S . Then, we can define a bijection from the subsets of S onto itself by taking complements. Since a subset T and its complement $S \setminus T$ are disjoint, we have $|S \setminus T| = |S| - |T|$. This means that every 0-subset is mapped bijectively to every n -subset, every 1-subset to every $(n - 1)$ -subset, and every r -subset to every $(n - r)$ -subset.

However, if we can bijectively map r -subsets to $(n - r)$ -subsets, there must be equally many such subsets. Since there are $\binom{n}{r}$ subsets of the first kind and

$\binom{n}{n-r}$ subsets of the second kind, they must be equal:

$$\binom{n}{r} = \binom{n}{n-r}$$

More intuitively, our reasoning is basically “choosing what r elements to include in a set is the same as choosing what $n - r$ elements to exclude”. This is very useful in our example of computing $\binom{10^9}{10^9-1}$, since this equals $\binom{10^9}{1} = 10^9$. More generally, this enables us to compute binomial coefficients in $O(\min\{r, n-r\})$ instead of $O(r)$.

Problem 17.6

Binomial Coefficients – binomial

Sjecista**Croatian Olympiad in Informatics 2006/2007, Contest #2**

In a convex polygon with N sides, line segments are drawn between all pairs of vertices in the polygon, so that no three line segments intersect in the same point. Some pairs of these inner segments intersect, however.

For $N = 6$, this number is 15.




Figure 17.1: A polygon with 4 vertices.

Given N , determine how many pairs of segments intersect.

Input

The integer $3 \leq N \leq 100$.

Output

The number of pairs of segments that intersect.

The problem is a classical counting problem. If we compute the answer by

hand starting at $N = 0$, we get 0, 0, 0, 0, 1, 5, 15, 35. A quick lookup on OEIS² suggests that the answer is the binomial coefficient $\binom{N}{4}$. While this certainly is a legit strategy when solving problems on your own, this approach is usually not applicable at contests where access to the Internet tend to be restricted.




Figure 17.2: Four points taken from Figure 17.1.

Instead, let us find some kind of bijection between the objects we count (intersections of line segments) with something easier to count. This strategy is one of the basic principles of combinatorial counting. An intersection is defined by two line segments, of which there are $\binom{N}{2}$. Does every pair of segments intersect? In Figure 17.2, two segments (the solid segments) do not intersect. However, two other segments which together have the same four endpoints *do* intersect with each other. This suggests that line segments was the wrong level of abstraction when finding a bijection. On the other hand, if we choose a set of four points, the segments formed by the two diagonals in the convex quadrilateral given by those four points will intersect at some point (the dashed segments in Figure 17.2).

Conversely, any intersection of two segments give rise to such a quadrilateral – the one given by the four endpoints of the segments that intersect. Thus there exists a bijection between intersections and quadrilaterals, meaning that there must be an equal number of both. There are $\binom{N}{4}$ such choices of quadrilaterals, meaning there are also $\binom{N}{4}$ points of intersection.

Exercise 17.5. Prove that

- 1) $\binom{n}{k} = \binom{n-1}{k-1} + \binom{n-1}{k}$
- 2) $\sum_{k=0}^n \binom{n}{k} = 2^n$
- 3) $\sum_{k=0}^n (-1)^k \binom{n}{k} = 0$
- 4) $\sum_{k=0}^n \binom{n}{k} 2^k = 3^n$
- 5) $\sum_{k=0}^n \left[\binom{n}{k} \left(\sum_{l=0}^k \binom{k}{l} 2^l \right) \right] = 4^n$

²<https://oeis.org/A000332>

Dyck Paths

In a grid of width W and height H , we stand in the lower left corner at coordinates $(0, 0)$, wanting to venture to the upper right corner at (W, H) . To do this, we are only allowed two different moves – we can either move one unit north, from (x, y) to $(x, y + 1)$ or one unit east, to $(x + 1, y)$. Such a path is called a **Dyck path**.




Figure 17.3: A Dyck path on a grid of width 8 and height 5.

As is the this spirit of this chapter, we ask how many Dyck paths there are in a grid of size $W \times H$. The solution is based on two facts: a Dyck path consists of exactly $H + W$ moves, and exactly H of those should be northbound moves, and W eastbound. Conversely, any path consisting of exactly $H + W$ moves where exactly H of those are northbound moves is a Dyck path.

If we consider e.g. the Dyck path in Figure 17.3, we can write down the sequence of moves we made, with the symbol N for northbound moves and E for eastbound moves:

$EENENNEEENEEN$

Such a sequence must consist of all $H + W$ moves, with exactly H “N”-moves. There are exactly $\binom{H+W}{H}$ such sequences, since this is the number of ways we can choose the subset of positions which should contain the N moves.




Figure 17.4: The two options for the last possible move in a Dyck path.

If we look at Figure 17.3, we can find another way to arrive at the same

answer. Letting $D(W, H)$ be the number of Dyck paths in a $W \times H$ grid, some case work on the last move gives us the recurrence

$$D(W, H) = D(W - 1, H) + D(W, H - 1)$$

with base cases

$$D(0, H) = D(W, 0) = 1$$

We introduce a new function D' , defined by $D'(W + H, W) = D(W, H)$. This gives us the recurrence

$$D'(W + H, H) = D'(W - 1 + H, W - 1) + D'(W + H - 1, H - 1)$$

with base cases

$$D'(0, 0) = D'(H, H) = 0$$

These relations are satisfied by the binomial coefficients (Exercise 17.5).

Exercise 17.6. Prove that $\sum_{i=0}^n \binom{n}{i} \binom{n}{n-i} = \binom{2n}{n}$.

While Dyck paths sometimes do appear directly in problems, they are also a useful tool to find bijections to other objects.

Sums

In how many ways can the numbers $0 \leq a_1, a_2, \dots, a_k$ be chosen such that

$$\sum_{i=1}^k a_i = n$$

Input

The integers $0 \leq n \leq 10^6$ and $0 \leq k \leq 10^6$.

Output

Output the number of ways modulo $10^9 + 7$.

Given a Dyck path such as the one in Figure 17.3, what happens if we count the number of northbound steps we take at each x -coordinate? There are a total of $W + 1$ coordinates and H northbound steps, so we expect this to be a sum of $W + 1$ (non-negative) variables with a sum of H . This is indeed similar to what we are counting, and Figure 17.5 shows this connection explicitly.

This mapping gives us a bijective mapping between sums of k terms with a sum of n , to Dyck paths on a grid of size $(k - 1) \times n$. We already know how many such Dyck paths there are: $\binom{n+k-1}{n}$.




Figure 17.5: A nine-term sum as a Dyck path.

Catalan Numbers

A special case of the Dyck paths are the paths on a square grid that *do not cross the diagonal* of the grid. See Figure 17.6 for an example.




Figure 17.6: A valid path (left) and an invalid path (right).

We are now going to count the number of such paths, the most complex counting problem we have encountered so far. It turns out there is a straightforward bijection between the *invalid* Dyck paths, i.e. those who do cross the diagonal of the grid, to Dyck paths in a grid of different dimensions. In Figure 17.6, the right grid contained a path that cross the diagonal. If we take the part of the grid just after the first segment that crossed the diagonal and mirror it in the diagonal translated one unit upwards, we get the situation in Figure 17.7.

We claim that when mirroring the remainder of the path in this translated diagonal, we will get a new Dyck path on the grid of size $(n - 1) \times (n + 1)$. Assume that the first crossing is at the point (c, c) . Then, after taking one step up in order to cross the diagonal, the remaining path goes from $(c, c + 1)$ to (n, n) . This needs $n - c$ steps to the right and $n - c - 1$ steps up. When mirroring, this instead turns into $n - c - 1$ steps up and $n - c$ steps right. Continuing from $(c, c + 1)$,




Figure 17.7: Mirroring the part of the Dyck path after its first diagonal crossing.

the new path must thus end at $(c + (n - c - 1), c + 1 + (n - c)) = (n - 1, n + 1)$. This mapping is also bijective.

This bijection lets us count the number of paths that do cross the diagonal: they are $\binom{2n}{n+1}$. The numbers of paths that does not cross the diagonal is then $\binom{2n}{n} - \binom{2n}{n+1}$.

Definition 17.3 — Catalan Numbers

The number of Dyck paths in an $n \times n$ grid is called the n 'th *Catalan number*

$$C_n = \binom{2n}{n} - \binom{2n}{n+1} = \binom{2n}{n} - \frac{n}{n+1} \binom{2n}{n} = \frac{1}{n+1} \binom{2n}{n}$$

The first few Catalan numbers³ are 1, 1, 2, 5, 14, 42, 132, 429, 1430.

Problem 17.7

Catalan Numbers – catalan

Catalan numbers count many other objects, most notably the number of *balanced parentheses* expressions. A balanced parentheses expression is a string of $2n$ characters $s_1 s_2 \dots s_{2n}$ of letters (and), such that every prefix $s_1 s_2 \dots s_k$ contain *at least* as many letters (as). Given such a string, like $((())())$ we can interpret it as a Dyck path, where (is a step to the right, and) is a step upwards. Then, the condition that the string is balanced is that, for every partial Dyck path, we have taken at least as many right steps as we have taken up steps. This is equivalent to the Dyck path never crossing the diagonal, giving us a bijection between parentheses expressions and Dyck paths. The number of such parentheses expressions are thus also C_n .

³<https://oeis.org/A000108>

17.5 The Principle of Inclusion and Exclusion

Often, we wish to compute the size of the union of a collection of sets S_1, S_2, \dots, S_n , where these sets are not pairwise disjoint. For this, the *principle of inclusion and exclusion* was developed.




Figure 17.8: The union of two sets A and B .

Let us consider the most basic case of the principle, using two sets A and B . If we wish to compute the size of their union $|A \cup B|$, we *at least* need to count every element in A and every set in B , i.e. $|A| + |B|$. The problem with this formula is that whenever an element is in *both* A and B , we count it twice. Fortunately, this is easily mitigated: the number of elements in both sets equals $|A \cap B|$ (Figure 17.8). Thus, we see that $|A \cup B| = |A| + |B| - |A \cap B|$.

Similarly, we can determine a formula for the union of three sets $|A \cup B \cup C|$. We begin by including every element: $|A| + |B| + |C|$. Again, we have included the pairwise intersections too many times, so we remove those and get

$$|A| + |B| + |C| - |A \cap B| - |A \cap C| - |B \cap C|$$

This time, however, we are not done. While we have counted the elements which are in exactly one of the sets correctly (using the first three terms), and the elements which are in exactly two of the sets correctly (by removing the double-counting using the three latter terms), we currently do not count the elements which are in all three sets at all! Thus, we need to add them back, which gives us the final formula:

$$|A \cup B \cup C| = |A| + |B| + |C| - |A \cap B| - |A \cap C| - |B \cap C| + |A \cap B \cap C|$$

Exercise 17.7. Compute the number of integers between 1 and 1000 that are divisible by 2, 3 or 5.

From the two examples, you can probably guess formula in the general case, which we write in the following way:

$$\left| \bigcup_{i=1}^n S_i \right| = \sum_i |S_i| - \sum_{i < j} |S_i \cap S_j| + \sum_{i < j < k} |S_i \cap S_j \cap S_k| - \cdots + (-1)^{n+1} |S_1 \cap S_2 \cap \cdots \cap S_n|$$

From this formula, we see the reason behind the naming of the principle. We *include* every element, *exclude* the ones we double-counted, *include* the ones we removed too many times, and so on. The principle is based on a very important assumption – that it is easier to compute intersections of sets than their unions. Whenever this is the case, you might want to consider if the principle is applicable.

Derangements

Compute the number of permutations π of length N such that $\pi(i) \neq i$ for every $i = 1 \dots N$.

This is a typical application of the principle. We will use it on those sets of permutations where the condition is false for *at least* a particular index i . If we let these sets be D_i , the set of all permutations where the condition is false is $D_1 \cup D_2 \cup \cdots \cup D_N$. This means we seek $N! - |D_1 \cup \cdots \cup D_N|$. To apply the inclusion and exclusion formula, we must be able to compute the size of intersections of the subsets of D_i . This task is simplified greatly since the intersection of k such subsets is entirely symmetrical (it does not matter for which elements the condition is false, only the number).

If we want to compute the intersection of k such subsets, this means that there are k indices i where $\pi(i) = i$. There are $N - k$ other elements, which can be arranged in $(N - k)!$ ways, so the intersection of these sets have size $(N - k)!$. Since we can choose which k elements that should be fixed in $\binom{N}{k}$ ways, the term in the formula where we compute all k -way intersections will evaluate to $\binom{N}{k}(N - k)! = \frac{N!}{k!}$. Thus, the formula can be simplified to

$$\frac{N!}{1!} - \frac{N!}{2!} + \frac{N!}{3!} - \cdots$$

Subtracting this from $N!$ means that there are

$$N! \left(1 - 1 + \frac{1}{2!} - \frac{1}{3!} + \cdots \right)$$

This gives us a $\Theta(N)$ algorithm to compute the answer.

It is possible to simplify this further, using some insights from calculus. We have that

$$e^{-1} = 1 - 1 + \frac{1}{2} - \frac{1}{3} + \dots$$

Then, we expect that the answer should converge to $\frac{N!}{e}$. As it happens, the answer will *always* be $\frac{N!}{e}$ rounded to the nearest integer.

Exercise 17.8. 8 persons are to be seated around a circular table. The company is made up of 4 married couples, where the two members of a couple prefer not to be seated next to each other. How many possible seating arrangements are possible, assuming the cyclic rotations of an arrangement are considered equivalent?

17.6 The Pigeon Hole Principle

The Pigeon Hole Principle is a principle as intuitive as it is useful. Its statement is short and simple: if we place $N + 1$ different objects into N boxes, some box must contain at least 2 objects.

17.7 Invariants

Many problems deal with processes which consist of many steps. During such processes, we are often interested in certain properties that never change. We call such a property an *invariant*. For example, consider the binary search algorithm to find a value in a sorted array. During the execution of the algorithm, we maintain the invariant that the value we are searching for must be contained in some given segment of the array indexed by $[lo, hi)$ at any time. The fact that this property is invariant basically constitutes the entire proof of correctness of binary search. Invariants are tightly attached to greedy algorithms, and is a common tool used in proving correctness of various greedy algorithms. They are also one of the main tools in proving impossibility results (for example when to answer *NO* in decision problems).

Permutation Swaps

Given is a permutation a_i of $\langle 1, 2, \dots, N \rangle$. Can you perform exactly K swaps, i.e. exchanging pairs of elements of the permutation, to obtain the identity permutation $\langle 1, 2, \dots, N \rangle$?

Input

The first line of input contains the size of the permutation $1 \leq N \leq 100\,000$. The next line contains N integers separated, the permutation a_1, a_2, \dots, a_N .

Output

Output YES if it is possible, and NO if it is impossible.

First, we need to compute the minimum number of swaps needed.

Assume the cycle decomposition of the permutation consists of C cycles (see 17.2.1 for a reminder of this concept), with lengths b_1, b_2, \dots, b_C . Then, we need at least

$$S = \sum_{i=1}^C b_i - 1$$

swaps to return it to the identity permutation, a fact you will be asked to prove in the next section on monovariants. This gives us one necessary condition: $K \geq S$. However, this is not sufficient. A single additional condition is needed – that S and K have the same parity! To prove this, we will look at the number of *inversions* of a permutation, one of the common invariant properties of permutations.

Given a permutation a_i , we say that the pair (i, j) is an inversion if $i < j$, but $a_i > a_j$. Intuitively, it is the number of pairs of elements that are “out of place” in relation to each other.

3	5	2	1	4	6 inversions
1	5	2	3	4	3 inversions
1	5	3	2	4	4 inversions
1	5	3	4	2	5 inversions
1	2	3	4	5	0 inversions

Figure 17.9: The number of inversions for permutations differing only by a single swap.

If we look at Figure 17.9, where we started out with a permutation and performed a number of swaps (transforming it to the identity permutation), we can spot a simple invariant. The parity of the number of swaps and the number of inversions seems to always be the same. This characterization of a permutation is called *odd* and *even* permutations depending on whether the number of inversions is odd or even. Let us prove that this invariant actually

holds.

If this is the case, it is obvious why S and K must have the same parity. Since S is the number of swaps needed to transform the identity permutation to the given permutation, it must have the same parity as the number of inversions. By performing K swaps, K must have the same parity as the number of inversions. As K and S must have the same parity as the number of inversions, they must have the same parity as each other.

To see why these two conditions are sufficient, we can, after performing S swaps to obtain the identity permutation, simply swap two numbers with each other the remaining swaps. This can be done since $K - S$ will be an even number due to their equal parity.

17.8 Monovariants

Another similar tool (sometimes called a *monovariant*) instead define some kind of value $p(v)$ to the state v of each step of the process. We choose p such that it is strictly increasing or decreasing. They are mainly used to prove the finiteness of a process, in which either:

- The value function assume e.g. integer values, and is easily bounded in the direction of monotonicity (e.g. an increasing function would have an upper bound).
- The value function assume can assume any real, but there are only finitely many states the process can be in. In this case, the monovariant is used to prove that the process will never return to a previous state since this would contradict the monotonicity of p .

Let us begin with a famous problem of the first kind.

Majority Graph Bipartitioning

Given is a graph G . Find a bipartition of this graph into parts U and V , such that every vertex v has at most $\frac{|N(v)|}{2}$ neighbors in the same part as v itself.

Input

The first line of input contains integers $1 \leq V \leq 100$ and $0 \leq E \leq \frac{V(V-1)}{2}$ – the number of vertices and edges respectively. The next E lines contain two integers $0 \leq a \neq b < V$, the zero-indexed numbers of two vertices that are connected by an edge. No pair of vertices will have two edges.

Output

Output N integers, one for each vertex. The i 'th integer should be 1 or 2 if the i 'th vertex is in the first or the second part of the partition, respectively.

As an example, consider the valid and invalid partitionings in Figure 17.10. The vertices which does not fulfill the neighbor condition are marked in gray.




Figure 17.10: An invalid bipartitioning, where vertices B, D, G break the condition.

Problems generally considered greedy algorithms and pure monovariant problems usually differ in that the choice of next action usually has less thought behind it in the monovariant problems. We will often focus not on optimally short sequences of choices as we do with greedy algorithms, but merely finding *any* valid configuration. For example, in the problem above, one might try to construct a greedy algorithm based on for example the degrees of the vertices, which seems reasonable. However, it turns out there is not enough structure in the problem to find any simple greedy algorithm to solve the problem.

Instead, we will attempt to use the most common monovariant attack. Roughly, the process follows these steps:

1. Start with any arbitrary state s .
2. Look for some kind of modification to this state, which is possible if and only if the state is not admissible. Generally, the goal of this modification is to “fix” whatever makes the state inadmissible.
3. Prove that there is some value $p(s)$ that must decrease whenever such a modification is done.
4. Prove that this value cannot decrease infinitely many times.

Using these four rules, we prove the existence of an admissible state. If (and only if) s is not admissible, by step 2 we can perform some specified action on

it, which by step 3 will decrease the value $p(s)$. Step 4 usually follows from one of the two value functions discussed previously. Hence, by performing finitely many such actions, we must (by rule 4) reach a state where no such action is possible. This happens only when the state is admissible, meaning such a state must exist. The process might seem a bit abstract, but will become clear once we walk you through the bipartitioning step.

Our algorithm will work as follows. First, consider any bipartition of the graph. Assume that this graph does not fulfill the neighbor condition. Then, there must exist a vertex v which has more than $\frac{|N(v)|}{2}$ vertices in the same part as v itself. Whenever such a vertex exists, we move any of them to the other side of the partition. See Figure 17.11 of the this process.




Figure 17.11: Two iterations of the algorithm, which brings the graph to a valid state.

One question remains – why does this move guarantee a finite process? We now have a general framework to prove such things, which suggests that perhaps we should look for a value function $p(s)$ which is either strictly increasing or decreasing as we perform an action. By studying the algorithm in action in Figure 17.11 we might notice that more and more edges tend to go between the

two parts. In fact, this number never decreased in our example, and it turns out this is always the case.

If a vertex v has a neighbors in the same part, b neighbors in the other part, and violates the neighbor condition, this means that $a > b$. When we move v to the other part, the b edges from v to its neighbors in the other part will no longer be between the two parts, while the a edges to its neighbors in the same part will. This means the number of edges between the parts will change by $a - b > 0$. Thus, we can choose this as our value function. Since this is an integer function with the obvious upper bound of E , we complete step 4 of our proof technique and can thus conclude the final state must be admissible.

In mathematical problem solving, monovariants are usually used to prove that an admissible state exists. However, such problems are really algorithmic problems in disguise, since they actually provide an algorithm to construct such an admissible state.

Let us complete our study of monovariants, by also showing a problem using the second value function rule.

Water Pistols

N girls and N boys stand on a large field, with no line going through three different children.

Each girl is equipped with a water pistol, and wants to pick a boy to fire at. While the boys probably will not appreciate being drenched in water, at least the girls are a fair menace – they will only fire at a single boy each. Unfortunately, it may be the case that two girls choose which boys to fire at in such a way that the water from their pistols will cross at some point. If this happens, they will cancel each other out, never hitting their targets.

Help the girls choose which boys to fire at, in such a way that no two girls fire at the same boy, and the water fired by two girls will not cross.




Figure 17.12: An assignment where some beams intersect (left), and an assignment where no beams intersect (right).

Input

The first line contains the integer $N \leq 200$. The next N lines contain two real numbers $-10^6 \leq x, y \leq 10^6$, separated by a space. Each line is the coordinate (x, y) of a girl. The next and final N lines contain the coordinates of the boys, in the same format.

Output

Output N lines. The i 'th line should contain the zero-indexed number of the boy which the i 'th girl should fire at.

After seeing the solution to the previous problem, the solution should not come as a surprise. We start by randomly assigning the girls to one boy each, with no two girls shooting at the same boy. If this assignment contains two girls firing water beams which cross, we simply swap their targets.

Unless you are geometrically minded, it may be hard to figure out an appropriate value function. The naive value function of counting the current number of water beams crossing unfortunately fails – and might even increase after a move.

Instead, let us look closer at what happens when we switch the targets of two girls. In Figure 17.13, we see the before and after of such an example, as well as the two situations interposed. If we consider the sum of the two lengths of the water beams before the swap $((C + D) + (E + F))$ versus the lengths after the swap $(A + B)$, we see that the latter must be less than the first. Indeed, we have $A < C + D$ and $B < E + F$ by the triangle inequality, which by summing the two inequalities give the desired result. Thus the sum of all water beam lengths




Figure 17.13: Swapping the targets of two intersecting beams.

will decrease whenever we perform such a move. As students of algorithmics, we can make the additional note that this means the minimum-cost matching of the complete bipartite graph of girls and boys, with edges given as cost the distance between a particular girl and boy, is a valid assignment. If this was not the case, we would be able to swap two targets and decrease the cost of the matching, contradicting the assumption that it was minimum-cost. Thus, this rather mathematical proof actually ended up giving us a very simple reduction to min-cost matching.

Problem 17.8

Army Division – armydivision

Bread Sorting – breadsorting

17.9 Chapter Notes

18 Game Theory

In ordinary life, most of us are familiar with the concept of a game. We play video games, sports, card games, board games or any of the other many kinds of games. As algorithmists, we primarily focus on a subset of competitive games tegic aspects and well-defined rules, where determining who won is simple. Games such as chess, poker, tic-tac-toe or or yatzee belong to this category, unlike soccer (running humans and the behaviour of rolling balls are not sufficiently well-defined) or most real-time video games where reaction speed counts. The mathematical area analyzing this kind of game is called *game theory*.

The games we deal with in algorithmic problem solving is a small subset of this category. We are often given a position in some kind of abstract turn-based game, tasked to determine if the player at turn wins if both players play as good as possible. Players never make mistakes in the games we analyze. For examples, the game of tic-tac-toe is considered to be a drawn game, since perfect play in the game always results in a draw.

In this chapter, we learn some basic techniques for determining who wins a certain game. Occasionally a problem also asks us to construct an optimal strategy (for example by making the problem interactive and playing against us). This is often the case when it is “obvious” who wins the game, or at least when it’s very easy to guess a winner but harder to prove why. Proofs often present us with an optimal strategy, so we will generally aim to prove who wins even when one in for example a contest situation just would guess.

18.1 Mathematical Techniques

Before we dive into analysis of games that will require programming, we start with some of the more basic techniques that one might use to solve games given in mathematical rather than algorithmic problems. They are useful in algorithmic problem solving as well, while also serving as an introduction to the kind of games we try to solve.

Symmetry

When children first to play chess, an early attempted strategy is playing that of playing symmetrically. White moves its E pawn two squares forward, the child responds with the same, and so on. Of course, this is not a very good strategy – once white plays the winning move, it is very difficult for black to reply with the symmetric move. In variants of chess, this strategy works better.

Knight Packing

On a $n \times n$ chess board, two players alternate placing a *knight* on the board. A knight can only be placed if there is no other knight which would be either 1 row and 2 columns or 2 rows and 1 column away from it. The first player who cannot place a knight on the board loses. Given n ($1 \leq n \leq 10^9$), determine if the first or second player to move wins.

Solution. In all problems of this kind, where one is supposed to make a move by e.g. choosing a square to place something at, the first question should be – if my opponent makes the first move, can I pick a symmetric move that is always possible? This symmetry can manifest in several ways, such as mirroring a move along an axis, rotating it 180° around a center, or even taking the complement of a subset, where a move consists of choosing a subset of something. Games on grids are great candidates for a mirroring strategy using the first two transformations.

□

Coloring Game 1

A graph consists of n vertices and n edges, connecting the vertices into a single cycle of length n . Two players play the following game on this graph. Initially, all vertices are colored white. A move consists of coloring one of the white vertexes red or blue, where player 1 colors vertexes red, and player 2 colors them blue. A vertex can only be colored red or blue if neither it nor its two neighbours have that same color.

For a given n , if both players play optimally, who wins?

Solution. The first steps in most games should be to try and solve it for a few smaller instances, to see if a pattern emerges.

□

18.2 The Graph Game

Most of the games we study can be reduced to the following abstraction, which we will call the **Graph Game**. Given a directed graph G , a game piece is placed on one of the vertices (generally called *positions*) in the graph. Two players (♔ and ♚) alternate in making moves, where a move consists of taking the piece from its current position v and moving it along an outgoing edge from v . Whenever a player is unable to make such a move, i.e. when the outdegree of v is 0, that player loses the game (called a normal game)¹. We call these states *terminal positions*.




Figure 18.1: An example of a graph game with 6 positions. Player ♔ starts and has three possible moves A , B and C . ♚ chooses to move to A , whereupon ♔ responds with the only available move D . Finally, ♔ ends the game with the move E , leaving ♚ with no possible moves who therefore lose the game.

We will study three variants of the game:

- The Graph Game on acyclic graphs.

¹The opposite kind of game, where the player unable to make a move wins, is called a *Misère* game.

- The Graph Game on general graphs.
- The Graph Game where a position may only be visited once.

Problems relating to games often ask us whether a certain position v is a **winning position**. This means that, if the game piece is currently placed on v , the current player can always play in such a way that the other player will eventually be unable to move (thus losing). Conversely, a **losing position** is one where the current player will eventually lose if the other player moves optimally.

Acyclic Games

An acyclic game admit a simple classification of winning and losing positions.




Figure 18.2: The winning and losing positions of the game in Figure 18.1

A position is winning if and only if **at least one** of its neighbouring positions is a losing position, and losing if and only if **all** of its neighbouring positions are winning positions. This can be proved by a short proof by contradiction. If a position has a losing position as neighbour, we can ensure our opponent has a losing position. This means the opponent will lose if we play optimally. On the other hand, if all the neighbours are winning positions, we are forced to give our opponent a winning position, meaning we will lose if the opponent plays optimally.

```
1 bool isWinning(graph G, vertex current) {  
2     vector<vertex> moves = G.adj(current);  
3     for (vertex pos : moves) {  
4         if (!isWinning(pos)) return true;  
5     }  
6     return false;  
7 }
```


If we memoize this algorithm, it is linear in the number of edges of the graph.

General Games**Non-Repetitive Games**

19 Number Theory

Number theory is the study of certain properties of integers. It makes an occasional appearance within algorithmic problem solving, in the form of its subfield *computational number theory*. It is within number theory topics such as divisibility and prime numbers belong.

In competitions, number theory problems range from simple applications of the main theorems you learn in the chapter, to trickier tasks where you must combine hard number theoretical insights with other algorithmic techniques. The latter can involve required insights that are difficult mathematical problems themselves. This should not come as a surprise. Most content in this chapter is essentially about learning efficient methods of computing the standard number theoretical objects, such as primes, modular inverses, divisors, and becoming well acquainted with time complexities and other asymptotic approximations that tend to arise in number theoretical problems.

19.1 Divisibility

All of the number theory in this chapter relate to a single property of integers, divisibility.

Definition 19.1 — Divisibility

An integer n is **divisible** by an integer d if there exists an integer q such that $n = dq$. We then call d a **divisor** of n .

We use the notation $d \mid n$ to state that d is a divisor of n , and $d \nmid n$ when it is not.

Dividing both sides of the equality $n = dq$ with d gives us an almost equivalent definition, namely that $\frac{n}{d}$ is an integer. The difference is that the first definition admits the divisibility of 0 by 0, while the second one does not (zero division is undefined). When we speak of the divisors of a number in most contexts (as in Example 19.1), we generally consider only the non-negative divisors. Since d is a divisor of n if and only if $-d$ is a divisor of n , this sloppiness loses little information.

Example 19.1 — Divisors of 12

The number 12 has 6 divisors – 1 ($1 \cdot 12 = 12$), 2 ($2 \cdot 6 = 12$), 3 ($3 \cdot 4 = 12$), 4 ($4 \cdot 3 = 12$), 6 ($6 \cdot 2 = 12$) and 12 ($12 \cdot 1 = 12$).

12 is not divisible by e.g. 5 – we have $\frac{12}{5} = 2 + \frac{2}{5}$, which is clearly not an integer.

Exercise 19.1. Compute the divisors of 7, 18 and 39.

The concept of divisibility raises many questions. First and foremost – how do we check if a number is divisible by another? This question has one short and one long answer. For small numbers, i.e. those that fit inside the native integer types of a language, checking for divisibility is as simple as using the modulo operator (%) of your favorite programming language. An integer n is divisible by d if and only if $n \% d = 0$, since this means $\frac{n}{d}$ has no remainder and is therefore an integer.

For large numbers, checking divisibility is more difficult. Some programming languages, such as Java and Python, have built-in support for dealing with large integers, but e.g. C++ does not. In Section 19.4 on modular arithmetic, we discuss the implementation of the modulo operator on large integers.

Dual Divisibility

Given two positive integers a and b with the same number of digits ($1 \leq b \leq a \leq 10^{18}$), compute the number of divisors of a that have b as a divisor.

For example, with $a = 96$ and $b = 12$, there are 5 such numbers: 12, 24, 36, 48 and 96.

Solution. Assume that c is such a number. The solution falls out from some applications of the definition of divisibility. We have $a = cq$ and $c = bq'$ for some positive integers q, q' .

The value of q' is at most 9 by the following argument. If $q' \geq 10$, we have $a = cq \geq c \geq 10b$, but then a has more digits than b , a contradiction. Thus, we can simply test all the values of c by letting $q' = 1, 2, \dots, 9$ and verifying that the two conditions hold using the modulo operator. \square

Problem 19.1

Dual Divisibility – dualdivisibility

Evening Out 1 – eveningout1

Multiplication Table – multtable

Note: Solve it for 1/2 points.

Divisor Shuffle – divisorshuffle

Next, how do we compute all the divisors of an integer n ?

Divisors

Given an integer n , compute all the positive divisors of n .

Every integer has at least two particular divisors called the *trivial divisors*, namely 1 and n itself. If we exclude the divisor n , we get the *proper divisors*. To find the remaining divisors, we can use the fact that any divisor d of n must satisfy $|d| \leq |n|$ (why?). This means that we can limit ourselves to testing whether the integers between 1 and n are divisors of n , a $\Theta(n)$ algorithm. We can do a bit better though, by exploiting a nice symmetry.

Hidden in Example 19.1 lies the key insight to speeding this up. It seems that whenever we had a divisor d , we were immediately given another divisor q . For example, when claiming 3 was a divisor of 12 since $3 \cdot 4 = 12$, we found another divisor, 4. This is not a surprise, given that the definition of divisibility (Definition 19.1) – the existence of the integer q in $n = dq$ – is symmetric in d and q , meaning divisors come in pairs $(d, \frac{n}{d})$.

Exercise 19.2. Prove that a positive integer has an odd number of divisors if and only if it is a perfect square.

Since divisors come in pairs, we can limit ourselves to finding one member of each such pair. Furthermore, one of the elements in each such pair must be bounded by \sqrt{n} . Otherwise, we would have that $n = d \cdot \frac{n}{d} > \sqrt{n} \cdot \sqrt{n} = n$, a contradiction (again, 0 is a special case here where we always have $\frac{0}{d} = 0$). This limit helps us reduce the time it takes to find the divisors of a number to $\Theta(\sqrt{n})$, which allows us to solve the problem sufficiently fast.

```

1: procedure DIVISORS( $N$ )
2:    $divisors \leftarrow$  new list
3:   for  $i$  from 1 up to  $i^2 \leq N$  do
4:     if  $N \bmod i = 0$  then
5:        $divisors.add(i)$ 
6:       if  $i \neq N/i$  then
7:          $divisors.add(\frac{N}{i})$ 

```

8: **return** *divisors*

Problem 19.2

Divisors – divisors

Let us look at an application of this algorithm.

Subcommittees

In a parliament of $P \leq 10^{16}$ people, the speaker wants to divide the parliament into (at least two) disjoint subcommittees of equal size. Of course, the chair of such a subcommittee furthermore wants to divide their subcommittee into (at least two) subsubcommittees of equal size, and so on, until no further divisions can be performed.

What is the maximum number of levels of subcommittees can be created?

Solution. What different sizes may the first level of subcommittees have? Well, if we perform a split into groups of size k , we get $\frac{P}{k}$ such groups. Of course, this must be an integer – i.e. k must be a divisor of P . This means we are looking for a sequence of numbers $c_0, c_1, c_2, \dots, c_n$ such that $c_0 = P$, $c_{i+1} \mid c_i$ and $c_n = 1$.

A simple solution would be to generate all divisors of P (the possible values of c_1), attempt a split into those group sizes, and then recursively solve the problem for them. However, this would be too slow. As an example, if we take $P = 8\,086\,598\,962\,041\,600$, the sum of the square roots of its divisors is $6\,636\,882\,083$, so even finding only the ways to split the parliament into 2-level committees would be too expensive.

Instead, we will use the following lemma:

Exercise 19.3. Prove that divisibility is a transitive relationship; if $b \mid a$ and $c \mid b$, then $c \mid a$.

This means that the possible values of c_i , i.e. the transitive closure of divisibility of P , are the divisors of P . Essentially, we are looking for the longest sequence of divisors of P such that each divisor is also a divisor of the previous divisor. By constructing the directed graph of all the divisors a with edges from a to its own divisors, we reduce the problem to finding the longest path in a DAG. Unfortunately, this too is slow – the above P has 41 472 divisors, leaving us with about $\binom{41472}{2} = 859\,942\,656$ modulo operations to construct the graph.

What if we instead try to look at the entire sequence at once? We have $P = c_0$, $c_0 = c_1 q_1$, $c_1 = c_2 q_2$, \dots , $c_{n-1} = c_n q_n$, $c_n = 1$ where $q_i > 1$. Inserting every equation into the previous one gives us $P = q_1 q_2 \cdots q_n$. Conversely, for each such choice of q_i , we can construct a valid sequence of c_i . Note that all q_i have only trivial divisors, or we could replace it with two numbers and still preserve the product, yielding a longer sequence.

Now comes the key insight – we can WLOG choose some q_i to be the smallest divisor of P – let us call this k . Pick an i such that $k \nmid Q = q_1 q_2 \cdots q_{i-1}$ and $k \mid Q q_i$ (such an i exists since $k \mid P$ and $k \nmid 1$). By the following theorem, $q_i = k$

Theorem 19.1 – Euclid’s Lemma

If $p > 1$ and $b > 1$ have only trivial divisors, $p \mid ab$ and $p \nmid a$, then $p = b$.

Proof. We prove a slightly stronger statement (called Euclid’s Lemma) instead; if $p > 1$ have only trivial divisors and $p \mid ab$, then either $p \mid a$ or $p \mid b$. This implies the original statement, since if $p \mid b$ and b only have trivial divisors, $p = 1$ or $p = b$ (but $p > 1$).

Consider the smallest p for which there exists a (smallest) a for which there exists a b where the theorem is false. We now prove that this minimal counterexample gives rise to an even smaller counterexample.

First, a lacks non-trivial divisors. Otherwise, we can pick n, m such that $a = nm$ where $0 < n \leq m < a$. Substitution gives us $p \mid (nm)b = n(mb)$. Since $n < a$, we have either $p \mid n$ or $p \mid mb$, since we otherwise find a smaller counterexample. We know that $p \nmid n$. Otherwise, as $n \mid a$, we get $p \mid a$. Therefore, $p \mid mb$. Because $m < a$, we again find that $p \mid m$ or $p \mid b$ (or we have a smaller counterexample). As $p \nmid b$ (by assumption), we get $p \mid m$. Again, as $m \mid a$ we get the corresponding contradiction $p \mid a$.

Next, we WLOG assume $a < p$. Otherwise, consider $c = a - p > 0$. Since $p \mid ab$, we have $p \mid ab - pb = cb$ (see Exercise 19.4). Thus, $p \mid c$ or $p \mid b$ (c would otherwise be a smaller counterexample). We have assumed $p \nmid b$, so $p \mid c$, i.e. $p \mid a - p$. By the same exercise, this implies $p \mid a$, contrary to our assumption.

Finally, let n be such that $pn = ab$. Since a is smaller than p and have

only trivial divisors, we have that $a \mid pn$ implies $a \mid p$ or $a \mid n$. As p lacks non-trivial divisors, the latter must be true. This means there exists m such that $n = ma$. Inserting this gives us that $ab = pma$, or $b = pm$. But this means $p \mid b$, a contradiction.

Since the assumption of a smallest counterexample only lead to contradictions, we find that that no such counterexample exists, meaning the theorem is true. \square

Exercise 19.4. Prove that if $a \mid b$ and $a \mid c$ then $a \mid b + c$.

With this, we are nearly there. The sequence q_i is independent of order – we can let $q_1 = k$. This means choosing the largest possible divisor of P as c_1 . How do we choose the remaining ones? Well, for c_2 , the same argument says we should choose greatest possible divisor of $\frac{P}{c_1}$, and so on. How do we find it? Since it must also be a divisor of P , we can iterate through the smaller divisors (in descending order) and pick the first one that was also a divisor of $\frac{P}{c_1}$. Eventually, we reach $c_n = 1$.

A solution could look something like the following.

```

1: procedure SUBCOMMITTEES( $P$ )
2:   list divisors  $\leftarrow$  Divisors( $P$ )
3:   sort divisors in descending order
4:    $ans \leftarrow 0$ 
5:   for each  $d$  in divisors do
6:     if  $d$  divides  $P$  then
7:        $ans \leftarrow ans + 1$ 
8:        $P \leftarrow d$ 
9:   return  $ans$ 

```

This solution only requires computing and iterating through all divisors of P , giving us a $\Theta(\sqrt{P})$ solution. \square

Problem 19.3

Subcommittees – subcommittees

Evening Out 2 – eveningout2

Multiplication Table – multtable

Note: Solve for 2 points.

This result that divisors comes in pairs happens to give us some help in answering our next question, regarding the plurality of divisors. The above result gives us an upper bound of $2\sqrt{n}$ divisors of an integer n . We can do a little better, with $\approx n^{\frac{1}{3}}$ being a commonly used approximation for the number of divisors when dealing with integers which fit in the native integer types.¹ For example, the maximal number of divisors of a number less than 10^3 is 32, 10^6 is 240, 10^9 is 1 344, 10^{18} is 103 680.²

A bound we will find more useful when solving problems concerns the *average* number of divisors of the integers between 1 and n .

Theorem 19.2 — Average Number of Divisors

Let $d(i)$ be the number of divisors of i . Then,

$$\sum_{i=1}^n d(i) = \Theta(n \ln n)$$

Proof. There are between $\frac{n-i+1}{i}$ and $\frac{n}{i}$ integers between 1 and n divisible by i , since every i 'th integer is divisible by i . Thus, the number of divisors of all those integers is bounded by

$$\sum_{j=1}^n \frac{n}{j} = n \sum_{j=1}^n \frac{1}{j} = O(n \ln n)$$

from above and

$$\sum_{j=1}^n \frac{n-j+1}{j} = n \sum_{j=1}^n \frac{1}{j} - n + \sum_{j=1}^n \frac{1}{j} \geq n \ln n - n + \ln n = \Omega(n \ln n)$$

from below. □

This proof also suggest a way to compute the divisors of all the integers 1, 2, ..., N .

¹In reality, the maximal number of divisors of the interval $[1, n]$ grows sub-polynomially, i.e., as $O(n^\epsilon)$ for every $\epsilon > 0$.

²Sequence A066150 from OEIS: <http://oeis.org/A066150>.

Divisor Counts

For every integer between 1 and N , compute the number of positive divisors it has.

Solving the problem with the previous algorithm, computing the divisors for every single integer, would yield a $\Theta(N\sqrt{N})$ algorithm. Instead, we inverse the problem. For each integer i , we find all the numbers divisible by i (in $\Theta(\frac{N}{i})$ time), which are $0i, 1i, 2i, \dots, \lfloor \frac{N}{i} \rfloor i$. In total, this takes $\Theta(N \ln N)$ time, a quite decent improvement.

Problem 19.4

Divisor Counts – divisorcounts

Organizator – organizer

This technique, called *sieving* is a quite common number theoretical method – we revisit it in the next section.

19.2 Prime Numbers

When talking about divisibility, we regularly used it as a tool to describe factorizations of an integer in various ways. For example, given the number 12, we could factor it as $2 \cdot 6$, or $3 \cdot 4$, or even further into $2 \cdot 2 \cdot 3$. This last factorization is special, in that no matter how hard we try, it cannot be factored further since 2 and 3 lack non-trivial divisors. It consists only of factors that are *prime numbers*.

Definition 19.2 – Prime Number

An integer $p \geq 2$ is called a **prime number** if its only positive divisors are 1 and p .

The integers $n \geq 2$ that are *not* prime numbers are called **composite numbers**.

Note that 1 is neither prime nor composite.

Example 19.2 The first 10 prime numbers are 2, 3, 5, 7, 11, 13, 17, 19, 23, 29. Which is the next one?

Problem 19.5

Longest Prime Sum – longprimesum

Shortest Composite Sum – shortcompositesum

Let us start with the simple questions – how do we determine if a number is a prime? Using the knowledge from the previous section, this is simple. A number is prime if it has only trivial divisors, so we can use the same approach that we used when computing all the divisors of a number to see if it has *any* divisor instead in $O(\sqrt{N})$ time.

Problem 19.6

Primality – primality

Note: Solve for 1 point.

In a similar manner, we can extend the algorithm used to count all divisors of all numbers up to some limit to counting primes up to a limit. Instead of counting divisors, we simply mark those numbers that had non-trivial divisors as non-prime. The running time is the same, $\Theta(N \ln N)$.

Problem 19.7

Prime Count – primecount

Note: Solve for 2 points.

Exercise 19.5. Let $\pi(N)$ be the number of primes up to N . Given a list of those primes, show how to determine the primality of any integer up to N^2 in $O(\pi(N)^2)$ time.

One might wonder if the algorithm in Exercise 19.5 is faster than testing all possible divisors. To answer this, we need to know more about the number of primes.

There are an infinite number of primes. This can be proven by a simple proof by contradiction. If p_1, p_2, \dots, p_q are the only primes, then $P = p_1 p_2 \dots p_q + 1$ is not divisible by any prime number (and by extension has no divisors but the trivial ones), so it is not composite. However, P is larger than any prime, so it is not a prime number either, a contradiction.

More relevant is instead the density of primes, since that is what determines how $\pi(N)$ relates to N .

Theorem 19.3 — Prime Number Theorem

The density of prime numbers in the interval $[1, N]$ is $\approx \frac{1}{\ln N}$ for large N .

The proof requires a lot of deep number theory, so we will not show it here. This means that precomputing primes and then using that list to check primality only gains you a logarithmic factor. More specifically, the number of primes below 10^3 is 168, below 10^6 is 78 498, below 10^9 is approximately $51 \cdot 10^6$.

Based on the Prime Number Theorem, one might have the reasonable thought that prime numbers shouldn't be that far apart. The Prime Number Theorem states that the average distance between primes is $\frac{1}{\ln N}$, but of course the maximum gap may be longer. For all integers up to 10^9 , the maximum gap is 282, and for 10^{18} it is 1442.

A very trivial upper bound on the gaps that is occasionally useful is the following one.

Theorem 19.4 — Bertrand's Postulate

For all $n \geq 2$, there exists a prime p where $n < p < 2n$.

Prime Time

Jon Marius Venstad, Nordic Collegiate Programming Contest 2011, CC BY-SA 3.0
Odd, Even and Ingmaria are playing a game. They start with an arbitrary positive integer and take turns either adding 1 or dividing by a prime (assuming the result is still an integer). Once they reach 1, they each gets points corresponding to the smallest of the numbers their move resulted in. If a player could make no move, their score is instead equal to the starting integer. They all play such that they minimize their own moves. If several possible moves would result in the same score for the player, they would choose to produce the lowest number that they can. They play in the order Odd \rightarrow Even \rightarrow Ingmaria $\rightarrow \dots$, but alternate who starts the round.

Given a list of starting integers of the rounds they played of the game, determine the final scores of the players.

Input

The first line contains n ($1 \leq n \leq 1000$), the number of rounds of the game. The next n lines contains the name of the starting player of the round and the starting integer (between 1 and 10 000).

Output

Output the final scores of the three players.

Note: The problem statement has been shortened.

Solution. The game theoretic solution of the problem would be to construct the graph of all the integers with edges between possible transitions. One could then compute the score a player would get for moving to a certain integer in the graph in the way described in Section 18.2. Unfortunately, the game graph contains possible loops such as $2 \rightarrow 3 \rightarrow 4 \rightarrow 2$. We can eliminate those loops with some additional insights using the tie-breaking rules the players use when picking moves.

Let us investigate the behaviour of the players more closely. If a player is presented with a prime number, they clearly will divide with it to end the game and get 1 point. In other cases, they may either add 1 or divide away a prime. This means we never want to move from a prime p to $p + 1$.

It turns out that removing these moves makes the game acyclic. Assume to the contrary that we currently are at a number a between two consecutive primes p_k and p_{k+1} and have a sequence of moves that takes us back to a . The next sequence of moves will be to add 1 until we hit an integer $p_k < b \leq p_{k+1}$ (remember we never want to go to $p_{k+1} + 1$) and dividing some prime p_i away. But $\frac{b}{p_i} \leq \frac{b}{2} \leq \frac{p_{k+1}}{2}$. By Bertrand's Postulate, $\frac{p_{k+1}}{2} < p_k$, so the new result will be less than p_k . However, since we are not allowed to make the transition $p_k \rightarrow p_k + 1$ we can never reach a again.

What remains is to compute the transitions from each integer $1 \leq a \leq N$. By precomputing all the primes up to 10 000, we can afford to test whether all of them are divisors of each a . In total, this is on the order of $\frac{10^4 \cdot 10^4}{\ln 10^4} \approx 10^7$ edges, which is a reasonable number. \square

Since the prime numbers have no other divisors besides the trivial ones, a factorization consisting only of prime numbers is special.

Definition 19.3 — Prime Factorization

The **prime factorization** of a positive integer n is a factorization of the form

$$p_1^{e_1} \cdot p_2^{e_2} \cdot \cdots \cdot p_k^{e_k}$$

where p_i are all distinct primes.

We define

$$\omega(n) = k$$

to be the number of primes dividing n , and

$$\Omega(n) = \sum_{i=1}^k e_i$$

to be the number of primes with multiplicity in the product of n .

Example 19.3 The prime factorization of 228 is $2 \cdot 2 \cdot 3 \cdot 19$.

Note that in the definition, we spoke of *the* prime factorization. It turns out that this factorization is indeed unique, except for a reordering of the p_i . It may be “intuitively obvious” that this is the case, but that is misguided. A proof can be constructed using Euclid’s Lemma from the previous section (p. 307).

Theorem 19.5 — Existence and Uniqueness of Prime Factorizations

There exists a unique factorization of a positive integer N into prime numbers.

Proof. First, the existence part, through a simple proof by induction. Assume that all integers up to $N - 1$ has a prime factorization. If N is a prime, then N is a prime factorization of itself. Otherwise, it has a non-trivial divisor, so we can write $N = ab$ with $1 < a \leq b < N$. By the induction hypothesis, a and b have prime factorizations. Concatenating the two factorizations gives us a prime factorization of N . Thus, by induction, all positive integers have prime factorizations.

Next, the uniqueness. We prove this too by induction. Our base case is $N = 1$ which has the empty product as unique prime factorization. Assume that N has two distinct prime factorizations $N = p_1 p_2 \cdots p_k = q_1 q_2 \cdots q_l$, but all integers up to $N - 1$ has only one. Consider the prime p_k . Since it divides the left side, it must also divide the right side. Let i be such that

$Q = q_1 \cdots q_{i-1}$ is not divisible by p_k , but Qq_i is. Such an i exists since when $i = 1$, the product Q is 1 (the empty product) which is not divisible by p_1 , and with $i = l$ we get $Q = N$ which is divisible by p_1 .

Then, by our version of Euclid's Lemma, since p_1 and q_i are primes, $p_1 \mid Qq_i$ but $p_1 \nmid Q$, we have $p_1 = q_i$. WLOG we can reorder the factors q and assume $q_i = q_1$. If we divide away this factor, we get that $p_2 \cdots p_k = q_2 \cdots q_l$ are both prime factorizations of $\frac{N}{p_1} < N$, so by the induction hypothesis they are the same. That means the original prime factorizations were also the same, a contradiction. Thus, N too has a unique prime factorization, completing the proof. \square

Factorization

Given an integer N , compute its prime factorization.

The simplest solution is to extend the algorithm the method used to test primality. An integer N can have at most one prime in its factorization that exceeds \sqrt{N} , since their product otherwise would exceed N . Looping over all possible prime divisors up to \sqrt{N} and factoring them out from N is thus sure to find all prime factors, except for possibly a single one that was larger than \sqrt{N} . The algorithm is called *trial division*.

```

1: procedure FACTOR( $N$ )
2:    $primes \leftarrow$  new list
3:   for  $i$  from 2 up to  $i^2 \leq N$  do
4:     while  $N \bmod i = 0$  do
5:        $primes.add(i)$ 
6:        $N \leftarrow \frac{N}{i}$ 
7:   if  $N \neq 1$  then
8:      $primes.add(N)$             $\triangleright N$  may have had a single prime factor  $> \sqrt{N}$ 
9:   return  $primes$ 

```

Exercise 19.6. In Algorithm ??, N is being modified in the loop when a new prime is found. Is it a problem to use the new, updated N in the $i^2 \leq N$ check in the loop? Is the time complexity the same?

Problem 19.8

Factorization – factorization

Note: Solve for 1 point.

For the next problem, we define some useful notation:

Definition 19.4 For integers a, b, k , we say that a^k **divides** b **exactly** if $a^k \mid b$ but $a^{k+1} \nmid b$. We use the notation $a^k \parallel b$ for this.

Factorial Power

Given integers n and m ($1 \leq n, m \leq 10^{14}$), determine the k for which $n^k \parallel m!$.

Solution. To start, we must first connect the prime factorization with divisibility. If n has the prime factorization $n = p_1^{e_1} \cdot p_2^{e_2} \cdots p_l^{e_l}$, a divisor of n must be of the form $d = p_1^{e'_1} \cdot p_2^{e'_2} \cdots p_l^{e'_l}$, where $0 \leq e'_i \leq e_i$. This can be proven using the uniqueness of the prime factorization, and the fact that $n = dq$ for some integer q . Any number of this form is also a divisor of n .

The exponent laws gives us that $n^k = p_1^{ke_1} \cdot p_2^{ke_2} \cdots p_l^{ke_l}$, so we are looking for the largest k such that ke_i does not exceed the power of p_i in $m!$. Thus, after factoring n the problem is reduced to determining how many times all the p_i divides $m!$. This equals

$$\left\lfloor \frac{m}{p_i} \right\rfloor + \left\lfloor \frac{m}{p_i^2} \right\rfloor + \left\lfloor \frac{m}{p_i^3} \right\rfloor + \dots$$

Exercise 19.7. Prove the above formula.

Since n can have at most $\log_2(n)$ prime factors, and all terms after the $\log_2(m)$ first ones are zero, the complexity of the computation is $O(\sqrt{n} + \log(n) \log(m))$. \square

Problem 19.9

Factorial Power – factorialpower

Perfect Pth Powers – perfectpowers

Divisor Guessing Game – divisorguessing

For the next problem, we will show a slightly faster version of the prime sieve we've previously seen, using it to factor all integers in an interval.

Product Divisors

Given a sequence of integers a_1, a_2, \dots, a_n , compute the number of divisors of $A = \prod_{i=1}^n a_i$.

Input

The length of the sequence $1 \leq n \leq 1\,000\,000$, and the sequence $1 \leq a_1, \dots, a_n \leq 10^6$.

Output

The number of divisors of n modulo $10^9 + 7$.

Solution. Again, we use the prime factorization interpretation of divisors. Let $A = \prod_{i=1}^k p_i^{e_i}$ where p_i are distinct primes. We can now use a simple combinatorial argument to count the number of divisors of A . A divisor of A is of the form $\prod_{i=1}^k p_i^{e'_i}$ where $e'_i \leq e_i$. All the e'_i can take any integer value between 0 and e_i to fulfill this condition. This gives us $e_i + 1$ choices for the value of e'_i . Since each e'_i is independent, there are a total of $(e_1 + 1)(e_2 + 1) \dots (e_k + 1)$ numbers of this form, and thus divisors of A by the multiplication principle.

We are left with the problem of determining the prime factorization of A . Essentially, this is tantamount to computing the prime factorization of *every* integer between 1 and 10^6 , since we could have $a_i = i$ for $i = 1 \dots 10^6$. Once this is done, we can go through the sequence a_i and tally up all primes in their factorization. Since an integer m has at most $\log_2 m$ prime factors, this step is bounded by approximately $n \log_2 10^6$ operations. Then, how do we factor all integers in $[1..10^6]$? We could obviously adapt the algorithm we used to count primes before, but we will now improve it a bit. The general idea was to loop over multiples of all numbers, and mark them as non-prime. When dealing with primes however, we only care about doing this for *primes*. After all, any non-prime must have a prime divisor, so we don't lose correctness by only sieving on primes. With this improvement, the sieve is called the Sieve of Eratosthenes.

This results in the following solution:

```
1: procedure PRODUCTDIVISORS(sequence A)
2:   counts  $\leftarrow$  new list[ $10^6 + 1$ ]
3:   for each  $a$  in  $A$  do
4:     counts[ $a$ ]  $\leftarrow$  counts[ $a$ ] + 1
```

```

5:  isPrime  $\leftarrow$  new list[ $10^6 + 1$ ] of true
6:  exponents  $\leftarrow$  new map
7:  for i from 2 to  $i^2 \leq N$  do
8:      if isPrime[i] then
9:          for j  $\in \{2i, 3i, \dots\}$  up to  $N - 1$  do
10:             if tmpj mod i = 0 then
11:                 isPrime[j]  $\leftarrow$  false
12:                 while tmpj mod i = 0 do
13:                     exponents[i] = exponents[i] + counts[j]
14:                     tmpj = tmpj / i
15:  ans  $\leftarrow$  1
16:  for each value in exponents do
17:      ans  $\leftarrow$  (ans · (value + 1)) mod ( $10^9 + 7$ )
18:  return ans
    
```

The complexity of this solution is a bit tricky to analyze. The important part of the sieve is the inner loop, which computes the actual factors. Let us count the number of times a prime p is pushed in this loop. First of all, every p 'th integer is divisible by p , which totals $\frac{n}{p}$ iterations. However, every p^2 'th integer is divisible by p yet again, contributing an additional $\frac{n}{p^2}$ iterations, and so on. Summing this over every p which is used in the sieve gives us the bound

$$\sum_{p \leq \sqrt{n}} \left(\frac{n}{p} + \frac{n}{p^2} + \frac{n}{p^3} + \dots \right) = n \sum_{p \leq \sqrt{n}} \left(\frac{1}{p} + \frac{1}{p^2} + \frac{1}{p^3} + \dots \right)$$

Using the formula for the sum of a geometric series ($\frac{1}{p} + \frac{1}{p^2} + \dots = \frac{p}{p-1}$) gives us the simplification

$$n \sum_{p \leq \sqrt{n}} \frac{1}{p-1} = \Theta \left(n \sum_{p \leq \sqrt{n}} \frac{1}{p} \right)$$

It is known that $\sum_{p \leq n} \frac{1}{p} = O(\ln \ln n)$. With this, the final complexity becomes a simple $O(n \ln \ln \sqrt{n}) = O(n \ln \ln n)$. \square

Competitive Tip

When using the Sieve of Eratosthenes, we can save quite a bit of memory by using a bitset instead since we only store a boolean state per number (whether it is prime or not). This gives us slightly better cache behaviour, improving the performance in real terms.

Competitive Tip**Problem 19.10**

Prime Count – primecount

Note: Solve for 3 points.

Problem 19.11

19.3 The Euclidean Algorithm

The Euclidean algorithm is one of the oldest known algorithms, dating back to Greek mathematician Euclid who wrote of it in his mathematical treatise *Elements*. It regards those numbers which are divisors of two different integers, and extends into integer equations of the form $ax + by = c$.

Definition 19.5 – Greatest Common Divisor

We call an integer d dividing both of the integers a and b a **common divisor** of a and b .

The greatest such integer is called the **greatest common divisor**, or *GCD* of a and b . This number is denoted $\gcd(a, b)$, or with the shorthand (a, b) when clear from the context.

Example 19.4 12 and 42 have divisors 1, 2, 3, 4, 6, 12 and 1, 2, 3, 6, 7, 14, 21, 42. Their shared divisors are 1, 2, 3, 6, so their greatest common divisor is 6.

We warm up with some simple properties of the GCD.

Theorem 19.6 – Properties of the GCD

Let a, b, c be non-negative integers. Then

$$(a, 0) = a \quad (19.1)$$

$$(a, a) = a \quad (19.2)$$

$$(a, b) \leq \max(a, b) \quad (19.3)$$

$$(ac, bc) = c \cdot (a, b) \quad (19.4)$$

$$(a, b) \mid (a, bc) \quad (19.5)$$

If $(a, c) = 1$, then

$$(a, bc) = (a, b) \quad (19.6)$$

Proof. We give a proof for the last equation – the others are good exercises to get acquainted with the GCD.

We can WLOG assume that c is prime. If it is not, we can assume we have a smallest counterexample, perform the substitution $c = pc'$ where p is prime, and prove it using $(a, (bc')p) = (a, bc')$ instead. As $(a, c') \mid (a, c) = 1$, the premises still hold for showing that $(a, bc') = (a, b)$.

First, let $d = (a, b)$. Then, $(a, bc) = d(\frac{a}{d}, \frac{bc}{d})$ by Equation 19.4. We are then trying to prove that $(\frac{a}{d}, c\frac{b}{d}) \neq 1$. Assume otherwise. Then there must be some prime p that divides $\frac{a}{d}$ and $c\frac{b}{d}$. The first condition implies $p \mid a$. However, since $(a, c) = 1$ we can not have $p \mid c$, or else they would share a divisor greater than 1. This means $p \mid \frac{b}{d}$. But then $p \mid (\frac{a}{d}, \frac{b}{d}) = 1$, which is impossible. Thus, we must have that $(\frac{a}{d}, c\frac{b}{d}) = 1$ and $(a, cb) = d = (a, b)$. \square

Exercise 19.8. Prove Equations 19.1-19.6.

Before we start looking into how to actually compute the greatest common divisor, we take a detour into the land of number theoretic sums to also get some practice and understanding of what the GCD actually means.

GCD Sum 1

Compute

$$\sum_{i|N} \sum_{j|N} \gcd(i, j)$$

where N ($1 \leq N \leq 10^{14}$) is a given integer.

Solution. Now and then problems consist of computing some number theoretic sum. There is a number of different techniques involved in this, so we will show two different solutions.

Let us first try to transform the sum into something simpler. In our case, we don't even know how to compute the gcd of two numbers quickly yet, so it makes sense to attempt to simplify that term. This approach is also supported by the gcd(i, j) being a non-trivial term that requires computation to figure out, but *we know all the values it will assume*. Just from the definition, we understand that gcd(i, j) $\mid i$, and $i \mid N$, so gcd(i, j) is a divisor of N too. Picking $i = j = d$, gives us gcd(i, j) = d for any divisor d .

By fixing the values of gcd(i, j) one at a time, the problem is instead transformed to the following: for each $d \mid N$, compute the number of pairs $i \mid N$, $j \mid N$ such that gcd(i, j) = d . If we let this number be $k(d)$, the sum in the problem simplifies to $\sum_{d|N} k(d) \cdot d$. Evaluating $k(d)$ is slightly tricky. We first need the following: if $a \mid b$ and a and b have a common divisor d , then $\frac{a}{d} \mid \frac{b}{d}$. This can be proved by looking at the prime factorization of the three numbers.

A consequence of this fact is that the condition $i, j \mid N$ reduces to $\frac{i}{d}, \frac{j}{d} \mid \frac{N}{d}$. Thus, we can instead try to count how many pair of divisors (i', j') that $\frac{N}{d}$ have. Can we choose any such pair of divisors? No – i' and j' cannot share any divisor $d' > 1$. If they did, i and j would share the divisor dd' , meaning d was not the *greatest* common divisor of i and j .

The remainder of the proof is now a combinatorial argument. Let $\frac{N}{d} = \prod_{p_i|N'} p_i^{e_i}$ be the prime factorization of N' . Since i' and j' are divisors of N' but themselves share no factor, there are three cases for each prime factor: it divides i' to some power, j' to some power, or neither (a consequence of Euclid's Lemma). In the first two cases there are e_i possibilities respectively, and in the second 1. Thus, there are in total $2e_i + 1$ choices for each factor, resulting in the product $\prod (2e_i + 1)$ for the total number of such pairs (i', j').

A note on implementation: when computing all the $d \mid N$ and their prime factorizations, we should not do this in an explicit manner. It is sufficient to

prime factor N and then use a recursive procedure to compute all divisors, one prime factor at a time.

The second approach we show gives us a quite different way of computing the sum. It involves first looking at the function we are computing, and figuring out how it is affected if we isolate one of the prime powers dividing the argument of the function (N). This is a common theme in number theory, where many sums and functions are easy to compute for prime powers, and hopefully easy to combine!

Let p be a prime where $p^k \parallel N$ and $N' = \frac{N}{p^k}$. Then, we can rewrite our sum as

$$\sum_{a=0}^k \sum_{b=0}^k \left(\sum_{i|N'} \sum_{j|N'} \gcd(i \cdot p^a, j \cdot p^b) \right)$$

By Equation 19.4, we can factor out $\min(p^a, p^b)$ from the innermost term. Assume for the purpose of demonstration that $a \leq b$. Then, the sum simplifies to

$$\sum_{a=0}^k \sum_{b=0}^k \left(\sum_{i|N'} \sum_{j|N'} p^a \gcd(i, j \cdot p^{b-a}) \right)$$

However, since $p \nmid i$, Equation 19.6, $\gcd(i, j \cdot p^k) = \gcd(i, j)$ since i is a divisor of the p -less N' . Further simplification becomes possible:

$$\left(\sum_{a=0}^k \sum_{b=0}^k p^{\min(a,b)} \right) \left(\sum_{i|N'} \sum_{j|N'} \gcd(i, j) \right)$$

The observant reader may notice that the left factor in the above product happens to be same sum you'd get if $N = p^k$, since $\gcd(p^a, p^b) = p^{\min(a,b)}$! It's apparently enough to compute the sum for all prime factors of N and multiply answers together. This is not uncommon – in Section 19.6 we study more functions like this.

Now, the only thing that remains is to evaluate this particular sum for each prime divisor. Luckily k and the number of primes in a number is both very small (*le* $\log_2(N)$), so it can be computed with nested loops. \square

Finally, we get to the big question. How do we compute the greatest common divisor of two integers?

Greatest Common Divisor

Given to integers a and b , compute (a, b) .

We already know of a $\Theta(\sqrt{a} + \sqrt{b})$ algorithm to compute (a, b) , namely to enumerate *all* divisors of a and b . A new identity is key to the much faster Euclidean algorithm.

$$(a, b) = (a, b - a) \quad (19.7)$$

We can prove the equality by proving an even stronger result – that *all* common divisors of a and b are also common divisors of a and $b - a$. Assume d is a common divisor of a and b , so that $a = da'$ and $b = db'$ for integers a', b' . Then $b - a = db' - da' = d(b' - a')$, with $b' - a'$ being an integer, is sufficient for d also being a divisor of $b - a$. The converse is shown in a similar way. Hence the divisors of a and b are also divisors of a and $b - a$. In particular, their largest common divisor is the same. The application of these identities yield a recursive solution to the problem. If we wish to compute (a, b) where a, b are positive and $a \leq b$, we reduce the problem to a smaller one by instead computing (a, b) , we compute $(a, b - a)$. This gives us a smaller problem, in the sense that $a + b$ decreases. Since both a and b are non-negative, this means we must at some point arrive at the situation where $b = 0$. Equation 19.1 tells us the GCD is then a .

One simple but important step remains before the algorithm is useful. Note how computing $(1, 10^9)$ requires about 10^9 steps right now, since we will do the reductions $(1, 10^9 - 1), (1, 10^9 - 2), (1, 10^9 - 3) \dots$ The fix is easy – the repeated application of subtraction of a number a from b while $b \geq a$ is exactly the modulo operation, meaning

$$(a, b) = (a, b \bmod a)$$

This last piece of our Euclidean puzzle complete our algorithm, and gives us a remarkably short algorithm, as seen in Algorithm ???. Note the recursive invocation to $(b \bmod a, a)$ to ensure that $a \leq b$.

```
1: procedure GCD( $A, B$ )
2:   if  $B = 0$  then
3:     return  $A$ 
4:   return GCD( $B \bmod A, A$ )
```

Competitive Tip

The Euclidean algorithm exists as the built-in function `__gcd(a, b)` in C++.

Whenever dealing with divisors in a problem, the greatest common divisor may be useful. This is the case in the next problem, where we also look closer at the prime factorization of the GCD.

Granica**Croatian Open Competition in Informatics 2007/2008, Contest #6**

Given integers a_1, a_2, \dots, a_n , find all those numbers d such that upon division by d , all of the numbers a_i leave the same remainder.

Input

The first line contains the integer $2 \leq n \leq 100$, the length of the sequence a_i . The second line contains the integers n integers $1 \leq a_1, a_2, \dots, a_n \leq 10^9$.

Output

Output all such integers d , separated by spaces.

Solution. What does it mean for two numbers a_i and a_j to have the same remainder when dividing by d ? Letting this remainder be r we can write $a_i = dn + r$ and $a_j = dm + r$ for integers n and m . Thus, $a_i - a_j = d(n - m)$ so that d is divisor of $a_i - a_j$! This gives us a necessary condition for our numbers d . Is it sufficient? If $a_i = dn + r$ and $a_j = dm + r'$, we have $a_i - a_j = d(n - m) + (r - r')$. Since d is a divisor of $a_i - a_j$ it must be a divisor of $d(n - m) + (r - r')$ too, meaning $d \mid r - r'$. As $0 \leq r, r' < d$, we have that $-d < r - r' < d$, implying $r - r' = 0$ so that $r = r'$ and both remainders were the same after all.

The question then is how we compute the set of common divisors of all numbers $a_i - a_j$. We claim that this set is (even for the case of only two numbers) the *divisors of their greatest common divisor*. Intuitively true for some, but to prove it we take aid in the prime factorizations of divisors. A divisor of some integer

$$n = p_1^{e_1} \cdots p_k^{e_k}$$

is of the form

$$d = p_1^{e'_1} \cdots p_k^{e'_k}$$

where $0 \leq e'_i \leq e_i$. Then, the requirement for d to be a **common** divisor of n and another number

$$m = p_1^{f_1} \cdots p_k^{f_k}$$

is that $0 \leq e'_i \leq \min(f_i, e_i)$. It should be clear that a number with this property is indeed a common divisor of n and m .

The largest such number is attained when $e'_i = \min(f_i, e_i)$ giving us the GCD. This also explains why all common divisors must be divisors of GCD.

Using this interpretation of the GCD, we can extend the result to finding the GCD d of a sequence b_1, b_2, \dots . Consider any prime p , such that $p^{q_i} \parallel b_i$. Then, we must have $p^{\min(q_1, q_2, \dots)} \parallel d$. This operation is exactly what the GCD algorithm does for two numbers. Since $\min(q_1, q_2, \dots) = \min(q_1, \min(q_2, \dots))$, we can use the recursion formula $d = \gcd(b_1, b_2, \dots) = \gcd(b_1, \gcd(b_2, \dots))$, simplest implemented in a loop:

```

1: procedure MULTIGCD(sequence A)
2:    $gcd \leftarrow 0$ 
3:   for each  $a \in A$  do
4:      $gcd \leftarrow \text{GCD}(gcd, a)$ 
5:   return  $gcd$ 
    
```

Finally, we need to find all the divisors of the GCD to arrive at the answer. \square

A complementary concept is the *least common multiple*.

Definition 19.6 — Least Common Multiple

The *least common multiple* of integers a and b is the smallest positive integer m such that $a \mid m$ and $b \mid m$.

Example 19.5 The multiples of 12 are 12, 24, 36, 48, 60, \dots . The multiples of 10 are 10, 20, 30, 40, 50, 60, \dots . The least common multiple of the numbers is thus 60.

Given a, b , ab is clearly a common multiple, but it doesn't have to be the smallest. Since $a \mid m$, we have that $m = ak$. The question is basically what extra factors we *must* add to a in the form of k in order to have $b \mid m$. Previously, we have determined that the condition for being a divisor is the following. If $p^k \parallel b$ is one of the primes dividing b , then $p^k \mid m$ has to hold. This basically means

that we have to add whatever factors to a that are additionally present in b . For example, since $10 = 2 \cdot 5$ and $12 = 2 \cdot 2 \cdot 3$, we need to add an additional factor 2 and 3 to 10 to make a common multiple – $2 \cdot 2 \cdot 3 \cdot 5 = 60$

To compute the LCM easily, note that a multiple m of an integer a with prime factorization

$$a = p_1^{e_1} \cdots p_k^{e_k}$$

must be of the form

$$m = p_1^{e'_1} \cdots p_k^{e'_k}$$

where $e_i \leq e'_i$.

Thus, if m is to be a common multiple of a and another integer

$$b = p_1^{f_1} \cdots p_k^{f_k}$$

it must hold that $\max(f_i, e_i) \leq e'_i$, with $e'_i = \max(f_i, e_i)$ giving us the smallest such multiple. Since $\max(e_i, f_i) + \min(e_i, f_i) = e_i + f_i$, we get that $\text{lcm}(a, b) \cdot \text{gcd}(a, b) = ab$. This gives us the formula $\text{lcm}(a, b) = \frac{a}{\text{gcd}(a, b)} b$ to compute the LCM. The order of operations is chosen to avoid overflows in computing the product ab .

As for the GCD of multiple integers, it holds that

$$\text{lcm}(a, b, c, \dots) = \text{lcm}(a, \text{lcm}(b, \text{lcm}(c, \dots)))$$

GCD and LCM

Given that $\text{gcd}(a, b) = x$ and $\text{lcm}(a, b) = y$, where $1 \leq x, y \leq 10^{14}$ determine the possible pairs (a, b) .

Solution. Consider a prime $p^k \parallel y$. By the above results, we get that $p^k \parallel a$ or $p^k \parallel b$, since the exponent of each prime in the LCM is the maximum of the corresponding exponents in the two. Similarly, if $p^l \parallel x$, either $p^l \parallel a$ or $p^l \parallel b$. There are then two cases; either $p^k \parallel a$ and $p^l \parallel b$, or the other way around. The possible (a, b) can be generated by checking both options for each prime dividing either y . \square

The Extended Euclidean Algorithm

Next up is the *extended Euclidean algorithm*. It is a way to solve certain integer equations.

Linear Diophantine Equation

Given integers a, b , find an integer solution x, y to

$$ax + by = (a, b)$$

It is not obvious that a solution exists. Let $S = \{ax + by \mid x, y \text{ integers}\}$. These numbers are called the *linear combinations* of a and b . S is closed under addition and negation (and thus also subtraction and multiplication). All numbers of the form $ax + by$ are multiples of (a, b) , and we claim that the set (a, b) itself (and thus all its multiples). Assume $d = (a, b)$ is the smallest positive member of S . Then $a - d\lceil \frac{a}{d} \rceil = a \bmod d \in S$, since it is closed under subtraction and multiplication. Similarly, $b \bmod d \in S$. As $0 \leq a \bmod d < d$, we must have $a \bmod d = 0$ and $b \bmod d = 0$ as d was the smallest element of S . However, this is equivalent to $d \mid a$ and $d \mid b$, so $d \mid (a, b)$ and $d = (a, b)$ since d was a multiple of (a, b) .

This proof might remind you somewhat of the Euclidean algorithm. The proof and the algorithm hide within them a method to write (a, b) as a linear combination of a and b . Remember that during the computation of the GCD, we repeatedly used that $(a, b) = (b, a \bmod b)$. Since $a \bmod b$ is a linear combination of a and b , it seems as if the numbers (a, b) during the computation of the GCD *always* are linear combinations of a and b . The algorithm concludes at $(d, 0)$, at which point $d = (a, b)$. If we only kept track of which linear combination that was equal to d , we would be able to construct a solution to $ax + by = (a, b)$. Let us try this with an example, where we use $[x, y]$ to denote the number $ax + by$.

Example 19.6 — Extended Euclidean algorithm

Consider the equation $15x + 11y = (15, 11) = 1$.

Performing the Euclidean algorithm on these numbers we find that

$$(15, 11) = ([1, 0], [0, 1]) = (11, 15 \bmod 11) =$$

$$(11, 15 - 1 \cdot 11) = ([0, 1], [1, 0] - 1 \cdot [0, 1]) =$$

$$(11, 4) = ([0, 1], [1, -1]) = (4, 11 \bmod 4) =$$

$$(4, 11 - 2 \cdot 4) = ([1, -1], [0, 1] - 2[1, -1]) =$$

$$(4, 3) = ([1, -1], [-2, 3]) = (3, 4 \bmod 3) =$$

$$(3, 4 - 1 \cdot 3) = ([-2, 3], [1, -1] - [-2, 3]) =$$

$$(3, 1) = ([-2, 3], [3, -4]) = (1, 3 \bmod 1) =$$

$$(1, 3 - 3 \cdot 1) = ([3, -4], [-2, 3] - 3[3, -4]) =$$

$$(1, 0) = ([3, -4], [-11, 15])$$

Verifying the results, $15 \cdot 3 + 11 \cdot (-4) = 45 - 44 = 1$.

Exercise 19.9. Find an integer solution to the equation $24x + 52y = 2$.

A short solution can be written recursively:

1: **procedure** EXTENDED_EUCLIDEAN(a, b)

This gives us a single solution. Finding the others is not much harder. Let $a' = \frac{a}{(a,b)}$ and $b' = \frac{b}{(a,b)}$. Given two solutions

$$ax_1 + by_1 = (a, b)$$

$$ax_2 + by_2 = (a, b)$$

we can first factor out (a, b) to get

$$a'x_1 + b'y_1 = 1$$

$$a'x_2 + b'y_2 = 1$$

A simple subtraction gives us that

$$a'(x_1 - x_2) + b'(y_1 - y_2) = 0$$

$$a'(x_1 - x_2) = b'(y_2 - y_1)$$

Because $(a', b') = 1$, $b' \mid x_1 - x_2$. Then there exists k such that $x_1 - x_2 = kb'$, so $x_1 = x_2 + kb'$. Inserting this gives us

$$a'(x_2 + kb' - x_2) = b'(y_2 - y_1)$$

$$a'kb' = b'(y_2 - y_1)$$

$$a'k = y_2 - y_1$$

$$y_1 = y_2 - ka'$$

Thus, any solution must be of the form

$$(x_1 + k \frac{b}{(a, b)}, y_1 - k \frac{a}{(a, b)}) \text{ for } k \in \mathbb{Z}$$

It is easily verified that any k also gives us a solution to this. This result is called **Bezout's identity**.

Generalized Knights

A *generalized knight* is a special chess piece. It moves by first choosing one of the four cardinal directions and moves a steps, and then chooses one of the two orthogonal cardinal directions and moves b steps (for example first up and then left or right, or first left and then up or down). Compute the minimum number of moves the knight needs to move from $(0, 0)$ to (x, y) .

Input

The four integers $1 \leq a, b, x, y \leq 10^8$, where $a \neq b$.

Solution. Let us split up all the 8 moves the knight can make into the following subcomponents: □

19.4 Modular Arithmetic

When first learning division, one is often introduced to the concept of *remainders*. For example, when dividing 7 by 3, you would get “2 with a remainder of 1”. In general, when dividing a number a with a number n , you would get a **quotient** q and a **remainder** r . These numbers would satisfy the identity $a = nq + r$, with $0 \leq r < b$.

Example 19.7 — Division with remainders

Consider division (with remainders) by 4 of the numbers $0, \dots, 6$. We have that

$$0 \div 4 = 0, \text{ remainder } 0$$

$$1 \div 4 = 0, \text{ remainder } 1$$

$$\text{frac}24 = 0, \text{ remainder } 2$$

$$\text{frac}34 = 0, \text{ remainder } 3$$

$$\text{frac}44 = 1, \text{ remainder } 0$$

$$\text{frac}54 = 1, \text{ remainder } 1$$

$$\text{frac}64 = 1, \text{ remainder } 2$$

Note how the remainder always increase by 1 when the nominator increased. As you might remember from Chapter 2 on C++ (or from your favorite programming language), there is an operator which compute this remainder called the *modulo operator*. Modular arithmetic is then the computation on numbers, where every number is taken modulo some integer n . Under such a scheme, we have that e.g. 3 and 7 are basically the same if computing modulo 4, since $3 \bmod 4 = 3 = 7 \bmod 4$. This concept, where numbers with the same remainder are treated as if they are equal is called *congruence*.

Definition 19.7 – Congruence

If a and b have the same remainder when divided by n , we say that a and b are ***congruent modulo n*** , written

$$a \equiv b \pmod{n}$$

An equivalent and in certain applications more useful definition is that $a \equiv b \pmod{n}$ if and only if $n \mid a - b$.

Exercise 19.10. What does it mean for a number a to be congruent to 0 modulo n ?

When counting modulo something, the laws of addition and multiplication are somewhat altered:

+	0	1	2
0	0	1	2
1	1	2	$3 \equiv 0$
2	2	$3 \equiv 0$	$4 \equiv 1$

*	0	1	2
0	0	0	0
1	0	1	2
2	0	2	$4 \equiv 1$

When we wish to perform arithmetic of this form, we use the *integers modulo n* rather than the ordinary integers. These has a special set notation as well: \mathbb{Z}_n .

While addition and multiplication is quite natural (i.e. performing the operation as usual and then taking the result modulo n), division is a more complicated story. For real numbers, the *inverse* x^{-1} of a number x is defined as the number which satisfy the equation $xx^{-1} = 1$. For example, the inverse of 4 is 0.25, since $4 \cdot 0.25 = 1$. The division $\frac{a}{b}$ is then simply a multiplied with the inverse of b . The same definition is applicable to modular arithmetic:

Definition 19.8 — Modular Inverse

The *modular inverse* of a modulo n is the integer a^{-1} such that $aa^{-1} \equiv 1 \pmod{n}$, if such an integer exists.

Considering our multiplication table of \mathbb{Z}_3 , we see that 0 has no inverse and 1 is its own inverse (just as with the real numbers). However, since $2 \cdot 2 = 4 \equiv 1 \pmod{3}$, 2 is actually its own inverse. If we instead consider multiplication in \mathbb{Z}_4 , the situation is quite different.

*	0	1	2	3
0	0	0	0	0
1	0	1	2	3
2	0	2	0	2
3	0	3	2	1

Now, 2 does not even have an inverse! To determine when an inverse exists – and if so, computing the inverse – we will make use of the extended Euclidean algorithm. If $aa^{-1} \equiv 1 \pmod{n}$, we have $n \mid aa^{-1} - 1$, meaning $aa^{-1} - 1 = nx$ for some integer x . Rearranging this equation gives us $aa^{-1} - nx = 1$. We know from Section 19.3 that this has a solution if and only if $(a, n) = 1$. In this case, we can use the extended Euclidean algorithm to compute a^{-1} . Note that be Bezout's identity, a^{-1} is actually unique modulo n .

Just like the reals, modular arithmetic has a cancellation law regarding

Theorem 19.7

Assume $a \perp n$. Then $ab \equiv ac \pmod{n}$ implies $b \equiv c \pmod{n}$.

Proof. Since $a \perp n$, there is a number a^{-1} such that $aa^{-1} \equiv 1 \pmod{n}$.

$$ab \equiv ac \pmod{n}$$

with a^{-1} results in

$$aa^{-1}b \equiv aa^{-1}c \pmod{n}$$

Simplifying aa^{-1} gives us

$$b \equiv c \pmod{n}$$

□

Another common modular operation is exponentiation, i.e. computing $a^m \pmod{n}$. While this can be computed easily in $\Theta(m)$, we can actually do better using a method called **exponentiation by squaring**. It is essentially based on the recursion

$$a^m \bmod n = \begin{cases} 1 \bmod n & \text{if } m = 0 \\ a \cdot (a^{m-1} \bmod n) \bmod n & \text{if } m \text{ odd} \\ (a^{\frac{m}{2}} \bmod n)^2 \bmod n & \text{if } m \text{ even} \end{cases}$$

This procedure is clearly $\Theta(\log_2 m)$, since applying the recursive formula for even numbers halve the m to be computed, while applying it an odd number will first make it even and then halve it in the next iteration. It is very important that $a^{\frac{m}{2}} \bmod n$ is computed only once, even though it is squared! Computing it twice causes the complexity to degrade to $\Theta(m)$ again.

19.5 Chinese Remainder Theorem

The Chinese Remainder Theorem is probably the most useful theorem in algorithmic problem solving. It gives us a way of solving certain systems of linear equations.

Theorem 19.8 — Chinese Remainder Theorem

Given a system of equations

$$x \equiv a_1 \pmod{m_1}$$

$$x \equiv a_2 \pmod{m_2}$$

...

$$x \equiv a_m \pmod{m_n}$$

where the numbers m_1, \dots, m_n are pairwise relatively prime, there is a *unique* integer $x \pmod{\prod_{i=1}^n m_i}$ that satisfy such a system.

Proof. We will prove the theorem inductively. The theorem is clearly true for $n = 1$, with the unique solution $x = a_1$. Now, consider the two equations

$$x \equiv a_1 \pmod{m_1}$$

$$x \equiv a_2 \pmod{m_2}$$

Let $x = a_1 \cdot m_2 \cdot (m_2^{-1} \bmod m_1) + a_2 \cdot m_1 \cdot (m_1^{-1} \bmod m_2)$, where $m_1^{-1} \bmod m_2$ is taken to be a modular inverse of m_1 modulo m_2 . These inverses exist, since $m_1 \perp m_2$ by assumption. We then have that $x = a_1 \cdot m_2 \cdot (m_2^{-1} \bmod m_1) + a_2 \cdot m_1 \cdot (m_1^{-1} \bmod m_2) \equiv a_1 \cdot m_2 \cdot (m_2^{-1} \bmod m_1) \equiv a_1 \pmod{m_1}$.

Since a solution exist for every a_1, a_2 , this solution must be unique by the pigeonhole principle – there are $m_1 \cdot m_2$ possible values for a_1, a_2 , and $m_1 \cdot m_2$ possible values for x . Thus, the theorem is also true for $n = 2$.

Assume the theorem is true for $k - 1$ equations. Then, we can replace the equations

$$x \equiv a_1 \pmod{m_1}$$

$$x \equiv a_2 \pmod{m_2}$$

with another equation

$$x \equiv x * \pmod{m_1 m_2}$$

where x^* is the solution to the first two equations. We just proved those two equations are equivalent with regards to x . This reduces the number of equations to $k - 1$, which by assumption the theorem holds for. Thus, it also holds for k equations. \square

Note that the theorem used an explicit construction of the solution, allowing us to find what the unique solution to such a system is.

Radar

KTH Challenge 2014

We say that an integer z is within distance y of an integer x modulo an integer m if

$$z \equiv x + t \pmod{m}$$

where $|t| \leq y$.

Find the smallest non-negative integer z such that it is:

- within distance y_1 of x_1 modulo m_1
- within distance y_2 of x_2 modulo m_2
- within distance y_3 of x_3 modulo m_3

Input

The integers $0 \leq m_1, m_2, m_3 \leq 10^6$. The integers $0 \leq x_1, x_2, x_3 \leq 10^6$. The integers $0 \leq y_1, y_2, y_3 \leq 300$.

Output

The integer z .

The problem gives rise to three linear equations of the form

$$z \equiv x_i + t_i \pmod{m_i}$$

where $-y_i \leq t_i \leq y_i$. If we fix all the variables t_i , the problem reduces to solving the system of equations using CRT. We could then find all possible values of z , and choose the minimum one. This requires applying the CRT construction about $2 \cdot 600^3 = 432\,000\,000$ times. Since the modulo operation involved is quite expensive, this approach would use too much time. Instead, let us exploit a useful greedy principle in finding minimal solutions.

Assume that z is the minimal answer to an instance. There are only two

situations where $z - 1$ cannot be a solution as well:

- $z = 0$ – since z must be non-negative, this is the smallest possible answer
- $z \equiv x_i - y_i$ – then, decreasing z would violate one of the constraints

In the first case, we only need to verify whether $z = 0$ is a solution to the three inequalities. In the second case, we managed to change an inequality to a linear equation. By testing which of the i this equation holds for, we only need to test the values of t_i for the two other equations. This reduce the number of times we need to use the CRT to $600^2 = 360\,000$ times, a modest amount well within the time limit.

19.6 Euler's totient function

Now that we have talked about modular arithmetic, we can give the numbers which are *not* divisors to some integer n their well-deserved attention. This discussion will start with the ϕ -function.

Definition 19.9 Two integers a and b are said to be **relatively prime** if their only (and thus greatest) common divisor is 1. If a and b are relatively prime, we write that $a \perp b$.

Example 19.8 The numbers 74 and 22 are not relatively prime, since they are both divisible by 2.

The numbers 72 and 65 are relatively prime. The prime factorization of 72 is $2 \cdot 2 \cdot 2 \cdot 3 \cdot 3$, and the factorization of 65 is $5 \cdot 13$. Since these numbers have no prime factors in common, they have no divisors other than 1 in common.

Given an integer n , we ask ourselves how many of the integers $1, 2, \dots, n$ which are relatively prime to n .

Definition 19.10 — Euler's totient function

Euler's totient function $\phi(n)$ is defined as the number of integers $k \in [1, n]$ such that $(k, n) = 1$, i.e. those positive integers less than n which are co-prime to n .

Example 19.9 What is $\phi(12)$? The numbers 2, 4, 6, 8, 10 all have the factor 2 in common with 12 and the numbers 3, 6, 9 all have the factor 3 in common with 12.

This leaves us with the integers 1, 5, 7, 11 which are relatively prime to 12. Thus, $\phi(12) = 4$.

For prime powers, $\phi(p^k)$ is easy to compute. The only integers which are not relatively prime to $\phi(p^k)$ are the multiples of p , which there are $\frac{p^k}{p} = p^{k-1}$ of, meaning

$$\phi(p^k) = p^k - p^{k-1} = p^{k-1}(p - 1)$$

It turns out $\phi(n)$ has a property which is highly useful in computing certain number theoretical functions – it is **multiplicative**, meaning

$$\phi(ab) = \phi(a)\phi(b) \text{ if } a \perp b$$

For multiplicative functions, we can reduce the problem of computing arbitrary values of the function to finding a formula only for prime powers. The reasoning behind the multiplicativity of ϕ is quite simple. Let $a' = a - \phi(a)$, i.e. the number of integers which do share a factor with a , and similarly $b' = b - \phi(b)$. Then, there will be ab' numbers between 1 and ab which share a factor with b . If x is one of the b' numbers sharing a factor with b , then so are $x, x + b, x + 2b, \dots, x + (a - 1)b$. Similarly, there will be $a'b$ numbers between 1 and ab sharing a factor with a . However, there may be some numbers sharing both a factor with a and b . Consider two such numbers $x + ib = y + ja$, which gives $ib - ja = y - x$. By Bezout's identity, this has a single solution (i, j) modulo ab , meaning every number $x + ib$ equals exactly one number $y + ja$. Thus, there were $a'b'$ numbers sharing a factor with both a and b . This means there are $ab' + a'b - a'b'$ numbers sharing a factor with either a and b , so

$$\phi(ab) = ab - ab' - a'b + a'b' = (a - a')(b - b') = \phi(a)\phi(b)$$

and we are done.

Using the multiplicativity of ϕ we get the simple formula

$$\phi(p_1^{e_1} \dots p_k^{e_k}) = \phi(p_1^{e_1}) \dots \phi(p_k^{e_k}) = p_1^{e_1-1}(p_1 - 1) \dots p_k^{e_k-1}(p_k - 1)$$

Computing ϕ for a single value can thus be done as quickly as factoring the number. If we wish to compute ϕ for an interval $[1, n]$ we can use the Sieve of Eratosthenes.

This seemingly convoluted function might seem useless, but is of great importance via the following theorem:

Theorem 19.9 — Euler's theorem

If a and n are relatively prime and $n \geq 1$,

$$a^{\phi(n)} \equiv 1 \pmod{n}$$

Proof. The proof of this theorem isn't trivial, but it is number theoretically interesting and helps to build some intuition for modular arithmetic. The idea behind the proof will be to consider the product of the $\phi(n)$ positive integers less than n which are relatively prime to n . We will call these $x_1, x_2, \dots, x_{\phi(n)}$. Since these are all distinct integers between 1 and n , they are incongruent modulo n . We call such a set of $\phi(n)$ numbers, all incongruent modulo n a **complete residue system** (CRS) modulo n .

Next, we will prove that $ax_1, ax_2, \dots, ax_{\phi(n)}$ also form a CRS modulo n . We need to show two properties for this:

1. All numbers are relatively prime to n
2. All numbers are incongruent modulo n

We will start with the first property. Since both a and x_i are relatively prime to n , neither number have a prime factor in common with n . This means ax_i have no prime factor in common with n either, meaning the two numbers are relatively prime. The second property requires us to make use of the cancellation property of modular arithmetic (Theorem 19.7). If $ax_i \equiv ax_j \pmod{n}$, the cancellation law gives us $x_i \equiv x_j \pmod{n}$. Since all x_i are incongruent modulo n , we must have $i = j$, meaning all the numbers ax_i are incongruent as well. Thus, these numbers did indeed form a complete residue system modulo n .

If $ax_1, \dots, ax_{\phi(n)}$ form a CRS, we know every ax_i must be congruent to some x_j , meaning

$$ax_1 \cdots ax_{\phi(n)} \equiv x_1 \cdots x_{\phi(n)} \pmod{n}$$

Factoring the left hand side turns this into

$$a^{\phi(n)} x_1 \cdots x_{\phi(n)} \equiv x_1 \cdots x_{\phi(n)} \pmod{n}$$

Since all the x_i are relatively prime to n , we can again use the cancellation law, leaving

$$a^{\phi(n)} \equiv 1 \pmod{n}$$

completing our proof of Euler's theorem. \square

For primes p we get a special case of Euler's theorem when since $\phi(p) = p - 1$.

Corollary 19.1

Fermat's Theorem For a prime p and an integer $a \not\equiv 0 \pmod{p}$, we have

$$a^{p-1} \equiv 1 \pmod{p}$$

Competitive Tip

By Fermat's Theorem, we also have $a^{p-2} \equiv a^{-1} \pmod{p}$ when p is prime (and $a \not\equiv 0 \pmod{p}$). This is often an easier way of computing modular inverses modulo primes than using the extended Euclidean algorithm, in particular if you already coded modular exponentiation.

Exponial

Nordic Collegiate Programming Contest 2016

Define the *exponial* of n as the function

$$opexponial(n) = n^{(n-1)^{(n-2)^{\dots^{2^1}}}}$$

Compute $exponial(n) \pmod{m}$.

Input

The input contains the integers $1 \leq n, m \leq 10^9$.

Output

Output a single integer, the value of $exponial(n) \pmod{m}$.

Euler's theorem suggests a recursive approach. Since $n^e \pmod{m}$ is periodic

(in e), with a period of $\phi(m)$, maybe when computing $n^{(n-1)\cdots}$ we could compute $e = (n-1)\cdots$ modulo $\phi(m)$ and only then compute $n^e \pmod{m}$? Alas, this is only useful when $n \perp m$, since this is a necessary precondition for Euler's theorem. When working modulo some integer m with a prime factorization of $p_1^{e_1} \cdots p_k^{e_k}$, a helpful approach is to instead work modulo its prime powers $p_i^{e_i}$ and then combine the results using the Chinese remainder theorem. Since the prime powers of a prime factorization is relatively prime, the remainder theorem applies.

Let us apply this principle to Euler's theorem. When computing $n^e \pmod{p^k}$ we have two cases. Either $p \mid n$, in which case $n^e \equiv 0 \pmod{p^k}$ whenever $e \geq k$. Otherwise, $p \perp n$, and $n^e \equiv n^{e \bmod \phi(p^k)} \pmod{p^k}$ by Euler's theorem.

This suggests that if $e \geq \max(e_1, \dots, e_k)$, we have that n^e is indeed periodic. Furthermore, e_i is bounded by $\log_2 n$, since

$$\begin{aligned} p_i^{e_i} &\leq n \Rightarrow \\ e_i \log_2(p_i) &\leq \log_2(N) \Rightarrow \\ e_i &\leq \frac{\log_2(N)}{\log_2(p_i)} \leq \log_2(N) \end{aligned}$$

As $p_i \geq 2$, we know that $\log_2(p_i) \geq 1$, which we used in the final inequality. Since $\log_2(10^9) \approx 30$ and $4^{32} \geq 30$, we can use the periodicity of n^e whenever $n \geq 5$.

$$n^{\phi(m) + \text{exponential}(n-1) \bmod \phi(m)} \bmod m$$

For $n = 4$, the exponential equals only 262144, meaning we can compute it immediately.

One final insight remains. If we use the recursive formula, i.e. first computing $e = (n-1)^{(n-2)\cdots} \bmod \phi(m)$ and then $n^{\phi(m) + e \bmod \phi(m)} \bmod m$, we still have the problem that n can be up to 10^9 . We would need to perform a number of exponentiations that is linear in n , which is slow for such large n . However, our modulo will actually very quickly converge to 1. While the final result is taken modulo m , the first recursive call is taken modulo $\phi(m)$. The recursive call performed at the next level will thus be modulo $\phi(\phi(m))$, and so on. That this sequence decrease very quickly is based on two facts. For even m , $\phi(m) = \phi(2)\phi(\frac{m}{2}) = \phi(\frac{m}{2}) \leq \frac{m}{2}$. For odd m , $\phi(m)$ is even. Any odd m consists only of odd prime factors, but since $\phi(p) = p-1$ (i.e. an even number for odd primes p) and ϕ is multiplicative, $\phi(m)$ must be even. Thus $\phi(\phi(m)) \leq \frac{m}{2}$ for

$m > 1$ (1 is neither even nor contains an odd prime factor). This means the modulo will become 1 in a logarithmic number of iterations, completing our algorithm.

Chapter Exercises

Problem 19.12

Longest Composite Sum – longcompositesum

Inheritance – inheritance

Evening Out 3 – eveningout3

Indivisible Sequence – indivisibleseq

Let S be the set of all positive integer divisors of k . How many numbers are the product of two distinct elements of S ?

Happy Happy Prime Prime – happyprime

Prime Path – primepathc

Chapter Notes

A highly theoretical introduction to classical number theory can be found in *An Introduction to the Theory of Numbers*[13] While devoid of exercises and examples, it is very comprehensive.

A Computational Introduction to Number Theory and Algebra[24] instead takes a more applied approach, and is freely available under a Creative Commons license at the authors home page.³

³<http://www.shoup.net/ntb/>

20 Competitive Programming Strategy

Competitive programming is what we call the mind sport of solving algorithmical problems and coding their solutions, often under the pressure of time. Most programming competitions are performed online, at your own computer through some kind of online judge system. For students of either high school or university, there are two main competitions. High school students compete in the *International Olympiad in Informatics* (IOI), and university students go for the *International Collegiate Programming Contest* (ICPC).

Different competition styles have different difficulty, problem types and strategies. In this chapter, we will discuss some basic strategy of programming competitions, and give tips on how to improve your competitive skills.

20.1 IOI

The IOI is an international event where a large number of countries send teams of up to 4 high school students to compete individually against each other during two days of competition. Every participating country has its own national selection olympiad first.

During a standard IOI contest, contestants are given 5 hours to solve 3 problems, each worth at most 100 points. These problems are not given in any particular order, and the scores of the other contestants are hidden until the end of the contest. Generally none of the problems are “easy” in the sense that it is immediately obvious how to solve the problem in the same way the first 1-2 problems of most other competitions are. This poses a large problem, in particular for the amateur. Without any trivial problems nor guidance from other contestants on what problems to focus on, how does an IOI competitor prioritize? The problem is further exacerbated by problems not having a simple binary scoring, with a submission being either accepted or rejected. Instead, IOI problems contain many so-called *subtasks*. These subtasks give partial credit for the problem, and contain additional restrictions and limits on either

input or output. Some problems do not even use discrete subtasks. In these tasks, scoring is done on some scale which determines how “good” the output produced by your program is.

Strategy

Very few contestants manage to solve every problem fully during an IOI contest. There is a very high probability you are not one of them, which leaves you with two options – you either skip a problem entirely, or you solve some of its subtasks. At the start of the competition, you should read through every problem and *all of the subtasks*. In the IOI you do not get extra points for submitting faster. Thus, it does not matter if you read the problems at the beginning instead of rushing to solve the first problem you read. Once you have read all the subtasks, you will often see the solutions to some of the subtasks immediately. Take note of the subtasks which you know how to solve!

Deciding on which order you should solve subtasks in is probably one of the most difficult parts of the IOI for contestants at or below the silver medal level. In IOI 2016, the difference between receiving a gold medal and a silver medal was a mere 3 points. On one of the problems, with subtasks worth 11, 23, 30 and 36 points, the first silver medalist solved the third subtask, worth 30 points (a submission that possibly was a failed attempt at 100 points). Most competitors instead solved the first two subtasks, together worth 34 points. If the contestant had solved the first two subtasks instead, he would have gotten a gold medal.

The problem basically boils down to the question *when should I solve subtasks instead of focusing on a 100 point solution?* There is no easy answer to this question, due to the lack of information about the other contestants’ performances. First of all, you need to get a good sense of how difficult a solution will be to implement correctly before you attempt it. If you only have 30 minutes left of a competition, it might not be a great idea to go for a 100 point solution on a very tricky problem. Instead, you might want to focus on some of the easier subtasks you have left on this or other problems. If you fail your 100 point solution which took over an hour to code, it is nice to know you did not have some easy subtasks worth 30-60 points which could have given you a medal.

Problems without discrete scoring (often called *heuristic* problems) are almost always the hardest ones to get a full score on. These problems tend to

be very fun, and some contestants often spend way too much time on these problems. They are treacherous in that it is often easy to increase your score by *something*. However, those 30 minutes you spent to gain one additional point may have been better spent coding a 15 point subtask on another problem. As a general rule, go for the heuristic problem last during a competition. This does not mean to skip the problem unless you completely solve the other two, just to focus on them until you decide that the heuristic problem is worth more points if given the remaining time.

In IOI, you are allowed to submit solution attempts a large number of times, without any penalty. Use this opportunity! When submitting a solution, you will generally be told the results of your submission on each of the secret test cases. This provides you with much details. For example, you can get a sense of how correct or wrong your algorithm is. If you only fail 1-2 cases, you probably just have a minor bug, but your algorithm in general is probably correct. You can also see if your algorithm is fast enough, since you will be told the execution time of your program on the test cases. Whenever you make a change to your code which you think affect correctness or speed – submit it again! This gives you a sense of your progress, and also works as a good regression test. If your change introduced more problems, you will know.

Whenever your solution should pass a subtask, submit it. These subtask results will help you catch bugs earlier when you have less code to debug.

Getting Better

The IOI usually tend to have pretty hard problems. Some areas get rather little attention. For example, there are basically no pure implementation tasks and very little geometry.

First and foremost, make sure you are familiar with all the content in the IOI syllabus¹. This is an official document which details what areas are allowed in IOI tasks. This book deals with most, if not all of the topics in the IOI syllabus.

In the Swedish IOI team, most of the top performers tend to also be good mathematical problem solvers (also getting IMO medals). Combinatorial problems from mathematical competitions tend to be somewhat similar to the algorithmic frame of mind, and can be good practice for the difficult IOI problems.

When selecting problems to practice on, there are a large number of national

¹<https://people.ksp.sk/~misof/ioi-syllabus/>

olympiads with great problems. The Croatian Open Competition in Informatics² is a good source. Their competitions are generally a bit easier than solving IOI with full marks, but are good practice. Additionally, they have a final round (the Croatian Olympiad in Informatics) which are of high quality and difficulty. COCI publishes solutions for all of their contests. These solutions help a lot in training.

One step up in difficulty from COCI is the Polish Olympiad in Informatics³. This is one of the most difficult European national olympiad published in English, but unfortunately they do not publish solutions in English for their competitions.

There are also many regional olympiads, such as the Baltic, Balkan, Central European, Asia-Pacific Olympiads in Informatics. Their difficulty is often higher than that of national olympiads, and of the same format as an IOI contest (3 problems, 5 hours). These, and old IOI problems, are probably the best sources of practice.

20.2 ICPC

In ICPC, you compete in teams of three to solve about 10-12 problems during 5 hours. A twist in the ICPC-style competitions is that the team shares a single computer. This makes it a bit harder to prioritize tasks in ICPC competitions than in IOI competitions. You will often have multiple problems ready to be coded, and wait for the computer. In ICPC, you see the progress of every other team as well, which gives you some suggestions on what to solve. As a beginner or medium-level team, this means you will generally have a good idea on what to solve next, since many better teams will have prioritized tasks correctly for you.

ICPC scoring is based on two factors. First, teams are ranked by the number of solved problems. As a tie breaker, the penalty time of the teams are used. The penalty time of a single problem is the number of minutes into the contest when your first successful attempt was submitted, plus a 20 minute penalty for any rejected attempts. Your total penalty time is the sum of penalties for every problem.

Strategy

In general, teams will be subject to the penalty tie-breaking. In the 2016 ICPC World Finals, both the winners and the team in second place solved 11 problems.

²<http://hsin.hr/coci/>

³<http://main.edu.pl/en/archive/oi>

Their penalty time differed by a mere 7 minutes! While such a small penalty difference in the very top is rather unusual, it shows the importance of taking your penalty into account.

Minimizing penalties generally comes down to a few basic strategic points:

- Solving the problems in the right order.
- Solving each problem quickly.
- Minimizing the number of rejected attempts.

In the very beginning of an ICPC contest, the first few problems will be solved quickly. In 2016, the first accepted submissions to five of the problems came in after 11, 15, 18, 32, 44 minutes. On the other hand, after 44 minutes no team had solved all of those problems. Why does not every team solve the problems in the same order? Teams are of different skill in different areas, make different judgment calls regarding difficulty or (especially early in the contest) simply read the problem in a different order. The better you get, the harder it is to distinguish between the “easy” problems of a contest – they are all “trivial” and will take less than 10-15 minutes to solve and code.

Unless you are a very good team or have very significant variations in skill among different areas (e.g., you are graph theory experts but do not know how to compute the area of a triangle), you should probably follow the order the other teams choose in solving the problems. In this case, you will generally always be a few problems behind the top teams.

The better you get, the harder it is to exploit the scoreboard. You will more often be tied in the top with teams who have solved the exact same problems. The problems which teams above you have solved but you have not may only be solved by 1-2 teams, which is not a particularly significant indicator in terms of difficulty. Teams who are very strong at math might prioritize a hard maths problem before an easier (on average for most teams) dynamic programming problem. This can risk confusing you into solving the wrong problems for the particular situation of your team.

The amount of cooperation during a contest is difficult to decide upon. The optimal amount varies a lot between different teams. In general, the amount of cooperation should increase within a single contest from the start to the end. In the beginning, you should work in parallel as much as possible, to quickly read all the problems, pick out the easy-medium problems and start solving

them. Once you have competed in a few contests, you will generally know the approximate difficulty of the simplest tasks, so you can skim the problem set for problems of this difficulty. Sometimes, you find an even easier problem in the beginning than the one the team decided to start coding.

If you run out of problems to code, you waste computer time. Generally, this should not happen. If it does, you need to become faster at solving problems.

Towards the end of the contest, it is a common mistake to parallelize on several of the hard problems at the same time. This carries a risk of not solving any of the problems in the end, due to none of the problems getting sufficient attention. Just as with subtasks in IOI, this is the hardest part of prioritizing tasks. During the last hour of an ICPC contest, the previously public scoreboard becomes frozen. You can still see the number of attempts other teams make, but not whether they were successful. Hence, you can not really know how many problems you have to solve to get the position that you want. Learning your own limits and practicing a lot as a team – especially on difficult contests – will help you get a feeling for how likely you are to get in all of your problems if you parallelize.

Read all the problems! You do not want to be in a situation where you run out of time during a competition, just to discover there was some easy problem you knew how to solve but never read the statement of. ICPC contests are made more complex by the fact that you are three different persons, with different skills and knowledge. Just because you can not solve a problem does not mean your team mates will not find the problem trivial, have seen something similar before or are just better at solving this kind of problem.

The scoreboard also displays failed attempts. If you see a problem where many teams require extra attempts, be more careful in your coding. Maybe you can perform some extra tests before submitting, or make a final read-through of the problem and solution to make sure you did not miss any details.

If you get Wrong Answer, you may want to spend a few minutes to code up your own test case generators. Prefer generators which create cases where you already know the answers. Learning e.g. Python for this helps, since it usually takes under a minute to code a reasonable complex input generator.

If you get Time Limit Exceeded, or even suspect time might be an issue – code a test case generator. Losing a minute on testing your program on the worst case, versus a risk of losing 20 minutes to penalty is a trade-off worth considering on some problems.

You are allowed to ask questions to the judges about ambiguities in the problems. Do this the moment you think something is ambiguous (judges generally take a few valuable minutes in answering). Most of the time they give you a “No comment” response, in which case the perceived ambiguity probably was not one.

If neither you nor your team mates can find a bug in a rejected solution, consider coding it again from scratch. Often, this can be done rather quickly when you have already coded a solution.

Getting Better

- Practice a lot with your team. Having a good team dynamic and learning what problems the other team members excel at can be the difference that helps you to solve an extra problem during a contest.
- Learn to debug on paper. Wasting computer time for debugging means not writing code! Whenever you submit a problem, print the code. This can save you a few minutes in getting your print-outs when judging is slow (in case your submission will need debugging). If your attempt was rejected, you can study your code on paper to find bugs. If you fail on the sample test cases and it takes more than a few minutes to fix, add a few lines of debug output and print it as well (or display it on half the computer screen).
- Learn to write code on paper while waiting for the computer. In particular, tricky subroutines and formulas are great to hammer out on paper before occupying valuable computer time.
- Focus your practice on your weak areas. If you write buggy code, learn your programming language better and code many complex solutions. If your team is bad at geometry, practice geometry problems. If you get stressed during contests, make sure you practice under time pressure. For example, Codeforces (<http://codeforces.com>) has an excellent gym feature, where you can compete retroactively in a contest using the same amount of time as in the original contest. The scoreboard will then show the corresponding scoreboard from the original contest during any given time.

21 Papers

This chapter contains a series of problems not associated to any particular topics. They are a way for you to practice problem solving without being primed with the techniques that will come up, and a way for us to show techniques, tricks and combinations that does not fit naturally into the text of any particular chapter.

The problems are divided into shorter “papers”. Each paper includes a suggested time that could be used if the problems were posed in a real contest. Some of the papers are actually real contests that was given in the past.

After each contest, you can find solution descriptions of the problem.

21.1 Paper 3

Divisor Solitaire

Nicolaas likes playing a solitaire game about divisors. First, he picks an integer N ($1 \leq N \leq 10^{14}$). Then, in each round of the game he picks a divisor of N , such that it neither divides or is a divisor of any previously picked number. Given N , determine the maximum number of rounds he can play.

Problem 21.1

Divisor Solitaire – divisorsolitaire

Solutions

Divisor Solitaire A reasonable guess guided by combinatorial intuition is to pick as solution exactly those divisors where $\Omega(d) = l$ for some l . Two such number cannot be divisors of each other (why?). It is also difficult to come up with something better in the case where all $e_i = 1$ (i.e. N is the product of distinct primes). In particular, we prove that it is optimal to let $l = \lfloor \frac{\sum_{i=1}^k e_i}{2} \rfloor$.

Roughly, we show that all divisors of N can be partitioned into chains of divisors of the form d_1, d_2, \dots, d_m such that $d_i \mid d_{i+1}$ and $\Omega(d_i) = l$ for some i . All the integers in a chain divide each other, so we can only pick a single divisor from each chain, making the number of such chains an upper bound to the answer. As each chain contains a d_i with $\Omega(d_i) = l$ (there can only be one per chain), then the number of divisors of this kind must also be an upper bound. Since that upper bound is attainable (by picking the set of all such divisors), it is the largest possible such set.

More specifically, every chain d_1, \dots, d_h , will be such that $d_{i+1} = d_i p$ for some prime p , and $\Omega(d_1) + \Omega(d_h) = N$. The first condition implies $\Omega(d_{i+1}) = \Omega(d_i) + 1$ and the second implies $\Omega(d_1) \leq l$ and $\Omega(d_h) \geq l$, so that $\Omega(d_i) = l$ for some i .

Such a partition exists by the following construction. Assume that all divisors of $N' = \prod_{i=2}^k p_i^{e_i}$ can be partitioned into such chains. If d_1, \dots, d_h , is such a chain, we can also partition the he_1 integers $d_i \cdot p_1^\alpha$ into chains. First, take $d_1, d_1 p_1, \dots, d_1 p_1^{e_1-1}, d_1 p_1^{e_1}, d_2 p_1^{e_1}, \dots, d_h p_1^{e_1}$. This is a valid chain, since $\Omega(d_1) + \Omega(d_h p_1^{e_1}) = \Omega(d_1) + \Omega(d_h) + e_1 = \Omega(N') + e_1 = \Omega(N' p_1^{e_1}) = \Omega(N)$. Similarly, $d_2, \dots, d_2 p_1^{e_1-1}, \dots, d_h p_1^{e_1-1}$ is a chain, and so on. Eventually, all numbers will be in one such chain. If we repeat this for *all* chains of N' , every divisor of N will also belong to a chain.

Constructing a partition for N' can be done in the exact same manner. The base case where we are partitioning the divisors of a single prime power p^k is straightforward – $1, p, \dots, p^k$ is exactly a chain.

Finally, computing the number of such divisors is then a straightforward exercise in brute force after factoring N to compute the exponents e_1, \dots, e_k .

The proof is due to De Bruijn et al[9].

Part III

Advanced Topics

22 Data Structures

- 22.1 Self-Balancing Trees
- 22.2 Persistent Data Structures
- 22.3 Heavy-Light Decomposition

23 Combinatorics

23.1 Convolutions

Fast Fourier Transform

Number Theoretic Transform

24 Strings

24.1 Hashing

Hashing is a concept most familiar from the hash table data structure. The idea behind the structure is to compress a set S of elements from a large set to a smaller set, in order to quickly determine memberships of S by having a direct indexing of the smaller set into an array (which has $\Theta(1)$ look-ups). In this section, we are going to look at hashing in a different light, as a way of speeding up comparisons of data. When comparing two pieces of data a and b of size n for equality, we need to use $\Theta(n)$ time in the worst case since every bit of data must be compared. This is fine if we perform only a single comparison. If we instead wish to compare many pieces of data, this becomes an unnecessary bottleneck. We can use the same kind of hashing as with hash tables, by defining a “random” function $H(x) : S \rightarrow \mathbb{Z}_n$ such that $x \neq y$ implies $H(x) \neq H(y)$ with high probability. Such a function allows us to perform comparisons in $\Theta(1)$ time (with linear preprocessing), by reducing the comparison of arbitrary data to small integers (we often choose n to be on the order of 2^{32} or 2^{64} to get constant-time comparisons). The trade-off lies in correctness, which is compromised in the unfortunate event that we perform a comparison $H(x) = H(y)$ even though $x \neq y$.

FriendBook

Swedish Olympiad in Informatics 2011, Finals

FriendBook is a web site where you can chat with your friends. For a long time, they have used a simple “friend system” where each user has a list of which other users are their “friends”. Recently, a somewhat controversial feature was added, namely a list of your “enemies”. While the friend relation will always be mutual (two users must confirm that they wish to be friends), enmity is sometimes one-way – a person A can have an enemy B , who – by plain animosity – refuse to accept A as an enemy.

Being a poet, you have lately been pondering the following quote.

A friend is someone who dislike the same people as yourself.

Given a FriendBook network, you wonder to what extent this quote applies. More specifically, for how many pairs of users is it the case that they are either friends with identical enemy lists, or are not friends and does not have identical enemy lists?

Input

The first line contains an integer $2 \leq N \leq 5000$, the number of friends on FriendBook. N lines follow, each containing n characters. The c 'th character on the r 'th line S_{rc} species what relation person r has to person c . This character is either

- V** – in case they are friends.
- F** – if r thinks of c as an enemy.
- .** – r has a neutral attitude towards c .

S_{ii} is always **.**, and S_{ij} is **V** if and only if S_{ji} is **V**.

Output

Output a single integer, the number of pairs of persons for which the quote holds.

This problem lends itself very well to hashing. It is clear that the problem is about comparisons – indeed, we are to count the number of pairs of persons who are either friends and have equal enemy lists or are not friends and have unequal enemy lists. The first step is to extract the enemy lists E_i for each person i . This will be a N -length string, where the j 'th character is **F** if person j is an enemy of person i , and **.** otherwise. Basically, we remove all the friendships from the input matrix. Performing naive comparisons on these strings would only give us a $O(N^3)$ time bound, since we need to perform N^2 comparisons of enemy lists of length N bounded only by $O(N)$ in the worst case. Here, hashing comes to our aid. By instead computing $h_i = H(E_i)$ for every i , comparisons of enemy lists instead become comparisons of the integers h_i – a $\Theta(1)$ operation – thereby reducing the complexity to $\Theta(N^2)$.

Alternative solutions exist. For example, we could instead have sorted all the enemy lists, after which we can perform a partitioning of the lists by equality in $\Theta(N^2)$ time. However, this takes $O(N^2 \log N)$ time with naive sorting (or $O(N^2)$ if radix sort is used, but it is more complex) and is definitely more

complicated to code than the hashing approach. Another option is to insert all the strings into a trie, simplifying this partitioning and avoiding the sorting altogether. This is better, but still more complex. While it would have the same complexity, the constant factor would be significantly worse compared to the hashing approach.

This is a common theme among string problems. While most string problems can be solved without hashes, solutions using them tend to be simpler.

The true power of string hashing is not this basic preprocessing step where we can only compare two strings. Another hashing technique allows us to compare arbitrary substring of a string in constant time.

Definition 24.1 — Polynomial Hash

Let $S = s_1 s_2 \dots s_n$ be a string. The *polynomial hash* $H(S)$ of S is the number

$$H(S) = (s_1 p^{n-1} + s_2 p^{n-2} + \dots + s_{n-1} p + s_n) \bmod M$$

As usual when dealing with strings in arithmetic expressions, we take s_i to be some numeric representation of the character, like its ASCII encoding. In C++, `char` is actually a numeric type and is thus usable as a number when using polynomial hashes.

Polynomial hashes have many useful properties.

Theorem 24.1 — Properties of the Polynomial Hash

If $S = s_1 \dots s_n$ is a string and c is a single character, we have that

1. $H(S||c) = (pH(S) + H(c)) \bmod M$
2. $H(c||S) = (H(S) + H(c)p^n) \bmod M$
3. $H(s_2 \dots s_n) = (H(S) - H(s_1)p^{n-1}) \bmod M$
4. $H(s_1 \dots s_{n-1}) = (H(S) - H(s_n))p^{-1} \bmod M$
5. $H(s_{l+1} \dots s_{r-2} s_{r-1}) = (H(s_1 \dots s_{R-1}) - H(s_1 - s_{L-1})p^{R-L}) \bmod M$

Exercise 24.1. Prove the properties of Theorem 24.1

Exercise 24.2. How can we compute the hash of $S||T$ in $O(1)$ given the hashes of the strings S and T ?

Properties 1-4 alone allow us to append and remove characters from the

beginning and end of a hash in constant time. We refer to this property as polynomial hashes being *rolling*. This property allows us to String Matching problem with a single pattern (Section 16.2) with the same complexity as KMP, by computing the hash of the pattern P and then rolling a $|P|$ -length hash through the string we are searching in. This algorithm is called the **Rabin-Karp** algorithm.

Property 5 allows us to compute the hash of any substring of a string in constant time, provided we have computed the hashes $H(S_1), H(s_1s_2), \dots, H(s_1s_2 \dots s_n)$ first. Naively this computation would be $\Theta(n^2)$, but property 1 allows us to compute them recursively, resulting in $\Theta(n)$ precomputation.

Radio Transmission

Baltic Olympiad in Informatics 2009

Given is a string S . Find the shortest string L , such that S is a substring of the infinite string $T = \dots LLLL \dots$

Input

The first and only line of the input contains the string S , with $1 \leq |S| \leq 10^6$.

Output

Output the string L . If there are multiple strings L of the shortest length, you can output any of them.

Assume that L has a particular length l . Then, since T is periodic with length l , S must be too (since it is a substring of T). Conversely, if S is periodic with some length l , we can choose as $L = s_1s_2 \dots s_l$. Thus, we are actually seeking the smallest l such that S is periodic with length l . The constraints this puts on S are simple. We must have that

$$s_1 = s_{l+1} = s_{2l+1} = \dots$$

$$s_2 = s_{l+2} = s_{2l+2} = \dots$$

$$\dots$$

$$s_l = s_{2l} = s_{3l} = \dots$$

Using this insight as-is gives us a $O(|S|^2)$ algorithm, where we first fix l and then verify if those constraints hold. The idea is sound, but a bit slow. Again, the problematic step is that we need to perform many slow, linear-time comparisons.

If we look at what comparisons we actually perform, we are actually comparing two substrings of S with each other:

$$s_1 s_2 \dots s_{n-l+1} = s_{l+1} s_{l+2} \dots s_n$$

Thus we are actually performing a linear number of substring comparisons, which we now know are actually constant-time operations after linear preprocessing. Hashes thus gave us a $\Theta(N)$ algorithm.

Radio Transmission

```

1  H lh = 0, Rh = 0;
2  int l = 0;
3  for (int i = 1; i <= n; ++i) {
4      Lh = (Lh * p + S[i]) % M;
5      Rh = (S[n - i + 1] * p^(i - 1) + Rh) % M;
6      if (Lh == Rh) {
7          l = i;
8      }
9  }
10 cout << n - l << endl;
```

Polynomial hashes are also a powerful tool to compare something against a large number of strings using hash sets. For example, we could actually use hashing as a replacement for Aho-Corasick. However, we would have to perform one pass of rolling hash for each different pattern length. If the string we are searching in is N and the sum of pattern lengths are P , this is not $O(N + P)$ however. If we have k different pattern lengths, their sum must be at least $1 + 2 + \dots + k = \Theta(k^2)$, so $k = O(\sqrt{P})$.

Substring Range Matching

Petrozavodsk Winter Training Camp 2015

Given N strings s_1, s_2, \dots, s_N and a list of queries of the form L, R, S , answer for each such query the number of strings in s_L, s_{L+1}, \dots, s_R which contain S as a substring.

Input

The first line contains $1 \leq N \leq 50\,000$ and the number of queries $0 \leq Q \leq 100\,000$. The next N lines contains the strings s_1, s_2, \dots, s_N , one per line. The next Q lines contains one query each. A query is given by the integers $1 \leq L \leq R \leq N$ and a string S .

The sum of $|S|$ over all queries is at most 20 000. The lengths $|s_1| + |s_2| +$

$\dots + |s_N|$ is at most 50 000.

Output

For each query L, R, S , output a line with the answer to the query.

Let us focus on how to solve the problem where every query has the same string S . In this case, we would first find which of the strings s_i that S is contained in using polynomial hashing. To respond to a query, could for example keep a set of all the i where s_i was an occurrence together with how many smaller s_i contained the string (i.e. some kind of partial sum). This would allow us to respond to a query where $L = 1$ using a upper bound in our set. Solving queries of the form $[1, R]$ is equivalent to general intervals however, since the interval $[L, R]$ is simply the interval $[1, R]$ with the interval $[1, L - 1]$ removed. This procedure would take $\Theta(\sum |s_i|)$ time to find the occurrences of S , and $O(Q \log N)$ time to answer the queries.

When extending this to the general case where our queries may contain different S , we do the same thing but instead find the occurrences of all the patterns of the same length p simultaneously. This can be done by keeping the hashes of those patterns in a map, to allow for fast look-up of our rolling hash. Since there can only be at most $\sqrt{20\,000} \approx 140$ different pattern lengths, we must perform about $140 \cdot 50\,000 \approx 7\,000\,000$ set look-ups, which is feasible.

Substring Range Matching

```
1  int countInterval(int upTo, const set<pii>& s) {
2      auto it = s.lower_bound(pii(upTo + 1, 0));
3      if (it == s.begin()) return 0;
4      return (--it)->second;
5  }
6
7  int main() {
8      int N, Q;
9      cin >> N >> Q;
10     vector<string> s(N);
11     rep(i,0,N) cin >> s[i];
12
13     map<int, set<string>> patterns;
14
15     vector<tuple<int, int, string>> queries;
16     rep(i,0,Q) {
17         int L, R;
18         string S;
19         cin >> L >> R >> S;
20         queries.emplace_back(L, R, S);
```

```

21     patterns[sz(s)].insert(S);
22 }
23
24 map<H, set<pii>> hits;
25 trav(pat, patterns) {
26     rep(i,0,N) {
27         vector<H> hashes = rollHash(s[i], pat.first);
28         trav(h, hashes)
29             if (pat.second.count(h))
30                 hits[h].emplace(i, sz(hits[h]) + 1);
31     }
32 }
33
34 trav(query, queries) {
35     H h = polyHash(get<2>(query));
36     cout << countInterval(R, hits[h]) - countInterval(L-1, hits[h]) << endl;
37 }
38 }

```

Exercise 24.3. Hashing can be used to determine which of two substrings are the lexicographically smallest one. How? Extend this result to a simple $\Theta(n \log S + S)$ construction of a suffix array, where n is the number of strings and S is the length of the string.

The Parameters of Polynomial Hashes

Until now, we have glossed over the choice of M and p in our polynomial hashing. These choices happen to be important. First of all, we want M and p to be relatively prime. This ensures p has an inverse modulo M , which we use when erasing characters from the end of a hash. Additionally, $p^i \bmod M$ have a smaller period when p and M share a factor.

We wish M to be sufficiently large, to avoid hash collisions. If we compare the hashes of c strings, we want $M = \Omega(\sqrt{c})$ to get a reasonable chance at avoiding collisions. However, this depends on how we use hashing. p must be somewhat large as well. If p is smaller than the alphabet, we get trivial collisions such as $H(10) = H(p)$.

Whenever we perform rolling hashes, we must have $(M - 1)p < 2^{64}$ if we use 64-bit unsigned integers to implement hashes. Otherwise, the addition of a character would overflow. If we perform substring hashes, we instead need that $(M - 1)^2 < 2^{64}$, since we perform multiplication of a hash and an arbitrary power of p . When using 32-bit or 128-bit hashes, these limits change correspondingly. Note that the choice of hash size depends on how large an M we can choose, which affect collision rates.

One might be tempted to choose $M = 2^{64}$ and use the overflow of 64-bit integers as a cheap way of using hashes modulo 2^{64} . This is a bad idea, since it is possible to construct strings which are highly prone to collisions.

Definition 24.2 — Thue-Morse Sequence

Let the binary sequence τ_i be defined as

$$\tau_i = \begin{cases} 0 & \text{if } i = 0 \\ \tau_{i-1} \overline{\tau_{i-1}} & \text{if } i > 0 \end{cases}$$

The *Thue-Morse* sequence is the infinite sequence τ_i as $i \rightarrow \infty$.

This sequence is well-defined since τ_i is a prefix of τ_{i-1} , meaning each recursive step only append a string to the sequence. It starts 0, 01, 0110, 01101001, 0110100110010110.

Exercise 24.4. Prove that τ_{2i} is a palindrome.

Theorem 24.2

For a polynomial hash H with an odd p , $2^{\frac{n(n+1)}{2}} \mid H(\overline{\tau_n}) - H(\tau_n)$.

Proof. We will prove this by induction on n . For $n = 0$, we have $1 \mid H(\overline{\tau_n}) - H(\tau_n)$ which is vacuously true.

In our inductive step, we have that

$$H(\tau_n) = H(\tau_{n-1} \parallel \overline{\tau_{n-1}}) = p^{2^{n-1}} \cdot H(\overline{\tau_{n-1}}) + H(\tau_{n-1})$$

and

$$H(\overline{\tau_n}) = H(\overline{\tau_{n-1}} \parallel \tau_{n-1}) = p^{2^{n-1}} \cdot H(\tau_{n-1}) + H(\overline{\tau_{n-1}})$$

Then,

$$\begin{aligned} H(\overline{\tau_n}) - H(\tau_n) &= p^{2^{n-1}} (H(\tau_{n-1}) - H(\overline{\tau_{n-1}})) + (H(\overline{\tau_{n-1}}) - H(\tau_{n-1})) \\ &= (p^{2^{n-1}} - 1)(H(\tau_{n-1}) - H(\overline{\tau_{n-1}})) \end{aligned}$$

Note that $p^{2^{n-1}} - 1 = (p^{2^{n-2}} - 1)(p^{2^{n-2}} + 1)$. If p is odd, the second factor is divisible by 2. By expanding $p^{2^{n-2}}$, we can prove that $p^{2^{n-1}}$ is divisible by 2^n .

Using our induction assumption, we have that

$$2^n \cdot 2^{\frac{(n-1)n}{2}} \mid (p^{2^{n-1}} - 1)(H(\tau_{n-1}) - H(\overline{\tau_{n-1}}))$$

But $2^n \cdot 2^{\frac{(n-1)n}{2}} = 2^{\frac{n(n+1)}{2}}$, proving our statement. \square

This means that we can construct a string of length linear in the bit size of M that causes hash collisions if we choose M as a power of 2, explaining why it is a bad choice.

2D Polynomial Hashing

Polynomial hashing can also be applied to pattern matching in grids, by first performing polynomial hashing on all rows of the grid (thus reducing the grid to a sequence) and then on the columns.

Surveillance

Swedish Olympiad in Informatics 2016, IOI Qualifiers

Given a matrix of integers $A = (a_{r,c})$ find all occurrences of another matrix $P = (p_{r,c})$ in A which may differ by a constant C . An occurrence (i, j) means that $a_{i+r,j+c} = p_{r,c} + C$ where C is a constant.

If we assume that $C = 0$, the problem is reduced to simple 2D pattern matching, which is easily solved by hashing. The requirement that such a pattern should be invariant to addition by a constant is a bit more complicated.

How would we solve this problem in one dimension, i.e. when $r = 1$? In this case, we have that a match on column j would imply

$$a_{1,j} - p_{1,1} = c$$

$$\dots$$

$$a_{1,j+n-1} - p_{1,n} = c$$

Since c is arbitrary, this means the only condition is that

$$a_{1,j} - p_{1,1} = \dots = a_{1,j+n-1} - p_{1,n} = c$$

Rearranging this gives us that

$$a_{1,j} - a_{1,j+1} = p_{1,1} - p_{1,2}$$

$$a_{1,j+1} - a_{1,j+2} = p_{1,2} - p_{1,3}$$

...

By computing these two sequences of the adjacent differences of elements $a_{1,i}$ and $r_{1,j}$, we have reduced the problem to substring matching and can apply hashing. In 2D, we can do something similar. For a match (i, j) , it is sufficient that this property holds for every line and every column in the match. We can then find matches using two 2D hashes.

Problem 24.1

Chasing Subs – chasingsubs

24.2 Dynamic Hashing

A Discrete Mathematics

This appendix reviews some basic discrete mathematics. Without a good grasp on the foundations of mathematics, algorithmic problem solving is basically impossible. When we analyze the efficiency of algorithms, we use sums, recurrence relations and a bit of algebra. Some basic topics, such as set theory, are essential to even understand some of the proofs and problems in this book.

This mathematical preliminary touches lightly upon these topics and is meant to complement a high school education in mathematics in preparation for the remaining text. While you can probably get by with the mathematics from this chapter, we highly recommend that you (at some point) delve deeper into discrete mathematics.

We do assume that you are familiar with induction or contradiction, and mathematics that is part of a pre-calculus course (trigonometry, polynomials, etc). Some more mathematically advanced parts of this book will go beyond these assumptions, but this is only the case in very few places.

A.1 Logic

In mathematics, we often deal with truths and falsehoods in the form of theorems, proofs, counter-examples and so on. Mathematical **logic** is a very exact discipline, and a precise language has been developed to help us deal with logical statements.

For example, consider the statements

1. an integer is either odd or even,
2. programming is more fun than mathematics,
3. x is negative if and only if x^3 is negative,
4. every apple is blue,
5. there exists an integer.
6. if there exists an odd number divisible by 6, every integer is even,

The first statement uses the logical connective *or*. It connects two statements, and requires only one of them to be true in order for the whole statement to be true. Since any integer is either odd or even, the statement is true.

The second statement is not really a logical statement. While we might have a personal conviction regarding the entertainment value of programming and maths, it is hard to consider the statements as having a truth value.

The third statement tells us that two statements are equivalent – one is true exactly when the other is. This is also a true statement by some simple algebraic manipulations.

The fourth statement concerns *every* object of some kind. It is a false statement, a fact that can be proved by exhibiting e.g., a green apple.

The fifth statement is true. It asks whether something exists, a statement we can prove by presenting an integer such as 42.

The sixth and last statement complicates matters by introducing an *implication*. It is a two-part statement, which only makes a claim regarding the second part if the first part is true. Since no odd number divisible by 6 exists, it makes no statement about the evenness of every integer. Thus, this implication is true.

To express such statements, a language has been developed where all these logical operations such as existence, implication and so on have symbols assigned to them. This enables us to remove the ambiguity inherent in the

English language, which is of utmost importance when dealing with the exactness required by logic.

The **disjunction** (a is true **or** b is true) is a common logical connective. It is given the symbol \vee , so that the above statement is written as $a \vee b$. Another common connective, the **conjunction** (a is true **and** b is true) is assigned the symbol \wedge . For example, we write that $a \wedge b$ for the statement that both a and b are true.

The third statement introduced the **equivalence**, a statement of the form “ a is true if, and only if, b is true”. This is the same as $a \rightarrow b$ (the only if part) and $b \rightarrow a$ (the if part). We use the symbol \leftrightarrow , which follows naturally for this reason. The statement would then be written as

$$x < 0 \leftrightarrow x^3 < 0$$

Logic also contains **quantifiers**. The fourth statement, that every apple is blue, actually makes a large number of statements – one for each apple. This concept is captured using the **universal quantifier** \forall , read as “for every”. For example, we could write the statement as

$$\forall \text{ apple } a : a \text{ is blue}$$

In the fifth statement, another quantifier was used, which speaks of the existence of something; the **existential quantifier** \exists , which we read as “there exists”. We would write the second statement as

$$\exists x : x \text{ is an integer}$$

An **implication** is a statement of the form “if a is true, then b must also be true”. This is a statement on its own, which is true whenever a is false (meaning it does not say anything of b), or when a is true and b is true. We use the symbol \rightarrow for this, writing the statement as $a \rightarrow b$. The sixth statement would hence be written as

$$(\exists p : p \text{ is prime} \wedge p \text{ is divisible by } 6) \rightarrow \forall \text{ prime } p : p \text{ is even}$$

The **negation** operator \neg inverts a statement. The statement “no penguin can fly” would thus be written as

$$\neg(\exists \text{ penguin } p : p \text{ can fly})$$

or, equivalently

$$\forall \text{ penguin } p : \neg p \text{ can fly}$$

Exercise A.1. Write the following statements using the logical symbols, and determine whether they are true or false:

- 1) If a and b are odd integers, $a + b$ is an even integer,
- 2) a and b are odd integers if and only if $a + b$ is an even integer,
- 3) Whenever it rains, the sun does not shine,
- 4) ab is 0 if and only if a or b is 0

Our treatment of logic ends here. Note that much is left unsaid – it is a most rudimentary walk-through. This section is mainly meant to give you some familiarity with the basic symbols used in logic, since they will appear later. If you wish to gain a better understanding of logic, you can follow the references in the chapter notes.

A.2 Sets and Sequences

A **set** is an unordered collection of **distinct** objects, such as numbers, letters, other sets, and so on. The objects contained within a set are called its **elements**, or **members**. Sets are written as a comma-separated list of its elements, enclosed by curly brackets:

$$A = \{2, 3, 5, 7\}$$

In this example, A contains four elements: the integers 2, 3, 5 and 7.

Because a set is unordered and only contains distinct objects, the set $\{1, 2, 2, 3\}$ is the exact same set as $\{3, 2, 1, 1\}$ and $\{1, 2, 3\}$.

If x is an element in a set S , we write that $x \in S$. For example, we have that $2 \in A$ (referring to our example A above). Conversely, we use the notation $x \notin S$ when the opposite holds. We have e.g., that $11 \notin A$.

Another way of describing the elements of a set uses the **set builder** notation, in which a set is constructed by explaining what properties its elements should have. The general syntax is

$$\{\text{element} \mid \text{properties that the element must have}\}$$

To construct the set of all even integers, we would use the syntax

$$\{2i \mid i \text{ is an integer}\}$$

which is read as “the set containing all numbers of the form $2i$ where i is an integer. To construct the set of all primes, we would write

$$\{p \mid p \text{ is prime}\}$$

Certain sets are used often enough to be assigned their own symbols:

- \mathbb{Z} – the set of integers $\{\dots, -2, -1, 0, 1, 2, \dots\}$,
- \mathbb{Z}_+ – the set of *positive* integers $\{1, 2, 3, \dots\}$,
- \mathbb{N} – the set of *non-negative* integers $\{0, 1, 2, \dots\}$,
- \mathbb{Q} – the set of all rational numbers $\{\frac{p}{q} \mid p, q \text{ integers where } q \neq 0\}$,
- \mathbb{R} – the set of all real numbers,
- $[n]$ – the set of the first n positive integers $\{1, 2, \dots, n\}$,
- \emptyset – the empty set.

Exercise A.2. 1) Use the set builder notation to describe the set of all odd integers.

2) Use the set builder notation to describe the set of all negative integers.

3) Compute the elements of the set $\{k \mid k \text{ is prime and } k^2 \leq 30\}$.

A set A is a **subset** of a set S if, for every $x \in A$, we also have $x \in S$ (i.e., every member of A is a member of S). We denote this with $A \subseteq S$. For example

$$\{2, 3\} \subseteq \{2, 3, 5, 7\}$$

and

$$\left\{\frac{2}{4}, 2, \frac{-1}{7}\right\} \subseteq \mathbb{Q}$$

For any set S , we have that $\emptyset \subseteq S$ and $S \subseteq S$. Whenever a set A is not a subset of another set B , we write that $A \not\subseteq B$. For example,

$$\{2, \pi\} \not\subseteq \mathbb{Q}$$

since π is not a rational number.

We say that sets A and B are **equal** whenever $x \in A$ if and only if $x \in B$. This is equivalent to $A \subseteq B$ and $B \subseteq A$. Sometimes, we will use the latter condition when proving set equality, i.e., first proving that every element of A must also be an element of B and then the other way round.

Exercise A.3. 1) List all subsets of the set $\{1, 2, 3\}$.

2) How many subsets does a set containing n elements have?

3) Determine which of the following sets are subsets of each other:

- \emptyset
- \mathbb{Z}
- \mathbb{Z}_+
- $\{2k \mid k \in \mathbb{Z}\}$
- $\{2, 16, 12\}$

Sets also have many useful operations defined on them. The **intersection** $A \cap B$ of two sets A and B is the set containing all the elements which are members of **both** sets, i.e.,

$$x \in A \cap B \Leftrightarrow x \in A \wedge x \in B$$

If the intersection of two sets is the empty set, we call the sets **disjoint**. A similar concept is the **union** $A \cup B$ of A and B , defined as the set containing those elements which are members of **either** set.

For example, if

$$X = \{1, 2, 3, 4\}, Y = \{4, 5, 6, 7\}, Z = \{1, 2, 6, 7\}$$

Then,

$$X \cap Y = \{4\}$$

$$X \cap Y \cap Z = \emptyset$$

$$X \cup Y = \{1, 2, 3, 4, 5, 6, 7\}$$

$$X \cup Z = \{1, 2, 3, 4, 6, 7\}$$

Exercise A.4. Compute the intersection and union of:

- 1) $A = \{1, 4, 2\}, B = \{4, 5, 6\}$
- 2) $A = \{a, b, c\}, B = \{d, e, f\}$
- 3) $A = \{\text{apple, orange}\}, B = \{\text{pear, orange}\}$

A *sequence* is an **ordered** collection of values (predominantly numbers) such as 1, 2, 1, 3, 1, 4, Sequences will mostly be a list of sub-scripted variables, such as a_1, a_2, \dots, a_n . A shorthand for this is $(a_i)_{i=1}^n$, denoting the sequence of variables a_i where i ranges from 1 to n . An infinite sequence is given ∞ as its upper bound: $(a_i)_{i=1}^\infty$.

A.3 Sums and Products

The most common mathematical expressions we deal with are sums of sequences of numbers, such as $1 + 2 + \cdots + n$. Such sums often have a variable number of terms and complex summands, such as $1 \cdot 3 \cdot 5 + 3 \cdot 5 \cdot 7 + \cdots + (2n+1)(2n+3)(2n+5)$. In these cases, sums given in the form of a few leading and trailing terms, with the remaining part hidden by \dots is too imprecise. Instead, we use a special syntax for writing sums in a formal way – the *sum operator*:

$$\sum_{i=j}^k a_i$$

The symbol denotes the sum of the $j - k + 1$ terms $a_j + a_{j+1} + a_{j+2} + \cdots + a_k$, which we read as “the sum of a_i from j to k ”.

For example, we can express the sum $2 + 4 + 6 + \cdots + 12$ of the 6 first even numbers as

$$\sum_{i=1}^6 2i$$

Exercise A.5. Compute the sum

$$\sum_{i=-2}^4 2 \cdot i - 1$$

Many useful sums have closed forms – expressions in which we do not need sums of a variable number of terms.

Exercise A.6. Prove the following identities:

$$\begin{aligned} \sum_{i=1}^n c &= cn \\ \sum_{i=1}^n i &= \frac{n(n+1)}{2} \\ \sum_{i=1}^n i^2 &= \frac{n(n+\frac{1}{2})(n+1)}{3} \\ \sum_{i=0}^n 2^i &= 2^{n+1} - 1 \end{aligned}$$

The sum of the inverses of the first n natural numbers happen to have a very neat approximation, which we will occasionally make use of later on:

$$\sum_{i=1}^n \frac{1}{i} \approx \ln n$$

This is a reasonable approximation, since $\int_1^l \frac{1}{x} dx = \ln(l)$

There is an analogous notation for products, using the *product operator* \prod :

$$\prod_{i=j}^k a_i$$

denotes the product of the $j - k + 1$ terms $a_j \cdot a_{j+1} \cdot a_{j+2} \cdots a_k$, which we read as “the product of a_i from j to k ”.

In this way, the product $1 \cdot 3 \cdot 5 \cdots (2n - 1)$ of the first n odd integers can be written as

$$\prod_{i=1}^n 2i - 1$$

Exercise A.7. Prove that

$$(n+2) \prod_{i=1}^n i + \frac{n}{n+1} \prod_{i=1}^{n+2} i = \prod_{i=1}^{n+2} i$$

Chapter Notes

If you need a refresher on some more basic mathematics, such as single-variable calculus, *Calculus* [26] by Michael Spivak is a solid textbook. It is not the easiest book, but one the best undergraduate text on single-variable calculus if you take the time to work it through.

For a gentle introduction to discrete mathematics, *Discrete and Combinatorial Mathematics: An Applied Introduction* [12] by Ralph Grimaldi is a nice book with a lot of breadth.

Logic in Computer Science [14] is an introduction to formal logic, with many interesting computational applications. The first chapter on propositional logic is sufficient for most algorithmic problem solving, but the remaining chapters shows many non-obvious applications that makes logic relevant to computer science.

One of the best works on discrete mathematics ever produced for the aspiring algorithmic problem solver is *Concrete Mathematics* [15], co-authored by famous computer scientist Donald Knuth. It is rather heavy-weight, and probably serves better as a more in-depth study of the foundations of discrete mathematics rather than an introductory text.

Graph Theory [10] by Reinhard Diestel is widely acknowledged as the go-to book on more advanced graph theory concepts. The author is freely available for viewing at the book's home page¹.

¹<http://diestel-graph-theory.com/>

Hints

1.1 Try dividing cards into smaller piles that can be sorted separately.

1.6 The optimal number of questions is 6.

2.12 Try solving it for the special case $y = 2$ first.

5.1 In the best case, line 4 of the insertion sort pseudo code never executes.

5.3 When is $\log^2 n < n$?

5.4 $c = 2$ for the upper bound.

5.5

1) Yes.

2) No.

5.6 Binomial expansion.

5.7

6.10

1. Sum the maximum number of steps each element can move.

2. What is the limit of an geometric series?

7.4 Use that $1.61 + 1 > 1.61^2$ and $1.62 + 1 < 1.62^2$.

7.5 The positive root of the equation $x^3 = x^2 + x + 1$ lies between 1.83 and 1.84.

7.6

3) The n choices are which of the two letters to put on each position in the string.

4) The n choices are whether to include each element or not.

7.7 Since the three recursions are structurally identical, they will have the same time complexity $T(n)$.

Solutions

1.1 One possible solution is to first divide the cards into separate piles by values $1 - 100\,000$, $100\,001 - 200\,000$, \dots . If we sort each such pile, the entire stack of cards is sorted by putting the piles together. Each such pile can be sorted the same way by instead dividing the cards up based on their $10\,000$ digits, and so on.

1.2

5) The input consists of two integers a and b , not both 0. The output should be the greatest common divisor of a and b .

6) The input consists of a sequence of real numbers, the coefficients x_i of a polynomial. The output should be a real number that is a root of the polynomial.

7) The input consists of two integers a and b . The output should be the product ab .

1.6 One can achieve 6 questions by always asking about the midpoint of the range of possible numbers. For example, by asking about the number 50, one knows if the correct number is between $1 - 49$ or $51 - 100$.

1.8 Given an algorithm that is correct with a probability $0.5 + \alpha$ for some $\alpha > 0$, we can find the correct answer by running it many times and choosing the answer that was most common.

1.12

Palindrome If we let the input n -letter word S have the letters s_0, s_1, \dots, s_{n-1} , it reads the same backwards and forwards if $s_0s_1 \dots s_{n-1} = s_{n-1}s_{n-2} \dots s_0$. We thus need to check all the letter pairs (s_0, s_{n-1}) , (s_1, s_{n-2}) and so forth for equality.

```
1: procedure PALINDROME(string  $S$ )
2:   for  $i$  from 0 to  $n - 1$  do
3:     if  $S_i \neq S_{n-1-i}$  then
4:       return false
5:   return true
```

Primality The problem can be solved by checking all the numbers between 2 and $n - 1$ to see if any of them are divisors of n . If not, then it is prime. This follows from the definition, and the fact that a (positive) divisor of a positive integer can not be greater than the integer itself.

```

1: procedure PRIMALITY(integer  $n$ )
2:   for  $i$  from 2 to  $n - 1$  do
3:     if  $i$  divides  $n$  then
4:       return false
5:   return true

```

2.12 We only analyze the case where x and y are positive. Assume that $0 \leq a < \frac{x}{y} \leq a + 1$, so that the result when rounded to an integer away from zero is $a + 1$. Multiplying by y gives us $ay < x \leq ay + y$, so that $ay \leq x - 1 < ay + y$ (since all values are now integers). Finally, adding y to both inequalities gives us $ay + y \leq x - 1 + y < ay + 2y$. After dividing by y , we get $a + 1 \leq \frac{x-1+y}{y} < a + 2$. This means that the result of $\frac{x-1+y}{y}$ rounded towards zero is $a + 1$, which is what we wanted.

Analysis for negative x and y is similar.

5.1 Consider the case when the array is already sorted. In this case, the inner loop on line 4 never executes, since $A[j] \geq A[j - 1]$ for all j . Thus, only the lines that take linear time in total are executed, making $O(n)$ an upper bound on the base case. On the other hand, the loop on line 2 always executes a linear number of times no matter the case, so $\Omega(n)$ is also a lower bound. Thus, the algorithm has a $\Theta(n)$ best-case running time.

5.2 To compute the sum in $\Theta(n)$ time, we can add all the variables to a counter using a for loop, one at a time.

To solve the problem in constant time, the formula $1 + 2 + \cdots + n = \frac{n(n+1)}{2}$ can be used.

5.3 Let $n_0 = 7$. For any $n \geq 1$, we have $\log^2 n < n$ as $n < 2^n$ (which can be proved using either induction or simple calculus). In this case, $10n^2 + 7n - 5 + \log^2 n \leq 10n^2 + n^2 + n^2 = 12n^2$. Thus, with $c = 12$ we get the required statement.

5.4 Clearly $\max\{f(n), g(n)\} \leq f(n) + g(n)$ since the maximum of the two functions is always equal to one of the functions. This means that $f(n) + g(n) = \Omega(\max\{f(n), g(n)\})$ with $c = 1$. Similarly, $f(n) + g(n) \leq 2 \max\{f(n), g(n)\}$

by the fact that each function individually is smaller than their maximum. Thus $f(n) + g(n) = O(\max\{f(n), g(n)\})$ with $c = 2$. Together this proves the statement.

5.5

8) This is clear with $c = 2$.

9) For any c , picking n such that $2^n > c$ gives us $2^{2n} = 2^n \cdot 2^n > c2^n$, so no c can satisfy the definition.

5.6 First, note that polynomials of higher powers are always greater than polynomials of lower powers eventually:

$$an^k < n^{k+1}$$

is true whenever $n > a$.

Next, we can write $(n + a)^b$ as the sum of n^b plus a lot of terms of lower powers of n using the formula for the binomial expansion. However, this means that $\max\{n^b, (n + a)^b - n^b\} = n^b$ for sufficiently large n . Thus, $(n + a)^b = n^b + ((n + a)^b - n^b) = \Theta(\max\{n^b, (n + a)^b - n^b\}) = \Theta(n^b)$ by a previous result.

5.7

6.10 Assume that the tree has $n = 2^k$ elements. The bottom n elements will move 0 steps. The $\frac{n}{2}$ elements of the next layer moves at most 1 step. The $\frac{n}{4}$ elements of the next layer moves at most 2 step, and so on. In total, this means that there are most

$$\frac{n}{2} + 2\frac{n}{4} + 3\frac{n}{8} + \cdots + k\frac{n}{2^k}$$

movements. Note that

$$\frac{n}{2} + \frac{n}{4} + \cdots \leq n$$

$$\frac{n}{4} + \frac{n}{8} + \cdots \leq \frac{n}{2}$$

$$\frac{n}{8} + \frac{n}{16} + \cdots \leq \frac{n}{4}$$

and so on. If we sum up all of these inequalities, we get that the original sum

$$\frac{n}{2} + 2\frac{n}{4} + 3\frac{n}{8} + \cdots + k\frac{n}{2^k} \leq$$

$$n + \frac{n}{2} + \frac{n}{4} + \cdots \leq 2n$$

proving that the elements are at most moved a linear number of time.

7.1 They are 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233, 377.

7.4 Assume that $T(n) \geq 1.61^n$ for all $n < n'$, and for $n = 1$. Then $T(n') \geq 1.61^{n'-1} + 1.61^{n'-2} = 1.61^{n'-2}(1.61 + 1) \geq 1.61^{n'}$ so that $T(n) \geq 1.61^n$ also for $n = n'$. By the principle of induction, $T(n) \geq 1.61^n$ holds for all $n \geq 0$.

The proof is similar for the upper bound.

7.5 We have that the time function $T(n) \geq T(n-1) + T(n-2) + T(n-3)$. If, by induction, $T(k) > 1.83^k$ for all $k < n$ we get

$$\begin{aligned} T(n) &\geq 1.83^{n-1} + 1.83^{n-2} + 1.83^{n-3} \\ &= 1.83^{n-3}(1 + 1.83 + 1.83^2) \\ &\geq 1.83^{n-3} 1.83^3 \\ &= 1.83^n \end{aligned}$$

so the claim holds for $T(n)$ too. Proving the upper bound is similar

7.6

10) Let $A(n)$ be the number of such strings. If the last character of the string was a B, the remaining string can be formed in $A(n-1)$ ways. If the last character of the string was an A, the second last character must have been a B (to avoid two consecutive A's). There are $A(n-2)$ ways in which the remaining string can be formed after fixing these two letters, so that $A(n) = A(n-1) + A(n-2)$. The base cases are $A(0) = 1$ and $A(1) = 2$.

11) Let $B(n)$ be the number of such subsets. If the element n is to be included in the subset, we can choose the remaining $n-1$ elements in $B(n-1)$ ways. If the element n is to be excluded from the subset, the element $n-1$ must be according to the problem. The remaining $n-2$ elements can then be chosen in $B(n-2)$ ways, for the recursion $B(n) = B(n-1) + B(n-2)$. The base cases are $A(0) = 1$ and $A(1) = 2$.

7.7 The time complexity fulfills $T(n) = 2T(n-1) + O(1)$. By induction, we get $T(n) = \Theta(2^n)$.

Bibliography

- [1] Exact Exponential Algorithms. *Fedor V. Fomin and Dieter Kratsch*. Springer, 2010.
- [2] Noga Alon, Raphy Yuster, and Uri Zwick. Color-coding: A new method for finding simple paths, cycles and other small subgraphs within large graphs. In *Proceedings of the Twenty-Sixth Annual ACM Symposium on Theory of Computing, STOC '94*, page 326–335, New York, NY, USA, 1994. Association for Computing Machinery.
- [3] Sanjeev Arora and Boaz Barak. *Computational Complexity: A Modern Approach*. Cambridge University Press, 2009.
- [4] David Beazley and Brian K. Jones. *Python Cookbook*. O'Reilly, 2013.
- [5] Joshua Bloch. *Effective Java*. Pearson Education, 2008.
- [6] Xuan Cai. Canonical coin systems for change-making problems. In *2009 Ninth International Conference on Hybrid Intelligent Systems*, volume 1, pages 499–504, Aug 2009.
- [7] Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms*. The MIT Press, 3rd edition, 2009.
- [8] Marek Cygan, Fedor V. Fomin, Łukasz Kowalik, Daniel Lokshtanov, Dániel Marx, Marcin Pilipczuk, Michał Pilipczuk, and Saket Saurabh. *Parameterized Algorithms*. Springer, 2015.
- [9] Nicolaas G. de Bruijn, Cornelia van Ebbenhorst Tengbergen, and Dirk Kruyswijk. On the set of divisors of a number. *Nieuw Archief voor Wiskunde*, serie 2, 23:191–193, 1951.
- [10] Reinhard Diestel. *Graph Theory*. Springer, 2016.
- [11] Philippe Flajolet and Robert Sedgewick. *An Introduction to the Analysis of Algorithms*. Addison-Wesley, 2013.

- [12] Ralph P. Grimaldi. *Discrete and Combinatorial Mathematics: An Applied Introduction*. Pearson Education, 2003.
- [13] Godfrey H. Hardy and Edward M. Wright. *An Introduction to the Theory of Numbers*. Oxford University Press, 2008.
- [14] Michael Huth. *Logic in Computer Science*. Cambridge University Press, 2004.
- [15] Donald E. Knuth, Oren Patashnik, and Ronald Graham. *Concrete Mathematics: A Foundation for Computer Science*. Addison-Wesley, 1994.
- [16] George S. Lueker. *Two NP-complete Problems in Nonnegative Integer Programming*. Princeton University. Department of Electrical Engineering, 1975.
- [17] Robert C. Martin. *Clean Code: A Handbook of Agile Software Craftsmanship*. Pearson Education, 2009.
- [18] Steve McConnell. *Code Complete: A Practical Handbook of Software Construction*. Microsoft Press, 2004.
- [19] Scott Meyers. *Effective STL*. O'Reilly, 2001.
- [20] Scott Meyers. *Effective C++*. O'Reilly, 2005.
- [21] Scott Meyers. *Effective Modern C++*. O'Reilly, 2014.
- [22] Christos Papadimitriou. *Computational Complexity*. Addison-Wesley, 1994.
- [23] Charles Petzold. *CODE*. Microsoft Press, 2000.
- [24] Victor Shoup. *A Computational Introduction to Number Theory and Algebra*. Cambridge University Press, 2008.
- [25] Brett Slatkin. *Effective Python*. Addison-Wesley, 2015.
- [26] Michael Spivak. *Calculus*. Springer, 1994.
- [27] Bjarne Stroustrup. *The C++ Programming Language*. Addison-Wesley, 2013.

- [28] Jeffrey Ullman and John Hopcroft. *Introduction to Automata Theory, Languages, and Computation*. Pearson Education, 2014.
- [29] Mark A. Weiss. *Data Structures and Algorithm Analysis in C++*. Pearson, 2013.

Index

K_n , 132

addition principle, 267

adjacency lists, 137

adjacency matrix, 136

algorithm, 5

amortized complexity, 90

and, 373

and operator, 29

array, 40

assignment operator, 21

auto, 24

BFS, 139, 232

bijection, 271

binary search, 210

binary tree, 104

binomial coefficient, 278

bipartite matching, 253

boolean, 24

breadth-first search, 139, 232

break statement, 32

char, 22

closed trail, 143

closed walk, 143

combinatorics, 267

comment, 19

comparison operators, 28

compiler, 16

complete graph, 132

component, 143

composite number, 310

computational problem, 3

conjunction, 373

connected, 143

connected component, 143

continue statement, 32

correctness, 7

cycle, 143

cycle decomposition, 273

data structure, 97

degree, of vertex, 133

Dijkstra's Algorithm, 239

directed graph, 135

disjoint sets, 376

disjunction, 373

divide and conquer, 201

divides exactly, 316

divisibility, 303

divisor, 303

double, 23

Dyck path, 282

dynamic array, 98

edge, 131

element, 374

equivalence, 373

existential quantifier, 373

factorial, 270

fixed-size array, 97

float, 23

- flow network, 247
- for loop, 31
- generate and test, 150
- graph, 131
- graph game, 299
- heap, 105
- identity permutation, 272
- if statements, 29
- implication, 373
- independent set, 162
- input description, 3
- insertion sort, 83
- instance, 4
- int, 22
- intersection
 - of sets, 376
- judgment, 11
- KMP, 264
- Knuth-Morris-Pratt, 264
- KS.Dev, 10
- lambda, 42
- length
 - of path, 138
- logic, 372
- long long, 22
- losing position, 300
- main function, 19
- maximum matching, 253
- member, 374
- memory complexity, 92
- modular inverse, 331
- modulo, 26
- multiplication principle, 267
- negation, 29, 373
- neighbour, 133
- next_permutation, 58
- NP-complete, 92
- online judge, 10
- operator, 26
- optimization problem, 149
- or, 373
- or operator, 29
- oracle, 92
- order
 - of a permutation, 274
- output description, 3
- partial correctness, 7
- path, 138
- permutation, 270, 271
 - cycles, 273
 - identity, 272
 - inverse of, 272
 - multiplication, 272
 - order, 274
- position
 - of game, 299
- prime number, 310
- priority queue, 103
- problem, 3
- product operator, 378
- programming language, 8
- pseudo code, 9
- quantifier, 373
- query complexity, 92
- queue, 51, 102

quotient, 329

Rabin-Karp, 364

recursion, 117

recursive definition, 117

remainder, 329

sequence, 376

set, 374

simple graph, 131

stable sort, 58

stack, 52, 101

string, 22

structure, 37

subset, 375

sum operator, 377

test data, 10

time complexity, 83

total correctness, 7

travelling salesman problem, 149

tree, 146

trial division, 315

typedef, 24

union, 376

universal quantifier, 373

variable declaration, 21

vertex, 131

Visual Studio Code, 16

weighted graph, 135

winning position, 300